

D Programming Language Specification

This is the specification for the D Programming Language. For more information see dlang.org.

Contents

Contents	1
1 Introduction	3
2 Lexical	5
3 Grammar	27
4 Modules	71
5 Declarations	85
6 Types	100
7 Properties	107
8 Attributes	113
9 Pragmas	127
10 Expressions	131
11 Statements	169
12 Arrays	201
13 Associative Arrays	219
14 Structs and Unions	229
15 Classes	241

16 Interfaces	265
17 Enums	273
18 Type Qualifiers	279
19 Functions	287
20 Operator Overloading	333
21 Templates	353
22 Template Mixins	383
23 Contract Programming	391
24 Conditional Compilation	397
25 Traits	409
26 Error Handling	429
27 Unit Tests	433
28 Garbage Collection	437
29 Floating Point	445
30 D x86 Inline Assembler	449
31 Embedded Documentation	463
32 Interfacing to C	483
33 Interfacing to C++	491
34 Interfacing to Objective-C	507
35 Portability Guide	513
36 Named Character Entities	515
37 Memory Safety	523

<i>CONTENTS</i>	3
38 Application Binary Interface	525
39 Vector Extensions	545

Chapter 1

Introduction

The D programming language is a general purpose systems programming language. To that end, a D program is a collection of modules that can be compiled separately to native code that is combined with libraries and compiled C code by a linker to create a native executable.

Phases of Compilation

The process of compiling is divided into multiple phases. Each phase has no dependence on subsequent phases. For example, the scanner is not perturbed by the semantic analyzer. This separation of the passes makes language tools like syntax directed editors relatively easy to produce. It also is possible to compress D source by storing it in ‘tokenized’ form.

1. **source character set**

The source file is checked to see what character set it is, and the appropriate scanner is loaded. ASCII and UTF formats are accepted.

2. **script line**

If the first line starts with "#!", then that line is ignored.

3. **lexical analysis**

The source file is divided up into a sequence of tokens. Special tokens are replaced with other tokens. *SpecialTokenSequences* are processed and removed.

4. **syntax analysis**

The sequence of tokens is parsed to form syntax trees.

5. **semantic analysis**

The syntax trees are traversed to declare variables, load symbol tables, assign types, and in general determine the meaning of the program.

6. **optimization**

Optimization is an optional pass that tries to rewrite the program in a semantically equivalent, but faster executing, version.

7. code generation

Instructions are selected from the target architecture to implement the semantics of the program. The typical result will be an object file, suitable for input to a linker.

Chapter 2

Lexical

The lexical analysis is independent of the syntax parsing and the semantic analysis. The lexical analyzer splits the source text up into tokens. The lexical grammar describes the syntax of those tokens. The grammar is designed to be suitable for high speed scanning and to make it easy to write a correct scanner for it. It has a minimum of special case rules and there is only one phase of translation. The tokens are readily recognizable by those familiar with C and C++.

Source Text

D source text can be in one of the following formats:

- ASCII
- UTF-8
- UTF-16BE
- UTF-16LE
- UTF-32BE
- UTF-32LE

UTF-8 is a superset of traditional 7-bit ASCII. One of the following UTF BOMs (Byte Order Marks) can be present at the beginning of the source text:

UTF Byte Order Marks	
Format	BOM
UTF-8	EF BB BF
UTF-16BE	FE FF
UTF-16LE	FF FE
UTF-32BE	00 00 FE FF
UTF-32LE	FF FE 00 00
ASCII	no BOM

If the source file does not start with a BOM, then the first character must be less than or equal to U+0000007F.

There are no digraphs or trigraphs in D.

The source text is decoded from its source representation into Unicode *Characters*. The *Characters* are further divided into: *WhiteSpace*, *EndOfLine*, *Comments*, *SpecialTokenSequences*, *Tokens*, all followed by *EndOfFile*.

The source text is split into tokens using the maximal munch technique, i.e., the lexical analyzer tries to make the longest token it can. For example >> is a right shift token, not two greater than tokens. There are two exceptions to this rule:

- A .. embedded inside what looks like two floating point literals, as in 1..2, is interpreted as if the .. was separated by a space from the first integer.
- A 1.a is interpreted as the three tokens 1, ., and a, while 1. a is interpreted as the two tokens 1. and a

Character Set

Character:

any Unicode character

End of File

EndOfFile:

physical end of the file

\u0000

\u001A

The source text is terminated by whichever comes first.

End of Line

```

EndOfLine:
    \u000D
    \u000A
    \u000D \u000A
    \u2028
    \u2029
    EndOfFile

```

There is no backslash line splicing, nor are there any limits on the length of a line.

White Space

```

WhiteSpace:
    Space
    Space WhiteSpace

```

```

Space:
    \u0020
    \u0009
    \u000B
    \u000C

```

Comments

```

Comment:
    BlockComment
    LineComment
    NestingBlockComment

```

```

BlockComment:
    /* Characters */

```

```

LineComment:
    // Characters EndOfLine

```

```

NestingBlockComment:
    /* NestingBlockCommentCharacters */

```

```

NestingBlockCommentCharacters:
    NestingBlockCommentCharacter
    NestingBlockCommentCharacter NestingBlockCommentCharacters

```

```

NestingBlockCommentCharacter:
    Character
    NestingBlockComment

```

```

Characters:
    Character
    Character Characters

```

D has three kinds of comments:

1. Block comments can span multiple lines, but do not nest.
2. Line comments terminate at the end of the line.
3. Nesting block comments can span multiple lines and can nest.

The contents of strings and comments are not tokenized. Consequently, comment openings occurring within a string do not begin a comment, and string delimiters within a comment do not affect the recognition of comment closings and nested `"+"` comment openings. With the exception of `"+"` occurring within a `"+"` comment, comment openings within a comment are ignored.

```

a = /* // + 1;    // parses as if 'a = 1;'
a = /* "+""_+/_1"/*; // parses as if 'a = " +/ 1;"'
a = /* /* + */ 3; // parses as if 'a = */ 3;'

```

Comments cannot be used as token concatenators, for example, `abc/**/def` is two tokens, `abc` and `def`, not one `abcdef` token.

Tokens

Token:

<i>Identifier</i>	<i>FloatLiteral</i>	.	<code>&=</code>
<i>StringLiteral</i>	<i>Keyword</i>	<code>..</code>	<code>&&</code>
<i>CharacterLiteral</i>	<code>/</code>	<code>...</code>	<code> </code>
<i>IntegerLiteral</i>	<code>/=</code>	<code>&</code>	<code>= </code>

	>=)	%
-	>>=	[%=
-=	>>>=]	^
--	>>	{	^=
+	>>>	}	^^
+=	!	?	^^=
++	!=	,	~
<	!<>	;	~=
<=	!<>=	:	@
<<	!<	\$	=>
<<=	!<=	=	#
<>	!>	==	
<>=	!>=	*	
>	(*=	

Identifiers

Identifier:

IdentifierStart

IdentifierStart IdentifierChars

IdentifierChars:

IdentifierChar

IdentifierChar IdentifierChars

IdentifierStart:

-

Letter

UniversalAlpha

IdentifierChar:

IdentifierStart

0

NonZeroDigit

Identifiers start with a letter, `_`, or universal alpha, and are followed by any number of letters,

_, digits, or universal alphas. Universal alphas are as defined in ISO/IEC 9899:1999(E) Appendix D. (This is the C99 Standard.) Identifiers can be arbitrarily long, and are case sensitive. Identifiers starting with `__` (two underscores) are reserved.

String Literals

StringLiteral:

WysiwygString
AlternateWysiwygString
DoubleQuotedString

HexString
DelimitedString
TokenString

WysiwygString:

r" *WysiwygCharacters* " *StringPostfix*_{opt}

AlternateWysiwygString:

' *WysiwygCharacters* ' *StringPostfix*_{opt}

WysiwygCharacters:

WysiwygCharacter
WysiwygCharacter WysiwygCharacters

WysiwygCharacter:

Character
EndOfLine

DoubleQuotedString:

" *DoubleQuotedCharacters* " *StringPostfix*_{opt}

DoubleQuotedCharacters:

DoubleQuotedCharacter
DoubleQuotedCharacter DoubleQuotedCharacters

DoubleQuotedCharacter:

Character
EscapeSequence
EndOfLine

EscapeSequence:

```

\ '
\ "
\ ?
\ \
\ 0
\ a
\ b
\ f
\ n
\ r
\ t
\ v
\ x HexDigit HexDigit
\ OctalDigit
\ OctalDigit OctalDigit
\ OctalDigit OctalDigit OctalDigit
\ u HexDigit HexDigit HexDigit HexDigit
\ U HexDigit HexDigit HexDigit HexDigit HexDigit HexDigit HexDigit HexDigit
\ NamedCharacterEntity

```

HexString:

```

x" HexStringChars " StringPostfixopt

```

HexStringChars:

```

HexStringChar
HexStringChar HexStringChars

```

HexStringChar:

```

HexDigit
WhiteSpace
EndOfLine

```

StringPostfix:

```

c
w
d

```

DelimitedString:

```
q" Delimiter WysiwygCharacters MatchingDelimiter "
```

```
TokenString:
  q{ Tokens }
```

A string literal is either a double quoted string, a wysiwyg quoted string, a delimited string, a token string, or a hex string.

In all string literal forms, an *EndOfLine* is regarded as a single `\n` character.

Wysiwyg Strings

Wysiwyg ("what you see is what you get") quoted strings are enclosed by `r"` and `"`. All characters between the `r"` and `"` are part of the string. There are no escape sequences inside `r" "`:

```
r"hello"
r"c:\root\foo.exe"
r"ab\n" // string is 4 characters,
        // 'a', 'b', '\', 'n'
```

An alternate form of wysiwyg strings are enclosed by backquotes, the ``` character. The ``` character is not available on some keyboards and the font rendering of it is sometimes indistinguishable from the regular `'` character. Since, however, the ``` is rarely used, it is useful to delineate strings with `"` in them.

```
`hello`
`c:\root\foo.exe`
`The "lazy" dog`
`a"b\n'//string is 5 characters,
      // 'a', 'b', '\', 'n'
```

Double Quoted Strings

Double quoted strings are enclosed by `"`. Escape sequences can be embedded into them with the typical `\` notation.

```
"hello"
"c:\\root\\foo.exe"
"ab\n" // string is 3 characters,
        // 'a', 'b', and a linefeed
"ab"
```



```
"           // string is 3 characters,
           // 'a', 'b', and a newline
```

Hex Strings

Hex strings allow string literals to be created using hex data. The hex data need not form valid UTF characters.

```
x"0A"           // same as "\x0A"
x"00_FBCD_32FD_0A" // same as
                  // "\x00\xFB\xCD\x32\xFD\x0A"
```

Whitespace and newlines are ignored, so the hex data can be easily formatted. The number of hex characters must be a multiple of 2.

Adjacent strings are concatenated with the `~` operator:

```
"hello_" ~ "world" ~ "\n" // forms the string
           // 'h','e','l','l','o',' ','\n',
           // 'w','o','r','l','d',newline
```

The following are all equivalent:

```
"ab" ~ "c"
r"ab" ~ r"c"
r"a" ~ "bc"
"a" ~ "b" ~ "c"
```

The optional *StringPostfix* character gives a specific type to the string, rather than it being inferred from the context. This is useful when the type cannot be unambiguously inferred, such as when overloading based on string type. The types corresponding to the postfix characters are:

String Literal Postfix Characters		
Postfix	Type	Aka
c	<code>immutable(char)[]</code>	string
w	<code>immutable(wchar)[]</code>	wstring
d	<code>immutable(dchar)[]</code>	dstring

```
"hello"c // string
"hello"w // wstring
"hello"d // dstring
```

The string literals are assembled as UTF-8 char arrays, and the postfix is applied to convert to `wchar` or `dchar` as necessary as a final step.

String literals are read only. Writes to string literals cannot always be detected, but cause undefined behavior.

Delimited Strings

Delimited strings use various forms of delimiters. The delimiter, whether a character or identifier, must immediately follow the `"` without any intervening whitespace. The terminating delimiter must immediately precede the closing `"` without any intervening whitespace. A *nesting delimiter* nests, and is one of the following characters:

Nesting Delimiters	
Delimiter	Matching Delimiter
[]
()
<	>
{	}

```
q"(foo(xxx))"    // "foo(xxx)"
q"[foo{"         // "foo{"
```

If the delimiter is an identifier, the identifier must be immediately followed by a newline, and the matching delimiter is the same identifier starting at the beginning of the line:

```
writeln(q"EOS
This
is_a_multi-line
heredoc_string
EOS"
);
```

The newline following the opening identifier is not part of the string, but the last newline before the closing identifier is part of the string. The closing identifier must be placed on its own line at the leftmost column.

Otherwise, the matching delimiter is the same as the delimiter character:

```
q"/foo]/"        // "foo]"
// q"/abc/def/"  // error
```

Token Strings

Token strings open with the characters `q{` and close with the token `}`. In between must be valid D tokens. The `{` and `}` tokens nest. The string is formed of all the characters between the opening and closing of the token string, including comments.

```
q{foo}           // "foo"
q{ /* */ }       // " /* */ "
q{ foo(q{hello}); } // " foo(q{hello}); "
q{ __TIME__ }     // " __TIME__ "
    // i.e. it is not replaced with the time
// q{ __EOF__ }   // error
    // __EOF__ is not a token, it's end of file
```

Escape Sequences

The following table explains the meaning of the escape sequences listed in *EscapeSequence*:

Escape Sequences	
Sequence	Meaning
<code>\'</code>	Literal single-quote: <code>'</code>
<code>\"</code>	Literal double-quote: <code>"</code>
<code>\?</code>	Literal question mark: <code>?</code>
<code>\\</code>	Literal backslash: <code>\</code>
<code>\0</code>	Binary zero (NUL, U+0000).
<code>\a</code>	BEL (alarm) character (U+0007).
<code>\b</code>	Backspace (U+0008).
<code>\f</code>	Form feed (FF) (U+000C).
<code>\n</code>	End-of-line (U+000A).
<code>\r</code>	Carriage return (U+000D).
<code>\t</code>	Horizontal tab (U+0009).
<code>\v</code>	Vertical tab (U+000B).
<code>\xnn</code>	Byte value in hexadecimal, where <i>nn</i> is specified as two hexadecimal digits. For example: <code>\xFF</code> represents the character with the value 255.
<code>\n\nn\nnn</code>	Byte value in octal. For example: <code>\775</code> represents the character with the value 509.
<code>\unnnn</code>	Unicode character U+ <i>nnnn</i> , where <i>nnnn</i> are four hexadecimal digits. For example, <code>\u042F</code> represents the character with the value 1071.
<code>\Unnnnnnnn</code>	Unicode character U+ <i>nnnnnnnn</i> , where <i>nnnnnnnn</i> are 8 hexadecimal digits. For example, <code>\U0000000A</code> represents the character with the value 10.
<code>\name</code>	Named character entity from the HTML5 specification. See <i>NamedCharacterEntity</i> for more details.

Character Literals

CharacterLiteral:

' SingleQuotedCharacter '

SingleQuotedCharacter:

Character

EscapeSequence

Character literals are a single character or escape sequence enclosed by single quotes, *'␣'*.

Integer Literals

IntegerLiteral:

Integer

Integer IntegerSuffix

Integer:

DecimalInteger

BinaryInteger

HexadecimalInteger

IntegerSuffix:

L

u

U

Lu

LU

uL

UL

DecimalInteger:

0

NonZeroDigit

NonZeroDigit DecimalDigitsUS

BinaryInteger:

BinPrefix BinaryDigitsUS

BinPrefix:

0b

0B

HexadecimalInteger:

HexPrefix HexDigitsNoSingleUS

NonZeroDigit:

1

2

3

4

5

6

7

8

9

DecimalDigits:

DecimalDigit

DecimalDigit DecimalDigits

DecimalDigitsUS:

DecimalDigitUS

DecimalDigitUS DecimalDigitsUS

DecimalDigitsNoSingleUS:

DecimalDigit

DecimalDigit DecimalDigitsUS

DecimalDigitsUS DecimalDigit

DecimalDigitsNoStartingUS:

DecimalDigit

DecimalDigit DecimalDigitsUS

DecimalDigit:

0

NonZeroDigit

DecimalDigitUS:

DecimalDigit

-

BinaryDigitsUS:
 BinaryDigitUS
 BinaryDigitUS BinaryDigitsUS

BinaryDigit:
 0
 1

BinaryDigitUS:
 BinaryDigit
 -

OctalDigit:
 0
 1
 2
 3
 4
 5
 6
 7

HexDigits:
 HexDigit
 HexDigit HexDigits

HexDigitsUS:
 HexDigitUS
 HexDigitUS HexDigitsUS

HexDigitsNoSingleUS:
 HexDigit
 HexDigit HexDigitsUS
 HexDigitsUS HexDigit

HexDigitsNoStartingUS:
 HexDigit
 HexDigit HexDigitsUS

HexDigit:

```

    DecimalDigit
    HexLetter

HexDigitUS:
    HexDigit
    -

```

```

HexLetter:
    a
    b
    c
    d
    e
    f
    A
    B
    C
    D
    E
    F

```

Integers can be specified in decimal, binary, or hexadecimal.

Decimal integers are a sequence of decimal digits.

Binary integers are a sequence of binary digits preceded by a ‘0b’ or ‘0B’.

C-style octal integer notation was deemed too easy to mix up with decimal notation; it is only fully supported in string literals. D still supports octal integer literals interpreted at compile time through the `std.conv.octal` template, as in `octal!167`.

Hexadecimal integers are a sequence of hexadecimal digits preceded by a ‘0x’ or ‘0X’.

Integers can have embedded ‘_’ characters, which are ignored. The embedded ‘_’ are useful for formatting long literals, such as using them as a thousands separator:

```

123_456          // 123456
1_2_3_4_5_6_    // 123456

```

Integers can be immediately followed by one ‘L’ or one of ‘u’ or ‘U’ or both. Note that there is no ‘l’ suffix.

The type of the integer is resolved as follows:

Decimal Literal Types		
Literal	Type	
<i>Usual decimal notation</i>		
0 .. 2_147_483_647	int	
2_147_483_648 .. 9_223_372_036_854_775_807	long	
<i>Explicit suffixes</i>		
0L .. 9_223_372_036_854_775_807L	long	
0U .. 4_294_967_295U	uint	
4_294_967_296U .. 18_446_744_073_709_551_615U		ulong
0UL .. 18_446_744_073_709_551_615UL	ulong	
<i>Hexadecimal notation</i>		
0x0 .. 0x7FFF_FFFF	int	
0x8000_0000 .. 0xFFFF_FFFF	uint	
0x1_0000_0000 .. 0x7FFF_FFFF_FFFF_FFFF	long	
0x8000_0000_0000_0000 .. 0xFFFF_FFFF_FFFF_FFFF		ulong
<i>Hexadecimal notation with explicit suffixes</i>		
0x0L .. 0x7FFF_FFFF_FFFF_FFFFL	long	
0x8000_0000_0000_0000L .. 0xFFFF_FFFF_FFFF_FFFFL		ulong
0x0U .. 0xFFFF_FFFFU	uint	
0x1_0000_0000U .. 0xFFFF_FFFF_FFFF_FFFFU		ulong
0x0UL .. 0xFFFF_FFFF_FFFF_FFFFUL	ulong	

Floating Point Literals

FloatLiteral:

Float

Float Suffix

Integer ImaginarySuffix

Integer FloatSuffix ImaginarySuffix

Integer RealSuffix ImaginarySuffix

Float:

DecimalFloat

HexFloat

DecimalFloat:

LeadingDecimal .

LeadingDecimal . DecimalDigits

DecimalDigits . DecimalDigitsNoStartingUS DecimalExponent
. DecimalInteger
. DecimalInteger DecimalExponent
LeadingDecimal DecimalExponent

DecimalExponent

DecimalExponentStart DecimalDigitsNoSingleUS

DecimalExponentStart

e
E
e+
E+
e-
E-

HexFloat:

HexPrefix HexDigitsNoSingleUS . HexDigitsNoStartingUS HexExponent
HexPrefix . HexDigitsNoStartingUS HexExponent
HexPrefix HexDigitsNoSingleUS HexExponent

HexPrefix:

0x
0X

HexExponent:

HexExponentStart DecimalDigitsNoSingleUS

HexExponentStart:

p
P
p+
P+
p-
P-

Suffix:

FloatSuffix
RealSuffix

```

    ImaginarySuffix
    FloatSuffix ImaginarySuffix
    RealSuffix ImaginarySuffix

```

```

FloatSuffix:
    f
    F

```

```

RealSuffix:
    L

```

```

ImaginarySuffix:
    i

```

```

LeadingDecimal:
    DecimalInteger
    0 DecimalDigitsNoSingleUS

```

Floats can be in decimal or hexadecimal format.

Hexadecimal floats are preceded by a **0x** or **0X** and the exponent is a **p** or **P** followed by a decimal number serving as the exponent of 2.

Floating literals can have embedded ‘_’ characters, which are ignored. The embedded ‘_’ are useful for formatting long literals to make them more readable, such as using them as a thousands separator:

```

123_456.567_8      // 123456.5678
1_2_3_4_5_6_.5_6_7_8 // 123456.5678
1_2_3_4_5_6_.5e-6_ // 123456.5e-6

```

Floating literals with no suffix are of type double. Floats can be followed by one **f**, **F**, or **L** suffix. The **f** or **F** suffix means it is a float, and **L** means it is a real.

If a floating literal is followed by **i**, then it is an *ireal* (imaginary) type.

Examples:

```

0x1.FFFFFFFFFFFFFp1023 // double.max
0x1p-52                // double.epsilon
1.175494351e-38F       // float.min
6.3i                   // idouble 6.3
6.3fi                  // ifloat 6.3
6.3Li                  // ireal 6.3

```

It is an error if the literal exceeds the range of the type. It is not an error if the literal is rounded to fit into the significant digits of the type.

Complex literals are not tokens, but are assembled from real and imaginary expressions during semantic analysis:

```
4.5 + 6.2i  // complex number (phased out)
```

Keywords are reserved identifiers.

See Also: Globally Defined *Symbols*.

Keyword:

abstract		↪	module
alias	dchar	function	
align	debug		new
asm	default	goto	nothrow
assert	delegate		null
auto	delete↯	idouble	
	↪ (deprecated)	if	out
body	deprecated	ifloat	override
bool	do	immutable	
break	double	import	package
byte		in	pragma
	else	inout	private
case	enum	int	protected
cast	export	interface	public
catch	extern	invariant	pure
cdouble		ireal	
cent	false	is	real
cfloat	final		ref
char	finally	lazy	return
class	float	long	
const	for		scope
continue	foreach	macro (unused)	shared
creal	foreach_reverse↯	mixin	short

<code>static</code>	↳ (deprecated)	<code>version</code>	<code>__MODULE__</code>
<code>struct</code>	<code>typeid</code>	<code>void</code>	<code>__LINE__</code>
<code>super</code>	<code>typeof</code>	<code>volatile</code> ↗	<code>__FUNCTION__</code>
<code>switch</code>		↳ (deprecated)	<code>__PRETTY_FUNCTION__</code> ↗
<code>synchronized</code>	<code>ubyte</code>		↳
	<code>ucnt</code>	<code>wchar</code>	
<code>template</code>	<code>uint</code>	<code>while</code>	<code>__gshared</code>
<code>this</code>	<code>ulong</code>	<code>with</code>	<code>__traits</code>
<code>throw</code>	<code>union</code>		<code>__vector</code>
<code>true</code>	<code>unittest</code>	<code>__FILE__</code>	<code>__parameters</code>
<code>try</code>	<code>ushort</code>	<code>__FILE_FULL_PATH__</code> ↗	
<code>typedef</code> ↗		↳	

Globally Defined Symbols

These are defined in `object_.d`, which is automatically imported by the default implementation.

Symbols:

<code>string</code> ↗	↳ (alias to <code>immutable(wchar)[]</code>)
↳ (alias to <code>immutable(char)[]</code>)↗	<code>size_t</code>
<code>wstring</code> ↗	↳ (alias to <code>immutable(dchar)[]</code>)↗

Special Tokens

These tokens are replaced with other tokens according to the following table:

Special Tokens

Special Token	Replaced with
<code>__DATE__</code>	string literal of the date of compilation " <i>mmm dd yyyy</i> "
<code>__EOF__</code>	sets the scanner to the end of the file
<code>__TIME__</code>	string literal of the time of compilation " <i>hh:mm:ss</i> "
<code>__TIMESTAMP__</code>	string literal of the date and time of compilation " <i>www mmm dd hh:mm:ss yyyy</i> "
<code>__VENDOR__</code>	Compiler vendor string, such as "Digital Mars D"
<code>__VERSION__</code>	Compiler version as an integer, such as 2001

Special Token Sequences

SpecialTokenSequence:

```
# line IntegerLiteral EndOfLine
# line IntegerLiteral Filespec EndOfLine
```

Filespec:

```
" Characters "
```

Special token sequences are processed by the lexical analyzer, may appear between any other tokens, and do not affect the syntax parsing.

There is currently only one special token sequence, `#line`.

This sets the source line number to *IntegerLiteral*, and optionally the source file name to *Filespec*, beginning with the next line of source text. The source file and line number is used for printing error messages and for mapping generated code back to the source for the symbolic debugging output.

For example:

```
int #line 6 "foo\bar"
x;  // this is now line 6 of file foo\bar
```

Note that the backslash character is not treated specially inside *Filespec* strings.

Chapter 3

Grammar

Lexical Syntax

Refer to the page for lexical syntax.

Type

Type:

*TypeCtors*_{opt} *BasicType* *BasicType2*_{opt}

TypeCtors:

TypeCtor

TypeCtor *TypeCtors*

TypeCtor:

`const`

`immutable`

`inout`

`shared`

BasicType:

BasicTypeX

. IdentifierList

IdentifierList

Typeof

Typeof . IdentifierList

TypeCtor (Type)

TypeVector

BasicTypeX:

```

bool
byte
ubyte
short
ushort
int
uint
long
ulong
char
wchar
dchar
float
double
real
ifloat
idouble
ireal
cfloat
cdouble
creal
void

```

BasicType2:

```

BasicType2X BasicType2opt

```

BasicType2X:

```

*
[ ]
[ AssignExpression ]
[ AssignExpression .. AssignExpression ]
[ Type ]
delegate Parameters MemberFunctionAttributesopt
function Parameters FunctionAttributesopt

```

IdentifierList:

```

Identifier
Identifier . IdentifierList

```



```

TemplateInstance
TemplateInstance . IdentifierList
Identifier [ AssignExpression ]. IdentifierList

```

```

Typeof:
  typeof ( Expression )
  typeof ( return )

```

```

TypeVector:
  __vector ( Type )

```

Expression

```

Expression:
  CommaExpression

```

```

CommaExpression:
  AssignExpression
  AssignExpression , CommaExpression

```

```

AssignExpression:
  ConditionalExpression
  ConditionalExpression = AssignExpression
  ConditionalExpression += AssignExpression
  ConditionalExpression -= AssignExpression
  ConditionalExpression *= AssignExpression
  ConditionalExpression /= AssignExpression
  ConditionalExpression %= AssignExpression
  ConditionalExpression &= AssignExpression
  ConditionalExpression |= AssignExpression
  ConditionalExpression ^= AssignExpression
  ConditionalExpression ~= AssignExpression
  ConditionalExpression <=< AssignExpression
  ConditionalExpression >>= AssignExpression
  ConditionalExpression >>>= AssignExpression
  ConditionalExpression ^= AssignExpression

```

```

ConditionalExpression:

```

OrOrExpression
OrOrExpression ? *Expression* : *ConditionalExpression*

OrOrExpression:
 AndAndExpression
 OrOrExpression || *AndAndExpression*

AndAndExpression:
 OrExpression
 AndAndExpression && *OrExpression*
 CmpExpression
 AndAndExpression && *CmpExpression*

OrExpression:
 XorExpression
 OrExpression | *XorExpression*

XorExpression:
 AndExpression
 XorExpression ^ *AndExpression*

AndExpression:
 ShiftExpression
 AndExpression & *ShiftExpression*

CmpExpression:
 ShiftExpression
 EqualExpression
 IdentityExpression
 RelExpression
 InExpression

EqualExpression:
 ShiftExpression == *ShiftExpression*

ShiftExpression != *ShiftExpression*

IdentityExpression:

ShiftExpression is *ShiftExpression*
ShiftExpression !is *ShiftExpression*

RelExpression:

ShiftExpression < *ShiftExpression*
ShiftExpression <= *ShiftExpression*
ShiftExpression > *ShiftExpression*
ShiftExpression >= *ShiftExpression*
ShiftExpression !<= *ShiftExpression*
ShiftExpression !< > *ShiftExpression*
ShiftExpression <> *ShiftExpression*
ShiftExpression <=> *ShiftExpression*
ShiftExpression !> *ShiftExpression*
ShiftExpression !>= *ShiftExpression*
ShiftExpression !< *ShiftExpression*
ShiftExpression !<= *ShiftExpression*

InExpression:

ShiftExpression in *ShiftExpression*
ShiftExpression !in *ShiftExpression*

ShiftExpression:

AddExpression
ShiftExpression << *AddExpression*
ShiftExpression >> *AddExpression*
ShiftExpression >>> *AddExpression*

AddExpression:

MulExpression
AddExpression + *MulExpression*
AddExpression - *MulExpression*
CatExpression

CatExpression:

AddExpression ~ MulExpression

MulExpression:

UnaryExpression

*MulExpression * UnaryExpression*

MulExpression / UnaryExpression

MulExpression % UnaryExpression

UnaryExpression:

& UnaryExpression

++ UnaryExpression

-- UnaryExpression

** UnaryExpression*

- UnaryExpression

+ UnaryExpression

! UnaryExpression

ComplementExpression

(Type) . Identifier

(Type) . TemplateInstance

DeleteExpression

CastExpression

PowExpression

ComplementExpression:

~ UnaryExpression

NewExpression:

new AllocatorArguments_{opt} Type

NewExpressionWithArgs

NewExpressionWithArgs:

new AllocatorArguments_{opt} Type [AssignExpression]

new AllocatorArguments_{opt} Type (ArgumentList_{opt})

NewAnonClassExpression

AllocatorArguments:
 (*ArgumentList*_{opt})

ArgumentList:
 AssignExpression
 AssignExpression ,
 AssignExpression , *ArgumentList*

NewAnonClassExpression:
 new *AllocatorArguments*_{opt} class *ClassArguments*_{opt} *SuperClass*_{opt} *Interfaces*_{opt} *AggregateBody*

ClassArguments:
 (*ArgumentList*_{opt})

DeleteExpression:
 delete *UnaryExpression*

CastExpression:
 cast (*Type*) *UnaryExpression*
 cast (*TypeCtors*_{opt}) *UnaryExpression*

PowExpression:
 PostfixExpression
 PostfixExpression ^^ *UnaryExpression*

PostfixExpression:
 PrimaryExpression
 PostfixExpression . *Identifier*
 PostfixExpression . *TemplateInstance*
 PostfixExpression . *NewExpression*
 PostfixExpression ++
 PostfixExpression --
 PostfixExpression (*ArgumentList*_{opt})
 *TypeCtors*_{opt} *BasicType* (*ArgumentList*_{opt})
 IndexExpression
 SliceExpression

IndexExpression:

PostfixExpression [*ArgumentList*]

SliceExpression:

PostfixExpression []

PostfixExpression [*Slice* ,_{opt}]

Slice:

AssignExpression

AssignExpression , *Slice*

AssignExpression .. *AssignExpression*

AssignExpression .. *AssignExpression* , *Slice*

PrimaryExpression:

Identifier

. *Identifier*

TemplateInstance

. *TemplateInstance*

this

super

null

true

false

\$

IntegerLiteral

FloatLiteral

CharacterLiteral

StringLiterals

ArrayLiteral

AssocArrayLiteral

FunctionLiteral

AssertExpression

MixinExpression

ImportExpression

NewExpressionWithArgs

BasicTypeX . *Identifier*

Typeof

TypeidExpression
IsExpression
 (*Expression*)
TraitsExpression
SpecialKeyword

StringLiterals:
StringLiteral
StringLiterals StringLiteral

ArrayLiteral:
 [*ArgumentList*_{opt}]

AssocArrayLiteral:
 [*KeyValuePairs*]

KeyValuePairs:
KeyValuePair
KeyValuePair , *KeyValuePairs*

KeyValuePair:
KeyExpression : *ValueExpression*

KeyExpression:
AssignExpression

ValueExpression:
AssignExpression

FunctionLiteral:
function *Type*_{opt} *ParameterAttributes* _{opt} *FunctionLiteralBody*
delegate *Type*_{opt} *ParameterMemberAttributes* _{opt} *FunctionLiteralBody*
ParameterMemberAttributes *FunctionLiteralBody*
FunctionLiteralBody
Lambda

ParameterAttributes:

Parameters FunctionAttributes_{opt}

ParameterMemberAttributes:

Parameters MemberFunctionAttributes_{opt}

FunctionLiteralBody:

BlockStatement

FunctionContracts_{opt} BodyStatement

Lambda:

function Type_{opt} ParameterAttributes => AssignExpression

delegate Type_{opt} ParameterMemberAttributes => AssignExpression

ParameterMemberAttributes => AssignExpression

Identifier => AssignExpression

AssertExpression:

assert (AssignExpression ,_{opt})

assert (AssignExpression , AssignExpression ,_{opt})

MixinExpression:

mixin (AssignExpression)

ImportExpression:

import (AssignExpression)

TypeidExpression:

typeid (Type)

typeid (Expression)

IsExpression:

is (Type)

is (Type : TypeSpecialization)

is (Type == TypeSpecialization)

is (Type : TypeSpecialization , TemplateParameterList)

is (Type == TypeSpecialization , TemplateParameterList)


```

is ( Type Identifier )
is ( Type Identifier : TypeSpecialization )
is ( Type Identifier == TypeSpecialization )
is ( Type Identifier : TypeSpecialization , TemplateParameterList )
is ( Type Identifier == TypeSpecialization , TemplateParameterList )

```

TypeSpecialization:

```

Type
struct
union
class
interface
enum
function
delegate
super
const
immutable
inout
shared
return
__parameters

```

TraitsExpression:

```

__traits ( TraitsKeyword , TraitsArguments )

```

TraitsKeyword:

```

isAbstractClass
isArithmetic
isAssociativeArray
isFinalClass
isPOD
isNested
isFloating
isIntegral
isScalar
isStaticArray
isUnsigned
isVirtualFunction

```

isVirtualMethod
isAbstractFunction
isFinalFunction
isStaticFunction
isOverrideFunction
isRef
isOut
isLazy
hasMember
identifier
getAliasThis
getAttributes
getFunctionAttributes
getMember
getOverloads
getPointerBitmap
getProtection
getVirtualFunctions
getVirtualMethods
getUnitTests
parent
classInstanceSize
getVirtualIndex
allMembers
derivedMembers
isSame
compiles

TraitsArguments:

TraitsArgument

TraitsArgument , *TraitsArguments*

TraitsArgument:

AssignExpression

Type

SpecialKeyword:

__FILE__

__FILE_FULL_PATH__

```

__MODULE__
__LINE__
__FUNCTION__
__PRETTY_FUNCTION__

```

Statement

Statement:

```

;
NonEmptyStatement
ScopeBlockStatement

```

NoScopeNonEmptyStatement:

```

NonEmptyStatement
BlockStatement

```

NoScopeStatement:

```

;
NonEmptyStatement
BlockStatement

```

NonEmptyOrScopeBlockStatement:

```

NonEmptyStatement
ScopeBlockStatement

```

NonEmptyStatement:

```

NonEmptyStatementNoCaseNoDefault
CaseStatement
CaseRangeStatement
DefaultStatement

```

NonEmptyStatementNoCaseNoDefault:

```

LabeledStatement
ExpressionStatement
DeclarationStatement
IfStatement
WhileStatement
DoStatement
ForStatement

```

ForeachStatement
SwitchStatement
FinalSwitchStatement
ContinueStatement
BreakStatement
ReturnStatement
GotoStatement
WithStatement
SynchronizedStatement
TryStatement
ScopeGuardStatement
ThrowStatement
AsmStatement
PragmaStatement
MixinStatement
ForeachRangeStatement
ConditionalStatement
StaticAssert
TemplateMixin
ImportDeclaration

ScopeStatement:
 NonEmptyStatement
 BlockStatement

ScopeBlockStatement:
 BlockStatement

LabeledStatement:
 Identifier :
 Identifier : *NoScopeStatement*
 Identifier : *Statement*

BlockStatement:
 { }
 { *StatementList* }

StatementList:

Statement
Statement StatementList

ExpressionStatement:

Expression ;

DeclarationStatement:

StorageClasses_{opt} Declaration

IfStatement:

if (IfCondition) ThenStatement
if (IfCondition) ThenStatement else ElseStatement

IfCondition:

Expression
auto Identifier = Expression
TypeCtors Identifier = Expression
TypeCtors_{opt} BasicType Declarator = Expression

ThenStatement:

ScopeStatement

ElseStatement:

ScopeStatement

WhileStatement:

while (Expression) ScopeStatement

DoStatement:

do ScopeStatement while (Expression) ;

ForStatement:

for (Initialize Test_{opt} ; Increment_{opt}) ScopeStatement

```
Initialize:
    ;
    NoScopeNonEmptyStatement

Test:
    Expression

Increment:
    Expression

ForeachStatement:
    Foreach ( ForeachTypeList ; ForeachAggregate ) NoScopeNonEmptyStatement

Foreach:
    foreach
    foreach_reverse

ForeachTypeList:
    ForeachType
    ForeachType , ForeachTypeList

ForeachType:
    ForeachTypeAttributesopt BasicType Declarator
    ForeachTypeAttributesopt Identifier

ForeachTypeAttributes
    ForeachTypeAttribute
    ForeachTypeAttribute ForeachTypeAttributesopt

ForeachTypeAttribute:
    ref
    TypeCtor

ForeachAggregate:
    Expression

ForeachRangeStatement:
```

```

    Foreach ( ForeachType ; LwrExpression .. UprExpression ) ScopeStatement

LwrExpression:
    Expression

UprExpression:
    Expression

SwitchStatement:
    switch ( Expression ) ScopeStatement

CaseStatement:
    case ArgumentList : ScopeStatementList

CaseRangeStatement:
    case FirstExp : .. case LastExp : ScopeStatementList

FirstExp:
    AssignExpression

LastExp:
    AssignExpression

DefaultStatement:
    default : ScopeStatementList

ScopeStatementList:
    StatementListNoCaseNoDefault

StatementListNoCaseNoDefault:
    StatementNoCaseNoDefault
    StatementNoCaseNoDefault StatementListNoCaseNoDefault

StatementNoCaseNoDefault:
    ;
    NonEmptyStatementNoCaseNoDefault
    ScopeBlockStatement

```

FinalSwitchStatement:
`final switch` (*Expression*) *ScopeStatement*

ContinueStatement:
`continue` *Identifier*_{opt} ;

BreakStatement:
`break` *Identifier*_{opt} ;

ReturnStatement:
`return` *Expression*_{opt} ;

GotoStatement:
`goto` *Identifier* ;
`goto default` ;
`goto case` ;
`goto case` *Expression* ;

WithStatement:
`with` (*Expression*) *ScopeStatement*
`with` (*Symbol*) *ScopeStatement*
`with` (*TemplateInstance*) *ScopeStatement*

SynchronizedStatement:
`synchronized` *ScopeStatement*
`synchronized` (*Expression*) *ScopeStatement*

TryStatement:
`try` *ScopeStatement* *Catches*
`try` *ScopeStatement* *Catches* *FinallyStatement*
`try` *ScopeStatement* *FinallyStatement*

Catches:
Catch

Catch Catches

Catch:

catch (*CatchParameter*) *NoScopeNonEmptyStatement*

CatchParameter:

BasicType Identifier

FinallyStatement:

finally *NoScopeNonEmptyStatement*

ThrowStatement:

throw *Expression* ;

ScopeGuardStatement:

scope(exit) *NonEmptyOrScopeBlockStatement*

scope(success) *NonEmptyOrScopeBlockStatement*

scope(failure) *NonEmptyOrScopeBlockStatement*

AsmStatement:

asm *FunctionAttributes*_{opt} { *AsmInstructionList*_{opt} }

AsmInstructionList:

AsmInstruction ;

AsmInstruction ; *AsmInstructionList*

PragmaStatement:

Pragma *NoScopeStatement*

MixinStatement:

mixin (*AssignExpression*) ;

Iasm

AsmInstruction:

Identifier : AsmInstruction
align IntegerExpression
even
naked
db Operands
ds Operands
di Operands
dl Operands
df Operands
dd Operands
de Operands
db StringLiteral
ds StringLiteral
di StringLiteral
dl StringLiteral
dw StringLiteral
dq StringLiteral
Opcode
Opcode Operands

Operands:
Operand
Operand , Operands

IntegerExpression:
IntegerLiteral
Identifier

Register:
AL AH AX EAX
BL BH BX EBX
CL CH CX ECX
DL DH DX EDX
BP EBP
SP ESP
DI EDI
SI ESI
ES CS SS DS GS FS

CR0 CR2 CR3 CR4
 DR0 DR1 DR2 DR3 DR6 DR7
 TR3 TR4 TR5 TR6 TR7
 ST
 ST(0) ST(1) ST(2) ST(3) ST(4) ST(5) ST(6) ST(7)
 MM0 MM1 MM2 MM3 MM4 MM5 MM6 MM7
 XMM0 XMM1 XMM2 XMM3 XMM4 XMM5 XMM6 XMM7

Register64:

RAX RBX RCX RDX
 BPL RBP
 SPL RSP
 DIL RDI
 SIL RSI
 R8B R8W R8D R8
 R9B R9W R9D R9
 R10B R10W R10D R10
 R11B R11W R11D R11
 R12B R12W R12D R12
 R13B R13W R13D R13
 R14B R14W R14D R14
 R15B R15W R15D R15
 XMM8 XMM9 XMM10 XMM11 XMM12 XMM13 XMM14 XMM15
 YMM0 YMM1 YMM2 YMM3 YMM4 YMM5 YMM6 YMM7
 YMM8 YMM9 YMM10 YMM11 YMM12 YMM13 YMM14 YMM15

Operand:

AsmExp

AsmExp:

AsmLogOrExp

AsmLogOrExp ? *AsmExp* : *AsmExp*

AsmLogOrExp:

AsmLogAndExp

AsmLogOrExp || *AsmLogAndExp*

AsmLogAndExp:

AsmOrExp
AsmLogAndExp && *AsmOrExp*

AsmOrExp:
AsmXorExp
AsmOrExp | *AsmXorExp*

AsmXorExp:
AsmAndExp
AsmXorExp ^ *AsmAndExp*

AsmAndExp:
AsmEqualExp
AsmAndExp & *AsmEqualExp*

AsmEqualExp:
AsmRelExp
AsmEqualExp == *AsmRelExp*
AsmEqualExp != *AsmRelExp*

AsmRelExp:
AsmShiftExp
AsmRelExp < *AsmShiftExp*
AsmRelExp <= *AsmShiftExp*
AsmRelExp > *AsmShiftExp*
AsmRelExp >= *AsmShiftExp*

AsmShiftExp:
AsmAddExp
AsmShiftExp << *AsmAddExp*
AsmShiftExp >> *AsmAddExp*
AsmShiftExp >>> *AsmAddExp*

AsmAddExp:
AsmMulExp
AsmAddExp + *AsmMulExp*
AsmAddExp - *AsmMulExp*

AsmMulExp:
AsmBrExp

```

AsmMulExp * AsmBrExp
AsmMulExp / AsmBrExp
AsmMulExp % AsmBrExp

```

```

AsmBrExp:
  AsmUnaExp
  AsmBrExp [ AsmExp ]

```

```

AsmUnaExp:
  AsmTypePrefix AsmExp
  offsetof AsmExp
  seg AsmExp
  + AsmUnaExp
  - AsmUnaExp
  ! AsmUnaExp
  ~ AsmUnaExp
  AsmPrimaryExp

```

```

AsmPrimaryExp:
  IntegerLiteral
  FloatLiteral
  __LOCAL_SIZE
  $
  Register
  Register : AsmExp
  Register64
  Register64 : AsmExp
  DotIdentifier
  this

```

```

DotIdentifier:
  Identifier
  Identifier . DotIdentifier

```

```

AsmTypePrefix:
  near ptr
  far ptr
  byte ptr
  short ptr

```

```

int ptr
word ptr
dword ptr
qword ptr
float ptr
double ptr
real ptr

```

Declaration

Declaration:

```

FuncDeclaration
VarDeclarations
AliasDeclaration
AggregateDeclaration
EnumDeclaration
ImportDeclaration

```

AliasDeclaration:

```

alias StorageClassesopt BasicType Declarator ;
alias StorageClassesopt BasicType FuncDeclarator ;
alias AliasDeclarationX ;

```

AliasDeclarationX:

```

AliasDeclarationY
AliasDeclarationX , AliasDeclarationY

```

AliasDeclarationY:

```

Identifier TemplateParametersopt = StorageClassesopt Type
Identifier TemplateParametersopt = FunctionLiteral

```

AutoDeclaration:

```

StorageClasses AutoDeclarationX ;

```

AutoDeclarationX:

```

AutoDeclarationY
AutoDeclarationX , AutoDeclarationY

```

AutoDeclarationY:

Identifier TemplateParameters_{opt} = Initializer

VarDeclarations:

StorageClasses_{opt} BasicType Declarators ;

AutoDeclaration

Declarators:

DeclaratorInitializer

DeclaratorInitializer , DeclaratorIdentifierList

DeclaratorInitializer:

VarDeclarator

VarDeclarator TemplateParameters_{opt} = Initializer

AltDeclarator

AltDeclarator = Initializer

DeclaratorIdentifierList:

DeclaratorIdentifier

DeclaratorIdentifier , DeclaratorIdentifierList

DeclaratorIdentifier:

VarDeclaratorIdentifier

AltDeclaratorIdentifier

VarDeclaratorIdentifier:

Identifier

Identifier TemplateParameters_{opt} = Initializer

AltDeclaratorIdentifier:

BasicType2 Identifier AltDeclaratorSuffixes_{opt}

BasicType2 Identifier AltDeclaratorSuffixes_{opt} = Initializer

BasicType2_{opt} Identifier AltDeclaratorSuffixes

BasicType2_{opt} Identifier AltDeclaratorSuffixes = Initializer

Declarator:

VarDeclarator

AltDeclarator

VarDeclarator:

BasicType2_{opt} Identifier

AltDeclarator:

BasicType2_{opt} Identifier AltDeclaratorSuffixes

BasicType2_{opt} (AltDeclaratorX)

BasicType2_{opt} (AltDeclaratorX) AltFuncDeclaratorSuffix

BasicType2_{opt} (AltDeclaratorX) AltDeclaratorSuffixes

AltDeclaratorX:

BasicType2_{opt} Identifier

BasicType2_{opt} Identifier AltFuncDeclaratorSuffix

AltDeclarator

AltDeclaratorSuffixes:

AltDeclaratorSuffix

AltDeclaratorSuffix AltDeclaratorSuffixes

AltDeclaratorSuffix:

[]

[AssignExpression]

[Type]

AltFuncDeclaratorSuffix:

Parameters MemberFunctionAttributes_{opt}

StorageClasses:

StorageClass

StorageClass StorageClasses

StorageClass:

LinkageAttribute

AlignAttribute

deprecated

enum

static

extern

abstract


```

final
override
synchronized
auto
scope
const
immutable
inout
shared
__gshared
Property
nothrow
pure
ref

```

Initializer:

```

    VoidInitializer
    NonVoidInitializer

```

VoidInitializer:

```

    void

```

NonVoidInitializer:

```

    ExpInitializer:
    ArrayInitializer
    StructInitializer

```

ExpInitializer:

```

    AssignExpression

```

ArrayInitializer:

```

    [ ArrayMemberInitializationsopt ]

```

ArrayMemberInitializations:

```

    ArrayMemberInitialization
    ArrayMemberInitialization ,
    ArrayMemberInitialization , ArrayMemberInitializations

```

ArrayMemberInitialization:

NonVoidInitializer
AssignExpression : *NonVoidInitializer*

StructInitializer:
 { *StructMemberInitializers*_{opt} }

StructMemberInitializers:
StructMemberInitializer
StructMemberInitializer ,
StructMemberInitializer , *StructMemberInitializers*

StructMemberInitializer:
NonVoidInitializer
Identifier : *NonVoidInitializer*

Function

FuncDeclaration:
*StorageClasses*_{opt} *BasicType* *FuncDeclarator* *FunctionBody*
AutoFuncDeclaration

AutoFuncDeclaration:
StorageClasses *Identifier* *FuncDeclaratorSuffix* *FunctionBody*

FuncDeclarator:
*BasicType2*_{opt} *Identifier* *FuncDeclaratorSuffix*

FuncDeclaratorSuffix:
Parameters *MemberFunctionAttributes*_{opt}
TemplateParameters *Parameters* *MemberFunctionAttributes*_{opt} *Constraint*_{opt}

Parameters:
 (*ParameterList*_{opt})

ParameterList:
Parameter
Parameter , *ParameterList*
 ...

Parameter:

*InOut*_{opt} *BasicType* *Declarator*
*InOut*_{opt} *BasicType* *Declarator* ...
*InOut*_{opt} *BasicType* *Declarator* = *AssignExpression*
*InOut*_{opt} *Type*
*InOut*_{opt} *Type* ...

InOut:

InOutX
InOut InOutX

InOutX:

auto
TypeCtor
final
in
lazy
out
ref
return ref
scope

FunctionAttributes:

FunctionAttribute
FunctionAttribute FunctionAttributes

FunctionAttribute:

nothrow
pure
Property

MemberFunctionAttributes:

MemberFunctionAttribute
MemberFunctionAttribute MemberFunctionAttributes

MemberFunctionAttribute:

const
immutable
inout

```

return
shared
FunctionAttribute

```

```

FunctionBody:
    BlockStatement
    FunctionContractsopt BodyStatement
    FunctionContracts

```

```

FunctionContracts:
    InStatement OutStatementopt
    OutStatement InStatementopt

```

```

InStatement:
    in BlockStatement

```

```

OutStatement:
    out BlockStatement
    out ( Identifier ) BlockStatement

```

```

BodyStatement:
    body BlockStatement

```

```

Constructor:
    this Parameters MemberFunctionAttributesopt ;
    this Parameters MemberFunctionAttributesopt FunctionBody
    ConstructorTemplate

```

```

ConstructorTemplate:
    this TemplateParameters Parameters MemberFunctionAttributesopt Constraintopt ; ↵
↵
    this TemplateParameters Parameters MemberFunctionAttributesopt Constraintopt FunctionBody

```

```

Destructor:
    ~ this ( ) MemberFunctionAttributesopt ;
    ~ this ( ) MemberFunctionAttributesopt FunctionBody

```

```

Postblit:

```

```

    this ( this ) MemberFunctionAttributesopt ;
    this ( this ) MemberFunctionAttributesopt FunctionBody

```

Allocator:

```

    new Parameters ;
    new Parameters FunctionBody

```

Deallocator:

```

    delete Parameters ;
    delete Parameters FunctionBody

```

Invariant:

```

    invariant ( ) BlockStatement
    invariant BlockStatement

```

UnitTest:

```

    unittest BlockStatement

```

StaticConstructor:

```

    static this ( ) MemberFunctionAttributesopt ;
    static this ( ) MemberFunctionAttributesopt FunctionBody

```

StaticDestructor:

```

    static ~ this ( ) MemberFunctionAttributesopt ;
    static ~ this ( ) MemberFunctionAttributesopt FunctionBody

```

SharedStaticConstructor:

```

    shared static this ( ) MemberFunctionAttributesopt ;
    shared static this ( ) MemberFunctionAttributesopt FunctionBody

```

SharedStaticDestructor:

```

    shared static ~ this ( ) MemberFunctionAttributesopt ;
    shared static ~ this ( ) MemberFunctionAttributesopt FunctionBody

```

Aggregate

AggregateDeclaration:

```

    ClassDeclaration

```

InterfaceDeclaration
StructDeclaration
UnionDeclaration

ClassDeclaration:

class Identifier ;
class Identifier BaseClassList_{opt} AggregateBody
ClassTemplateDeclaration

ClassTemplateDeclaration:

class Identifier TemplateParameters Constraint_{opt} BaseClassList_{opt} AggregateBody
class Identifier TemplateParameters BaseClassList Constraint AggregateBody

InterfaceDeclaration:

interface Identifier ;
interface Identifier BaseInterfaceList_{opt} AggregateBody
InterfaceTemplateDeclaration

InterfaceTemplateDeclaration:

interface Identifier TemplateParameters Constraint_{opt} BaseInterfaceList_{opt} AggregateBody
interface Identifier TemplateParameters BaseInterfaceList Constraint AggregateBody

StructDeclaration:

struct Identifier ;
struct Identifier AggregateBody
StructTemplateDeclaration
AnonStructDeclaration

StructTemplateDeclaration:

struct Identifier TemplateParameters Constraint_{opt} AggregateBody

AnonStructDeclaration:

struct AggregateBody

UnionDeclaration:

union Identifier ;
union Identifier AggregateBody
UnionTemplateDeclaration
AnonUnionDeclaration

UnionTemplateDeclaration:

union Identifier TemplateParameters Constraint_{opt} AggregateBody

AnonUnionDeclaration:

union AggregateBody

AggregateBody:

{ DeclDefs_{opt} }

BaseClassList:

: SuperClass

: SuperClass , Interfaces

: Interfaces

BaseInterfaceList:

: Interfaces

SuperClass:

BasicType

Interfaces:

Interface

Interface , Interfaces

Interface:

BasicType

AliasThis:

alias Identifier this ;

Enum

EnumDeclaration:

enum Identifier EnumBody

enum Identifier : EnumBaseType EnumBody

AnonymousEnumDeclaration

EnumBaseType:
Type

EnumBody:
{ EnumMembers }
 ;

EnumMembers:
EnumMember
EnumMember ,
EnumMember , *EnumMembers*

EnumMember:
Identifier
Identifier = *AssignExpression*

AnonymousEnumDeclaration:
enum : *EnumBaseType* { *EnumMembers* }
enum { *EnumMembers* }
enum { *AnonymousEnumMembers* }

AnonymousEnumMembers:
AnonymousEnumMember
AnonymousEnumMember ,
AnonymousEnumMember , *AnonymousEnumMembers*

AnonymousEnumMember:
EnumMember
Type Identifier = *AssignExpression*

Template

TemplateDeclaration:
template *Identifier* *TemplateParameters* *Constraint*_{opt} { *DeclDefs*_{opt} }

TemplateParameters:
 (*TemplateParameterList*_{opt})

TemplateParameterList:

TemplateParameter
TemplateParameter ,
TemplateParameter , *TemplateParameterList*

TemplateParameter:

TemplateTypeParameter
TemplateValueParameter
TemplateAliasParameter
TemplateTupleParameter
TemplateThisParameter

Constraint:

if (*Expression*)

TemplateInstance:

Identifier TemplateArguments

TemplateArguments:

! (*TemplateArgumentList*_{opt})
 ! *TemplateSingleArgument*

TemplateArgumentList:

TemplateArgument
TemplateArgument ,
TemplateArgument , *TemplateArgumentList*

TemplateArgument:

Type
AssignExpression
Symbol

Symbol:

SymbolTail
 . *SymbolTail*

SymbolTail:

Identifier
Identifier . *SymbolTail*
TemplateInstance

TemplateInstance . SymbolTail

TemplateSingleArgument:

Identifier
BasicTypeX
CharacterLiteral
StringLiteral
IntegerLiteral
FloatLiteral
true
false
null
this
SpecialKeyword

TemplateTypeParameter:

Identifier
Identifier TemplateTypeParameterSpecialization
Identifier TemplateTypeParameterDefault
Identifier TemplateTypeParameterSpecialization TemplateTypeParameterDefault

TemplateTypeParameterSpecialization:

: Type

TemplateTypeParameterDefault:

= Type

TemplateThisParameter:

this TemplateTypeParameter

TemplateValueParameter:

BasicType Declarator
BasicType Declarator TemplateValueParameterSpecialization
BasicType Declarator TemplateValueParameterDefault
BasicType Declarator TemplateValueParameterSpecialization TemplateValueParameterDefault

TemplateValueParameterSpecialization:

: *ConditionalExpression*

TemplateValueParameterDefault:

= *AssignExpression*

= *SpecialKeyword*

TemplateAliasParameter:

alias *Identifier* *TemplateAliasParameterSpecialization*_{opt} *TemplateAliasParameterDefault*_{opt}

alias *BasicType* *Declarator* *TemplateAliasParameterSpecialization*_{opt} *TemplateAliasParameterD*

TemplateAliasParameterSpecialization:

: *Type*

: *ConditionalExpression*

TemplateAliasParameterDefault:

= *Type*

= *ConditionalExpression*

TemplateTupleParameter:

Identifier ...

TemplateMixinDeclaration:

mixin template *Identifier* *TemplateParameters* *Constraint*_{opt} { *DeclDefs*_{opt} }

TemplateMixin:

mixin *MixinTemplateName* *TemplateArguments*_{opt} *Identifier*_{opt} ;

MixinTemplateName:

. *QualifiedIdentifierList*

QualifiedIdentifierList

Typeof . *QualifiedIdentifierList*

QualifiedIdentifierList:

Identifier

Identifier . *QualifiedIdentifierList*

TemplateInstance . *QualifiedIdentifierList*

Attribute

AttributeSpecifier:

Attribute :
Attribute DeclarationBlock

Attribute:

LinkageAttribute
AlignAttribute
DeprecatedAttribute
ProtectionAttribute
Pragma
static
extern
abstract
final
override
synchronized
auto
scope
const
immutable
inout
shared
__gshared
Property
nothrow
pure
ref

DeclarationBlock:

DeclDef
{ *DeclDefs*_{opt} }

LinkageAttribute:

extern (*LinkageType*)
extern (C++, *IdentifierList*)
extern (C++, *CppAggregateType*)

LinkageType:

C
C++
D
Windows
Pascal
System
Objective-C

CppAggregateType: `class struct`

AlignAttribute:

`align`
`align (AssignExpression)`

DeprecatedAttribute:

`deprecated`
`deprecated (AssignExpression)`

ProtectionAttribute:

`private`
`package`
`package (IdentifierList)`
`protected`
`public`
`export`

Property:

`@ PropertyIdentifier`
`UserDefinedAttribute`

PropertyIdentifier:

`property`
`safe`
`trusted`
`system`
`disable`

nogc

UserDefinedAttribute:

```
@ ( ArgumentList )
@ Identifier
@ Identifier ( ArgumentListopt )
@ TemplateInstance
@ TemplateInstance ( ArgumentListopt )
```

Pragma:

```
pragma ( Identifier )
pragma ( Identifier , ArgumentList )
```

Conditional

ConditionalDeclaration:

```
Condition DeclarationBlock
Condition DeclarationBlock else DeclarationBlock
Condition : DeclDefsopt
Condition DeclarationBlock else : DeclDefsopt
```

ConditionalStatement:

```
Condition NoScopeNonEmptyStatement
Condition NoScopeNonEmptyStatement else NoScopeNonEmptyStatement
```

Condition:

```
VersionCondition
DebugCondition
StaticIfCondition
```

VersionCondition:

```
version ( IntegerLiteral )
version ( Identifier )
version ( unittest )
version ( assert )
```

DebugCondition:

```

debug
debug ( IntegerLiteral )
debug ( Identifier )

```

```

StaticIfCondition:
    static if ( AssignExpression )

```

```

VersionSpecification:
    version = Identifier ;
    version = IntegerLiteral ;

```

```

DebugSpecification:
    debug = Identifier ;
    debug = IntegerLiteral ;

```

```

StaticAssert:
    static assert ( AssignExpression ,opt );
    static assert ( AssignExpression , AssignExpression ,opt );

```

Module

```

Module:
    ModuleDeclaration DeclDefs
    DeclDefs

```

```

DeclDefs:
    DeclDef
    DeclDef DeclDefs

```

```

DeclDef:
    AttributeSpecifier
    Declaration
    Constructor
    Destructor
    Postblit
    Allocator
    Deallocator

```

Invariant
UnitTest
AliasThis
StaticConstructor
StaticDestructor
SharedStaticConstructor
SharedStaticDestructor
ConditionalDeclaration
DebugSpecification
VersionSpecification
StaticAssert
TemplateDeclaration
TemplateMixinDeclaration
TemplateMixin
MixinDeclaration
;

ModuleDeclaration:

`module` *ModuleFullyQualifiedName* ;

ModuleFullyQualifiedName:

ModuleName

Packages . *ModuleName*

ModuleName:

Identifier

Packages:

PackageName

Packages . *PackageName*

PackageName:

Identifier

ImportDeclaration:

`import` *ImportList* ;

`static import` *ImportList* ;


```

ImportList:
  Import
  ImportBindings
  Import , ImportList

Import:
  ModuleFullyQualifiedName
  ModuleAliasIdentifier = ModuleFullyQualifiedName

ImportBindings:
  Import : ImportBindList

ImportBindList:
  ImportBind
  ImportBind , ImportBindList

ImportBind:
  Identifier
  Identifier = Identifier

ModuleAliasIdentifier:
  Identifier

MixinDeclaration:
  mixin ( AssignExpression ) ;

)

```


Chapter 4

Modules

Module:

ModuleDeclaration DeclDefs
DeclDefs

DeclDefs:

DeclDef
DeclDef DeclDefs

DeclDef:

AttributeSpecifier
Declaration
Constructor
Destructor
Postblit
Allocator
Deallocator
Invariant
UnitTest
AliasThis
StaticConstructor
StaticDestructor
SharedStaticConstructor
SharedStaticDestructor
ConditionalDeclaration
DebugSpecification
VersionSpecification

```

StaticAssert
TemplateDeclaration
TemplateMixinDeclaration
TemplateMixin
MixinDeclaration
;

```

Modules have a one-to-one correspondence with source files. The module name is, by default, the file name with the path and extension stripped off, and can be set explicitly with the module declaration.

Modules automatically provide a namespace scope for their contents. Modules superficially resemble classes, but differ in that:

- There's only one instance of each module, and it is statically allocated.
- There is no virtual table.
- Modules do not inherit, they have no super modules, etc.
- Only one module per file.
- Module symbols can be imported.
- Modules are always compiled at global scope, and are unaffected by surrounding attributes or other modifiers.

Modules can be grouped together in hierarchies called *packages*.

Modules offer several guarantees:

- The order in which modules are imported does not affect the semantics.
- The semantics of a module are not affected by what imports it.
- If a module C imports modules A and B, any modifications to B will not silently change code in C that is dependent on A.

Module Declaration

The *ModuleDeclaration* sets the name of the module and what package it belongs to. If absent, the module name is taken to be the same name (stripped of path and extension) of the source file name.

```

ModuleDeclaration:
  ModuleAttributesopt module ModuleFullyQualifiedName ;

```

```

ModuleAttributes:

```

ModuleAttribute
ModuleAttribute ModuleAttributes

ModuleAttribute:
DeprecatedAttribute
UserDefinedAttribute

ModuleFullyQualifiedName:
ModuleName
Packages . ModuleName

ModuleName:
Identifier

Packages:
PackageName
Packages . PackageName

PackageName:
Identifier

The *Identifiers* preceding the rightmost are the *Packages* that the module is in. The packages correspond to directory names in the source file path. Package names cannot be keywords, hence the corresponding directory names cannot be keywords, either.

If present, the *ModuleDeclaration* appears syntactically first in the source file, and there can be only one per source file.

Example:

```
module c.stdio; // module stdio in the c package
```

By convention, package and module names are all lower case. This is because those names can have a one-to-one correspondence with the operating system's directory and file names, and many file systems are not case sensitive. All lower case package and module names will minimize problems moving projects between dissimilar file systems.

If the file name of a module is an invalid module name (e.g. `foo-bar.d`), you may use a module declaration to set a valid module name:

```
module foo_bar;
```

ModuleDeclaration can have an optional *DeprecatedAttribute*. The compiler will produce a message when the deprecated module is imported.

```
deprecated module foo;
```

```
module bar;
import foo; // Deprecated: module foo is deprecated
```

DeprecatedAttribute can have an optional string argument to provide a more expressive message.

```
deprecated("Please use foo2 instead.")
module foo;
```

```
module bar;
import foo; // Deprecated: module foo is deprecated - Please use foo2 instead.
```

Import Declaration

Symbols from one module are made available in another module by using the *ImportDeclaration*:

ImportDeclaration:

```
import ImportList ;
static import ImportList ;
```

ImportList:

```
Import
ImportBindings
Import , ImportList
```

Import:

```
ModuleFullyQualifiedName
ModuleAliasIdentifier = ModuleFullyQualifiedName
```

ImportBindings:

```
Import : ImportBindList
```

ImportBindList:

```
ImportBind
ImportBind , ImportBindList
```

ImportBind:

```
Identifier
Identifier = Identifier
```

ModuleAliasIdentifier:
Identifier

There are several forms of the *ImportDeclaration*, from generalized to fine-grained importing. The order in which *ImportDeclarations* occur has no significance.

ModuleFullyQualifiedNames in the *ImportDeclaration* must be fully qualified with whatever packages they are in. They are not considered to be relative to the module that imports them.

Symbol Name Lookup

The simplest form of importing is to just list the modules being imported:

```
module myapp.main;

import std.stdio; // import module stdio from package std

class Foo : BaseClass
{
    import myapp.foo; // import module myapp.foo in this class' scope
    void bar ()
    {
        import myapp.bar; // import module myapp.bar in this function' scope
        writeln("hello!"); // calls std.stdio.writeln
    }
}
```

When a symbol name is used unqualified, a two-phase lookup will happen. First, the module scope will be searched, starting from the innermost scope. For example, in the previous example, while looking for ‘writeln’, the order will be:

- Declarations inside ‘bar’.
- Declarations inside ‘Foo’.
- Declarations inside ‘BaseClass’.
- Declarations at module scope.

If the first lookup wasn’t successful, a second one is performed on imports. In the second lookup phase inherited scopes are ignored. This includes scope of base classes and interface (in this example, ‘BaseClass’'s imports would be ignored), as well as imports in mixed-in [template](#).

Symbol lookup stops as soon as a symbol is found. If two symbols with the same name are found at the same lookup phase, this ambiguity will result in a compilation error.

```
module A;
void foo();
void bar();

module B;
void foo();
void bar();

module C;
import A;
void foo();
void test()
{
    foo(); // C.foo() is called, it is found before imports are searched
    bar(); // A.bar() is called, since imports are searched
}

module D;
import A;
import B;
void test()
{
    foo(); // error, A.foo() or B.foo() ?
    A.foo(); // ok, call A.foo()
    B.foo(); // ok, call B.foo()
}

module E;
import A;
import B;
alias foo = B.foo;
void test()
{
    foo(); // call B.foo()
    A.foo(); // call A.foo()
    B.foo(); // call B.foo()
}
```


Public Imports

By default, imports are *private*. This means that if module A imports module B, and module B imports module C, then names from C are visible only from B and not from A. This prevents surprising name clashes from modules which haven't directly been imported.

An import can be however specifically declared *public*, which will cause names from the imported module to be visible to further imports. So in the above example where module A imports module B, if module B *publicly* imports module C, names from C will be visible in A as well.

All symbols from a publicly imported module are also aliased in the importing module. Thus in the above example if C contains the name `foo`, it will be accessible in A as `foo`, `B.foo` and `C.foo`.

For another example:

```
module W;
void foo() { }

module X;
void bar() { }

module Y;
import W;
public import X;
...
foo(); // calls W.foo()
bar(); // calls X.bar()

module Z;
import Y;
...
foo(); // error, foo() is undefined
bar(); // ok, calls X.bar()
X.bar(); // ditto
Y.bar(); // ok, Y.bar() is an alias to X.bar()
```

Static Imports

Basic imports work well for programs with relatively few modules and imports. If there are a lot of imports, name collisions can start occurring between the names in the various imported modules. One way to stop this is by using static imports. A static import requires one to use a fully qualified name to reference the module's names:

```
static import std.stdio;
```

```
void main()
{
    writeln("hello!");           // error, writeln is undefined
    std.stdio.writeln("hello!"); // ok, writeln is fully qualified
}
```

Renamed Imports

A local name for an import can be given, through which all references to the module's symbols must be qualified with:

```
import io = std.stdio;

void main()
{
    io.writeln("hello!");           // ok, calls std.stdio.writeln
    std.stdio.writeln("hello!");    // error, std is undefined
    writeln("hello!");              // error, writeln is undefined
}
```

Renamed imports are handy when dealing with very long import names.

Selective Imports

Specific symbols can be exclusively imported from a module and bound into the current namespace:

```
import std.stdio : writeln, foo = write;

void main()
{
    std.stdio.writeln("hello!"); // error, std is undefined
    writeln("hello!");           // ok, writeln bound into current namespace
    write("world");              // error, write is undefined
    foo("world");                // ok, calls std.stdio.write()
    fwritefln(stdout, "abc");    // error, fwritefln undefined
}
```

`static` cannot be used with selective imports.

Renamed and Selective Imports

When renaming and selective importing are combined:

```

import io = std.stdio : foo = writeln;

void main()
{
    writeln("bar");           // error, writeln is undefined
    std.stdio.foo("bar");     // error, foo is bound into current namespace
    std.stdio.writeln("bar"); // error, std is undefined
    foo("bar");               // ok, foo is bound into current namespace,
                             // FQN not required
    io.writeln("bar");        // ok, io=std.stdio bound the name io in
                             // the current namespace to refer to the entire ↗
                             // ↘ module
    io.foo("bar");            // error, foo is bound into current namespace,
                             // foo is not a member of io
}

```

Scoped Imports

Import declarations may be used at any scope. For example:

```

void main()
{
    import std.stdio;
    writeln("bar");
}

```

The imports are looked up to satisfy any unresolved symbols at that scope. Imported symbols may hide symbols from outer scopes.

In function scopes, imported symbols only become visible after the import declaration lexically appears in the function body. In other words, imported symbols at function scope cannot be forward referenced.

```

void main()
{
    void writeln(string) {}
    void foo()
    {
        writeln("bar"); // calls main.writeln
        import std.stdio;
        writeln("bar"); // calls std.stdio.writeln
        void writeln(string) {}
    }
}

```

```

        writeln("bar"); // calls main.foo.writeln
    }
    writeln("bar"); // calls main.writeln
    std.stdio.writeln("bar"); // error, std is undefined
}

```

Module Scope Operator

Sometimes, it's necessary to override the usual lexical scoping rules to access a name hidden by a local name. This is done with the global scope operator, which is a leading `'.'`:

```

int x;

int foo(int x)
{
    if (y)
        return x; // returns foo.x, not global x
    else
        return .x; // returns global x
}

```

The leading `'.'` means look up the name at the module scope level.

Static Construction and Destruction

Static constructors are code that gets executed to initialize a module or a class before the `main()` function gets called. Static destructors are code that gets executed after the `main()` function returns, and are normally used for releasing system resources.

There can be multiple static constructors and static destructors within one module. The static constructors are run in lexical order, the static destructors are run in reverse lexical order.

Static constructors and static destructors run on thread local data, and are run whenever threads are created or destroyed.

Shared static constructors and shared static destructors are run on global shared data, and constructors are run once on program startup and destructors are run once on program termination.

Order of Static Construction

Shared static constructors on all modules are run before any static constructors.

The order of static initialization is implicitly determined by the *import* declarations in each module. Each module is assumed to depend on any imported modules being statically constructed

first. Other than following that rule, there is no imposed order on executing the module static constructors.

Cycles (circular dependencies) in the import declarations are allowed as long as not both of the modules contain static constructors or static destructors. Violation of this rule will result in a runtime exception.

Overriding Cycle Detection Abort

You can override the cyclic detection behavior using the D Runtime switch `--DRT-oncycle=...`. The following behaviors are supported:

1. **deprecate** The default behavior. This functions just like **abort**, but will use the pre-2.072 algorithm to determine if the cycle was undetected. If so, the old algorithm is used, but a deprecation message is printed. After 2.073, the **abort** option will be the default.
2. **abort** The normal behavior as described in the previous section. After 2.073, this will be the default behavior.
3. **print** Print all cycles detected, but do not abort execution. Order of static construction is implementation defined, and not guaranteed to be valid.
4. **ignore** Do not abort execution or print any cycles. Order of static construction is implementation defined, and not guaranteed to be valid.

Order of Static Construction within a Module

Within a module, the static construction occurs in the lexical order in which they appear.

Order of Static Destruction

It is defined to be exactly the reverse order that static construction was performed in. Static destructors for individual modules will only be run if the corresponding static constructor successfully completed.

Shared static destructors are executed after static destructors.

Order of Unit tests

Unit tests are run in the lexical order in which they appear within a module.

Mixin Declaration

```
MixinDeclaration:
    mixin ( AssignExpression ) ;
```

The *AssignExpression* must evaluate at compile time to a constant string. The text contents of the string must be compilable as a valid *DeclDefs*, and is compiled as such.

Package Module

A package module can be used to publicly import other modules, while enabling a simpler import syntax. It enables converting a module into a package of modules, without breaking existing code which uses that module. Example of a set of library modules:

libweb/client.d:

```
module libweb.client;
```

```
void runClient() { }
```

libweb/server.d:

```
module libweb.server;
```

```
void runServer() { }
```

libweb/package.d:

```
module libweb;
```

```
public import libweb.client;
```

```
public import libweb.server;
```

The package module must have the file name `package.d`. The module name is declared to be the fully qualified name of the package. Package modules can be imported just like any other modules:

test.d:

```
module test;
```

```
// import the package module
```

```
import libweb;
```

```
void main()
```

```
{
```

```
    runClient();
```

```
    runServer();
```

```
}
```

A package module can be nested inside of a sub-package:

libweb/utils/package.d:

```
// must be declared as the fully qualified name of the package, not just 'utils'  
module libweb.utils;
```

```
// publicly import modules from within the 'libweb.utils' package.  
public import libweb.utils.conv;  
public import libweb.utils.text;
```

The package module can then be imported with the standard module import declaration:

test.d:

```
module test;
```

```
// import the package module  
import libweb.utils;
```

```
void main() { }
```


Chapter 5

Declarations

Declaration:

FuncDeclaration
VarDeclarations
AliasDeclaration
AggregateDeclaration
EnumDeclaration
ImportDeclaration

VarDeclarations:

*StorageClasses*_{opt} *BasicType* *Declarators* ;
AutoDeclaration

Declarators:

DeclaratorInitializer
DeclaratorInitializer , *DeclaratorIdentifierList*

DeclaratorInitializer:

VarDeclarator
VarDeclarator *TemplateParameters*_{opt} = *Initializer*
AltDeclarator
AltDeclarator = *Initializer*

DeclaratorIdentifierList:

DeclaratorIdentifier
DeclaratorIdentifier , *DeclaratorIdentifierList*

DeclaratorIdentifier:

VarDeclaratorIdentifier
AltDeclaratorIdentifier

VarDeclaratorIdentifier:

Identifier
*Identifier TemplateParameters*_{opt} = *Initializer*

AltDeclaratorIdentifier:

*BasicType2 Identifier AltDeclaratorSuffixes*_{opt}
*BasicType2 Identifier AltDeclaratorSuffixes*_{opt} = *Initializer*
*BasicType2*_{opt} *Identifier AltDeclaratorSuffixes*
*BasicType2*_{opt} *Identifier AltDeclaratorSuffixes* = *Initializer*

Declarator:

VarDeclarator
AltDeclarator

VarDeclarator:

*BasicType2*_{opt} *Identifier*

AltDeclarator:

*BasicType2*_{opt} *Identifier AltDeclaratorSuffixes*
*BasicType2*_{opt} (*AltDeclaratorX*)
*BasicType2*_{opt} (*AltDeclaratorX*) *AltFuncDeclaratorSuffix*
*BasicType2*_{opt} (*AltDeclaratorX*) *AltDeclaratorSuffixes*

AltDeclaratorX:

*BasicType2*_{opt} *Identifier*
*BasicType2*_{opt} *Identifier AltFuncDeclaratorSuffix*
AltDeclarator

AltDeclaratorSuffixes:

AltDeclaratorSuffix
AltDeclaratorSuffix AltDeclaratorSuffixes

AltDeclaratorSuffix:

[]
[*AssignExpression*]
[*Type*]

AltFuncDeclaratorSuffix:

*Parameters MemberFunctionAttributes*_{opt}

Type:

*TypeCtors*_{opt} *BasicType* *BasicType2*_{opt}

TypeCtors:

TypeCtor

TypeCtor TypeCtors

TypeCtor:

const

immutable

inout

shared

BasicType:

BasicTypeX

. IdentifierList

IdentifierList

Typeof

Typeof . IdentifierList

TypeCtor (Type)

TypeVector

BasicTypeX:

bool

byte

ubyte

short

ushort

int

uint

long

ulong

char

wchar

dchar

float

double

real

ifloat

idouble

ireal

cfloat

cdouble

creal

void

BasicType2:
BasicType2X BasicType2_{opt}

BasicType2X:
 *
 []
 [*AssignExpression*]
 [*AssignExpression* .. *AssignExpression*]
 [*Type*]
delegate *Parameters MemberFunctionAttributes_{opt}*
function *Parameters FunctionAttributes_{opt}*

IdentifierList:
Identifier
Identifier . *IdentifierList*
TemplateInstance
TemplateInstance . *IdentifierList*
Identifier [*AssignExpression*]. *IdentifierList*

StorageClasses:
StorageClass
StorageClass StorageClasses

StorageClass:

<i>LinkageAttribute</i>	<i>static</i>	<i>synchronized</i> ↗	<i>immutable</i>	<i>nothrow</i>
<i>AlignAttribute</i>	<i>extern</i>	↘	<i>inout</i>	<i>pure</i>
<i>deprecated</i> ↗	<i>abstract</i>	<i>auto</i>	<i>shared</i>	<i>ref</i>
↘	<i>final</i>	<i>scope</i>	<i>__gshared</i>	
<i>enum</i>	<i>override</i>	<i>const</i>	<i>Property</i>	

Initializer:
VoidInitializer

NonVoidInitializer

NonVoidInitializer:

ExpInitializer :

ArrayInitializer

StructInitializer

ExpInitializer:

AssignExpression

ArrayInitializer:

[*ArrayMemberInitializations*_{opt}]

ArrayMemberInitializations:

ArrayMemberInitialization

ArrayMemberInitialization ,

ArrayMemberInitialization , *ArrayMemberInitializations*

ArrayMemberInitialization:

NonVoidInitializer

AssignExpression : *NonVoidInitializer*

StructInitializer:

{ *StructMemberInitializers*_{opt} }

StructMemberInitializers:

StructMemberInitializer

StructMemberInitializer ,

StructMemberInitializer , *StructMemberInitializers*

StructMemberInitializer:

NonVoidInitializer

Identifier : *NonVoidInitializer*

Declaration Syntax

Declaration syntax generally reads right to left:

```
int x;      // x is an int
int* x;     // x is a pointer to int
```

```

int** x; // x is a pointer to a pointer to int
int[] x; // x is an array of ints
int*[] x; // x is an array of pointers to ints
int[]* x; // x is a pointer to an array of ints

```

Arrays read right to left as well:

```

int[3] x; // x is an array of 3 ints
int[3][5] x; // x is an array of 5 arrays of 3 ints
int[3]*[5] x; // x is an array of 5 pointers to arrays of 3 ints

```

Pointers to functions are declared using the `function` keyword:

```

int function(char) x; // x is a pointer to
                      // a function taking a char argument
                      // and returning an int
int function(char)[] x; // x is an array of
                      // pointers to functions
                      // taking a char argument
                      // and returning an int

```

C-style array, function pointer and pointer to array declarations are deprecated:

```

int x[3]; // x is an array of 3 ints
int x[3][5]; // x is an array of 3 arrays of 5 ints
int (*x[5])[3]; // x is an array of 5 pointers to arrays of 3 ints
int (*x)(char); // x is a pointer to a function taking a char argument
                // and returning an int
int (*[] x)(char); // x is an array of pointers to functions
                // taking a char argument and returning an int

```

In a declaration declaring multiple symbols, all the declarations must be of the same type:

```

int x,y; // x and y are ints
int* x,y; // x and y are pointers to ints
int x,*y; // error, multiple types
int[] x,y; // x and y are arrays of ints
int x[],y; // error, multiple types

```

Implicit Type Inference

AutoDeclaration:

StorageClasses AutoDeclarationX ;

```

AutoDeclarationX:
    AutoDeclarationY
    AutoDeclarationX , AutoDeclarationY

```

```

AutoDeclarationY:
    Identifier TemplateParametersopt = Initializer

```

If a declaration starts with a *StorageClass* and has a *NonVoidInitializer* from which the type can be inferred, the type on the declaration can be omitted.

```

static x = 3;          // x is type int
auto y = 4u;          // y is type uint

auto s = "string";    // s is type immutable(char)[]

class C { ... }

auto c = new C();     // c is a handle to an instance of class C

```

The *NonVoidInitializer* cannot contain forward references (this restriction may be removed in the future). The implicitly inferred type is statically bound to the declaration at compile time, not run time.

An *ArrayLiteral* is inferred to be a dynamic array type rather than a static array:

```

auto v = ["hello", "world"]; // type is string[], not string[2]

```

Alias Declarations

```

AliasDeclaration:
    alias StorageClassesopt BasicType Declarators ;
    alias StorageClassesopt BasicType FuncDeclarator ;
    alias AliasDeclarationX ;

```

```

AliasDeclarationX:
    AliasDeclarationY
    AliasDeclarationX , AliasDeclarationY

```

```

AliasDeclarationY:
    Identifier TemplateParametersopt = StorageClassesopt Type
    Identifier TemplateParametersopt = FunctionLiteral

```

AliasDeclarations create a symbol that is an alias for another type, and can be used anywhere that other type may appear.

```
alias myint = abc.Foo.bar;
```

Aliased types are semantically identical to the types they are aliased to. The debugger cannot distinguish between them, and there is no difference as far as function overloading is concerned. For example:

```
alias myint = int;

void foo(int x) { ... }
void foo(myint m) { ... } // error, multiply defined function foo
```

A symbol can be declared as an *alias* of another symbol. For example:

```
import string;

alias mylen = string.strlen;
...
int len = mylen("hello"); // actually calls string.strlen()
```

The following alias declarations are valid:

```
template Foo2(T) { alias t = T; }
alias t1 = Foo2!(int);
alias t2 = Foo2!(int).t;
alias t3 = t1.t;
alias t4 = t2;

t1.t v1; // v1 is type int
t2 v2;   // v2 is type int
t3 v3;   // v3 is type int
t4 v4;   // v4 is type int
```

Aliased symbols are useful as a shorthand for a long qualified symbol name, or as a way to redirect references from one symbol to another:

```
version (Win32)
{
    alias myfoo = win32.foo;
}
version (linux)
{
    alias myfoo = linux.bar;
}
```


Aliasing can be used to ‘import’ a symbol from an import into the current scope:

```
alias strlen = string.strlen;
```

Aliases can also ‘import’ a set of overloaded functions, that can be overloaded with functions in the current scope:

```
class A
{
    int foo(int a) { return 1; }
}

class B : A
{
    int foo( int a, uint b ) { return 2; }
}

class C : B
{
    int foo( int a ) { return 3; }
    alias foo = B.foo;
}

class D : C
{
}

void test()
{
    D b = new D();
    int i;

    i = b.foo(1, 2u);    // calls B.foo
    i = b.foo(1);        // calls C.foo
}
```

Note: Type aliases can sometimes look indistinguishable from alias declarations:

```
alias abc = foo.bar; // is it a type or a symbol?
```

The distinction is made in the semantic analysis pass.

Aliases cannot be used for expressions:

```
struct S { static int i; }
```

```

S s;

alias a = s.i; // illegal, s.i is an expression
alias b = S.i; // ok
b = 4;         // sets S.i to 4

```

Extern Declarations

Variable declarations with the storage class `extern` are not allocated storage within the module. They must be defined in some other object file with a matching name which is then linked in. The primary usefulness of this is to connect with global variable declarations in C files.

An `extern` declaration can optionally be followed by an `extern` linkage attribute. If there is no linkage attribute it defaults to `extern(D)`:

```

extern(C) int foo;           // variable allocated and initialized in this module ↗
    ↪ with C linkage
extern extern(C) int bar;    // variable allocated outside this module with C ↗
    ↪ linkage
                               // (e.g. in a statically linked C library or another ↗
                               ↪ module)

```

typeof

```

Typeof:
    typeof ( Expression )
    typeof ( return )

```

Typeof is a way to specify a type based on the type of an expression. For example:

```

void func(int i)
{
    typeof(i) j;           // j is of type int
    typeof(3 + 6.0) x;      // x is of type double
    typeof(1)* p;          // p is of type pointer to int
    int[sizeof(p)] a;       // a is of type int[int*]

    writefln("%d", sizeof(typeof('c'))); // prints 1
    double c = cast(typeof(1.0))j; // cast j to double
}

```

Expression is not evaluated, just the type of it is generated:

```
void func()
{
    int i = 1;
    typeof(++i) j; // j is declared to be an int, i is not incremented
    writefln("%d", i); // prints 1
}
```

There are three special cases:

1. `typeof(this)` will generate the type of what `this` would be in a non-static member function, even if not in a member function.
2. Analogously, `typeof(super)` will generate the type of what `super` would be in a non-static member function.
3. `typeof(return)` will, when inside a function scope, give the return type of that function.

```
class A { }

class B : A
{
    typeof(this) x; // x is declared to be a B
    typeof(super) y; // y is declared to be an A
}

struct C
{
    static typeof(this) z; // z is declared to be a C

    typeof(super) q; // error, no super struct for C
}

typeof(this) r; // error, no enclosing struct or class
```

If the expression is a Property Function, `typeof` gives its return type.

```
struct S
{
    @property int foo() { return 1; }
}

typeof(S.foo) n; // n is declared to be an int
```

Where *typeof* is most useful is in writing generic template code.

Void Initializations

```
VoidInitializer:
    void
```

Normally, variables are initialized either with an explicit *Initializer* or are set to the default value for the type of the variable. If the *Initializer* is `void`, however, the variable is not initialized. If its value is used before it is set, undefined program behavior will result.

```
void foo()
{
    int x = void;
    writeln(x); // will print garbage
}
```

Therefore, one should only use `void` initializers as a last resort when optimizing critical code.

Global and Static Initializers

The *Initializer* for a global or static variable must be evaluable at compile time. Whether some pointers can be initialized with the addresses of other functions or data is implementation defined. Runtime initialization can be done with static constructors.

Type Qualifiers vs. Storage Classes

D draws a distinction between a type qualifier and a storage class.

A *type qualifier* creates a derived type from an existing base type, and the resulting type may be used to create multiple instances of that type.

For example, the `immutable` type qualifier can be used to create variables of immutable type, such as:

```
immutable(int)    x; // typeof(x) == immutable(int)
immutable(int)[] y; // typeof(y) == immutable(int)[]
                  // typeof(y[0]) == immutable(int)

// Type constructors create new types that can be aliased:
alias ImmutableInt = immutable(int);
ImmutableInt z;    // typeof(z) == immutable(int)
```

A *storage class*, on the other hand, does not create a new type, but describes only the type of storage used by the variable or function being declared. For example, a member function can be declared with the `const` storage class to indicate that it does not modify its implicit `this` argument:

```

struct S
{
    int x;
    int method() const
    {
        //x++;    // Error: this method is const and cannot modify this.x
        return x; // OK: we can still read this.x
    }
}

```

Although some keywords can be used both as a type qualifier and a storage class, there are some storage classes that cannot be used to construct new types. One example is [ref](#):

```

// ref declares the parameter x to be passed by reference
void func(ref int x)
{
    x++; // so modifications to x will be visible in the caller
}

void main()
{
    auto x = 1;
    func(x);
    assert(x == 2);

    // However, ref is not a type qualifier, so the following is illegal:
    ref(int) y; // Error: ref is not a type qualifier.
}

// Functions can also be declared as 'ref', meaning their return value is
// passed by reference:
ref int func2()
{
    static int y = 0;
    return y;
}

void main()
{
    func2() = 2; // The return value of func2() can be modified.
    assert(func2() == 2);
}

```

```

// However, the reference returned by func2() does not propagate to
// variables, because the 'ref' only applies to the return value itself,
// not to any subsequent variable created from it:
auto x = func2();
static_assert(is(typeof(x) == int)); // N.B.: *not* ref(int);
                                     // there is no such type as ref(int).

x++;
assert(x == 3);
assert(func2() == 2); // x is not a reference to what func2() returned; it
                     // does not inherit the ref storage class from func2().
}

```

Due to the fact that some keywords, such as `const`, can be used both as a type qualifier and a storage class, it may sometimes result in ambiguous-looking code:

```

struct S
{
    // Is const here a type qualifier or a storage class?
    // Is the return value const(int), or is this a const function that returns
    // (mutable) int?
    const int func() { return 1; }
}

```

To avoid such confusion, it is recommended that type qualifier syntax with parentheses always be used for return types, and that function storage classes be written on the right-hand side of the declaration instead of the left-hand side where it may be visually confused with the return type:

```

struct S
{
    // Now it is clear that the 'const' here applies to the return type:
    const(int) func1() { return 1; }

    // And it is clear that the 'const' here applies to the function:
    int func2() const { return 1; }
}

```


Chapter 6

Types

Basic Data Types

Basic Data Types		
Keyword	Default Initializer (.init)	Description
<code>void</code>	-	no type
<code>bool</code>	<code>false</code>	boolean value
<code>byte</code>	0	signed 8 bits
<code>ubyte</code>	0	unsigned 8 bits
<code>short</code>	0	signed 16 bits
<code>ushort</code>	0	unsigned 16 bits
<code>int</code>	0	signed 32 bits
<code>uint</code>	0	unsigned 32 bits
<code>long</code>	0L	signed 64 bits
<code>ulong</code>	0L	unsigned 64 bits
<code>cent</code>	0	signed 128 bits (reserved for future use)
<code>ucent</code>	0	unsigned 128 bits (reserved for future use)
<code>float</code>	<code>float.nan</code>	32 bit floating point
<code>double</code>	<code>double.nan</code>	64 bit floating point
<code>real</code>	<code>real.nan</code>	largest FP size implemented in hardware ¹
<code>ifloat</code>	<code>float.nan*1.0i</code>	imaginary float
<code>idouble</code>	<code>double.nan*1.0i</code>	imaginary double
<code>ireal</code>	<code>real.nan*1.0i</code>	imaginary real
<code>cfloat</code>	<code>float.nan+float.nan*1.0i</code>	a complex number of two float values
<code>cdouble</code>	<code>double.nan+double.nan*1.0i</code>	complex double
<code>creal</code>	<code>real.nan+real.nan*1.0i</code>	complex real
<code>char</code>	0xFF	unsigned 8 bit (UTF-8 code unit)
<code>wchar</code>	0xFFFF	unsigned 16 bit (UTF-16 code unit)
<code>dchar</code>	0x0000FFFF	unsigned 32 bit (UTF-32 code unit)

¹Implementation Note: 80 bits for x86 CPUs or `double` size, whichever is larger

Derived Data Types

- pointer
- array
- associative array
- function
- delegate

Strings are a special case of arrays.

User Defined Types

- alias
- enum
- struct
- union
- class

Base Types

The *base type* of an enum is the type it is based on:

```
enum E : T { ... } // T is the base type of E
```

Pointer Conversions

Casting pointers to non-pointers and vice versa is allowed in D, however, do not do this for any pointers that point to data allocated by the garbage collector.

Implicit Conversions

Implicit conversions are used to automatically convert types as required.

A enum can be implicitly converted to its base type, but going the other way requires an explicit conversion.

For example:

```
int i;

enum Foo { E }
Foo f;
i = f;           // OK
f = i;           // error
```

```
f = cast(Foo)i;  // OK
f = 0;          // error
f = Foo.E;      // OK
```

Integer Promotions

Integer Promotions are conversions of the following types:

Integer Promotions

from	to
<code>bool</code>	<code>int</code>
<code>byte</code>	<code>int</code>
<code>ubyte</code>	<code>int</code>
<code>short</code>	<code>int</code>
<code>ushort</code>	<code>int</code>
<code>char</code>	<code>int</code>
<code>wchar</code>	<code>int</code>
<code>dchar</code>	<code>uint</code>

If an enum has as a base type one of the types in the left column, it is converted to the type in the right column.

Usual Arithmetic Conversions

The usual arithmetic conversions convert operands of binary operators to a common type. The operands must already be of arithmetic types. The following rules are applied in order, looking at the base type:

1. If either operand is real, the other operand is converted to real.
2. Else if either operand is double, the other operand is converted to double.
3. Else if either operand is float, the other operand is converted to float.
4. Else the integer promotions are done on each operand, followed by:
 - a) If both are the same type, no more conversions are done.
 - b) If both are signed or both are unsigned, the smaller type is converted to the larger.
 - c) If the signed type is larger than the unsigned type, the unsigned type is converted to the signed type.
 - d) The signed type is converted to the unsigned type.

If one or both of the operand types is an enum after undergoing the above conversions, the result type is:

1. If the operands are the same type, the result will be the that type.
2. If one operand is an enum and the other is the base type of that enum, the result is the base type.
3. If the two operands are different enums, the result is the closest base type common to both. A base type being closer means there is a shorter sequence of conversions to base type to get there from the original type.

Integer values cannot be implicitly converted to another type that cannot represent the integer bit pattern after integral promotion. For example:

```
ubyte  u1 = cast(byte)-1;  // error, -1 cannot be represented in a ubyte
ushort u2 = cast(short)-1; // error, -1 cannot be represented in a ushort
uint   u3 = cast(int)-1;   // ok, -1 can be represented in a uint
ulong  u4 = cast(long)-1;  // ok, -1 can be represented in a ulong
```

Floating point types cannot be implicitly converted to integral types.

Complex floating point types cannot be implicitly converted to non-complex floating point types.

Imaginary floating point types cannot be implicitly converted to float, double, or real types. Float, double, or real types cannot be implicitly converted to imaginary floating point types.

bool

The bool type is a 1 byte size type that can only hold the value true or false. The only operators that can accept operands of type bool are: `& | ^ &= |= ^= ! && || ?:`. A `bool` value can be implicitly converted to any integral type, with false becoming 0 and true becoming 1. The numeric literals 0 and 1 can be implicitly converted to the bool values false and true, respectively. Casting an expression to bool means testing for 0 or `!=0` for arithmetic types, and null or `!=null` for pointers or references.

Delegates

There are no pointers-to-members in D, but a more useful concept called *delegates* are supported. Delegates are an aggregate of two pieces of data: an object reference and a pointer to a non-static member function, or a pointer to a closure and a pointer to a nested function. The object reference forms the *this* pointer when the function is called.

Delegates are declared similarly to function pointers, except that the keyword `delegate` takes the place of `*`, and the identifier occurs afterwards:

```
int function(int) fp; // fp is pointer to a function
int delegate(int) dg; // dg is a delegate to a function
```

The C style syntax for declaring pointers to functions is deprecated:

```
int (*fp)(int); // fp is pointer to a function
```

A delegate is initialized analogously to function pointers:

```
int func(int);
fp = &func; // fp points to func

class OB
{
    int member(int);
}
OB o;
dg = &o.member; // dg is a delegate to object o and
                // member function member
```

Delegates cannot be initialized with static member functions or non-member functions.

Delegates are called analogously to function pointers:

```
fp(3); // call func(3)
dg(3); // call o.member(3)
```

The equivalent of member function pointers can be constructed using anonymous lambda functions:

```
class C
{
    int a;
    int foo(int i) { return i + a; }
}

// mfp is the member function pointer
auto mfp = function(C self, int i) { return self.foo(i); };
auto c = new C(); // create an instance of C
mfp(c, 1); // and call c.foo(1)
```

size_t

`size_t` is an alias to one of the unsigned integral basic types, and represents a type that is large enough to represent an offset into all addressible memory.

`ptrdiff_t`

`ptrdiff_t` is an alias to the signed basic type the same size as `size_t`.

Chapter 7

Properties

Every type and expression has properties that can be queried:

Property Examples	
Expression	Value
<code>int.sizeof</code>	yields 4
<code>float.nan</code>	yields the floating point nan (Not A Number) value
<code>(float).nan</code>	yields the floating point nan value
<code>(3).sizeof</code>	yields 4 (because 3 is an int)
<code>int.init</code>	default initializer for int's
<code>int.mangleof</code>	yields the string "i"
<code>int.stringof</code>	yields the string "int"
<code>(1+2).stringof</code>	yields the string "1 + 2"

Properties for All Types	
Property	Description
<code>.init</code>	initializer
<code>.sizeof</code>	size in bytes (equivalent to C's <code>sizeof(type)</code>)
<code>.alignof</code>	alignment size
<code>.mangleof</code>	string representing the ‘mangled’ representation of the type
<code>.stringof</code>	string representing the source representation of the type

Properties for Integral Types

Property	Description
<code>.init</code>	initializer (0)
<code>.max</code>	maximum value
<code>.min</code>	minimum value

Properties for Floating Point Types

Property	Description
<code>.init</code>	initializer (NaN)
<code>.infinity</code>	infinity value
<code>.nan</code>	NaN value
<code>.dig</code>	number of decimal digits of precision
<code>.epsilon</code>	smallest increment to the value 1
<code>.mant_dig</code>	number of bits in mantissa
<code>.max_10_exp</code>	maximum int value such that $10^{\text{max_10_exp}}$ is representable
<code>.max_exp</code>	maximum int value such that $2^{\text{max_exp}-1}$ is representable
<code>.min_10_exp</code>	minimum int value such that $10^{\text{min_10_exp}}$ is representable as a normalized value
<code>.min_exp</code>	minimum int value such that $2^{\text{min_exp}-1}$ is representable as a normalized value
<code>.max</code>	largest representable value that's not infinity
<code>.min_normal</code>	smallest representable normalized value that's not 0
<code>.re</code>	real part
<code>.im</code>	imaginary part

Properties for Class Types

Property	Description
<code>.classinfo</code>	Information about the dynamic type of the class

.init Property

`.init` produces a constant expression that is the default initializer. If applied to a type, it is the default initializer for that type. If applied to a variable or field, it is the default initializer for that variable or field's type. For example:

```
int a;
int b = 1;
typedef int t = 2;
t c;
```



```

t d = cast(t)3;

int.init // is 0
a.init  // is 0
b.init  // is 0
t.init  // is 2
c.init  // is 2
d.init  // is 2

struct Foo
{
    int a;
    int b = 7;
}

Foo.init.a // is 0
Foo.init.b // is 7

```

Note: `.init` produces a default initialized object, not default constructed. That means using `.init` is sometimes incorrect.

1. If `T` is a nested struct, the context pointer in `T.init` is `null`.

```

void main()
{
    int a;
    struct S
    {
        void foo() { a = 1; } // access a variable in enclosing scope
    }
    S s1;           // OK. S() correctly initialize its frame pointer.
    S s2 = S();     // OK. same as s1
    S s3 = S.init;  // Bad. the frame pointer in s3 is null
    s3.foo();       // Access violation
}

```

2. If `T` is a struct which has `@disable this();`, `T.init` might return a logically incorrect object.

```

struct S
{
    int a;
    @disable this();
}

```

```

    this(int n) { a = n; }
    invariant { assert(a > 0); }
    void check() {}
}
void main()
{
    //S s1;           // Error: variable s1 initializer required for type S
    //S s2 = S();     // Error: constructor S.this is not callable
                    // because it is annotated with @disable
    S s3 = S.init;    // Bad. s3.a == 0, and it violates the invariant of S.
    s3.check();       // Assertion failure.
}

```

.stringof Property

.stringof produces a constant string that is the source representation of its prefix. If applied to a type, it is the string for that type. If applied to an expression, it is the source representation of that expression. Semantic analysis is not done for that expression. For example:

```

module test;
import std.stdio;

struct Foo { }

enum Enum { RED }

typedef int myint;

void main()
{
    writeln((1+2).stringof);    // "1 + 2"
    writeln(Foo.stringof);      // "Foo"
    writeln(test.Foo.stringof); // "Foo"
    writeln(int.stringof);      // "int"
    writeln((int*[5][]).stringof); // "int*[5u][]"
    writeln(Enum.RED.stringof); // "cast(enum)0"
    writeln(test.myint.stringof); // "myint"
    writeln((5).stringof);      // "5"
}

```

Note: Using `.stringof` for code generation is not recommended, as the internal representation of a type or expression can change between different compiler versions.

Instead you should prefer to use the identifier trait, or one of the Phobos helper functions such as `fullyQualifiedName`.

.sizeof Property

`e.sizeof` gives the size in bytes of the expression `e`.

When getting the size of a member, it is not necessary for there to be a *this* object:

```
struct S
{
    int a;
    static int foo()
    {
        return a.sizeof; // returns 4
    }
}

void test()
{
    int x = S.a.sizeof; // sets x to 4
}
```

`.sizeof` applied to a class object returns the size of the class reference, not the class instantiation.

.alignof Property

`.alignof` gives the aligned size of an expression or type. For example, an aligned size of 1 means that it is aligned on a byte boundary, 4 means it is aligned on a 32 bit boundary.

.classinfo Property

`.classinfo` provides information about the dynamic type of a class object.

It returns a reference to type `object.TypeInfo_Class`.

`.classinfo` applied to an interface gives the information for the interface, not the class it might be an instance of.

User Defined Properties

User defined properties can be created using Property Functions.

Chapter 8

Attributes

AttributeSpecifier:

Attribute :

Attribute DeclarationBlock

Attribute:

LinkageAttribute

AlignAttribute

DeprecatedAttribute

VisibilityAttribute

Pragma

static

extern

abstract

final

override

synchronized

auto

scope

const

immutable

inout

shared

__gshared

Property

nothrow

pure

```

ref
return

```

```

Property:
  @ PropertyIdentifier
  UserDefinedAttribute

```

```

PropertyIdentifier:
  property
  safe
  trusted
  system
  disable
  nogc

```

```

DeclarationBlock:
  DeclDef
  { DeclDefsopt }

```

Attributes are a way to modify one or more declarations. The general forms are:

```

attribute declaration; // affects the declaration

```

```

attribute:      // affects all declarations until the end of
                // the current scope
  declaration;
  declaration;
  ...

```

```

attribute {    // affects all declarations in the block
  declaration;
  declaration;
  ...
}

```

Linkage Attribute

```

LinkageAttribute:
  extern ( LinkageType )

```

```
extern ( C++, IdentifierList )
```

LinkageType:

```
C
C++
D
Windows
Pascal
System
Objective-C
```

D provides an easy way to call C functions and operating system API functions, as compatibility with both is essential. The *LinkageType* is case sensitive, and is meant to be extensible by the implementation (they are **not** keywords). C and D must be supplied, the others are what makes sense for the implementation. C++ offers limited compatibility with C++. Objective-C offers limited compatibility with Objective-C, see the Interfacing to Objective-C documentation for more information. System is the same as Windows on Windows platforms, and C on other platforms. **Implementation Note:** for Win32 platforms, Windows and Pascal should exist.

C function calling conventions are specified by:

```
extern (C):
    int foo(); // call foo() with C conventions
```

D conventions are:

```
extern (D):
```

Windows API conventions are:

```
extern (Windows):
    void *VirtualAlloc(
        void *lpAddress,
        uint dwSize,
        uint flAllocationType,
        uint flProtect
    );
```

The Windows convention is distinct from the C convention only on Win32 platforms, where it is equivalent to the stdcall convention.

Note that a lone **extern** declaration is used as a storage class.

C++ Namespaces

The linkage form `extern (C++, IdentifierList)` creates C++ declarations that reside in C++ namespaces. The *IdentifierList* specifies the namespaces.

```
extern (C++, N) { void foo(); }
```

refers to the C++ declaration:

```
namespace N { void foo(); }
```

and can be referred to with or without qualification:

```
foo();  
N.foo();
```

Namespaces create a new named scope that is imported into its enclosing scope.

```
extern (C++, N) { void foo(); void bar(); }  
extern (C++, M) { void foo(); }
```

```
bar();    // ok  
foo();    // error - N.foo() or M.foo() ?  
M.foo();  // ok
```

Multiple identifiers in the *IdentifierList* create nested namespaces:

```
extern (C++, N.M) { extern (C++) { extern (C++, R) { void foo(); } } }  
N.M.R.foo();
```

refers to the C++ declaration:

```
namespace N { namespace M { namespace R { void foo(); } } }
```

align Attribute

AlignAttribute:

```
align  
align ( AssignExpression )
```

Specifies the alignment of:

1. variables
2. struct fields
3. union fields
4. class fields
5. struct, union, and class types

`align` by itself sets it to the default, which matches the default member alignment of the companion C compiler.

```
struct S
{
    align:
        byte a;    // placed at offset 0
        int b;     // placed at offset 4
        long c;    // placed at offset 8
}
```

```
auto sz = S.sizeof; // 16
```

AssignExpression specifies the alignment which matches the behavior of the companion C compiler when non-default alignments are used. It must be a positive power of 2.

A value of 1 means that no alignment is done; fields are packed together.

```
struct S
{
    align (1):
        byte a;    // placed at offset 0
        int b;     // placed at offset 1
        long c;    // placed at offset 5
}
```

```
auto sz = S.sizeof; // 16
```

The alignment for the fields of an aggregate does not affect the alignment of the aggregate itself - that is affected by the alignment setting outside of the aggregate.

```
align (2) struct S
{
    align (1):
        byte a;    // placed at offset 0
        int b;     // placed at offset 1
        long c;    // placed at offset 5
}
```

```
auto sz = S.sizeof; // 14
```

Setting the alignment of a field aligns it to that power of 2, regardless of the size of the field.

```
struct S
{
```

```

align (4):
    byte a;    // placed at offset 0
    byte b;    // placed at offset 4
    short c;   // placed at offset 8
}

auto sz = S.sizeof; // 12

```

Do not align references or pointers that were allocated using *NewExpression* on boundaries that are not a multiple of `size_t`. The garbage collector assumes that pointers and references to gc allocated objects will be on `size_t` byte boundaries. If they are not, undefined behavior will result.

The *AlignAttribute* is reset to the default when entering a function scope or a non-anonymous struct, union, class, and restored when exiting that scope. It is not inherited from a base class.

deprecated Attribute

```

DeprecatedAttribute:
    deprecated
    deprecated ( AssignExpression )

```

It is often necessary to deprecate a feature in a library, yet retain it for backwards compatibility. Such declarations can be marked as **deprecated**, which means that the compiler can be instructed to produce an error if any code refers to deprecated declarations:

```

deprecated
{
    void oldFoo();
}

oldFoo();    // Deprecated: function test.oldFoo is deprecated

```

Optionally a string literal or manifest constant can be used to provide additional information in the deprecation message.

```

deprecated("Don't use bar") void oldBar();
oldBar();    // Deprecated: function test.oldBar is deprecated - Don't use bar

```

Calling CTFE-able functions or using manifest constants is also possible.

```

import std.format;
enum Message = format("%s and all its members are obsolete", Foobar.stringof);
deprecated(Message) class Foobar {}

```

```

auto f = new Foobar();    // Deprecated: class test.Foobar is deprecated - Foobar ↯
    ↪ and all its members are obsolete
deprecated(format("%s is also obsolete", "This class")) class BarFoo {}
auto bf = new BarFoo();   // Deprecated: class test.BarFoo is deprecated - This ↯
    ↪ class is also obsolete

```

Implementation Note: The compiler should have a switch specifying if `deprecated` should be ignored, cause a warning, or cause an error during compilation.

Visibility Attribute

VisibilityAttribute:

```

private
package
package ( IdentifierList )
protected
public
export

```

Visibility is an attribute that is one of `private`, `package`, `protected`, `public` or `export`.

They may be referred to as protection attributes in documents written before DIP22 was implemented.

Symbols with `private` visibility can only be accessed from within the same module. Private member functions are implicitly final and cannot be overridden.

`package` extends private so that package members can be accessed from code in other modules that are in the same package. If no identifier is provided, this applies to the innermost package only, or defaults to `private` if a module is not nested in a package.

`package` may have an optional parameter, dot-separated identifier list which is resolved as the qualified package name. If this optional parameter is present, the symbol will be visible by this package and all its descendants.

`protected` only applies inside classes (and templates as they can be mixed in) and means that a symbol can only be seen by members of the same module, or by a derived class. If accessing a protected instance member through a derived class member function, that member can only be accessed for the object instance which can be implicitly cast to the same type as ‘this’. `protected` module members are illegal.

`public` means that any code within the executable can see the member. It is the default visibility attribute.

`export` means that any code outside the executable can access the member. `export` is analogous to exporting definitions from a DLL.

Visibility participates in symbol name lookup.

const Attribute

The `const` attribute changes the type of the declared symbol from `T` to `const(T)`, where `T` is the type specified (or inferred) for the introduced symbol in the absence of `const`.

```
const int foo = 7;
static assert(is(typeof(foo) == const(int)));

const
{
    double bar = foo + 6;
}
static assert(is(typeof(bar) == const(double)));

class C
{
    const void foo();
    const
    {
        void bar();
    }
    void baz() const;
}
#pragma(msg, typeof(C.foo)); // const void()
#pragma(msg, typeof(C.bar)); // const void()
#pragma(msg, typeof(C.baz)); // const void()
static assert(is(typeof(C.foo) == typeof(C.bar)) &&
              is(typeof(C.bar) == typeof(C.baz)));
```

immutable Attribute

The `immutable` attribute modifies the type from `T` to `immutable(T)`, the same way as `const` does.

inout Attribute

The `inout` attribute modifies the type from `T` to `inout(T)`, the same way as `const` does.

shared Attribute

The `shared` attribute modifies the type from `T` to `shared(T)`, the same way as `const` does.

`__gshared` Attribute

By default, non-immutable global declarations reside in thread local storage. When a global variable is marked with the `__gshared` attribute, its value is shared across all threads.

```
int foo;           // Each thread has its own exclusive copy of foo.
__gshared int bar; // bar is shared by all threads.
```

`__gshared` may also be applied to member variables and local variables. In these cases, `__gshared` is equivalent to `static`, except that the variable is shared by all threads rather than being thread local.

```
class Foo
{
    __gshared int bar;
}

int foo()
{
    __gshared int bar = 0;
    return bar++; // Not thread safe.
}
```

Unlike the `shared` attribute, `__gshared` provides no safe-guards against data races or other multi-threaded synchronization issues. It is the responsibility of the programmer to ensure that access to variables marked `__gshared` is synchronized correctly.

`__gshared` is disallowed in safe mode.

`@disable` Attribute

A reference to a declaration marked with the `@disable` attribute causes a compile time error. This can be used to explicitly disallow certain operations or overloads at compile time rather than relying on generating a runtime error.

```
@disable void foo() { }

void main()
{
    foo(); // error, foo is disabled
}
```

Disabling struct no-arg constructor disallows default construction of the struct.

Disabling struct postblit makes the struct not copyable.

@nobjc Attribute

`@nobjc` applies to functions, and means that that function does not allocate memory on the GC heap, either directly such as with *NewExpression* or indirectly through functions it may call, or through language features such as array concatenation and dynamic closures.

```
@nobjc void foo(char[] a)
{
    auto p = new int; // error, operator new allocates
    a ~= 'c';         // error, appending to arrays allocates
    bar();             // error, bar() may allocate
}

void bar() { }
```

`@nobjc` affects the type of the function. An `@nobjc` function is covariant with a non-`@nobjc` function.

```
void function() fp;
void function() @nobjc gp; // pointer to @nobjc function

void foo();
@nobjc void bar();

void test()
{
    fp = &foo; // ok
    fp = &bar; // ok, it's covariant
    gp = &foo; // error, not contravariant
    gp = &bar; // ok
}
```

@property Attribute

See Property Functions.

nothrow Attribute

See Nothrow Functions.

pure Attribute

See Pure Functions.

ref Attribute

See Ref Functions.

return Attribute

See Return Ref Parameters.

override Attribute

The **override** attribute applies to virtual functions. It means that the function must override a function with the same name and parameters in a base class. The override attribute is useful for catching errors when a base class's member function gets its parameters changed, and all derived classes need to have their overriding functions updated.

```
class Foo
{
    int bar();
    int abc(int x);
}

class Foo2 : Foo
{
    override
    {
        int bar(char c); // error, no bar(char) in Foo
        int abc(int x);  // ok
    }
}
```

static Attribute

The **static** attribute applies to functions and data. It means that the declaration does not apply to a particular instance of an object, but to the type of the object. In other words, it means there is no **this** reference. **static** is ignored when applied to other declarations.

```
class Foo
{
    static int bar() { return 6; }
    int foobar() { return 7; }
}
```

...

```

Foo f = new Foo;
Foo.bar();      // produces 6
Foo.fooBar();   // error, no instance of Foo
f.bar();        // produces 6;
f.fooBar();     // produces 7;

```

Static functions are never virtual.

Static data has one instance per thread, not one per object.

Static does not have the additional C meaning of being local to a file. Use the `private` attribute in D to achieve that. For example:

```

module foo;
int x = 3;      // x is global
private int y = 4; // y is local to module foo

```

auto Attribute

The `auto` attribute is used when there are no other attributes and type inference is desired.

```

auto i = 6.8;    // declare i as a double

```

scope Attribute

The `scope` attribute is used for local variables and for class declarations. For class declarations, the `scope` attribute creates a *scope* class. For local declarations, `scope` implements the RAI (Resource Acquisition Is Initialization) protocol. This means that the destructor for an object is automatically called when the reference to it goes out of scope. The destructor is called even if the scope is exited via a thrown exception, thus `scope` is used to guarantee cleanup.

If there is more than one `scope` variable going out of scope at the same point, then the destructors are called in the reverse order that the variables were constructed.

`scope` cannot be applied to globals, statics, data members, ref or out parameters. Arrays of `scopes` are not allowed, and `scope` function return values are not allowed. Assignment to a `scope`, other than initialization, is not allowed. **Rationale:** These restrictions may get relaxed in the future if a compelling reason to appears.

abstract Attribute

An abstract member function must be overridden by a derived class. Only virtual member functions may be declared abstract; non-virtual member functions and free-standing functions cannot be declared abstract.

Classes become abstract if any of its virtual member functions are declared abstract or if they are defined within an abstract attribute. Note that an abstract class may also contain non-virtual member functions.

Classes defined within an abstract attribute or with abstract member functions cannot be instantiated directly. They can only be instantiated as a base class of another, non-abstract, class.

Member functions declared as abstract can still have function bodies. This is so that even though they must be overridden, they can still provide ‘base class functionality’, e.g. through `super.foo()` in a derived class. Note that the class is still abstract and cannot be instantiated directly.

User Defined Attributes

User Defined Attributes (UDA) are compile time expressions that can be attached to a declaration. These attributes can then be queried, extracted, and manipulated at compile time. There is no runtime component to them.

Grammatically, a UDA is a StorageClass:

```
UserDefinedAttribute:
    @ ( ArgumentList )
    @ Identifier
    @ Identifier ( ArgumentListopt )
    @ TemplateInstance
    @ TemplateInstance ( ArgumentListopt )
```

And looks like:

```
@(3) int a;
@("string", 7) int b;

enum Foo;
@Foo int c;

struct Bar
{
    int x;
}

@Bar(3) int d;
```

If there are multiple UDAs in scope for a declaration, they are concatenated:

```
@(1)
{
    @(2) int a;           // has UDA's (1, 2)
    @("string") int b;   // has UDA's (1, "string")
}
```

UDA's can be extracted into an expression tuple using `__traits`:

```
@('c') string s;
pragma(msg, __traits(getAttributes, s)); // prints tuple('c')
```

If there are no user defined attributes for the symbol, an empty tuple is returned. The expression tuple can be turned into a manipulatable tuple:

```
template Tuple (T...)
{
    alias Tuple = T;
}

enum EEE = 7;
@("hello") struct SSS { }
@(3) { @(4) @EEE @SSS int foo; }

alias TP = Tuple!(__traits(getAttributes, foo));

pragma(msg, TP); // prints tuple(3, 4, 7, (SSS))
pragma(msg, TP[2]); // prints 7
```

Of course the tuple types can be used to declare things:

```
TP[3] a; // a is declared as an SSS
```

The attribute of the type name is not the same as the attribute of the variable:

```
pragma(msg, __traits(getAttributes, typeof(a))); // prints tuple("hello")
```

Of course, the real value of UDA's is to be able to create user defined types with specific values. Having attribute values of basic types does not scale. The attribute tuples can be manipulated like any other tuple, and can be passed as the argument list to a template.

Whether the attributes are values or types is up to the user, and whether later attributes accumulate or override earlier ones is also up to how the user interprets them.

Chapter 9

Pragmas

Pragma:

```
pragma ( Identifier )  
pragma ( Identifier , ArgumentList )
```

Pragmas are a way to pass special information to the compiler and to add vendor specific extensions to D. Pragmas can be used by themselves terminated with a ‘;’, they can influence a statement, a block of statements, a declaration, or a block of declarations.

Pragmas can appear as either declarations, *Pragma DeclarationBlock*, or as statements, *PragmaStatement*.

```
pragma(ident);           // just by itself  
  
pragma(ident) declaration; // influence one declaration  
  
pragma(ident): // influence subsequent declarations  
    declaration;  
    declaration;  
  
pragma(ident) // influence block of declarations  
{  
    declaration;  
    declaration;  
}  
  
pragma(ident) statement; // influence one statement
```

```
pragma(ident)    // influence block of statements
{
    statement;
    statement;
}
```

The kind of pragma it is determined by the *Identifier*. *ExpressionList* is a comma-separated list of *AssignExpressions*. The *AssignExpressions* must be parsable as expressions, but what they mean semantically is up to the individual pragma semantics.

Predefined Pragmas

All implementations must support these, even if by just ignoring them:

- `pragma inline`
- `pragma lib`
- `pragma mangle`
- `pragma msg`
- `pragma startaddress`

inline Affects whether functions are inlined or not. If at the declaration level, it affects the functions declared in the block it controls. If inside a function, it affects the function it is enclosed by. If there are multiple pragma inlines in a function, the lexically last one takes effect.

It takes three forms:

1. `pragma(inline)`

Sets the behavior to match the default behavior set by the compiler switch `-inline`.

2. `pragma(inline, false)`

Functions are never inlined.

3. `pragma(inline, true)`

If a function cannot be inlined with the `-inline` switch, an error message is issued. This is expected to be improved in the future to causing functions to always be inlined regardless of compiler switch settings. Whether a compiler can inline a particular function or not is implementation defined.

```

pragma(inline):
int foo(int x) // foo() is never inlined
{
    pragma(inline, true);
    ++x;
    pragma(inline, false); // supercedes the others
    return x + 3;
}

```

lib Inserts a directive in the object file to link in the library specified by the *AssignExpression*. The *AssignExpressions* must be a string literal:

```
pragma(lib, "foo.lib");
```

mangle Overrides the default mangling for a symbol. It's only effective when the symbol is a function declaration or a variable declaration. For example this allows linking to a symbol which is a D keyword, which would normally be disallowed as a symbol name:

```

pragma(mangle, "body")
extern(C) void body_func();

```

msg Constructs a message from the arguments and prints to the standard error stream while compiling:

```
pragma(msg, "compiling...", 1, 1.0);
```

startaddress Puts a directive into the object file saying that the function specified in the first argument will be the start address for the program:

```

void foo() { ... }
pragma(startaddress, foo);

```

This is not normally used for application level programming, but is for specialized systems work. For applications code, the start address is taken care of by the runtime library.

Vendor Specific Pragmas

Vendor specific pragma *Identifiers* can be defined if they are prefixed by the vendor's trademarked name, in a similar manner to version identifiers:

```
pragma(DigitalMars_funky_extension) { ... }
```

Compilers must diagnose an error for unrecognized *Pragmas*, even if they are vendor specific ones. This implies that vendor specific pragmas should be wrapped in version statements:

```
version (DigitalMars)
{
    pragma(DigitalMars_funky_extension)
    { ... }
}
```

Chapter 10

Expressions

C and C++ programmers will find the D expressions very familiar, with a few interesting additions.

Expressions are used to compute values with a resulting type. These values can then be assigned, tested, or ignored. Expressions can also have side effects.

Order Of Evaluation

Binary expressions and function arguments are evaluated in strictly left-to-right order. This is similar to Java but different to C and C++, where the evaluation order is unspecified. Thus, the following code is valid and well defined.

```
import std.conv;
int i = 0;
(i = 2) = ++i * i++ + i;
assert(i == 13); // left to right evaluation of side effects
assert(text(++i, ++i) == "1415"); // left to right evaluation of arguments
```

But even though the order of evaluation is well defined, writing code that depends on it is rarely recommended. **Note that dmd currently does not comply with left to right evaluation of function arguments and AssignExpression.**

Expressions

Expression:

CommaExpression

CommaExpression:

AssignExpression

AssignExpression , *CommaExpression*

The left operand of the `,` is evaluated, then the right operand is evaluated. The type of the expression is the type of the right operand, and the result is the result of the right operand.

Assign Expressions

AssignExpression:

```
ConditionalExpression
ConditionalExpression = AssignExpression
ConditionalExpression += AssignExpression
ConditionalExpression -= AssignExpression
ConditionalExpression *= AssignExpression
ConditionalExpression /= AssignExpression
ConditionalExpression %= AssignExpression
ConditionalExpression &= AssignExpression
ConditionalExpression |= AssignExpression
ConditionalExpression ^= AssignExpression
ConditionalExpression ~= AssignExpression
ConditionalExpression <<= AssignExpression
ConditionalExpression >>= AssignExpression
ConditionalExpression >>>= AssignExpression
ConditionalExpression ^^= AssignExpression
```

The right operand is implicitly converted to the type of the left operand, and assigned to it. The result type is the type of the left operand, and the result value is the value of the left operand after the assignment.

The left operand must be an lvalue.

Assignment Operator Expressions

Assignment operator expressions, such as:

```
a op= b
```

are semantically equivalent to:

```
a = cast(typeof(a))(a op b)
```

except that:

- operand `a` is only evaluated once

- overloading *op* uses a different function than overloading *op=* does
- the left operand of *>>=* does not undergo integral promotions before shifting

Conditional Expressions

ConditionalExpression:

OrOrExpression

OrOrExpression ? *Expression* : *ConditionalExpression*

The first expression is converted to **bool**, and is evaluated.

If it is **true**, then the second expression is evaluated, and its result is the result of the conditional expression.

If it is **false**, then the third expression is evaluated, and its result is the result of the conditional expression.

If either the second or third expressions are of type **void**, then the resulting type is **void** ↗ ↘ . Otherwise, the second and third expressions are implicitly converted to a common type which becomes the result type of the conditional expression.

OrOr Expressions

OrOrExpression:

AndAndExpression

OrOrExpression || *AndAndExpression*

The result type of an *OrOrExpression* is **bool**, unless the right operand has type **void**, when the result is type **void**.

The *OrOrExpression* evaluates its left operand.

If the left operand, converted to type **bool**, evaluates to **true**, then the right operand is not evaluated. If the result type of the *OrOrExpression* is **bool** then the result of the expression is **true**.

If the left operand is **false**, then the right operand is evaluated. If the result type of the *OrOrExpression* is **bool** then the result of the expression is the right operand converted to type **bool**.

AndAnd Expressions

AndAndExpression:

OrExpression

AndAndExpression && *OrExpression*

The result type of an *AndAndExpression* is `bool`, unless the right operand has type `void`, when the result is type `void`.

The *AndAndExpression* evaluates its left operand.

If the left operand, converted to type `bool`, evaluates to `false`, then the right operand is not evaluated. If the result type of the *AndAndExpression* is `bool` then the result of the expression is `false`.

If the left operand is `true`, then the right operand is evaluated. If the result type of the *AndAndExpression* is `bool` then the result of the expression is the right operand converted to type `bool`.

Bitwise Expressions

Bit wise expressions perform a bitwise operation on their operands. Their operands must be integral types. First, the default integral promotions are done. Then, the bitwise operation is done.

Or Expressions

OrExpression:

XorExpression

OrExpression | *XorExpression*

The operands are OR'd together.

Xor Expressions

XorExpression:

AndExpression

XorExpression ^ *AndExpression*

The operands are XOR'd together.

And Expressions

AndExpression:

CmpExpression

AndExpression & *CmpExpression*

The operands are AND'd together.

Compare Expressions

```
CmpExpression:
    ShiftExpression
    EqualExpression
    IdentityExpression
    RelExpression
    InExpression
```

Equality Expressions

```
EqualExpression:
    ShiftExpression == ShiftExpression
    ShiftExpression != ShiftExpression
```

Equality expressions compare the two operands for equality (==) or inequality (!=). The type of the result is `bool`. The operands go through the usual conversions to bring them to a common type before comparison.

If they are integral values or pointers, equality is defined as the bit pattern of the type matches exactly.

Equality for floating point types is more complicated. `-0` and `+0` compare as equal. If either or both operands are NAN, then both the `==` returns false and `!=` returns `true`. Otherwise, the bit patterns are compared for equality.

For complex numbers, equality is defined as equivalent to:

```
x.re == y.re && x.im == y.im
```

and inequality is defined as equivalent to:

```
x.re != y.re || x.im != y.im
```

Equality for struct objects means the logical product of all equality results of the corresponding object fields. If all struct fields use bitwise equality, the whole struct equality could be optimized to one memory comparison operation (the existence of alignment holes in the objects is accounted for, usually by setting them all to 0 upon initialization).

For class and struct objects, the expression `(a == b)` is rewritten as `a.opEquals(b)`, and `(a != b)` is rewritten as `!a.opEquals(b)`.

For class objects, the `==` and `!=` operators are intended to compare the contents of the objects, however an appropriate `opEquals` override must be defined for this to work. The default `opEquals` provided by the root `Object` class is equivalent to the `is` operator. Comparing against `null` is invalid, as `null` has no contents. Use the `is` and `!is` operators instead.

```

class C;
C c;
if (c == null) // error
    ...
if (c is null) // ok
    ...

```

For static and dynamic arrays, equality is defined as the lengths of the arrays matching, and all the elements are equal.

Identity Expressions

IdentityExpression:
ShiftExpression is ShiftExpression
ShiftExpression !is ShiftExpression

The **is** compares for identity. To compare for nonidentity, use **e1 !is e2**. The type of the result is **bool**. The operands go through the usual conversions to bring them to a common type before comparison.

For class objects, identity is defined as the object references are for the same object. Null class objects can be compared with **is**.

For struct objects and floating point values, identity is defined as the bits in the operands being identical.

For static and dynamic arrays, identity is defined as referring to the same array elements and the same number of elements.

For other operand types, identity is defined as being the same as equality.

The identity operator **is** cannot be overloaded.

Relational Expressions

RelExpression:
ShiftExpression < ShiftExpression
ShiftExpression <= ShiftExpression
ShiftExpression > ShiftExpression
ShiftExpression >= ShiftExpression
ShiftExpression !<=> ShiftExpression
ShiftExpression !<> ShiftExpression
ShiftExpression <> ShiftExpression
ShiftExpression <=> ShiftExpression
ShiftExpression !> ShiftExpression

```

ShiftExpression !>= ShiftExpression
ShiftExpression !< ShiftExpression
ShiftExpression !<= ShiftExpression

```

First, the integral promotions are done on the operands. The result type of a relational expression is `bool`.

For class objects, the result of `Object.opCmp()` forms the left operand, and 0 forms the right operand. The result of the relational expression (`o1 op o2`) is:

```
(o1.opCmp(o2) op 0)
```

It is an error to compare objects if one is `null`.

For static and dynamic arrays, the result of the relational op is the result of the operator applied to the first non-equal element of the array. If two arrays compare equal, but are of different lengths, the shorter array compares as "less" than the longer array.

Integer comparisons

Integer comparisons happen when both operands are integral types.

Integer comparison operators

Operator	Relation
<	less
>	greater
<=	less or equal
>=	greater or equal
==	equal
!=	not equal

It is an error to have one operand be signed and the other unsigned for a `<`, `<=`, `>` or `>=` expression. Use casts to make both operands signed or both operands unsigned.

Floating point comparisons

If one or both operands are floating point, then a floating point comparison is performed.

Useful floating point operations must take into account NAN values. In particular, a relational operator can have NAN operands. The result of a relational operation on float values is less, greater, equal, or unordered (unordered means either or both of the operands is a NAN). That means there are 14 possible comparison conditions to test for:

Floating point comparison operators						
Operator	Greater	Less	Equal	Unordered	Exception	Relation
<code>==</code>	F	F	T	F	no	equal
<code>!=</code>	T	T	F	T	no	unordered, less, or greater
<code>></code>	T	F	F	F	yes	greater
<code>>=</code>	T	F	T	F	yes	greater or equal
<code><</code>	F	T	F	F	yes	less
<code><=</code>	F	T	T	F	yes	less or equal
<code>!<>=</code>	F	F	F	T	no	unordered
<code><></code>	T	T	F	F	yes	less or greater
<code><>=</code>	T	T	T	F	yes	less, equal, or greater
<code>!<=</code>	T	F	F	T	no	unordered or greater
<code>!<</code>	T	F	T	T	no	unordered, greater, or equal
<code>!>=</code>	F	T	F	T	no	unordered or less
<code>!></code>	F	T	T	T	no	unordered, less, or equal
<code>!<></code>	F	F	T	T	no	unordered or equal

Notes:

1. For floating point comparison operators, `(a !op b)` is not the same as `!(a op b)`.
2. "Unordered" means one or both of the operands is a NAN.
3. "Exception" means the *Invalid Exception* is raised if one of the operands is a NAN. It does not mean an exception is thrown. The *Invalid Exception* can be checked using the functions in `core.stdc.fenv`.

Class comparisons

For class objects, the relational operators compare the contents of the objects. Therefore, comparing against `null` is invalid, as `null` has no contents.

```
class C;
C c;
if (c < null) // error
    ...
```

In Expressions

InExpression:

```
ShiftExpression in ShiftExpression
ShiftExpression !in ShiftExpression
```

An associative array can be tested to see if an element is in the array:

```
int foo[char[]];
...
if ("hello" in foo)
    ...
```

The `in` expression has the same precedence as the relational expressions `<`, `<=`, etc. The return value of the *InExpression* is `null` if the element is not in the array; if it is in the array it is a pointer to the element.

The `!in` expression is the logical negation of the `in` operation.

Shift Expressions

```
ShiftExpression:
    AddExpression
    ShiftExpression << AddExpression
    ShiftExpression >> AddExpression
    ShiftExpression >>> AddExpression
```

The operands must be integral types, and undergo the usual integral promotions. The result type is the type of the left operand after the promotions. The result value is the result of shifting the bits by the right operand's value.

`<<` is a left shift. `>>` is a signed right shift. `>>>` is an unsigned right shift.

It's illegal to shift by the same or more bits than the size of the quantity being shifted:

```
int c;
auto x = c << 33;           // error
```

Add Expressions

```
AddExpression:
    MulExpression
    AddExpression + MulExpression
    AddExpression - MulExpression
    CatExpression
```

If the operands are of integral types, they undergo integral promotions, and then are brought to a common type using the usual arithmetic conversions.

If either operand is a floating point type, the other is implicitly converted to floating point and they are brought to a common type via the usual arithmetic conversions.

If the operator is `+` or `-`, and the first operand is a pointer, and the second is an integral type, the resulting type is the type of the first operand, and the resulting value is the pointer plus (or minus) the second operand multiplied by the size of the type pointed to by the first operand.

If the second operand is a pointer, and the first is an integral type, and the operator is `+`, the operands are reversed and the pointer arithmetic just described is applied.

If both operands are pointers, and the operator is `+`, then it is illegal.

If both operands are pointers, and the operator is `-`, the pointers are subtracted and the result is divided by the size of the type pointed to by the operands. It is an error if the pointers point to different types. The type of the result is `ptrdiff_t`.

If both operands are of integral types and an overflow or underflow occurs in the computation, wrapping will happen. That is, `uint.max + 1 == uint.min` and `uint.min - 1 == uint.max`.

Add expressions for floating point operands are not associative.

Cat Expressions

CatExpression:

AddExpression ~ MulExpression

A *CatExpression* concatenates arrays, producing a dynamic array with the result. The arrays must be arrays of the same element type. If one operand is an array and the other is of that array's element type, that element is converted to an array of length 1 of that element, and then the concatenation is performed.

Mul Expressions

MulExpression:

UnaryExpression

*MulExpression * UnaryExpression*

MulExpression / UnaryExpression

MulExpression % UnaryExpression

The operands must be arithmetic types. They undergo integral promotions, and then are brought to a common type using the usual arithmetic conversions.

For integral operands, the `*`, `/`, and `%` correspond to multiply, divide, and modulus operations. For multiply, overflows are ignored and simply chopped to fit into the integral type.

For integral operands of the `/` and `%` operators, the quotient rounds towards zero and the remainder has the same sign as the dividend. If the divisor is zero, an `Exception` is thrown.

For floating point operands, the `*` and `/` operations correspond to the IEEE 754 floating point equivalents. the IEEE 754 remainder. For example, 15.0 for IEEE 754, `remainder(15.0,10.0) == -5.0`.

Mul expressions for floating point operands are not associative.

Unary Expressions

UnaryExpression:

```
& UnaryExpression
++ UnaryExpression
-- UnaryExpression
* UnaryExpression
- UnaryExpression
+ UnaryExpression
! UnaryExpression
ComplementExpression
( Type ) . Identifier
( Type ) . TemplateInstance
DeleteExpression
CastExpression
PowExpression
```

Complement Expressions

ComplementExpression:

```
~ UnaryExpression
```

ComplementExpressions work on integral types (except `bool`). All the bits in the value are complemented.

Note: unlike in C and C++, the usual integral promotions are not performed prior to the complement operation.

New Expressions

NewExpression:

```
new AllocatorArgumentsopt Type
NewExpressionWithArgs
```

```

NewExpressionWithArgs:
    new AllocatorArgumentsopt Type [ AssignExpression ]
    new AllocatorArgumentsopt Type ( ArgumentListopt )
    NewAnonClassExpression

```

```

AllocatorArguments:
    ( ArgumentListopt )

```

```

ArgumentList:
    AssignExpression
    AssignExpression ,
    AssignExpression , ArgumentList

```

NewExpressions are used to allocate memory on the garbage collected heap (default) or using a class or struct specific allocator.

To allocate multidimensional arrays, the declaration reads in the same order as the prefix array declaration order.

```

char[] [] foo;    // dynamic array of strings
...
foo = new char[] [30]; // allocate array of 30 strings

```

The above allocation can also be written as:

```

foo = new char[] [] (30); // allocate array of 30 strings

```

To allocate the nested arrays, multiple arguments can be used:

```

int[] [] [] bar;
...
bar = new int[] [] [] (5, 20, 30);

```

Which is equivalent to:

```

bar = new int[] [] [5];
foreach (ref a; bar)
{
    a = new int[] [20];
    foreach (ref b; a)
    {
        b = new int[30];
    }
}

```

If there is a `new (ArgumentList)`, then those arguments are passed to the class or struct specific allocator function after the size argument.

If a *NewExpression* is used as an initializer for a function local variable with `scope` storage class, and the *ArgumentList* to `new` is empty, then the instance is allocated on the stack rather than the heap or using the class specific allocator.

Delete Expressions

DeleteExpression:

`delete` *UnaryExpression*

If the *UnaryExpression* is a class object reference, and there is a destructor for that class, the destructor is called for that object instance.

Next, if the *UnaryExpression* is a class object reference, or a pointer to a struct instance, and the class or struct has overloaded operator delete, then that operator delete is called for that class object instance or struct instance.

Otherwise, the garbage collector is called to immediately free the memory allocated for the class instance or struct instance. If the garbage collector was not used to allocate the memory for the instance, undefined behavior will result.

If the *UnaryExpression* is a pointer or a dynamic array, the garbage collector is called to immediately release the memory. If the garbage collector was not used to allocate the memory for the instance, undefined behavior will result.

The pointer, dynamic array, or reference is set to `null` after the delete is performed. Any attempt to reference the data after the deletion via another reference to it will result in undefined behavior.

If *UnaryExpression* is a variable allocated on the stack, the class destructor (if any) is called for that instance. Neither the garbage collector nor any class deallocator is called.

Cast Expressions

CastExpression:

`cast (Type) UnaryExpression`

`cast (TypeCtorsopt) UnaryExpression`

A *CastExpression* converts the *UnaryExpression* to *Type*.

```
cast(foo) -p; // cast (-p) to type foo
(foo) - p;    // subtract p from foo
```

Any casting of a class reference to a derived class reference is done with a runtime check to make sure it really is a downcast. `null` is the result if it isn't.

Note: This is equivalent to the behavior of the `dynamic_cast` operator in C++.

```
class A { ... }
class B : A { ... }

void test(A a, B b)
{
    B bx = a;           // error, need cast
    B bx = cast(B) a;   // bx is null if a is not a B
    A ax = b;           // no cast needed
    A ax = cast(A) b;   // no runtime check needed for upcast
}
```

In order to determine if an object `o` is an instance of a class `B` use a cast:

```
if (cast(B) o)
{
    // o is an instance of B
}
else
{
    // o is not an instance of B
}
```

Casting a pointer type to and from a class type is done as a type paint (i.e. a reinterpret cast).

Casting a dynamic array to another dynamic array is done only if the array lengths multiplied by the element sizes match. The cast is done as a type paint, with the array length adjusted to match any change in element size. If there's not a match, a runtime error is generated.

```
import std.stdio;

int main()
{
    byte[] a = [1,2,3];
    auto b = cast(int[])a; // runtime array cast misalignment

    int[] c = [1, 2, 3];
    auto d = cast(byte[])c; // ok
    // prints:
    // [1, 0, 0, 0, 2, 0, 0, 0, 3, 0, 0, 0]
    writeln(d);
}
```

```
    return 0;
}
```

Casting a floating point literal from one type to another changes its type, but internally it is retained at full precision for the purposes of constant folding.

```
void test()
{
    real a = 3.40483L;
    real b;
    b = 3.40483;      // literal is not truncated to double precision
    assert(a == b);
    assert(a == 3.40483);
    assert(a == 3.40483L);
    assert(a == 3.40483F);
    double d = 3.40483; // truncate literal when assigned to variable
    assert(d != a);      // so it is no longer the same
    const double x = 3.40483; // assignment to const is not
    assert(x == a);      // truncated if the initializer is visible
}
```

Casting a floating point value to an integral type is the equivalent of converting to an integer using truncation.

```
void main()
{
    int a = cast(int) 0.8f;
    assert(a == 0);
    long b = cast(long) 1.5;
    assert(b == 1L);
    long c = cast(long) -1.5;
    assert(c == -1);
}
```

Casting a value v to a struct S , when value is not a struct of the same type, is equivalent to: $S(v)$

Casting to a *CastQual* replaces the qualifiers to the type of the *UnaryExpression*.

```
shared int x;
assert(is(typeof(cast(const)x) == const int));
```

Casting with no *Type* or *CastQual* removes any top level `const`, `immutable`, `shared` or `inout` type modifiers from the type of the *UnaryExpression*.

```
shared int x;
assert(is(typeof(cast())x) == int));
```

Casting an expression to `void` type is allowed to mark that the result is unused. On *ExpressionStatement*, it could be used properly to avoid "has no effect" error.

```
void foo(lazy void exp) {}
void main()
{
    foo(10);           // NG - has no effect in expression '10'
    foo(cast(void)10); // OK
}
```

Pow Expressions

```
PowExpression:
    PostfixExpression
    PostfixExpression ^^ UnaryExpression
```

PowExpression raises its left operand to the power of its right operand.

Postfix Expressions

```
PostfixExpression:
    PrimaryExpression
    PostfixExpression . Identifier
    PostfixExpression . TemplateInstance
    PostfixExpression . NewExpression
    PostfixExpression ++
    PostfixExpression --
    PostfixExpression ( ArgumentListopt )
    TypeCtorsopt BasicType ( ArgumentListopt )
    IndexExpression
    SliceExpression
```

Index Expressions

```
IndexExpression:
    PostfixExpression [ ArgumentList ]
```

PostfixExpression is evaluated.

If *PostfixExpression* is an expression of type static array or dynamic array, the symbol \$ is set to be the number of elements in the array.

If *PostfixExpression* is an *ExpressionTuple*, the symbol \$ is set to be the number of elements in the tuple.

A new declaration scope is created for the evaluation of the *ArgumentList* and \$ appears in that scope only.

If *PostfixExpression* is an *ExpressionTuple*, then the *ArgumentList* must consist of only one argument, and that must be statically evaluatable to an integral constant. That integral constant *n* then selects the *n*th expression in the *ExpressionTuple*, which is the result of the *IndexExpression*. It is an error if *n* is out of bounds of the *ExpressionTuple*.

Slice Expressions

SliceExpression:

PostfixExpression []

PostfixExpression [*Slice* ,_{opt}]

Slice:

AssignExpression

AssignExpression , *Slice*

AssignExpression .. *AssignExpression*

AssignExpression .. *AssignExpression* , *Slice*

PostfixExpression is evaluated. if *PostfixExpression* is an expression of type static array or dynamic array, the special variable \$ is declared and set to be the length of the array. A new declaration scope is created for the evaluation of the *AssignExpression..AssignExpression* and \$ appears in that scope only.

The first *AssignExpression* is taken to be the inclusive lower bound of the slice, and the second *AssignExpression* is the exclusive upper bound. The result of the expression is a slice of the *PostfixExpression* array.

If the [] form is used, the slice is of the entire array.

The type of the slice is a dynamic array of the element type of the *PostfixExpression*.

A *SliceExpression* is not a modifiable lvalue.

If the slice bounds can be known at compile time, the slice expression is implicitly convertible to an lvalue of static array. For example:

```
arr[a .. b]      // typed T[]
```

If both *a* and *b* are integers (may be constant-folded), the slice expression can be converted to a static array type *T*[*b* - *a*].

```

void foo(int[2] a)
{
    assert(a == [2, 3]);
}
void bar(ref int[2] a)
{
    assert(a == [2, 3]);
    a[0] = 4;
    a[1] = 5;
    assert(a == [4, 5]);
}
void baz(int[3] a) {}

void main()
{
    int[] arr = [1, 2, 3];

    foo(arr[1 .. 3]);
    assert(arr == [1, 2, 3]);

    bar(arr[1 .. 3]);
    assert(arr == [1, 4, 5]);

    //baz(arr[1 .. 3]); // cannot match length
}

```

Following forms of slice expression can be convertible to a static array type:

e An expression that contains no side effects.

a, *b* Integers (that may be constant-folded).

Form	The length calculated at compile time
<code>arr[]</code>	The compile time length of <code>arr</code> if it's known.
<code>arr[a .. b]</code>	$b - a$
<code>arr[\textit{e}-a .. \textit{e}]</code>	a
<code>arr[\textit{e} .. \textit{e}+b]</code>	b
<code>arr[\textit{e}-a .. \textit{e}+b]</code>	$a + b$
<code>arr[\textit{e}+a .. \textit{e}+b]</code>	$b - a$ <i>if</i> $a \leq b$
<code>arr[\textit{e}-a .. \textit{e}-b]</code>	$a - b$ <i>if</i> $a \geq b$

If *PostfixExpression* is an *ExpressionTuple*, then the result of the slice is a new *ExpressionTuple* formed from the upper and lower bounds, which must statically evaluate to integral constants. It is an error if those bounds are out of range.

Primary Expressions

PrimaryExpression:

```

  Identifier
  . Identifier
  TemplateInstance
  . TemplateInstance
  this
  super
  null
  true
  false
  $
  IntegerLiteral
  FloatLiteral
  CharacterLiteral
  StringLiterals
  ArrayLiteral
  AssocArrayLiteral
  FunctionLiteral
  AssertExpression
  MixinExpression
  ImportExpression
  NewExpressionWithArgs
  BasicTypeX . Identifier
  BasicTypeX ( ArgumentListopt )
  TypeCtor ( Type ) . Identifier
  TypeCtor ( Type ) ( ArgumentListopt )
  Typeof
  TypeidExpression
  IsExpression
  ( Expression )
  TraitsExpression
  SpecialKeyword
```

.Identifier

Identifier is looked up at module scope, rather than the current lexically nested scope.

this

Within a non-static member function, `this` resolves to a reference to the object for which the function was called. If the object is an instance of a struct, `this` will be a pointer to that instance. If a member function is called with an explicit reference to `typeof(this)`, a non-virtual call is made:

```
class A
{
    char get() { return 'A'; }

    char foo() { return typeof(this).get(); }
    char bar() { return this.get(); }
}

class B : A
{
    override char get() { return 'B'; }
}

void main()
{
    B b = new B();

    assert(b.foo() == 'A');
    assert(b.bar() == 'B');
}
```

Assignment to `this` is not allowed.

super

`super` is identical to `this`, except that it is cast to `this`'s base class. It is an error if there is no base class. It is an error to use `super` within a struct member function. (Only class `Object` has no base class.) If a member function is called with an explicit reference to `super`, a non-virtual call is made.

Assignment to `super` is not allowed.

null

null represents the null value for pointers, pointers to functions, delegates, dynamic arrays, associative arrays, and class objects. If it has not already been cast to a type, it is given the singular type `typeof(null)` and it is an exact conversion to convert it to the null value for pointers, pointers to functions, delegates, etc. After it is cast to a type, such conversions are implicit, but no longer exact.

true, false

These are of type `bool` and when cast to another integral type become the values 1 and 0, respectively.

Character Literals

Character literals are single characters and resolve to one of type `char`, `wchar`, or `dchar`. If the literal is a `\u` escape sequence, it resolves to type `wchar`. If the literal is a `\U` escape sequence, it resolves to type `dchar`. Otherwise, it resolves to the type with the smallest size it will fit into.

String Literals

```
StringLiterals:
    StringLiteral
    StringLiterals StringLiteral
```

String literals can implicitly convert to any of the following types, they have equal weight:

```
immutable(char)*
immutable(wchar)*
immutable(dchar)*
immutable(char)[]
immutable(wchar)[]
immutable(dchar)[]
```

By default, a string literal is typed as a dynamic array, but the element count is known at compile time. So all string literals can be implicitly converted to static array types.

```
void foo(char[2] a)
{
    assert(a == "bc");
}
void bar(ref const char[2] a)
{
```

```

    assert(a == "bc");
}
void baz(const char[3] a) {}

void main()
{
    string str = "abc";
    foo(str[1 .. 3]);
    bar(str[1 .. 3]);
    //baz(str[1 .. 3]); // cannot match length
}

```

String literals have a 0 appended to them, which makes them easy to pass to C or C++ functions expecting a `const char*` string. The 0 is not included in the `.length` property of the string literal.

Array Literals

ArrayLiteral:
`[ArgumentListopt]`

Array literals are a comma-separated list of *AssignExpressions* between square brackets [and]. The *AssignExpressions* form the elements of a dynamic array, the length of the array is the number of elements. The common type of the all elements is taken to be the type of the array element, and all elements are implicitly converted to that type.

```

auto a1 = [1,2,3]; // type is int[], with elements 1, 2 and 3
auto a2 = [1u,2,3]; // type is uint[], with elements 1u, 2u, and 3u

```

By default, an array literal is typed as a dynamic array, but the element count is known at compile time. So all array literals can be implicitly converted to static array types.

```

void foo(long[2] a)
{
    assert(a == [2, 3]);
}
void bar(ref long[2] a)
{
    assert(a == [2, 3]);
    a[0] = 4;
    a[1] = 5;
    assert(a == [4, 5]);
}

```

```

void baz(const char[3] a) {}

void main()
{
    long[] arr = [1, 2, 3];

    foo(arr[1 .. 3]);
    assert(arr == [1, 2, 3]);

    bar(arr[1 .. 3]);
    assert(arr == [1, 4, 5]);

    //baz(arr[1 .. 3]); // cannot match length
}

```

If any of the arguments in the *ArgumentList* are an *ExpressionTuple*, then the elements of the *ExpressionTuple* are inserted as arguments in place of the tuple.

Array literals are allocated on the memory managed heap. Thus, they can be returned safely from functions:

```

int[] foo()
{
    return [1, 2, 3];
}

```

When array literals are cast to another array type, each element of the array is cast to the new element type. When arrays that are not literals are cast, the array is reinterpreted as the new type, and the length is recomputed:

```

import std.stdio;

void main()
{
    // cast array literal
    const short[] ct = cast(short[]) [cast(byte)1, 1];
    // this is equivalent with:
    // const short[] ct = [cast(short)1, cast(short)1];
    writeln(ct); // writes [1, 1]

    // cast other array expression
    // --> normal behavior of CastExpression
    byte[] arr = [cast(byte)1, cast(byte)1];
}

```

```

    short[] rt = cast(short[]) arr;
    writeln(rt); // writes [257]
}

```

In other words, casting literal expression will change the literal type.

Associative Array Literals

```

AssocArrayLiteral:
    [ KeyValuePairs ]

KeyValuePairs:
    KeyValuePair
    KeyValuePair , KeyValuePairs

KeyValuePair:
    KeyExpression : ValueExpression

KeyExpression:
    AssignExpression

ValueExpression:
    AssignExpression

```

Associative array literals are a comma-separated list of *key:value* pairs between square brackets [and]. The list cannot be empty. The common type of the all keys is taken to be the key type of the associative array, and all keys are implicitly converted to that type. The common type of the all values is taken to be the value type of the associative array, and all values are implicitly converted to that type. An *AssocArrayLiteral* cannot be used to statically initialize anything.

```

[21u:"he", 38:"ho", 2:"hi"]; // type is string[uint],
                             // with keys 21u, 38u and 2u
                             // and values "he", "ho", and "hi"

```

If any of the keys or values in the *KeyValuePairs* are an *ExpressionTuple*, then the elements of the *ExpressionTuple* are inserted as arguments in place of the tuple.

Function Literals

```

FunctionLiteral:
    function Typeopt ParameterAttributesopt FunctionLiteralBody

```

```

delegate Typeopt ParameterMemberAttributesopt FunctionLiteralBody
ParameterMemberAttributes FunctionLiteralBody
FunctionLiteralBody
Lambda

```

ParameterAttributes:

```
Parameters FunctionAttributesopt
```

ParameterMemberAttributes:

```
Parameters MemberFunctionAttributesopt
```

FunctionLiteralBody:

```
BlockStatement
FunctionContractsopt BodyStatement
```

FunctionLiterals enable embedding anonymous functions and anonymous delegates directly into expressions. *Type* is the return type of the function or delegate, if omitted it is inferred from any *ReturnStatements* in the *FunctionLiteralBody*. (*ArgumentList*) forms the arguments to the function. If omitted it defaults to the empty argument list (). The type of a function literal is pointer to function or pointer to delegate. If the keywords **function** or **delegate** are omitted, it is inferred from whether *FunctionLiteralBody* is actually accessing to the outer context.

For example:

```
int function(char c) fp; // declare pointer to a function
```

```
void test()
{
    static int foo(char c) { return 6; }

    fp = &foo;
}
```

is exactly equivalent to:

```
int function(char c) fp;

void test()
{
    fp = function int(char c) { return 6;} ;
}
```

And:

```
int abc(int delegate(int i));

void test()
{
    int b = 3;
    int foo(int c) { return 6 + b; }

    abc(&foo);
}
```

is exactly equivalent to:

```
int abc(int delegate(int i));

void test()
{
    int b = 3;

    abc( delegate int(int c) { return 6 + b; } );
}
```

and the following where the return type `int` and `function/delegate` are inferred:

```
int abc(int delegate(int i));
int def(int function(int s));

void test()
{
    int b = 3;

    abc( (int c) { return 6 + b; } ); // inferred to delegate
    def( (int c) { return c * 2; } ); // inferred to function
    //def( (int c) { return c * b; } ); // error!
    // Because the FunctionLiteralBody accesses b, then the function literal type
    // is inferred to delegate. But def cannot receive delegate.
}
```

If the type of a function literal can be uniquely determined from its context, the parameter type inference is possible.

```
void foo(int function(int) fp);

void test()
```



```

{
    int function(int) fp = (n) { return n * 2; };
    // The type of parameter n is inferred to int.

    foo((n) { return n * 2; });
    // The type of parameter n is inferred to int.
}

```

Anonymous delegates can behave like arbitrary statement literals. For example, here an arbitrary statement is executed by a loop:

```

double test()
{
    double d = 7.6;
    float f = 2.3;

    void loop(int k, int j, void delegate() statement)
    {
        for (int i = k; i < j; i++)
        {
            statement();
        }
    }

    loop(5, 100, { d += 1; });
    loop(3, 10, { f += 3; });

    return d + f;
}

```

When comparing with nested functions, the **function** form is analogous to static or non-nested functions, and the **delegate** form is analogous to non-static nested functions. In other words, a delegate literal can access stack variables in its enclosing function, a function literal cannot.

Lambdas

Lambda:

```

function Typeopt ParameterAttributes => AssignExpression
delegate Typeopt ParameterMemberAttributes => AssignExpression
ParameterMemberAttributes => AssignExpression
Identifier => AssignExpression

```

Lambdas are a shorthand syntax for *FunctionLiterals*.

1. Just one *Identifier* is rewritten to *Parameters*:
(*Identifier*)
2. The following part \Rightarrow *AssignExpression* is rewritten to *FunctionLiteralBody*:
{ **return** *AssignExpression* ; }

Example usage:

```
import std.stdio;

void main()
{
    auto i = 3;
    auto twice = function (int x) => x * 2;
    auto square = delegate (int x) => x * x;

    auto n = 5;
    auto mul_n = (int x) => x * n;

    writeln(twice(i));    // prints 6
    writeln(square(i));   // prints 9
    writeln(mul_n(i));    // prints 15
}
```

Uniform construction syntax for built-in scalar types

The implicit conversions of built-in scalar types can be explicitly represented by using function call syntax. For example:

```
auto a = short(1);    // implicitly convert an integer literal '1' to short
auto b = double(a);   // implicitly convert a short variable 'a' to double
auto c = byte(128);   // error, 128 cannot be represented in a byte
```

If the argument is omitted, it means default construction of the scalar type:

```
auto a = ushort();    // same as: ushort.init
auto b = wchar();     // same as: wchar.init
auto c = creal();     // same as: creal.init
```

Assert Expressions

AssertExpression:

```
assert ( AssignExpression ,opt )
assert ( AssignExpression , AssignExpression ,opt )
```

The assert expression is used to declare conditions that the programmer asserts must hold at that point in the program if the program logic has been correctly implemented. It can be used both as a debugging tool and as a way of communicating to the compiler facts about the code that it may employ to produce more efficient code.

Programs for which *AssignExpression* is false are invalid. Subsequent to such a false result, the program is in an invalid, non-recoverable state.

As a debugging tool, the compiler may insert checks to verify that the condition indeed holds by evaluating *AssignExpression* at runtime. If it evaluates to a non-null class reference, the class invariant is run. Otherwise, if it evaluates to a non-null pointer to a struct, the struct invariant is run. Otherwise, if the result is false, an **AssertError** is thrown. If the result is true, then no exception is thrown. In this way, if a bug in the code causes the assertion to fail, execution is aborted, prompting the programmer to fix the problem.

It is implementation defined whether the *AssignExpression* is evaluated at run time or not. Programs that rely on side effects of *AssignExpression* are invalid.

The result type of an assert expression is **void**. Asserts are a fundamental part of the Contract Programming support in D.

The expression **assert(0)** is a special case; it signifies that it is unreachable code. Either **AssertError** is thrown at runtime if it is reachable, or the execution is halted (on the x86 processor, a HLT instruction can be used to halt execution). The optimization and code generation phases of compilation may assume that it is unreachable code.

The second *AssignExpression*, if present, must be implicitly convertible to type **const(char)[]**. It is evaluated if the result is false, and the string result is appended to the **AssertError**'s message.

```
void main()
{
    assert(0, "an" ~ "␣error␣message");
}
```

When compiled and run, it will produce the message:

```
Error: AssertError Failure test.d(3) an error message
```

Mixin Expressions

MixinExpression:

```
mixin ( AssignExpression )
```

The *AssignExpression* must evaluate at compile time to a constant string. The text contents of the string must be compilable as a valid *Expression*, and is compiled as such.

```
int foo(int x)
{
    return mixin("x_+1") * 7;  // same as ((x + 1) * 7)
}
```

Import Expressions

```
ImportExpression:
    import ( AssignExpression )
```

The *AssignExpression* must evaluate at compile time to a constant string. The text contents of the string are interpreted as a file name. The file is read, and the exact contents of the file become a string literal.

Implementations may restrict the file name in order to avoid directory traversal security vulnerabilities. A possible restriction might be to disallow any path components in the file name.

Note that by default an import expression will not compile unless you pass one or more paths via the **-J** switch. This tells the compiler where it should look for the files to import. This is a security feature.

```
void foo()
{
    // Prints contents of file foo.txt
    writeln(import("foo.txt"));
}
```

Typeid Expressions

```
TypeidExpression:
    typeid ( Type )
    typeid ( Expression )
```

If *Type*, returns an instance of class **TypeInfo** corresponding to *Type*.

If *Expression*, returns an instance of class **TypeInfo** corresponding to the type of the *Expression*. If the type is a class, it returns the **TypeInfo** of the dynamic type (i.e. the most derived type). The *Expression* is always executed.

```

class A { }
class B : A { }

void main()
{
    writeln(typeid(int));           // int
    uint i;
    writeln(typeid(i++));           // uint
    writeln(i);                     // 1
    A a = new B();
    writeln(typeid(a));              // B
    writeln(typeid(typeof(a)));      // A
}

```

IsExpression

IsExpression:

```

is ( Type )
is ( Type : TypeSpecialization )
is ( Type == TypeSpecialization )
is ( Type : TypeSpecialization , TemplateParameterList )
is ( Type == TypeSpecialization , TemplateParameterList )
is ( Type Identifier )
is ( Type Identifier : TypeSpecialization )
is ( Type Identifier == TypeSpecialization )
is ( Type Identifier : TypeSpecialization , TemplateParameterList )
is ( Type Identifier == TypeSpecialization , TemplateParameterList )

```

TypeSpecialization:

```

Type
struct
union
class
interface
enum
function
delegate
super
const

```

```

immutable
inout
shared
return
__parameters

```

IsExpressions are evaluated at compile time and are used for checking for valid types, comparing types for equivalence, determining if one type can be implicitly converted to another, and deducing the subtypes of a type. The result of an *IsExpression* is an int of type 0 if the condition is not satisfied, 1 if it is.

Type is the type being tested. It must be syntactically correct, but it need not be semantically correct. If it is not semantically correct, the condition is not satisfied.

Identifier is declared to be an alias of the resulting type if the condition is satisfied. The *Identifier* forms can only be used if the *IsExpression* appears in a *StaticIfCondition*.

TypeSpecialization is the type that *Type* is being compared against.

The forms of the *IsExpression* are:

1. `is (Type)` The condition is satisfied if *Type* is semantically correct (it must be syntactically correct regardless).

```

alias int func(int);    // func is a alias to a function type
void foo()
{
    if (is(func[]))      // not satisfied because arrays of
                        // functions are not allowed
        writeln("satisfied");
    else
        writeln("not_satisfied");

    if (is([] []))       // error, [][] is not a syntactically valid type
        ...
}

```

2. `is (Type : TypeSpecialization)` The condition is satisfied if *Type* is semantically correct and it is the same as or can be implicitly converted to *TypeSpecialization*. *TypeSpecialization* is only allowed to be a *Type*.

```

alias bar = short;
void foo(bar x)
{
    if (is(bar : int))   // satisfied because short can be

```

```

        // implicitly converted to int
        writeln("satisfied");
    else
        writeln("not_satisfied");
}

```

3. `is (Type == TypeSpecialization)`

The condition is satisfied if *Type* is semantically correct and is the same type as *TypeSpecialization*.

If *TypeSpecialization* is one of

`struct union class` interface `enum` function delegate `const` immutable shared
then the condition is satisfied if *Type* is one of those.

```

alias bar = short;

void test(bar x)
{
    if (is(bar == int))    // not satisfied because short is not
                          // the same type as int
        writeln("satisfied");
    else
        writeln("not_satisfied");
}

```

4. `is (Type Identifier)` The condition is satisfied if *Type* is semantically correct. If so, *Identifier* is declared to be an alias of *Type*.

```

alias bar = short;
void foo(bar x)
{
    static if (is(bar T))
        alias S = T;
    else
        alias S = long;

    writeln(typeid(S)); // prints "short"
    if (is(bar T))      // error, Identifier T form can
                        // only be in StaticIfConditions
        ...
}

```

5. `is (Type Identifier : TypeSpecialization)`

The condition is satisfied if *Type* is the same as *TypeSpecialization*, or if *Type* is a class and *TypeSpecialization* is a base class or base interface of it. The *Identifier* is declared to be either an alias of the *TypeSpecialization* or, if *TypeSpecialization* is dependent on *Identifier*, the deduced type.

```
alias bar = int;
alias abc = long*;
void foo(bar x, abc a)
{
    static if (is(bar T : int))
        alias S = T;
    else
        alias S = long;

    writeln(typeid(S)); // prints "int"

    static if (is(abc U : U*))
    {
        U u;
        writeln(typeid(typeof(u))); // prints "long"
    }
}
```

The way the type of *Identifier* is determined is analogous to the way template parameter types are determined by *TemplateTypeParameterSpecialization*.

6. `is (Type Identifier == TypeSpecialization)`

The condition is satisfied if *Type* is semantically correct and is the same as *TypeSpecialization*. The *Identifier* is declared to be either an alias of the *TypeSpecialization* or, if *TypeSpecialization* is dependent on *Identifier*, the deduced type.

If *TypeSpecialization* is one of `struct union class interface enum function delegate` ↪ `const immutable shared`

then the condition is satisfied if *Type* is one of those. Furthermore, *Identifier* is set to be an alias of the type:

keyword	alias type for <i>Identifier</i>
<code>struct</code>	<i>Type</i>
<code>union</code>	<i>Type</i>
<code>class</code>	<i>Type</i>
<code>interface</code>	<i>Type</i>
<code>super</code>	<i>TypeTuple</i> of base classes and interfaces
<code>enum</code>	the base type of the enum
<code>function</code>	<i>TypeTuple</i> of the function parameter types. For C- and D-style variadic functions, only the non-variadic parameters are included. For typesafe variadic functions, the ... is ignored.
<code>delegate</code>	the function type of the delegate
<code>return</code>	the return type of the function, delegate, or function pointer
<code>__parameters</code>	the parameter tuple of a function, delegate, or function pointer. This includes the parameter types, names, and default values.
<code>const</code>	<i>Type</i>
<code>immutable</code>	<i>Type</i>
<code>shared</code>	<i>Type</i>

```

alias bar = short;
enum E : byte { Emember }
void foo(bar x)
{
    static if (is(bar T == int))    // not satisfied, short is not int
        alias S = T;
    alias U = T;                    // error, T is not defined

    static if (is(E V == enum))    // satisfied, E is an enum
        V v;                       // v is declared to be a byte
}

```

7. `is (Type : TypeSpecialization , TemplateParameterList) is (Type == TypeSpecialization , TemplateParameterList) is (Type Identifier : TypeSpecialization , TemplateParameterList) is (Type Identifier == TypeSpecialization , TemplateParameterList)`

More complex types can be pattern matched; the *TemplateParameterList* declares symbols based on the parts of the pattern that are matched, analogously to the way implied template parameters are matched.

```
import std.stdio, std.typecons;
```

```

void main()
{
    alias Tup = Tuple!(int, string);
    alias AA = long[char[]];

    static if (is(Tup : TX!TL, alias TX, TL...))
    {
        writeln(is(TX!(int, long) == Tuple!(int, long))); // true
        writeln(typeid(TL[0])); // int
        writeln(typeid(TL[1])); // immutable(char)[]
    }

    static if (is(AA T : T[U], U : const char[]))
    {
        writeln(typeid(T)); // long
        writeln(typeid(U)); // const char[]
    }

    static if (is(AA A : A[B], B : int))
    {
        assert(0); // should not match, as B is not an int
    }

    static if (is(int[10] W : W[V], int V))
    {
        writeln(typeid(W)); // int
        writeln(V); // 10
    }

    static if (is(int[10] X : X[Y], int Y : 5))
    {
        assert(0); // should not match, Y should be 10
    }
}

```

Associativity and Commutativity

An implementation may rearrange the evaluation of expressions according to arithmetic associativity and commutativity rules as long as, within that thread of execution, no observable difference is possible.

This rule precludes any associative or commutative reordering of floating point expressions.

Chapter 11

Statements

C and C++ programmers will find the D statements very familiar, with a few interesting additions.

Statement:

;
NonEmptyStatement
ScopeBlockStatement

NoScopeNonEmptyStatement:

NonEmptyStatement
BlockStatement

NoScopeStatement:

;
NonEmptyStatement
BlockStatement

NonEmptyOrScopeBlockStatement:

NonEmptyStatement
ScopeBlockStatement

NonEmptyStatement:

NonEmptyStatementNoCaseNoDefault
CaseStatement
CaseRangeStatement
DefaultStatement

NonEmptyStatementNoCaseNoDefault:

LabeledStatement
ExpressionStatement
DeclarationStatement
IfStatement
WhileStatement
DoStatement
ForStatement
ForeachStatement
SwitchStatement
FinalSwitchStatement
ContinueStatement
BreakStatement
ReturnStatement
GotoStatement
WithStatement
SynchronizedStatement
TryStatement
ScopeGuardStatement
ThrowStatement
AsmStatement
PragmaStatement
MixinStatement
ForeachRangeStatement
ConditionalStatement
StaticAssert
TemplateMixin
ImportDeclaration

Any ambiguities in the grammar between *Statements* and *Declarations* are resolved by the declarations taking precedence. If a *Statement* is desired instead, wrapping it in parentheses will disambiguate it in favor of being a *Statement*.

Scope Statements

ScopeStatement:
 NonEmptyStatement
 BlockStatement

A new scope for local symbols is introduced for the *NonEmptyStatement* or *BlockStatement*.

Even though a new scope is introduced, local symbol declarations cannot shadow (hide) other local symbol declarations in the same function.

```
void func1(int x)
{
    int x;    // illegal, x shadows parameter x

    int y;

    { int y; } // illegal, y shadows enclosing scope's y

    void delegate() dg;
    dg = { int y; }; // ok, this y is not in the same function

    struct S
    {
        int y;    // ok, this y is a member, not a local
    }

    { int z; }
    { int z; } // ok, this z is not shadowing the other z

    { int t; }
    { t++; } // illegal, t is undefined
}
```

The idea is to avoid bugs in complex functions caused by scoped declarations inadvertently hiding previous ones. Local names should all be unique within a function.

Scope Block Statements

ScopeBlockStatement:
BlockStatement

A scope block statement introduces a new scope for the *BlockStatement*.

Labeled Statements

Statements can be labeled. A label is an identifier that precedes a statement.

LabeledStatement:

```

Identifier :
Identifier : NoScopeStatement
Identifier : Statement

```

Any statement can be labeled, including empty statements, and so can serve as the target of a goto statement. Labeled statements can also serve as the target of a break or continue statement.

A label can appear without a following statement at the end of a block.

Labels are in a name space independent of declarations, variables, types, etc. Even so, labels cannot have the same name as local declarations. The label name space is the body of the function they appear in. Label name spaces do not nest, i.e. a label inside a block statement is accessible from outside that block.

Block Statement

```

BlockStatement:
    { }
    { StatementList }

StatementList:
    Statement
    Statement StatementList

```

A block statement is a sequence of statements enclosed by `{ }`. The statements are executed in lexical order.

Expression Statement

```

ExpressionStatement:
    Expression ;

```

The expression is evaluated.

Expressions that have no effect, like `(x + x)`, are illegal in expression statements. If such an expression is needed, casting it to void will make it legal.

```

int x;
x++;           // ok
x;             // illegal
1+1;          // illegal
cast(void)(x + x); // ok

```


Declaration Statement

Declaration statements declare variables and types.

```
DeclarationStatement:
    StorageClassesopt Declaration
```

Some declaration statements:

```
int a;           // declare a as type int and initialize it to 0
struct S { }    // declare struct s
alias myint = int;
```

If Statement

If statements provide simple conditional execution of statements.

```
IfStatement:
    if ( IfCondition ) ThenStatement
    if ( IfCondition ) ThenStatement else ElseStatement
```

```
IfCondition:
    Expression
    auto Identifier = Expression
    TypeCtors Identifier = Expression
    TypeCtorsopt BaseType Declarator = Expression
```

```
ThenStatement:
    ScopeStatement
```

```
ElseStatement:
    ScopeStatement
```

Expression is evaluated and must have a type that can be converted to a boolean. If it's true the *ThenStatement* is transferred to, else the *ElseStatement* is transferred to.

The 'dangling else' parsing problem is solved by associating the else with the nearest if statement.

If an `auto Identifier` is provided, it is declared and initialized to the value and type of the *Expression*. Its scope extends from when it is initialized to the end of the *ThenStatement*.

If a *Declarator* is provided, it is declared and initialized to the value of the *Expression*. Its scope extends from when it is initialized to the end of the *ThenStatement*.

```

import std.regex;
...
if (auto m = std.regex.matchFirst("abcdef", "b(c)d"))
{
    writefln("[%s]", m.pre);    // prints [a]
    writefln("[%s]", m.post);   // prints [ef]
    writefln("[%s]", m[0]);     // prints [bcd]
    writefln("[%s]", m[1]);     // prints [c]
}
else
{
    writeln(m.post); // Error: undefined identifier 'm'
}
writeln(m.pre);     // Error: undefined identifier 'm'

```

While Statement

WhileStatement:

```
while ( Expression ) ScopeStatement
```

While statements implement simple loops.

Expression is evaluated and must have a type that can be converted to a boolean. If it's true the *ScopeStatement* is executed. After the *ScopeStatement* is executed, the *Expression* is evaluated again, and if true the *ScopeStatement* is executed again. This continues until the *Expression* evaluates to false.

```

int i = 0;
while (i < 10)
{
    foo(i);
    i++;
}

```

A *BreakStatement* will exit the loop. A *ContinueStatement* will transfer directly to evaluating *Expression* again.

Do Statement

DoStatement:

```
do ScopeStatement while ( Expression );
```

Do while statements implement simple loops.

ScopeStatement is executed. Then *Expression* is evaluated and must have a type that can be converted to a boolean. If it's true the loop is iterated again. This continues until the *Expression* evaluates to false.

```
int i = 0;
do
{
    foo(i);
} while (++i < 10);
```

A *BreakStatement* will exit the loop. A *ContinueStatement* will transfer directly to evaluating *Expression* again.

For Statement

For statements implement loops with initialization, test, and increment clauses.

ForStatement:

```
for ( Initialize Testopt ; Incrementopt ) ScopeStatement
```

Initialize:

```
;
NoScopeNonEmptyStatement
```

Test:

```
Expression
```

Increment:

```
Expression
```

Initialize is executed. *Test* is evaluated and must have a type that can be converted to a boolean. If it's true the statement is executed. After the statement is executed, the *Increment* is executed. Then *Test* is evaluated again, and if true the statement is executed again. This continues until the *Test* evaluates to false.

A *BreakStatement* will exit the loop. A *ContinueStatement* will transfer directly to the *Increment*.

A *ForStatement* creates a new scope. If *Initialize* declares a variable, that variable's scope extends through the end of the for statement. For example:

```
for (int i = 0; i < 10; i++)
    foo(i);
```

is equivalent to:

```
{
    int i;
    for (i = 0; i < 10; i++)
        foo(i);
}
```

Function bodies cannot be empty:

```
for (int i = 0; i < 10; i++)
    ;           // illegal
```

Use instead:

```
for (int i = 0; i < 10; i++)
{
}
}
```

The *Initialize* may be omitted. *Test* may also be omitted, and if so, it is treated as if it evaluated to true.

Foreach Statement

A foreach statement loops over the contents of an aggregate.

ForeachStatement:

Foreach (ForeachTypeList ; ForeachAggregate) NoScopeNonEmptyStatement

Foreach:

foreach
foreach_reverse

ForeachTypeList:

ForeachType
ForeachType , ForeachTypeList

ForeachType:

*ForeachTypeAttributes*_{opt} *BasicType Declarator*
*ForeachTypeAttributes*_{opt} *Identifier*

ForeachTypeAttributes

ForeachTypeAttribute
*ForeachTypeAttribute ForeachTypeAttributes*_{opt}

ForeachTypeAttribute:

ref
TypeCtor

ForeachAggregate:

Expression

ForeachAggregate is evaluated. It must evaluate to an expression of type static array, dynamic array, associative array, struct, class, delegate, or tuple. The *NoScopeNonEmptyStatement* is executed, once for each element of the aggregate. At the start of each iteration, the variables declared by the *ForeachTypeList* are set to be a copy of the elements of the aggregate. If the variable is **ref**, it is a reference to the contents of that aggregate.

The aggregate must be loop invariant, meaning that elements to the aggregate cannot be added or removed from it in the *NoScopeNonEmptyStatement*.

Foreach over Arrays

If the aggregate is a static or dynamic array, there can be one or two variables declared. If one, then the variable is said to be the *value* set to the elements of the array, one by one. The type of the variable must match the type of the array contents, except for the special cases outlined below. If there are two variables declared, the first is said to be the *index* and the second is said to be the *value*. The *index* must be of **int**, **uint** or **size_t** type, it cannot be *ref*, and it is set to be the index of the array element.

```
char[] a;
...
foreach (int i, char c; a)
{
    writefln("a[%d]= '%c'", i, c);
}
```

For **foreach**, the elements for the array are iterated over starting at index 0 and continuing to the maximum of the array. For **foreach_reverse**, the array elements are visited in the reverse order.

Note: The *ForeachTypeAttribute* is implicit, and when a type is not specified, it is inferred. In that case, **auto** is implied, and it is not necessary (and actually forbidden) to use it.

```
int[] arr;
...
foreach (n; arr) // ok, n is an int
```

```

        writeln(n);

foreach (auto n; arr) // error, auto is redundant
    writeln(n);

```

Foreach over Arrays of Characters

If the aggregate expression is a static or dynamic array of `chars`, `wchars`, or `dchars`, then the *Type* of the *value* can be any of `char`, `wchar`, or `dchar`. In this manner any UTF array can be decoded into any UTF type:

```

char[] a = "\xE2\x89\xA0".dup; // \u2260 encoded as 3 UTF-8 bytes

foreach (dchar c; a)
{
    writefln("a[]_=%x", c); // prints 'a[] = 2260'
}

dchar[] b = "\u2260".dup;

foreach (char c; b)
{
    writef("%x,_", c); // prints 'e2, 89, a0, '
}

```

Aggregates can be string literals, which can be accessed as `char`, `wchar`, or `dchar` arrays:

```

void test()
{
    foreach (char c; "ab")
    {
        writefln("'%s'", c);
    }
    foreach (wchar w; "xy")
    {
        writefln("'%s'", w);
    }
}

```

which would print:

```
'a' 'b' 'x' 'y'
```

Foreach over Associative Arrays

If the aggregate expression is an associative array, there can be one or two variables declared. If one, then the variable is said to be the *value* set to the elements of the array, one by one. The type of the variable must match the type of the array contents. If there are two variables declared, the first is said to be the *index* and the second is said to be the *value*. The *index* must be of the same type as the indexing type of the associative array. It cannot be *ref*, and it is set to be the index of the array element. The order in which the elements of the array are iterated over is unspecified for `foreach`. `foreach_reverse` for associative arrays is illegal.

```
double[string] a; // index type is string, value type is double
...
foreach (string s, double d; a)
{
    writefln("a['%s']=%g", s, d);
}
```

Foreach over Structs and Classes with opApply

If the aggregate expression is a struct or class object, the `foreach` is defined by the special `opApply` member function, and the `foreach_reverse` behavior is defined by the special `opApplyReverse` member function. These functions have the type:

```
int opApply(scope int delegate(ref Type [, ...]) dg);

int opApplyReverse(scope int delegate(ref Type [, ...]) dg);
```

where *Type* matches the *Type* used in the *ForeachType* declaration of *Identifier*. Multiple *ForeachTypes* correspond with multiple *Type*'s in the delegate type passed to `opApply` or `opApplyReverse`. There can be multiple `opApply` and `opApplyReverse` functions, one is selected by matching the type of *dg* to the *ForeachTypes* of the *ForeachStatement*. The body of the apply function iterates over the elements it aggregates, passing them each to the *dg* function. If the *dg* returns 0, then apply goes on to the next element. If the *dg* returns a nonzero value, apply must cease iterating and return that value. Otherwise, after done iterating across all the elements, apply will return 0.

For example, consider a class that is a container for two elements:

```
class Foo
{
    uint[2] array;

    int opApply(scope int delegate(ref uint) dg)
```

```

{
    int result = 0;

    for (int i = 0; i < array.length; i++)
    {
        result = dg(array[i]);
        if (result)
            break;
    }
    return result;
}

```

An example using this might be:

```

void test()
{
    Foo a = new Foo();

    a.array[0] = 73;
    a.array[1] = 82;

    foreach (uint u; a)
    {
        writefln("%d", u);
    }
}

```

which would print:

73 82 The `scope` storage class on the `dg` parameter means that the parameter's value does not escape the scope of the `opApply` function (an example would be assigning `dg` to a global). If it cannot be statically guaranteed that `dg` does not escape, a closure may be allocated for it on the heap instead of the stack. Best practice is to annotate delegate parameters with `scope` when possible.

`opApply` can also be a templated function, which will infer the types of parameters based on the *ForeachStatement*.

For example:

```

struct S
{
    import std.traits : ParameterTypeTuple; // introspection template

```



```

int opApply(Dg)(scope Dg dg)
    if (ParameterTypeTuple!Dg.length == 2) // foreach function takes 2 ↗
        ↪ parameters
    {
        return 0;
    }

int opApply(Dg)(scope Dg dg)
    if (ParameterTypeTuple!Dg.length == 3) // foreach function takes 3 ↗
        ↪ parameters
    {
        return 0;
    }
}

void main()
{
    foreach (int a, int b; S()) { } // calls first opApply function
    foreach (int a, int b, float c; S()) { } // calls second opApply function
}

```

It is important to make sure that, if `opApply` catches any exceptions, that those exceptions did not originate from the delegate passed to `opApply`. The user would expect exceptions thrown from a `foreach` body to both terminate the loop, and propagate outside the `foreach` body.

Foreach over Structs and Classes with Ranges

If the aggregate expression is a struct or class object, but the `opApply` for `foreach`, or `opApplyReverse` `foreach_reverse` do not exist, then iteration over struct and class objects can be done with range primitives. For `foreach`, this means the following properties and methods must be defined:

Foreach Range Properties	
Property	Purpose
<code>.empty</code>	returns true if no more elements
<code>.front</code>	return the leftmost element of the range
Foreach Range Methods	
Method	Purpose
<code>.popFront()</code>	move the left edge of the range right by one

Meaning:

```
foreach (e; range) { ... }
```

translates to:

```
for (auto __r = range; !__r.empty; __r.popFront())
{
    auto e = __r.front;
    ...
}
```

Similarly, for `foreach_reverse`, the following properties and methods must be defined:

Foreach_reverse Range Properties

Property	Purpose
<code>.empty</code>	returns true if no more elements
<code>.back</code>	return the rightmost element of the range

Foreach_reverse Range Methods

Method	Purpose
<code>.popBack()</code>	move the right edge of the range left by one

Meaning:

```
foreach_reverse (e; range) { ... }
```

translates to:

```
for (auto __r = range; !__r.empty; __r.popBack())
{
    auto e = __r.back;
    ...
}
```

Foreach over Delegates

If *ForeachAggregate* is a delegate, the type signature of the delegate is of the same as for `opApply`. This enables many different named looping strategies to coexist in the same class or struct.

For example:

```

void main()
{
    // Custom loop implementation, that iterates over powers of 2 with
    // alternating sign. The loop body is passed in dg.
    int myLoop(int delegate(ref int) dg)
    {
        for (int z = 1; z < 128; z *= -2)
        {
            auto ret = dg(z);

            // If the loop body contains a break, ret will be non-zero.
            if (ret != 0)
                return ret;
        }
        return 0;
    }

    // This example loop simply collects the loop index values into an array.
    int[] result;
    foreach (ref x; &myLoop)
    {
        result ~= x;
    }
    assert(result == [1, -2, 4, -8, 16, -32, 64, -128]);
}

```

Note: When *ForeachAggregate* is a delegate, the compiler does not try to implement reverse traversal of the results returned by the delegate when `foreach_reverse` is used. This may result in code that is confusing to read. Therefore, using `foreach_reverse` with a delegate is now deprecated, and will be rejected in the future.

Foreach over Tuples

If the aggregate expression is a tuple, there can be one or two variables declared. If one, then the variable is said to be the *value* set to the elements of the tuple, one by one. If the type of the variable is given, it must match the type of the tuple contents. If it is not given, the type of the variable is set to the type of the tuple element, which may change from iteration to iteration. If there are two variables declared, the first is said to be the *index* and the second is said to be the *value*. The *index* must be of `int` or `uint` type, it cannot be *ref*, and it is set to be the index of the tuple element.

If the tuple is a list of types, then the `foreach` statement is executed once for each type, and the value is aliased to that type.

```
import std.stdio;
import std.tupletuple; // for TypeTuple

void main()
{
    alias TL = TypeTuple!(int, long, double);

    foreach (T; TL)
    {
        writeln(typeid(T));
    }
}
```

Prints:

```
int long double
```

Foreach Ref Parameters

`ref` can be used to update the original elements:

```
void test()
{
    static uint[2] a = [7, 8];

    foreach (ref uint u; a)
    {
        u++;
    }
    foreach (uint u; a)
    {
        writefln("%d", u);
    }
}
```

which would print:

```
8 9 ref can not be applied to the index values.
```

If not specified, the *Types* in the *ForeachType* can be inferred from the type of the *ForeachAggregate*.

Foreach Restrictions

The aggregate itself must not be resized, reallocated, free'd, reassigned or destructed while the foreach is iterating over the elements.

```
int[] a;
int[] b;
foreach (int i; a)
{
    a = null;           // error
    a.length += 10;    // error
    a = b;              // error
}
a = null;              // ok
```

Foreach Range Statement

A foreach range statement loops over the specified range.

ForeachRangeStatement:

Foreach (ForeachType ; LwrExpression .. UprExpression) ScopeStatement

LwrExpression:

Expression

UprExpression:

Expression

ForeachType declares a variable with either an explicit type, or a type inferred from *LwrExpression* and *UprExpression*. The *ScopeStatement* is then executed *n* times, where *n* is the result of *UprExpression* - *LwrExpression*. If *UprExpression* is less than or equal to *LwrExpression*, the *ScopeStatement* is executed zero times. If *Foreach* is **foreach**, then the variable is set to *LwrExpression*, then incremented at the end of each iteration. If *Foreach* is **foreach_reverse**, then the variable is set to *UprExpression*, then decremented before each iteration. *LwrExpression* and *UprExpression* are each evaluated exactly once, regardless of how many times the *ScopeStatement* is executed.

```
import std.stdio;
```

```
int foo()
{
```

```

    write("foo");
    return 10;
}

void main()
{
    foreach (i; 0 .. foo())
    {
        write(i);
    }
}

Prints:
foo0123456789

```

Break and Continue out of Foreach

A *BreakStatement* in the body of the foreach will exit the foreach, a *ContinueStatement* will immediately start the next iteration.

Switch Statement

A switch statement goes to one of a collection of case statements depending on the value of the switch expression.

SwitchStatement:

```
switch ( Expression ) ScopeStatement
```

CaseStatement:

```
case ArgumentList : ScopeStatementList
```

CaseRangeStatement:

```
case FirstExp : .. case LastExp : ScopeStatementList
```

FirstExp:

```
AssignExpression
```

LastExp:

```
AssignExpression
```

DefaultStatement:

```

    default : ScopeStatementList

ScopeStatementList:
    StatementListNoCaseNoDefault

StatementListNoCaseNoDefault:
    StatementNoCaseNoDefault
    StatementNoCaseNoDefault StatementListNoCaseNoDefault

StatementNoCaseNoDefault:
    ;
    NonEmptyStatementNoCaseNoDefault
    ScopeBlockStatement

```

Expression is evaluated. The result type *T* must be of integral type or `char[]`, `wchar[]` or `dchar[]`. The result is compared against each of the case expressions. If there is a match, the corresponding case statement is transferred to.

The case expressions, *ArgumentList*, are a comma separated list of expressions.

A *CaseRangeStatement* is a shorthand for listing a series of case statements from *FirstExp* to *LastExp*.

If none of the case expressions match, and there is a default statement, the default statement is transferred to.

A switch statement must have a default statement.

The case expressions must all evaluate to a constant value or array, or a runtime initialized const or immutable variable of integral type.

They must be implicitly convertible to the type of the switch *Expression*.

Case expressions must all evaluate to distinct values. Const or immutable variables must all have different names. If they share a value, the first case statement with that value gets control. There must be exactly one default statement.

The *ScopeStatementList* introduces a new scope.

Case statements and default statements associated with the switch can be nested within block statements; they do not have to be in the outermost block. For example, this is allowed:

```

switch (i)
{
    case 1:
    {
        case 2:
        }
}

```

```

    break;
}

```

A *ScopeStatementList* must either be empty, or be ended with a *ContinueStatement*, *BreakStatement*, *ReturnStatement*, *GotoStatement*, *ThrowStatement* or `assert(0)` expression unless this is the last case. This is to set apart with C's error-prone implicit fall-through behavior. `goto case;` could be used for explicit fall-through:

```

int number;
string message;
switch (number)
{
    default:    // valid: ends with 'throw'
                throw new Exception("unknown_␣number");

    case 3:     // valid: ends with 'break' (break out of the 'switch' only)
                message ~= "three_␣";
                break;

    case 4:     // valid: ends with 'continue' (continue the enclosing loop)
                message ~= "four_␣";
                continue;

    case 5:     // valid: ends with 'goto' (explicit fall-through to next case.)
                message ~= "five_␣";
                goto case;

    case 6:     // ERROR: implicit fall-through
                message ~= "six_␣";

    case 1:     // valid: the body is empty
    case 2:     // valid: this is the last case in the switch statement.
                message = "one_␣or_␣two";
}

```

A `break` statement will exit the switch *BlockStatement*.

Strings can be used in switch expressions. For example:

```

char[] name;
...
switch (name)
{

```



```

    case "fred":
    case "sally":
        ...
}

```

For applications like command line switch processing, this can lead to much more straightforward code, being clearer and less error prone. `char`, `wchar` and `dchar` strings are allowed.

Implementation Note: The compiler's code generator may assume that the case statements are sorted by frequency of use, with the most frequent appearing first and the least frequent last. Although this is irrelevant as far as program correctness is concerned, it is of performance interest.

Final Switch Statement

```

FinalSwitchStatement:
    final switch ( Expression ) ScopeStatement

```

A final switch statement is just like a switch statement, except that:

- No *DefaultStatement* is allowed.
- No *CaseRangeStatements* are allowed.
- If the switch *Expression* is of enum type, all the enum members must appear in the *CaseStatements*.
- The case expressions cannot evaluate to a run time initialized value.

Continue Statement

```

ContinueStatement:
    continue Identifieropt ;

```

A continue aborts the current iteration of its enclosing loop statement, and starts the next iteration.

continue executes the next iteration of its innermost enclosing while, for, foreach, or do loop. The increment clause is executed.

If continue is followed by *Identifier*, the *Identifier* must be the label of an enclosing while, for, or do loop, and the next iteration of that loop is executed. It is an error if there is no such statement.

Any intervening finally clauses are executed, and any intervening synchronization objects are released.

Note: If a finally clause executes a throw out of the finally clause, the continue target is never reached.

```
for (i = 0; i < 10; i++)
{
    if (foo(i))
        continue;
    bar();
}
```

Break Statement

BreakStatement:
`break Identifieropt ;`

A break exits the enclosing statement.

break exits the innermost enclosing while, for, foreach, do, or switch statement, resuming execution at the statement following it.

If break is followed by *Identifier*, the *Identifier* must be the label of an enclosing while, for, do or switch statement, and that statement is exited. It is an error if there is no such statement.

Any intervening finally clauses are executed, and any intervening synchronization objects are released.

Note: If a finally clause executes a throw out of the finally clause, the break target is never reached.

```
for (i = 0; i < 10; i++)
{
    if (foo(i))
        break;
}
```

Return Statement

ReturnStatement:
`return Expressionopt ;`

A return exits the current function and supplies its return value.

Expression is required if the function specifies a return type that is not void. The *Expression* is implicitly converted to the function return type.

At least one return statement, throw statement, or assert(0) expression is required if the function specifies a return type that is not void, unless the function contains inline assembler code.

Before the function actually returns, any objects with scope storage duration are destroyed, any enclosing finally clauses are executed, any scope(exit) statements are executed, any scope(success) statements are executed, and any enclosing synchronization objects are released.

The function will not return if any enclosing finally clause does a return, goto or throw that exits the finally clause.

If there is an out postcondition (see Contract Programming), that postcondition is executed after the *Expression* is evaluated and before the function actually returns.

```
int foo(int x)
{
    return x + 3;
}
```

Goto Statement

```
GotoStatement:
    goto Identifier ;
    goto default ;
    goto case ;
    goto case Expression ;
```

A goto transfers to the statement labeled with *Identifier*.

```
if (foo)
    goto L1;
x = 3;
L1:
    x++;
```

The second form, `goto default;`, transfers to the innermost *DefaultStatement* of an enclosing *SwitchStatement*.

The third form, `goto case;`, transfers to the next *CaseStatement* of the innermost enclosing *SwitchStatement*.

The fourth form, `goto case Expression;`, transfers to the *CaseStatement* of the innermost enclosing *SwitchStatement* with a matching *Expression*.

```
switch (x)
{
    case 3:
        goto case;
    case 4:
```

```

        goto default;
    case 5:
        goto case 4;
    default:
        x = 4;
        break;
}

```

Any intervening finally clauses are executed, along with releasing any intervening synchronization mutexes.

It is illegal for a *GotoStatement* to be used to skip initializations.

With Statement

The with statement is a way to simplify repeated references to the same object.

```

WithStatement:
    with ( Expression ) ScopeStatement
    with ( Symbol ) ScopeStatement
    with ( TemplateInstance ) ScopeStatement

```

where *Expression* evaluates to a class reference or struct instance. Within the with body the referenced object is searched first for identifier symbols. The *WithStatement*

```

with (expression)
{
    ...
    ident;
}

```

is semantically equivalent to:

```

{
    Object tmp;
    tmp = expression;
    ...
    tmp.ident;
}

```

Note that *Expression* only gets evaluated once. The with statement does not change what **this** or **super** refer to.

For *Symbol* which is a scope or *TemplateInstance*, the corresponding scope is searched when looking up symbols. For example:

```

struct Foo
{
    alias Y = int;
}
...
Y y;           // error, Y undefined
with (Foo)
{
    Y y;       // same as Foo.Y y;
}

```

Use of with object symbols that shadow local symbols with the same identifier are not allowed. This is to reduce the risk of inadvertant breakage of with statements when new members are added to the object declaration.

```

struct S
{
    float x;
}

void main()
{
    int x;
    S s;
    with (s)
    {
        x++; // error, shadows the int x declaration
    }
}

```

Synchronized Statement

The synchronized statement wraps a statement with a mutex to synchronize access among multiple threads.

```

SynchronizedStatement:
    synchronized ScopeStatement
    synchronized ( Expression ) ScopeStatement

```

Synchronized allows only one thread at a time to execute *ScopeStatement* by using a mutex.

What mutex is used is determined by the *Expression*. If there is no *Expression*, then a global mutex is created, one per such synchronized statement. Different synchronized statements will have different global mutexes.

If there is an *Expression*, it must evaluate to either an Object or an instance of an *Interface*, in which case it is cast to the Object instance that implemented that *Interface*. The mutex used is specific to that Object instance, and is shared by all synchronized statements referring to that instance.

The synchronization gets released even if *ScopeStatement* terminates with an exception, goto, or return.

Example:

```
synchronized { ... }
```

This implements a standard critical section.

Synchronized statements support recursive locking; that is, a function wrapped in synchronized is allowed to recursively call itself and the behavior will be as expected: The mutex will be locked and unlocked as many times as there is recursion.

Try Statement

Exception handling is done with the try-catch-finally statement.

TryStatement:

```
try ScopeStatement Catches
try ScopeStatement Catches FinallyStatement
try ScopeStatement FinallyStatement
```

Catches:

```
LastCatch
Catch
Catch Catches
```

LastCatch:

```
catch NoScopeNonEmptyStatement
```

Catch:

```
catch ( CatchParameter ) NoScopeNonEmptyStatement
```

CatchParameter:

```
BasicType Identifier
```

FinallyStatement:
 finally NoScopeNonEmptyStatement

CatchParameter declares a variable *v* of type *T*, where *T* is

Throwable or derived from Throwable. *v* is initialized by the throw expression if *T* is of the same type or a base class of the throw expression. The catch clause will be executed if the exception object is of type *T* or derived from *T*.

If just type *T* is given and no variable *v*, then the catch clause is still executed.

It is an error if any *CatchParameter* type *T1* hides a subsequent *Catch* with type *T2*, i.e. it is an error if *T1* is the same type as or a base class of *T2*.

LastCatch catches all exceptions.

The *FinallyStatement* is always executed, whether the *try ScopeStatement* exits with a goto, break, continue, return, exception, or fall-through.

If an exception is raised in the *FinallyStatement* and is not caught before the original exception is caught, it is chained to the previous exception via the *next* member of *Throwable*. Note that, in contrast to most other programming languages, the new exception does not replace the original exception. Instead, later exceptions are regarded as 'collateral damage' caused by the first exception. The original exception must be caught, and this results in the capture of the entire chain.

Thrown objects derived from *Error* are treated differently. They bypass the normal chaining mechanism, such that the chain can only be caught by catching the first *Error*. In addition to the list of subsequent exceptions, *Error* also contains a pointer that points to the original exception (the head of the chain) if a bypass occurred, so that the entire exception history is retained.

```
import std.stdio;

int main()
{
    try
    {
        try
        {
            throw new Exception("first");
        }
        finally
        {
            writeln("finally");
            throw new Exception("second");
        }
    }
}
```

```

    catch (Exception e)
    {
        writeln("catch_␣%", e.msg);
    }
    writeln("done");
    return 0;
}

```

prints:

finally catch first done

A *FinallyStatement* may not exit with a goto, break, continue, or return; nor may it be entered with a goto.

A *FinallyStatement* may not contain any *Catches*. This restriction may be relaxed in future versions.

Throw Statement

Throw an exception.

ThrowStatement:

```
throw Expression ;
```

Expression is evaluated and must be a Throwable reference. The Throwable reference is thrown as an exception.

```
throw new Exception("message");
```

Scope Guard Statement

ScopeGuardStatement:

```

scope(exit) NonEmptyOrScopeBlockStatement
scope(success) NonEmptyOrScopeBlockStatement
scope(failure) NonEmptyOrScopeBlockStatement

```

The *ScopeGuardStatement* executes *NonEmptyOrScopeBlockStatement* at the close of the current scope, rather than at the point where the *ScopeGuardStatement* appears. **scope(exit)** executes *NonEmptyOrScopeBlockStatement* when the scope exits normally or when it exits due to exception unwinding. **scope(failure)** executes *NonEmptyOrScopeBlockStatement* when the scope exits due to exception unwinding. **scope(success)** executes *NonEmptyOrScopeBlockStatement* when the scope exits normally.

If there are multiple *ScopeGuardStatements* in a scope, they will be executed in the reverse lexical order in which they appear. If any scope instances are to be destroyed upon the close of the scope, their destructions will be interleaved with the *ScopeGuardStatements* in the reverse lexical order in which they appear.

```

write("1");
{
    write("2");
    scope(exit) write("3");
    scope(exit) write("4");
    write("5");
}
writeln();

writes:
12543

{
    scope(exit) write("1");
    scope(success) write("2");
    scope(exit) write("3");
    scope(success) write("4");
}
writeln();

writes:
4321

struct Foo
{
    this(string s) { write(s); }
    ~this() { write("1"); }
}

try
{
    scope(exit) write("2");
    scope(success) write("3");
    Foo f = Foo("0");
    scope(failure) write("4");
    throw new Exception("msg");
    scope(exit) write("5");
    scope(success) write("6");
}

```

```

        scope(failure) write("7");
    }
    catch (Exception e)
    {
    }
    writeln();

```

writes:

0412

A `scope(exit)` or `scope(success)` statement may not exit with a `throw`, `goto`, `break`, `continue`, or `return`; nor may it be entered with a `goto`.

Catching C++ Class Objects

On many platforms, catching C++ class objects is supported. Catching C++ objects and D objects cannot both be done in the same *TryStatement*. Upon exit from the *Catch*, any destructors for the C++ object will be run and the storage used for it reclaimed. C++ objects cannot be caught in `@safe` code.

Asm Statement

Inline assembler is supported with the `asm` statement:

AsmStatement:

```
asm FunctionAttributesopt { AsmInstructionListopt }
```

AsmInstructionList:

```
AsmInstruction ;
```

```
AsmInstruction ; AsmInstructionList
```

An `asm` statement enables the direct use of assembly language instructions. This makes it easy to obtain direct access to special CPU features without resorting to an external assembler. The D compiler will take care of the function calling conventions, stack setup, etc.

The format of the instructions is, of course, highly dependent on the native instruction set of the target CPU, and so is implementation defined. But, the format will follow the following conventions:

- It must use the same tokens as the D language uses.
- The comment form must match the D language comments.
- Asm instructions are terminated by a `;`, not by an end of line.

These rules exist to ensure that D source code can be tokenized independently of syntactic or semantic analysis.

For example, for the Intel Pentium:

```
int x = 3;
asm
{
    mov EAX,x; // load x and put it in register EAX
}
```

Inline assembler can be used to access hardware directly:

```
int gethardware()
{
    asm
    {
        mov EAX, dword ptr 0x1234;
    }
}
```

For some D implementations, such as a translator from D to C, an inline assembler makes no sense, and need not be implemented. The version statement can be used to account for this:

```
version (D_InlineAsm_X86)
{
    asm
    {
        ...
    }
}
else
{
    /* ... some workaround ... */
}
```

Semantically consecutive *AsmStatements* shall not have any other instructions (such as register save or restores) inserted between them by the compiler.

Pragma Statement

PragmaStatement:
Pragma NoScopeStatement

Mixin Statement

```
MixinStatement:
    mixin ( AssignExpression ) ;
```

The *AssignExpression* must evaluate at compile time to a constant string. The text contents of the string must be compilable as a valid *StatementList*, and is compiled as such.

```
import std.stdio;

void main()
{
    int j;
    mixin("
        int x=3;
        for(int i=0;i<3;i++)
            writeln(x+i,++j);
        ");    // ok

    const char[] s = "int y;";
    mixin(s);  // ok
    y = 4;     // ok, mixin declared y

    char[] t = "y=3;";
    mixin(t);  // error, t is not evaluatable at compile time

    mixin("y=") 4; // error, string must be complete statement

    mixin("y=" ~ "4;"); // ok
}
```

Chapter 12

Arrays

There are four kinds of arrays:

Kinds of Arrays	
Syntax	Description
<i>type</i> *	Pointers to data
<i>type</i> [<i>integer</i>]	Static arrays
<i>type</i> []	Dynamic arrays
<i>type</i> [<i>type</i>]	Associative arrays

Pointers

```
int* p;
```

These are simple pointers to data, analogous to C pointers. Pointers are provided for interfacing with C and for specialized systems work. There is no length associated with it, and so there is no way for the compiler or runtime to do bounds checking, etc., on it. Most conventional uses for pointers can be replaced with dynamic arrays, **out** and **ref** parameters, and reference types.

Static Arrays

```
int[3] s;
```

These are analogous to C arrays. Static arrays are distinguished by having a length fixed at compile time.

The total size of a static array cannot exceed 16Mb. A dynamic array should be used instead for such large arrays.

A static array with a dimension of 0 is allowed, but no space is allocated for it. It's useful as the last member of a variable length struct, or as the degenerate case of a template expansion.

Static arrays are value types. Unlike in C and D version 1, static arrays are passed to functions by value. Static arrays can also be returned by functions.

Dynamic Arrays

```
int[] a;
```

Dynamic arrays consist of a length and a pointer to the array data. Multiple dynamic arrays can share all or parts of the array data.

12.1 Array Declarations

Declarations appear before the identifier being declared and read right to left, so:

```
int[] a;           // dynamic array of ints
int[4][3] b;       // array of 3 arrays of 4 ints each
int[][5] c;        // array of 5 dynamic arrays of ints.
int*[][3] d;       // array of 3 pointers to dynamic arrays of pointers to ints
int[]* e;          // pointer to dynamic array of ints
```

12.2 Array Usage

There are two broad kinds of operations to do on an array - affecting the handle to the array, and affecting the contents of the array. C only has operators to affect the handle. In D, both are accessible.

The handle to an array is specified by naming the array, as in p, s or a:

```
int* p;
int[3] s;
int[] a;

int[] b;

p = s.ptr; // p points to the first element of the array s.
p = a.ptr; // p points to the first element of the array a.

a = p;     // error, since the length of the array pointed
           // to by p is unknown
a = s;     // a is initialized to point to the s array
a = b;     // a points to the same array as b does
```

12.3 Slicing

Slicing an array means to specify a subarray of it. An array slice does not copy the data, it is only another reference to it. For example:

```
int[10] a;    // declare array of 10 ints
int[] b;

b = a[1..3]; // a[1..3] is a 2 element array consisting of
              // a[1] and a[2]
foo(b[1]);   // equivalent to foo(0)
a[2] = 3;
foo(b[1]);   // equivalent to foo(3)
```

The `[]` is shorthand for a slice of the entire array. For example, the assignments to `b`:

```
int[10] a;
int[] b;

b = a;
b = a[];
b = a[0 .. a.length];
```

are all semantically equivalent.

Slicing is not only handy for referring to parts of other arrays, but for converting pointers into bounds-checked arrays:

```
int* p;
int[] b = p[0..8];
```

12.4 Array Copying

When the slice operator appears as the left-hand side of an assignment expression, it means that the contents of the array are the target of the assignment rather than a reference to the array. Array copying happens when the left-hand side is a slice, and the right-hand side is an array of or pointer to the same type.

```
int[3] s;
int[3] t;

s = t;           // the 3 elements of t are copied into s
s[] = t;         // the 3 elements of t are copied into s
s[] = t[];       // the 3 elements of t are copied into s
```

```

s[1..2] = t[0..1]; // same as s[1] = t[0]
s[0..2] = t[1..3]; // same as s[0] = t[1], s[1] = t[2]
s[0..4] = t[0..4]; // error, only 3 elements in s
s[0..2] = t;       // error, operands have different lengths

```

Overlapping Copying

Overlapping copies are an error:

```

s[0..2] = s[1..3]; // error, overlapping copy
s[1..3] = s[0..2]; // error, overlapping copy

```

Disallowing overlapping makes it possible for more aggressive parallel code optimizations than possible with the serial semantics of C.

If overlapping is required, use `std.algorithm.mutation.copy`:

```

import std.algorithm;
int[] s = [1, 2, 3, 4];

copy(s[1..3], s[0..2]);
assert(s == [2, 3, 3, 4]);

```

12.5 Array Setting

If a slice operator appears as the left-hand side of an assignment expression, and the type of the right-hand side is the same as the element type of the left-hand side, then the array contents of the left-hand side are set to the right-hand side.

```

int[3] s;
int* p;

s[] = 3; // same as s[0] = 3, s[1] = 3, s[2] = 3
p[0..2] = 3; // same as p[0] = 3, p[1] = 3

```

12.6 Array Concatenation

The binary operator `~` is the *cat* operator. It is used to concatenate arrays:

```

int[] a;
int[] b;
int[] c;

```



```
a = b ~ c; // Create an array from the concatenation
           // of the b and c arrays
```

Many languages overload the `+` operator to mean concatenation. This confusingly leads to, does:

```
"10" + 3 + 4
```

produce the number 17, the string "1034" or the string "107" as the result? It isn't obvious, and the language designers wind up carefully writing rules to disambiguate it - rules that get incorrectly implemented, overlooked, forgotten, and ignored. It's much better to have `+` mean addition, and a separate operator to be array concatenation.

Similarly, the `=` operator means append, as in:

```
a ~= b; // a becomes the concatenation of a and b
```

Concatenation always creates a copy of its operands, even if one of the operands is a 0 length array, so:

```
a = b;           // a refers to b
a = b ~ c[0..0]; // a refers to a copy of b
```

Appending does not always create a copy, see setting dynamic array length for details.

12.7 Array Operations

Many array operations, also known as vector operations, can be expressed at a high level rather than as a loop. For example, the loop:

```
T[] a, b;
...
for (size_t i = 0; i < a.length; i++)
    a[i] = b[i] + 4;
```

assigns to the elements of `a` the elements of `b` with 4 added to each. This can also be expressed in vector notation as:

```
T[] a, b;
...
a[] = b[] + 4;
```

A vector operation is indicated by the slice operator appearing as the left-hand side of an `=`, `+=`, `-=`, `*=`, `/=`, operator. The right-hand side can be an expression consisting either of an array slice of the same length and type as the left-hand side or an expression of the element type of the

left-hand side, in any combination. The operators supported for vector operations are the binary operators $+$, $-$, $*$, $/$, $-$ and $.$

The slice on the left and any slices on the right must not overlap. The vector assignment operators are evaluated right to left, and the other binary operators are evaluated left to right. All operands are evaluated exactly once, even if the array slice has zero elements in it.

The order in which the array elements are computed is implementation defined, and may even occur in parallel. An application must not depend on this order.

Implementation note: many of the more common vector operations are expected to take advantage of any vector math instructions available on the target computer.

12.8 Pointer Arithmetic

```
int[3] abc;           // static array of 3 ints
int[] def = [ 1, 2, 3 ]; // dynamic array of 3 ints

void dibb(int* array)
{
    array[2];          // means same thing as *(array + 2)
    *(array + 2);      // get 3rd element
}

void diss(int[] array)
{
    array[2];          // ok
    *(array + 2);      // error, array is not a pointer
}

void ditt(int[3] array)
{
    array[2];          // ok
    *(array + 2);      // error, array is not a pointer
}
```

12.9 Rectangular Arrays

Experienced FORTRAN numerics programmers know that multidimensional "rectangular" arrays for things like matrix operations are much faster than trying to access them via pointers to pointers resulting from "array of pointers to array" semantics. For example, the D syntax:

```
double[] [] matrix;
```

declares `matrix` as an array of pointers to arrays. (Dynamic arrays are implemented as pointers to the array data.) Since the arrays can have varying sizes (being dynamically sized), this is sometimes called "jagged" arrays. Even worse for optimizing the code, the array rows can sometimes point to each other! Fortunately, D static arrays, while using the same syntax, are implemented as a fixed rectangular layout:

```
double[3][3] matrix;
```

declares a rectangular matrix with 3 rows and 3 columns, all contiguously in memory. In other languages, this would be called a multidimensional array and be declared as:

```
double matrix[3,3];
```

12.10 Array Length

Within the `[]` of a static or a dynamic array, the symbol `$` represents the length of the array.

```
int[4] foo;
int[] bar = foo;
int* p = &foo[0];

// These expressions are equivalent:
bar[]
bar[0 .. 4]
bar[0 .. $]
bar[0 .. bar.length]

p[0 .. $]      // '$' is not defined, since p is not an array
bar[0]+$      // '$' is not defined, out of scope of [ ]

bar[$-1] // retrieves last element of the array
```

12.11 Array Properties

Static array properties are:

Static Array Properties

Property	Description
<code>.init</code>	Returns an array literal with each element of the literal being the <code>.init</code> property of the array element type.
<code>.sizeof</code>	Returns the array length multiplied by the number of bytes per array element.
<code>.length</code>	Returns the number of elements in the array. This is a fixed quantity for static arrays. It is of type <code>size_t</code> .
<code>.ptr</code>	Returns a pointer to the first element of the array.
<code>.dup</code>	Create a dynamic array of the same size and copy the contents of the array into it. The copy will have any immutability or <code>const</code> stripped. If this conversion is invalid the call will not compile.
<code>.idup</code>	Create a dynamic array of the same size and copy the contents of the array into it. The copy is typed as being immutable. If this conversion is invalid the call will not compile.
<code>.reverse</code>	Reverses in place the order of the elements in the array. Returns the array.
<code>.sort</code>	Sorts in place the order of the elements in the array. Returns the array.

Dynamic array properties are:

Dynamic Array Properties

Property	Description
<code>.init</code>	Returns <code>null</code> .
<code>.sizeof</code>	Returns the size of the dynamic array reference, which is 8 in 32-bit builds and 16 on 64-bit builds.
<code>.length</code>	Get/set number of elements in the array. It is of type <code>size_t</code> .
<code>.ptr</code>	Returns a pointer to the first element of the array.
<code>.dup</code>	Create a dynamic array of the same size and copy the contents of the array into it. The copy will have any immutability or <code>const</code> stripped. If this conversion is invalid the call will not compile.
<code>.idup</code>	Create a dynamic array of the same size and copy the contents of the array into it. The copy is typed as being immutable. If this conversion is invalid the call will not compile.
<code>.reverse</code>	Reverses in place the order of the elements in the array. Returns the array.
<code>.sort</code>	Sorts in place the order of the elements in the array. Returns the array.

For the `.sort` property to work on arrays of class objects, the class definition must define the function: `int opCmp(Object)`. This is used to determine the ordering of the class objects. Note that the parameter is of type `Object`, not the type of the class.

For the `.sort` property to work on arrays of structs or unions, the struct or union definition must define the function:

```
int opCmp(ref const S) const.
```

The type `S` is the type of the struct or union. This function will determine the sort ordering.

Examples:

```
int* p;
int[3] s;
int[] a;

p.length; // error, length not known for pointer
s.length; // compile time constant 3
a.length; // runtime value

p.dup;    // error, length not known
s.dup;    // creates an array of 3 elements, copies
           // elements s into it
a.dup;    // creates an array of a.length elements, copies
           // elements of a into it
```

Setting Dynamic Array Length

The `.length` property of a dynamic array can be set as the left-hand side of an `=` operator:

```
array.length = 7;
```

This causes the array to be reallocated in place, and the existing contents copied over to the new array. If the new array length is shorter, the array is not reallocated, and no data is copied. It is equivalent to slicing the array:

```
array = array[0..7];
```

If the new array length is longer, the remainder is filled out with the default initializer.

To maximize efficiency, the runtime always tries to resize the array in place to avoid extra copying.

It will always do a copy if the new size is larger and the array was not allocated via the new operator or resizing in place would overwrite valid data in the array.

For example:

```
char[] a = new char[20];
char[] b = a[0..10];
char[] c = a[10..20];
char[] d = a;

b.length = 15; // always reallocates because extending in place would
```

```

        // overwrite other data in a.
b[11] = 'x';    // a[11] and c[1] are not affected

d.length = 1;
d.length = 20; // also reallocates, because doing this will overwrite a and
               // c

c.length = 12; // may reallocate in place if space allows, because nothing
               // was allocated after c.
c[5] = 'y';    // may affect contents of a, but not b or d because those
               // were reallocated.

a.length = 25; // This always reallocates because if c extended in place,
               // then extending a would overwrite c. If c didn't
               // reallocate in place, it means there was not enough space,
               // which will still be true for a.
a[15] = 'z';    // does not affect c, because either a or c has reallocated.

```

To guarantee copying behavior, use the `.dup` property to ensure a unique array that can be resized. Also, one may use the `.capacity` property to determine how many elements can be appended to the array without reallocating.

These issues also apply to appending arrays with the `=` operator. Concatenation using the `+` operator is not affected since it always reallocates.

Resizing a dynamic array is a relatively expensive operation. So, while the following method of filling an array:

```

int[] array;
while (1)
{
    c = getinput();
    if (!c)
        break;
    ++array.length;
    array[array.length - 1] = c;
}

```

will work, it will be inefficient. A more practical approach would be to minimize the number of resizes:

```

int[] array;
array.length = 100;           // guess
for (i = 0; ; i++)

```

```

{
    c = getinput();
    if (!c)
        break;
    if (i == array.length)
        array.length *= 2;
    array[i] = c;
}
array.length = i;

```

Picking a good initial guess is an art, but you usually can pick a value covering 99% of the data. For example, when gathering user input from the console - it's unlikely to be longer than 80.

Also, you may wish to utilize the `reserve` function to pre-allocate array data to use with the append operator.

```

int[] array;
size_t cap = array.reserve(10); // request
array ~= [1, 2, 3, 4, 5];
assert(cap >= 10); // allocated may be more than request
assert(cap == array.capacity);

```

Functions as Array Properties

If the first parameter to a function is an array, the function can be called as if it were a property of the array:

```

int[] array;
void foo(int[] a, int x);

foo(array, 3);
array.foo(3); // means the same thing

```

12.12 Array Bounds Checking

It is an error to index an array with an index that is less than 0 or greater than or equal to the array length. If an index is out of bounds, a `RangeError` exception is raised if detected at runtime, and an error if detected at compile time. A program may not rely on array bounds checking happening, for example, the following program is incorrect:

```

try
{

```

```
    for (i = 0; ; i++)
    {
        array[i] = 5;
    }
}
catch (RangeError)
{
    // terminate loop
}
```

The loop is correctly written:

```
for (i = 0; i < array.length; i++)
{
    array[i] = 5;
}
```

Implementation Note: Compilers should attempt to detect array bounds errors at compile time, for example:

```
int[3] foo;
int x = foo[3]; // error, out of bounds
```

Insertion of array bounds checking code at runtime should be turned on and off with a compile time switch.

12.13 Array Initialization

Default Initialization

- Pointers are initialized to `null`.
- Static array contents are initialized to the default initializer for the array element type.
- Dynamic arrays are initialized to having 0 elements.
- Associative arrays are initialized to having 0 elements.

Void Initialization

Void initialization happens when the *Initializer* for an array is `void`. What it means is that no initialization is done, i.e. the contents of the array will be undefined. This is most useful as an efficiency optimization. Void initializations are an advanced technique and should only be used when profiling indicates that it matters.

Static Initialization of Statically Allocated Arrays

Static initializations are supplied by a list of array element values enclosed in `[]`. The values can be optionally preceded by an index and a `:`. If an index is not supplied, it is set to the previous index plus 1, or 0 if it is the first value.

```
int[3] a = [ 1:2, 3 ]; // a[0] = 0, a[1] = 2, a[2] = 3
```

This is most handy when the array indices are given by enums:

```
enum Color { red, blue, green };

int value[Color.max + 1] =
    [ Color.blue :6,
      Color.green:2,
      Color.red   :5 ];
```

These arrays are statically allocated when they appear in global scope. Otherwise, they need to be marked with `const` or `static` storage classes to make them statically allocated arrays.

12.14 Special Array Types

Strings

A string is an array of characters. String literals are just an easy way to write character arrays. String literals are immutable (read only).

```
char[] str1 = "abc";           // error, "abc" is not mutable
char[] str2 = "abc".dup;       // ok, make mutable copy
immutable(char)[] str3 = "abc"; // ok
immutable(char)[] str4 = str1;  // error, str4 is not mutable
immutable(char)[] str5 = str1.dup; // ok, make immutable copy
```

The name `string` is aliased to `immutable(char)[]`, so the above declarations could be equivalently written as:

```
char[] str1 = "abc";           // error, "abc" is not mutable
char[] str2 = "abc".dup;       // ok, make mutable copy
string str3 = "abc";           // ok
string str4 = str1;            // error, str4 is not mutable
string str5 = str1.dup;        // ok, make immutable copy
```

`char[]` strings are in UTF-8 format. `wchar[]` strings are in UTF-16 format. `dchar[]` strings are in UTF-32 format.

Strings can be copied, compared, concatenated, and appended:

```

str1 = str2;
if (str1 < str3) { ... }
func(str3 ~ str4);
str4 ~= str1;

```

with the obvious semantics. Any generated temporaries get cleaned up by the garbage collector (or by using `alloca()`). Not only that, this works with any array not just a special `String` array.

A pointer to a char can be generated:

```

char* p = &str[3]; // pointer to 4th element
char* p = str;     // pointer to 1st element

```

Since strings, however, are not 0 terminated in D, when transferring a pointer to a string to C, add a terminating 0:

```

str ~= "\0";

```

or use the function `std.string.toStringz`.

The type of a string is determined by the semantic phase of compilation. The type is one of: `char[]`, `wchar[]`, `dchar[]`, and is determined by implicit conversion rules. If there are two equally applicable implicit conversions, the result is an error. To disambiguate these cases, a cast or a postfix of `c`, `w` or `d` can be used:

```

cast(immutable(wchar) [])"abc" // this is an array of wchar characters
"abc"w                          // so is this

```

String literals that do not have a postfix character and that have not been cast can be implicitly converted between `string`, `wstring`, and `dstring` as necessary.

```

char c;
wchar w;
dchar d;

c = 'b'; // c is assigned the character 'b'
w = 'b'; // w is assigned the wchar character 'b'
//w = 'bc'; // error - only one wchar character at a time
w = "b"[0]; // w is assigned the wchar character 'b'
w = "\r"[0]; // w is assigned the carriage return wchar character
d = 'd'; // d is assigned the character 'd'

```

Strings and Unicode

Note that built-in comparison operators operate on a code unit basis. The end result for valid strings is the same as that of code point for code point comparison as long as both strings are in

the same normalization form. Since normalization is a costly operation not suitable for language primitives it's assumed to be enforced by the user.

The standard library lends a hand for comparing strings with mixed encodings (by transparently decoding, see `std.algorithm.cmp`), case-insensitive comparison and normalization.

Last but not least, a desired string sorting order differs by culture and language and is usually nothing like code point for code point comparison. The natural order of strings is obtained by applying the Unicode collation algorithm that should be implemented in the standard library.

C's `printf()` and Strings

`printf()` is a C function and is not part of D. `printf()` will print C strings, which are 0 terminated. There are two ways to use `printf()` with D strings. The first is to add a terminating 0, and cast the result to a `char*`:

```
str ~= "\0";
printf("the_string_is_%s\n", cast(char*)str);
```

or:

```
import std.string;
printf("the_string_is_%s\n", std.string.toStringz(str));
```

String literals already have a 0 appended to them, so can be used directly:

```
printf("the_string_is_%s\n", cast(char*)"string_literal");
```

So, why does the first string literal to `printf` not need the cast? The first parameter is prototyped as a `const(char)*`, and a string literal can be implicitly cast to a `const(char)*`. The rest of the arguments to `printf`, however, are variadic (specified by `...`), and a string literal is passed as a (length, pointer) combination to variadic parameters.

The second way is to use the precision specifier. The length comes first, followed by the pointer:

```
printf("the_string_is_%.s\n", str.length, str.ptr);
```

The best way is to use `std.stdio.writeln`, which can handle D strings:

```
import std.stdio;
writeln("the_string_is_%s", str);
```

Void Arrays

There is a special type of array which acts as a wildcard that can hold arrays of any kind, declared as `void[]`. Void arrays are used for low-level operations where some kind of array data is being handled, but the exact type of the array elements are unimportant. The `.length` of a void array

is the length of the data in bytes, rather than the number of elements in its original type. Array indices in indexing and slicing operations are interpreted as byte indices.

Arrays of any type can be implicitly converted to a void array; the compiler inserts the appropriate calculations so that the `.length` of the resulting array's size is in bytes rather than number of elements. Void arrays cannot be converted back to the original type without using a cast, and it is an error to convert to an array type whose element size does not evenly divide the length of the void array.

```
void main()
{
    int[] data1 = [1,2,3];
    long[] data2;

    void[] arr = data1;           // OK, int[] implicit converts to void[].
    assert(data1.length == 3);
    assert(arr.length == 12);     // length is implicitly converted to bytes.

    //data1 = arr;                // Illegal: void[] does not implicitly
                                // convert to int[].

    int[] data3 = cast(int[]) arr; // OK, can convert with explicit cast.
    data2 = cast(long[]) arr;      // Runtime error: long.sizeof == 8, which
                                // does not divide arr.length, which is 12
                                // bytes.
}
```

Void arrays can also be static if their length is known at compile-time. The length is specified in bytes:

```
void main()
{
    byte[2] x;
    int[2] y;

    void[2] a = x; // OK, lengths match
    void[2] b = y; // Error: int[2] is 8 bytes long, doesn't fit in 2 bytes.
}
```

While it may seem that void arrays are just fancy syntax for `ubyte[]`, there is a subtle distinction. The garbage collector generally will not scan `ubyte[]` arrays for pointers, `ubyte[]` being presumed to contain only pure byte data, not pointers. However, it *will* scan `void[]` arrays for pointers, since such an array may have been implicitly converted from an array of pointers or an array of elements that contain pointers. Allocating an array that contains pointers as `ubyte[]` may

run the risk of the GC collecting live memory if these pointers are the only remaining references to their targets.

12.15 Implicit Conversions

A pointer `T*` can be implicitly converted to one of the following:

- `void*`

A static array `T[dim]` can be implicitly converted to one of the following:

- `T[]`
- `const(U) []`
- `const(U[])`
- `void[]`

A dynamic array `T[]` can be implicitly converted to one of the following:

- `const(U) []`
- `const(U[])`
- `void[]`

Where `U` is a base class of `T`.

Chapter 13

Associative Arrays

Associative arrays have an index that is not necessarily an integer, and can be sparsely populated. The index for an associative array is called the *key*, and its type is called the *KeyType*.

Associative arrays are declared by placing the *KeyType* within the [] of an array declaration:

```
int[string] aa;    // Associative array of ints that are
                  // indexed by string keys.
                  // The KeyType is string.
aa["hello"] = 3;   // set value associated with key "hello" to 3
int value = aa["hello"]; // lookup value from a key
assert(value == 3);
```

Note: The built-in associative arrays do not preserve the order of the keys inserted into the array. In particular, in a `foreach` loop the order in which the elements are iterated is unspecified.

Removing Keys

Particular keys in an associative array can be removed with the `remove` function:

```
aa.remove("hello");
```

`remove(key)` does nothing if the given *key* does not exist and returns `false`. If the given *key* does exist, it removes it from the AA and returns `true`.

All keys can be removed by using the method `clear`.

Testing Membership

The *InExpression* yields a pointer to the value if the key is in the associative array, or `null` if not:

```
int* p;
```

```

p = ("hello" in aa);
if (p !is null)
{
    *p = 4; // update value associated with key
    assert(aa["hello"] == 4);
}

```

Neither the *KeyType*s nor the element types of an associative array can be function types or `void`.

Using Classes as the KeyType

Classes can be used as the *KeyType*. For this to work, the class definition must override the following member functions of class `Object`:

- `size_t toHash() @trusted nothrow`
- `bool opEquals(Object)`

Note that the parameter to `opEquals` is of type `Object`, not the type of the class in which it is defined.

For example:

```

class Foo
{
    int a, b;

    override size_t toHash() { return a + b; }

    override bool opEquals(Object o)
    {
        Foo foo = cast(Foo) o;
        return foo && a == foo.a && b == foo.b;
    }
}

```

Care should be taken that `toHash` should consistently be the same value when `opEquals` returns true. In other words, two objects that are considered equal should always have the same hash value. If this is not the case, the associative array will not function properly. Also note that `opCmp` is not used to check for equality by the associative array. However, since the actual `opEquals` or `opCmp` called is not decided until runtime, the compiler cannot always detect mismatched functions. Because of legacy issues, the compiler may reject an associative array key type that overrides `opCmp` but not `opEquals`. This restriction may be removed in future versions of D.

Using Structs or Unions as the KeyType

If the *KeyType* is a struct or union type, a default mechanism is used to compute the hash and comparisons of it based on the binary data within the struct value. A custom mechanism can be used by providing the following functions as struct members:

```
size_t toHash() const @safe pure nothrow;
bool opEquals(ref const typeof(this) s) @safe pure nothrow;
```

For example:

```
import std.string;

struct MyString
{
    string str;

    size_t toHash() const @safe pure nothrow
    {
        size_t hash;
        foreach (char c; str)
            hash = (hash * 9) + c;
        return hash;
    }

    bool opEquals(ref const MyString s) const @safe pure nothrow
    {
        return std.string.cmp(this.str, s.str) == 0;
    }
}
```

Care should be taken that `toHash` should consistently be the same value when `opEquals` returns true. In other words, two structs that are considered equal should always have the same hash value. If this is not the case, the associative array will not function properly.

If necessary the functions can use `@trusted` instead of `@safe`.

Also note that `opCmp` is not used to check for equality by the associative array. For this reason, and for legacy reasons, an associative array key is not allowed to define a specialized `opCmp`, but omit a specialized `opEquals`. This restriction may be removed in future versions of D.

Construction or Assignment on Setting AA Entries

When an AA indexing access appears on the left side of an assignment operator, it is specially handled for setting an AA entry associated with the key.

```
string[int] aa;
string s;
s = aa[1];           // throws RangeError in runtime
aa[1] = "hello";     // handled for setting AA entry
s = aa[1];           // succeeds to lookup
assert(s == "hello");
```

If the assigned value type is equivalent with the AA element type:

1. If the indexing key does not yet exist in AA, a new AA entry will be allocated, and it will be initialized with the assigned value.
2. If the indexing key already exists in the AA, the setting runs normal assignment.

```
struct S
{
    int val;
    void opAssign(S rhs) { this.val = rhs.val * 2; }
}
S[int] aa;
aa[1] = S(10); // first setting initializes the entry aa[1]
assert(aa[1].val == 10);
aa[1] = S(10); // second setting invokes normal assignment, and
               // operator-overloading rewrites it to member opAssign function.
assert(aa[1].val == 20);
```

If the assigned value type is **not** equivalent with the AA element type, the expression could invoke operator overloading with normal indexing access:

```
struct S
{
    int val;
    void opAssign(int v) { this.val = v * 2; }
}
S[int] aa;
aa[1] = 10; // is rewritten to: aa[1].opAssign(10), and
           // throws RangeError before opAssign is called
```

However, if the AA element type is a struct which supports an implicit constructor call from the assigned value, implicit construction is used for setting the AA entry:

```

struct S
{
    int val;
    this(int v) { this.val = v; }
    void opAssign(int v) { this.val = v * 2; }
}

S s = 1;    // OK, rewritten to: S s = S(1);
s = 1;      // OK, rewritten to: s.opAssign(1);

S[int] aa;
aa[1] = 10; // first setting is rewritten to: aa[1] = S(10);
assert(aa[1].val == 10);
aa[1] = 10; // second setting is rewritten to: aa[1].opAssign(10);
assert(aa[1].val == 20);

```

This is designed for efficient memory reuse with some value-semantics structs, eg. `std.bigint` ↪ `.BigInt`.

```

import std.bigint;
BigInt[string] aa;
aa["a"] = 10;    // construct BigInt(10) and move it in AA
aa["a"] = 20;    // call aa["a"].opAssign(20)

```

Static Initialization of AAs

```

immutable long[string] aa = [
    "foo": 5,
    "bar": 10,
    "baz": 2000
];

unittest
{
    assert(aa["foo"] == 5);
    assert(aa["bar"] == 10);
    assert(aa["baz"] == 2000);
}

```

Runtime Initialization of Immutable AAs

Immutable associative arrays are often desirable, but sometimes initialization must be done at runtime. This can be achieved with a constructor (static constructor depending on scope), a buffer

associative array and `assumeUnique`:

```
immutable long[string] aa;

static this()
{
    import std.exception : assumeUnique;
    import std.conv : to;

    long[string] temp; // mutable buffer
    foreach(i; 0 .. 10)
    {
        temp[to!string(i)] = i;
    }
    temp.rehash; // for faster lookups

    aa = assumeUnique(temp);
}

unittest
{
    assert(aa["1"] == 1);
    assert(aa["5"] == 5);
    assert(aa["9"] == 9);
}
```

Construction and Reference Semantics

An Associative Array defaults to `null`, and is constructed upon assigning the first key/value pair. However, once constructed, an associative array has *reference semantics*, meaning that assigning one array to another does not copy the data. This is especially important when attempting to create multiple references to the same array.

```
int[int] aa;           // defaults to null
int[int] aa2 = aa;     // copies the null reference
aa[1] = 1;
assert(aa2.length == 0); // aa2 still is null
aa2 = aa;
aa2[2] = 2;
assert(aa[2] == 2);     // now both refer to the same instance
```

Properties

Properties for associative arrays are:

Associative Array Properties	
Property	Description
<code>.sizeof</code>	Returns the size of the reference to the associative array; it is 4 in 32-bit builds and 8 on 64-bit builds.
<code>.length</code>	Returns number of values in the associative array. Unlike for dynamic arrays, it is read-only.
<code>.dup</code>	Create a new associative array of the same size and copy the contents of the associative array into it.
<code>.keys</code>	Returns dynamic array, the elements of which are the keys in the associative array.
<code>.values</code>	Returns dynamic array, the elements of which are the values in the associative array.
<code>.rehash</code>	Reorganizes the associative array in place so that lookups are more efficient. <code>rehash</code> is effective when, for example, the program is done loading up a symbol table and now needs fast lookups in it. Returns a reference to the reorganized array.
<code>.clear</code>	Removes all remaining keys and values from an associative array. The array is not rehashed after removal, to allow for the existing storage to be reused. This will affect all references to the same instance and is not equivalent to <code>destroy(aa)</code> ↯ ↵ which only sets the current reference to <code>null</code>
<code>.byKey()</code>	Returns a forward range suitable for use as a <i>ForEachAggregate</i> to a <i>ForEachStatement</i> which will iterate over the keys of the associative array.
<code>.byValue()</code>	Returns a forward range suitable for use as a <i>ForEachAggregate</i> to a <i>ForEachStatement</i> which will iterate over the values of the associative array.
<code>.byKeyValue()</code>	Returns a forward range suitable for use as a <i>ForEachAggregate</i> to a <i>ForEachStatement</i> which will iterate over key-value pairs of the associative array. The returned pairs are represented by an opaque type with <code>.key</code> and <code>.value</code> properties for accessing the key and value of the pair, respectively.
<code>.get(Key key, lazy Value defVal)</code>	Looks up <code>key</code> ; if it exists returns corresponding value else evaluates and returns <code>defVal</code> .

Associative Array Example: word count

Let's consider the file is ASCII encoded with LF EOL. In general case we should use *dchar c* for iteration over code points and functions from `std.uni`.

```
import std.file;           // D file I/O
import std.stdio;
import std.ascii;

void main (string[] args)
{
    ulong totalWords, totalLines, totalChars;
    ulong[string] dictionary;

    writeln("lineswordsbytesfile");
    foreach (arg; args[1 .. $]) // for each argument except the first one
    {
        ulong wordCount, lineCount, charCount;

        foreach(line; File(arg).byLine())
        {
            bool inWord;
            size_t wordStart;

            void tryFinishWord(size_t wordEnd)
            {
                if (inWord)
                {
                    auto word = line[wordStart .. wordEnd];
                    ++dictionary[word.idup]; // increment count for word
                    inWord = false;
                }
            }

            foreach (i, char c; line)
            {
                if (std.ascii.isDigit(c))
                {
                    // c is a digit (0..9)
                }
            }
        }
    }
}
```

```

    }
    else if (std::ascii::isAlpha(c))
    {
        // c is an ASCII letter (A..Z, a..z)
        if (!inWord)
        {
            wordStart = i;
            inWord = true;
            ++wordCount;
        }
    }
    else
        tryFinishWord(i);
    ++charCount;
}
tryFinishWord(line.length);
++lineCount;
}

writefln("%8s%8s%8s□%s", lineCount, wordCount, charCount, arg);
totalWords += wordCount;
totalLines += lineCount;
totalChars += charCount;
}

if (args.length > 2)
{
    writefln("-----\n%8s%8s%8s□total",
            totalLines, totalWords, totalChars);
}

writeln("-----");
foreach (word; dictionary.keys.sort)
{
    writefln("%3s□%s", dictionary[word], word);
}
}

```


Chapter 14

Structs and Unions

Whereas classes are reference types, structs are value types. Any C struct can be exactly represented as a D struct, except non-static function-nested D structs which access the context of their enclosing scope. Structs and unions are meant as simple aggregations of data, or as a way to paint a data structure over hardware or an external type. External types can be defined by the operating system API, or by a file format. Object oriented features are provided with the class data type.

A struct is defined to not have an identity; that is, the implementation is free to make bit copies of the struct as convenient.

Struct, Class Comparison Table

Feature	<code>struct</code>	<code>class</code>	C <code>struct</code>	C++ <code>struct</code>	C++ <code>class</code>
value type	✓		✓	✓	✓
reference type		✓			
data members	✓	✓	✓	✓	✓
hidden members	✓	✓		✓	✓
static members	✓	✓		✓	✓
default member initializers	✓	✓			
bit fields			✓	✓	✓
non-virtual member functions	✓	✓		✓	✓
virtual member functions		✓		✓	✓
constructors	✓	✓		✓	✓
postblit/copy constructors	✓			✓	✓
destructors	✓	✓		✓	✓
<i>SharedStaticConstructors</i>	✓	✓			
<i>SharedStaticDestructors</i>	✓	✓			
RAII	✓	✓		✓	✓
identity assign overload	✓			✓	✓
literals	✓				
operator overloading	✓	✓		✓	✓
inheritance		✓		✓	✓
invariants	✓	✓			
unit tests	✓	✓			
synchronizable		✓			
parameterizable	✓	✓		✓	✓
alignment control	✓	✓			
member protection	✓	✓		✓	✓
default public	✓	✓	✓	✓	
tag name space			✓	✓	✓
anonymous	✓		✓	✓	✓
static constructor	✓	✓			
static destructor	✓	✓			
const/immutable/shared	✓	✓			
inner nesting	✓	✓			

AggregateDeclaration:

ClassDeclaration

InterfaceDeclaration

StructDeclaration

UnionDeclaration

StructDeclaration:

```
struct Identifier ;
struct Identifier AggregateBody
StructTemplateDeclaration
AnonStructDeclaration
```

AnonStructDeclaration:

```
struct AggregateBody
```

UnionDeclaration:

```
union Identifier ;
union Identifier AggregateBody
UnionTemplateDeclaration
AnonUnionDeclaration
```

AnonUnionDeclaration:

```
union AggregateBody
```

AggregateBody:

```
{ DeclDefsopt }
```

They work like they do in C, with the following exceptions:

- no bit fields
- alignment can be explicitly specified
- no separate tag name space - tag names go into the current scope
- declarations like:

```
struct ABC x;
```

are not allowed, replace with:

```
ABC x;
```

- anonymous structs/unions are allowed as members of other structs/unions
- Default initializers for members can be supplied.
- Member functions and static members are allowed.

Opaque Structs and Unions

Opaque struct and union declarations do not have a *AggregateBody*:

```
struct S;
union U;
```

The members are completely hidden to the user, and so the only operations on those types are ones that do not require any knowledge of the contents of those types. For example:

```
struct S;
S.sizeof; // error, size is not known
S s;      // error, cannot initialize unknown contents
S* p;     // ok, knowledge of members is not necessary
```

They can be used to implement the PIMPL idiom.

Static Initialization of Structs

Static struct members are by default initialized to whatever the default initializer for the member is, and if none supplied, to the default initializer for the member's type. If a static initializer is supplied, the members are initialized by the member name, colon, expression syntax. The members may be initialized in any order. Initializers for statics must be evaluatable at compile time. Members not specified in the initializer list are default initialized.

```
struct S { int a; int b; int c; int d = 7;}
static S x = { a:1, b:2};           // c is set to 0, d to 7
static S z = { c:4, b:5, a:2 , d:5}; // z.a = 2, z.b = 5, z.c = 4, z.d = 5
```

C-style initialization, based on the order of the members in the struct declaration, is also supported:

```
static S q = { 1, 2 }; // q.a = 1, q.b = 2, q.c = 0, q.d = 7
```

Struct literals can also be used to initialize statics, but they must be evaluable at compile time.

```
static S q = S( 1, 2+3 ); // q.a = 1, q.b = 5, q.c = 0, q.d = 7
```

The static initializer syntax can also be used to initialize non-static variables, provided that the member names are not given. The initializer need not be evaluatable at compile time.

```
void test(int i)
{
    S q = { 1, i }; // q.a = 1, q.b = i, q.c = 0, q.d = 7
}
```

Static Initialization of Unions

Unions are initialized explicitly.

```
union U { int a; double b; }
static U u = { b : 5.0 }; // u.b = 5.0
```

Other members of the union that overlay the initializer, but occupy more storage, have the extra storage initialized to zero.

Dynamic Initialization of Structs

Structs can be dynamically initialized from another value of the same type:

```
struct S { int a; }
S t;      // default initialized
t.a = 3;
S s = t;  // s.a is set to 3
```

If `opCall` is overridden for the struct, and the struct is initialized with a value that is of a different type, then the `opCall` operator is called:

```
struct S
{
    int a;

    static S opCall(int v)
    {
        S s;
        s.a = v;
        return s;
    }

    static S opCall(S v)
    {
        S s;
        s.a = v.a + 1;
        return s;
    }
}

S s = 3; // sets s.a to 3
S t = s; // sets t.a to 3, S.opCall(s) is not called
```

Struct Literals

Struct literals consist of the name of the struct followed by a parenthesized argument list:

```
struct S { int x; float y; }
```

```
int foo(S s) { return s.x; }
```

```
foo( S(1, 2) ); // set field x to 1, field y to 2
```

Struct literals are syntactically like function calls. If a struct has a member function named `opCall`, then struct literals for that struct are not possible. See also `opCall` operator overloading for the issue workaround. It is an error if there are more arguments than fields of the struct. If there are fewer arguments than fields, the remaining fields are initialized with their respective default initializers. If there are anonymous unions in the struct, only the first member of the anonymous union can be initialized with a struct literal, and all subsequent non-overlapping fields are default initialized.

Struct Properties

Struct Properties	
Name	Description
<code>.sizeof</code>	Size in bytes of struct
<code>.alignof</code>	Size boundary struct needs to be aligned on
<code>.tupleof</code>	Gets type tuple of fields

Struct Field Properties

Struct Field Properties	
Name	Description
<code>.offsetof</code>	Offset in bytes of field from beginning of struct

Const, Immutable and Shared Structs

A struct declaration can have a storage class of `const`, `immutable` or `shared`. It has an equivalent effect as declaring each member of the struct as `const`, `immutable` or `shared`.

```
const struct S { int a; int b = 2; }
```

```
void main()
{
```

```

    S s = S(3); // initializes s.a to 3
    S t;        // initializes t.a to 0
    t = s;      // error, t.a and t.b are const, so cannot modify them.
    t.a = 4;    // error, t.a is const
}

```

Struct Constructors

Struct constructors are used to initialize an instance of a struct. The *ParameterList* may not be empty. Struct instances that are not instantiated with a constructor are default initialized to their `.init` value.

```

struct S
{
    int x, y;

    this() // error, cannot implement default ctor for structs
    {
    }

    this(int a, int b)
    {
        x = a;
        y = b;
    }
}

void main()
{
    S a = S(4, 5);
    auto b = S(); // same as auto b = S.init;
}

```

A constructor qualifier allows the object to be constructed with that specific qualifier.

```

struct S1
{
    int[] a;
    this(int n) { a = new int[] (n); }
}

struct S2

```

```

{
    int[] a;
    this(int n) immutable { a = new int[] (n); }
}
void main()
{
    // Mutable constructor creates mutable object.
    S1 m1 = S1(1);

    // Constructed mutable object is implicitly convertible to const.
    const S1 c1 = S1(1);

    // Constructed mutable object is not implicitly convertible to immutable.
    // immutable i1 = S1(1);

    // Mutable constructor cannot construct immutable object.
    // auto x1 = immutable S1(1);

    // Immutable constructor cannot construct mutable object.
    // auto x2 = S2(1);

    // Constructed immutable object is not implicitly convertible to mutable.
    // S2 m2 = immutable S2(1);

    // Constructed immutable object is implicitly convertible to const.
    const S2 c2 = immutable S2(1);

    // Immutable constructor creates immutable object.
    immutable i2 = immutable S2(1);
}

```

If struct constructor is annotated with `@disable` and has empty parameter, the struct is disabled construction without calling other constructor.

```

struct S
{
    int x;

    // Disables default construction, function body can be empty.
    @disable this();
}

```



```

    this(int v) { x = v; }
}
void main()
{
    //S s;          // default construction is disabled
    //S s = S();    // also disabled
    S s = S(1);     // construction with calling constructor
}

```

Struct Postblits

Postblit:

```

this ( this ) MemberFunctionAttributesopt ;
this ( this ) MemberFunctionAttributesopt FunctionBody

```

Copy construction is defined as initializing a struct instance from another struct of the same type. Copy construction is divided into two parts:

1. blitting the fields, i.e. copying the bits
2. running *postblit* on the result

The first part is done automatically by the language, the second part is done if a postblit function is defined for the struct. The postblit has access only to the destination struct object, not the source. Its job is to ‘fix up’ the destination as necessary, such as making copies of referenced data, incrementing reference counts, etc. For example:

```

struct S
{
    int[] a;    // array is privately owned by this instance
    this(this)
    {
        a = a.dup;
    }
}

```

Disabling struct postblit makes the object not copyable.

```

struct T
{
    @disable this(this); // disabling makes T not copyable
}

```

```

struct S
{
    T t;    // uncopyable member makes S also not copyable
}

void main()
{
    S s;
    S t = s; // error, S is not copyable
}

```

Unions may not have fields that have postblits.

Struct Destructors

Destructors are called when an object goes out of scope. Their purpose is to free up resources owned by the struct object.

Unions may not have fields that have destructors.

Identity Assignment Overload

While copy construction takes care of initializing an object from another object of the same type, or elaborate destruction is needed for the type, assignment is defined as copying the contents of one object over another, already initialized, type:

```

struct S { ... } // S has postblit or destructor
S s;             // default construction of s
S t = s;         // t is copy-constructed from s
t = s;           // t is assigned from s

```

Struct assignment `t=s` is defined to be semantically equivalent to:

```
t.opAssign(s);
```

where `opAssign` is a member function of `S`:

```

ref S opAssign(ref S s)
{
    S tmp = this; // bitcopy this into tmp
    this = s;     // bitcopy s into this
    tmp.__dtor(); // call destructor on tmp
    return this;
}

```

An identity assignment overload is required for a struct if one or more of these conditions hold:

- it has a destructor
- it has a postblit
- it has a field with an identity assignment overload

If an identity assignment overload is required and does not exist, an identity assignment overload function of the type `ref S opAssign(ref S)` will be automatically generated.

A user-defined one can implement the equivalent semantics, but can be more efficient.

One reason a custom `opAssign` might be more efficient is if the struct has a reference to a local buffer:

```
struct S
{
    int[] buf;
    int a;

    ref S opAssign(ref const S s)
    {
        a = s.a;
        return this;
    }

    this(this)
    {
        buf = buf.dup;
    }
}
```

Here, `S` has a temporary workspace `buf[]`. The normal postblit will pointlessly free and reallocate it. The custom `opAssign` will reuse the existing storage.

14.1 Nested Structs

A *nested struct* is a struct that is declared inside the scope of a function or a templated struct that has aliases to local functions as a template argument. Nested structs have member functions. It has access to the context of its enclosing scope (via an added hidden field).

```
void foo()
{
    int i = 7;
    struct SS
    {
```

```
    int x,y;
    int bar() { return x + i + 1; }
}
SS s;
s.x = 3;
s.bar(); // returns 11
}
```

A struct can be prevented from being nested by using the static attribute, but then of course it will not be able to access variables from its enclosing scope.

```
void foo()
{
    int i = 7;
    static struct SS
    {
        int x, y;
        int bar()
        {
            return i; // error, SS is not a nested struct
        }
    }
}
```

14.2 Unions and Special Member Functions

Unions may not have postblits, destructors, or invariants.

Chapter 15

Classes

The object-oriented features of D all come from classes. The class hierarchy has as its root the class `Object`. `Object` defines a minimum level of functionality that each derived class has, and a default implementation for that functionality.

Classes are programmer defined types. Support for classes are what make D an object oriented language, giving it encapsulation, inheritance, and polymorphism. D classes support the single inheritance paradigm, extended by adding support for interfaces. Class objects are instantiated by reference only.

A class can be exported, which means its name and all its non-private members are exposed externally to the DLL or EXE.

A class declaration is defined:

ClassDeclaration:

```
class Identifier ;  
class Identifier BaseClassListopt AggregateBody  
ClassTemplateDeclaration
```

BaseClassList:

```
: SuperClass  
: SuperClass , Interfaces  
: Interfaces
```

SuperClass:

```
BasicType
```

Interfaces:

```
Interface  
Interface , Interfaces
```

Interface:
BasicType

Classes consist of:

- a super class
- interfaces
- dynamic fields
- static fields
- types
- an optional synchronized attribute
- member functions
 - static member functions
 - Virtual Functions
 - Constructors
 - Destructors
 - Static Constructors
 - Static Destructors
 - *SharedStaticConstructors*
 - *SharedStaticDestructors*
 - Class Invariants
 - Unit Tests
 - Class Allocators
 - Class Deallocators
 - Alias This

A class is defined:

```
class Foo
{
    ... members ...
}
```

Note that there is no trailing ; after the closing } of the class definition. It is also not possible to declare a variable var like:

```
class Foo { } var;
```

Instead:

```
class Foo { }
Foo var;
```

Access Control

Access to class members is controlled using *ProtectionAttributes*. The default protection attribute is `public`. Access control does not affect visibility.

Fields

Class members are always accessed with the `.` operator.

Members of a base class can be accessed by prepending the name of the base class followed by a dot:

```
class A { int a; int a2;}
class B : A { int a; }

void foo(B b)
{
    b.a = 3;    // accesses field B.a
    b.a2 = 4;   // accesses field A.a2
    b.A.a = 5; // accesses field A.a
}
```

The D compiler is free to rearrange the order of fields in a class to optimally pack them in an implementation-defined manner. Consider the fields much like the local variables in a function - the compiler assigns some to registers and shuffles others around all to get the optimal stack frame layout. This frees the code designer to organize the fields in a manner that makes the code more readable rather than being forced to organize it according to machine optimization rules. Explicit control of field layout is provided by struct/union types, not classes.

Field Properties

The `.offsetof` property gives the offset in bytes of the field from the beginning of the class instantiation. `.offsetof` can only be applied to expressions which produce the type of the field itself, not the class type:

```
class Foo
{
    int x;
}
...
```

```

void test(Foo foo)
{
    size_t o;

    o = Foo.x.offsetof; // error, Foo.x needs a 'this' reference
    o = foo.x.offsetof; // ok
}

```

Class Properties

The `.tupleof` property returns an *ExpressionTuple* of all the fields in the class, excluding the hidden fields and the fields in the base class.

```

class Foo { int x; long y; }
void test(Foo foo)
{
    foo.tupleof[0] = 1; // set foo.x to 1
    foo.tupleof[1] = 2; // set foo.y to 2
    foreach (x; foo.tupleof)
        write(x);      // prints 12
}

```

The properties `__vptr` and `__monitor` give access to the class object's `vtbl[]` and `monitor`, respectively, but should not be used in user code.

Super Class

All classes inherit from a super class. If one is not specified, it inherits from `Object`. `Object` forms the root of the D class inheritance hierarchy.

Member Functions

Non-static member functions have an extra hidden parameter called *this* through which the class object's other members can be accessed.

Non-static member functions can have, in addition to the usual *FunctionAttributes*, the attributes `const`, `immutable`, `shared`, or `inout`. These attributes apply to the hidden *this* parameter.

```

class C
{
    int a;
    const void foo()
    {

```



```

        a = 3; // error, 'this' is const
    }
    void foo() immutable
    {
        a = 3; // error, 'this' is immutable
    }
}

```

Synchronized Classes

All member functions of synchronized classes are synchronized. A static member function is synchronized on the *classinfo* object for the class, which means that one monitor is used for all static member functions for that synchronized class. For non-static functions of a synchronized class, the monitor used is part of the class object. For example:

```

synchronized class Foo
{
    void bar() { ...statements... }
}

```

is equivalent to (as far as the monitors go):

```

synchronized class Foo
{
    void bar()
    {
        synchronized (this) { ...statements... }
    }
}

```

Member functions of non-synchronized classes cannot be individually marked as synchronized. The synchronized attribute must be applied to the class declaration itself:

```

class Foo
{
    synchronized void foo() { } // disallowed!
}

synchronized class Bar
{
    void bar() { } // bar is synchronized
}

```

Member fields of a synchronized class cannot be public:

```
synchronized class Foo
{
    int foo;  // disallowed: public field
}
```

```
synchronized class Bar
{
    private int bar;  // ok
}
```

The `synchronized` attribute can only be applied to classes, structs cannot be marked to be synchronized.

Constructors

Constructor:

```
this Parameters MemberFunctionAttributesopt ;
this Parameters MemberFunctionAttributesopt FunctionBody
ConstructorTemplate
```

Members are always initialized to the default initializer for their type, which is usually 0 for integer types and NAN for floating point types. This eliminates an entire class of obscure problems that come from neglecting to initialize a member in one of the constructors. In the class definition, there can be a static initializer to be used instead of the default:

```
class Abc
{
    int a;      // default initializer for a is 0
    long b = 7; // default initializer for b is 7
    float f;    // default initializer for f is NAN
}
```

This static initialization is done before any constructors are called.

Constructors are defined with a function name of `this` and having no return value:

```
class Foo
{
    this(int x)  // declare constructor for Foo
    {
        ...
    }
}
```

```

    this()
    {
        ...
    }
}

```

Base class construction is done by calling the base class constructor by the name **super**:

```

class A { this(int y) { } }

class B : A
{
    int j;
    this()
    {
        ...
        super(3); // call base constructor A.this(3)
        ...
    }
}

```

Constructors can also call other constructors for the same class in order to share common initializations (this is called delegating constructors):

```

class C
{
    int j;
    this()
    {
        ...
    }
    this(int i)
    {
        this();
        j = i;
    }
}

```

If no call to constructors via **this** or **super** appear in a constructor, and the base class has a constructor, a call to **super()** is inserted at the beginning of the constructor.

If there is no constructor for a class, but there is a constructor for the base class, a default constructor of the form:

```

this() { }

```

is implicitly generated.

Class object construction is very flexible, but some restrictions apply:

1. It is illegal for constructors to mutually call each other, although the compiler is not required to detect it. It will result in undefined behavior.

```
this() { this(1); }
this(int i) { this(); } // illegal, cyclic constructor calls
```

2. If any constructor call appears inside a constructor, any path through the constructor must make exactly one constructor call:

```
this() { a || super(); } // illegal

this() { (a) ? this(1) : super(); } // ok

this()
{
    for (...)
    {
        super(); // illegal, inside loop
    }
}
```

3. It is illegal to refer to `this` implicitly or explicitly prior to making a constructor call.
4. Constructor calls cannot appear after labels (in order to make it easy to check for the previous conditions in the presence of goto's).

Instances of class objects are created with *NewExpressions*:

```
A a = new A(3);
```

The following steps happen:

1. Storage is allocated for the object. If this fails, rather than return `null`, an `OutOfMemoryError` is thrown. Thus, tedious checks for null references are unnecessary.
2. The raw data is statically initialized using the values provided in the class definition. The pointer to the `vtbl[]` (the array of pointers to virtual functions) is assigned. This ensures that constructors are passed fully formed objects for which virtual functions can be called. This operation is equivalent to doing a memory copy of a static version of the object onto the newly allocated one, although more advanced compilers may be able to optimize much of this away.
3. If there is a constructor defined for the class, the constructor matching the argument list is called.

4. If class invariant checking is turned on, the class invariant is called at the end of the constructor.

Constructors can have one of these member function attributes: `const`, `immutable`, and `shared`. Construction of qualified objects will then be restricted to the implemented qualified constructors.

```
class C
{
    this();    // non-shared mutable constructor
}

// create mutable object
C m = new C();

// create const object using by mutable constructor
const C c2 = new const C();

// a mutable constructor cannot create an immutable object
// immutable C i = new immutable C();

// a mutable constructor cannot create a shared object
// shared C s = new shared C();
```

Constructors can be overloaded with different attributes.

```
class C
{
    this();           // non-shared mutable constructor
    this() shared;    // shared mutable constructor
    this() immutable; // immutable constructor
}

C m = new C();
shared s = new shared C();
immutable i = new immutable C();
```

If the constructor can create unique object (e.g. if it is pure), the object can be implicitly convertible to any qualifiers.

```
class C
{
    this() pure;
```

```

// Based on the definition, this creates a mutable object. But the
// created object cannot contain any mutable global data.
// Then compiler can guarantee that the created object is unique.

this(int[] arr) immutable pure;
// Based on the definition, this creates an immutable object. But
// the argument int[] never appears in the created object so it
// isn't implicitly convertible to immutable. Also, it cannot store
// any immutable global data.
// Therefore the compiler can guarantee that the created object is
// unique.
}

immutable i = new immutable C();           // this() pure is called
shared s = new shared C();                 // this() pure is called
C m = new C([1,2,3]);                       // this(int[]) immutable pure is called

```

Field initialization inside constructor

Inside constructor, the first instance field assignment is specially handled for its initialization.

```

class C
{
    int num;
    this()
    {
        num = 1; // initialize
        num = 2; // assignment
    }
}

```

If the field type has opAssign method, it won't be used for initialization.

```

struct A
{
    this(int n) {}
    void opAssign(A rhs) {}
}

class C
{
    A val;
}

```

```

    this()
    {
        val = A(1);  // A(1) is moved in this.val for initializing
        val = A(2);  // rewritten to val.opAssign(A(2))
    }
}

```

If the field type is not modifiable, multiple initialization will be rejected.

```

class C
{
    immutable int num;
    this()
    {
        num = 1;  // OK
        num = 2;  // Error: multiple field initialization
    }
}

```

If the assignment expression for the field initialization may be invoked multiple times, it would also be rejected.

```

class C
{
    immutable int num;
    immutable string str;
    this()
    {
        foreach (i; 0..2)
        {
            num = 1;  // Error: field initialization not allowed in loops
        }
        size_t i = 0;
    Label:
        str = "hello";  // Error: field initialization not allowed after labels
        if (i++ < 2)
            goto Label;
    }
}

```

Destructors

Destructor:

```
~ this ( ) MemberFunctionAttributesopt ;
~ this ( ) MemberFunctionAttributesopt FunctionBody
```

The garbage collector calls the destructor function when the object is deleted. The syntax is:

```
class Foo
{
    ~this() // destructor for Foo
    {
    }
}
```

There can be only one destructor per class, the destructor does not have any parameters, and has no attributes. It is always virtual.

The destructor is expected to release any resources held by the object.

The program can explicitly inform the garbage collector that an object is no longer referred to (with the delete expression), and then the garbage collector calls the destructor immediately, and adds the object's memory to the free storage. The destructor is guaranteed to never be called twice.

The destructor for the super class automatically gets called when the destructor ends. There is no way to call the super destructor explicitly.

The garbage collector is not guaranteed to run the destructor for all unreferenced objects. Furthermore, the order in which the garbage collector calls destructors for unreference objects is not specified. This means that when the garbage collector calls a destructor for an object of a class that has members that are references to garbage collected objects, those references may no longer be valid. This means that destructors cannot reference sub objects. This rule does not apply to auto objects or objects deleted with the *DeleteExpression*, as the destructor is not being run by the garbage collector, meaning all references are valid.

Objects referenced from the data segment never get collected by the gc.

Static Constructors

StaticConstructor:

```
static this ( ) MemberFunctionAttributesopt ;
static this ( ) MemberFunctionAttributesopt FunctionBody
```

A static constructor is a function that performs initializations of thread local data before the `main()` function gets control for the main thread, and upon thread startup.

Static constructors are used to initialize static class members with values that cannot be computed at compile time.

Static constructors in other languages are built implicitly by using member initializers that can't be computed at compile time. The trouble with this stems from not having good control over exactly when the code is executed, for example:

```
class Foo
{
    static int a = b + 1;
    static int b = a * 2;
}
```

What values do a and b end up with, what order are the initializations executed in, what are the values of a and b before the initializations are run, is this a compile error, or is this a runtime error? Additional confusion comes from it not being obvious if an initializer is static or dynamic.

D makes this simple. All member initializations must be determinable by the compiler at compile time, hence there is no order-of-evaluation dependency for member initializations, and it is not possible to read a value that has not been initialized. Dynamic initialization is performed by a static constructor, defined with a special syntax `static this()`.

```
class Foo
{
    static int a;           // default initialized to 0
    static int b = 1;
    static int c = b + a; // error, not a constant initializer

    static this()          // static constructor
    {
        a = b + 1;         // a is set to 2
        b = a * 2;         // b is set to 4
    }
}
```

If `main()` or the thread returns normally, (does not throw an exception), the static destructor is added to the list of functions to be called on thread termination.

Static constructors have empty parameter lists.

Static constructors within a module are executed in the lexical order in which they appear. All the static constructors for modules that are directly or indirectly imported are executed before the static constructors for the importer.

The `static` in the static constructor declaration is not an attribute, it must appear immediately before the `this`:

```
class Foo
```

```

{
    static this() { ... } // a static constructor
    static private this() { ... } // not a static constructor
    static
    {
        this() { ... } // not a static constructor
    }
    static:
        this() { ... } // not a static constructor
}

```

Static Destructors

```

StaticDestructor:
    static ~ this ( ) MemberFunctionAttributesopt ;
    static ~ this ( ) MemberFunctionAttributesopt FunctionBody

```

A static destructor is defined as a special static function with the syntax `static ~this()`.

```

class Foo
{
    static ~this() // static destructor
    {
    }
}

```

A static destructor gets called on thread termination, but only if the static constructor completed successfully. Static destructors have empty parameter lists. Static destructors get called in the reverse order that the static constructors were called in.

The `static` in the static destructor declaration is not an attribute, it must appear immediately before the `~this`:

```

class Foo
{
    static ~this() { ... } // a static destructor
    static private ~this() { ... } // not a static destructor
    static
    {
        ~this() { ... } // not a static destructor
    }
    static:

```

```

    ~this() { ... } // not a static destructor
}

```

Shared Static Constructors

```

SharedStaticConstructor:
    shared static this ( ) MemberFunctionAttributesopt ;
    shared static this ( ) MemberFunctionAttributesopt FunctionBody

```

Shared static constructors are executed before any *StaticConstructors*, and are intended for initializing any shared global data.

Shared Static Destructors

```

SharedStaticDestructor:
    shared static ~ this ( ) MemberFunctionAttributesopt ;
    shared static ~ this ( ) MemberFunctionAttributesopt FunctionBody

```

Shared static destructors are executed at program termination in the reverse order that *SharedStaticConstructors* were executed.

Class Invariants

```

Invariant:
    invariant ( ) BlockStatement
    invariant BlockStatement

```

Class invariants are used to specify characteristics of a class that always must be true (except while executing a member function). They are described in *Invariants*.

Class Allocators

Note: Class allocators are deprecated in D2.

```

Allocator:
    new Parameters ;
    new Parameters FunctionBody

```

A class member function of the form:

```
new(uint size)
{
    ...
}
```

is called a class allocator. The class allocator can have any number of parameters, provided the first one is of type `uint`. Any number can be defined for a class, the correct one is determined by the usual function overloading rules. When a new expression:

```
new Foo;
```

is executed, and `Foo` is a class that has an allocator, the allocator is called with the first argument set to the size in bytes of the memory to be allocated for the instance. The allocator must allocate the memory and return it as a `void*`. If the allocator fails, it must not return a `null`, but must throw an exception. If there is more than one parameter to the allocator, the additional arguments are specified within parentheses after the `new` in the *NewExpression*:

```
class Foo
{
    this(char[] a) { ... }

    new(uint size, int x, int y)
    {
        ...
    }
}

...

new(1,2) Foo(a);           // calls new(Foo.sizeof,1,2)
```

Derived classes inherit any allocator from their base class, if one is not specified.

The class allocator is not called if the instance is created on the stack.

See also Explicit Class Instance Allocation.

Class Deallocators

Note: Class deallocators and the delete operator are deprecated in D2. Use the `destroy` function to finalize an object by calling its destructor. The memory of the object is **not** immediately deallocated, instead the GC will collect the memory of the object at an undetermined point after finalization:

```
class Foo { int x; this() { x = 1; } }
Foo foo = new Foo;
destroy(foo);
assert(foo.x == int.init); // object is still accessible
```

```
Deallocator:
    delete Parameters ;
    delete Parameters FunctionBody
```

A class member function of the form:

```
delete(void *p)
{
    ...
}
```

is called a class deallocator. The deallocator must have exactly one parameter of type `void*`. Only one can be specified for a class. When a delete expression:

```
delete f;
```

is executed, and `f` is a reference to a class instance that has a deallocator, the deallocator is called with a pointer to the class instance after the destructor (if any) for the class is called. It is the responsibility of the deallocator to free the memory.

Derived classes inherit any deallocator from their base class, if one is not specified.

The class allocator is not called if the instance is created on the stack.

See also Explicit Class Instance Allocation.

Alias This

```
AliasThis:
    alias Identifier this ;
```

An *AliasThis* declaration names a member to subtype. The *Identifier* names that member.

A class or struct can be implicitly converted to the *AliasThis* member.

```
struct S
{
    int x;
    alias x this;
}
```

```

int foo(int i) { return i * 2; }

void test()
{
    S s;
    s.x = 7;
    int i = -s;    // i == -7
    i = s + 8;     // i == 15
    i = s + s;     // i == 14
    i = 9 + s;     // i == 16
    i = foo(s);    // implicit conversion to int
}

```

If the member is a class or struct, undefined lookups will be forwarded to the *AliasThis* member.

```

struct Foo
{
    int baz = 4;
    int get() { return 7; }
}

class Bar
{
    Foo foo;
    alias foo this;
}

void test()
{
    auto bar = new Bar;
    int i = bar.baz; // i == 4
    i = bar.get();  // i == 7
}

```

If the *Identifier* refers to a property member function with no parameters, conversions and undefined lookups are forwarded to the return value of the function.

```

struct S
{
    int x;
    @property int get()
}

```

```

    {
        return x * 2;
    }
    alias get this;
}

void test()
{
    S s;
    s.x = 2;
    int i = s; // i == 4
}

```

Multiple *AliasThis* are allowed. For implicit conversions and forwarded lookups, all *AliasThis* declarations are attempted; if more than one *AliasThis* is eligible, the ambiguity is disallowed by raising an error. Note: Multiple *AliasThis* is currently unimplemented.

Scope Classes

Note: Scope classes have been recommended for deprecation.

A scope class is a class with the `scope` attribute, as in:

```
scope class Foo { ... }
```

The `scope` characteristic is inherited, so any classes derived from a scope class are also scope.

A scope class reference can only appear as a function local variable. It must be declared as being `scope`:

```
scope class Foo { ... }
```

```

void func()
{
    Foo f;    // error, reference to scope class must be scope
    scope Foo g = new Foo(); // correct
}

```

When a scope class reference goes out of scope, the destructor (if any) for it is automatically called. This holds true even if the scope was exited via a thrown exception.

Final Classes

Final classes cannot be subclassed:

```
final class A { }  
class B : A { } // error, class A is final
```

15.1 Nested Classes

A *nested class* is a class that is declared inside the scope of a function or another class. A nested class has access to the variables and other symbols of the classes and functions it is nested inside:

```
class Outer  
{  
    int m;  
  
    class Inner  
    {  
        int foo()  
        {  
            return m; // Ok to access member of Outer  
        }  
    }  
}  
  
void func()  
{  
    int m;  
  
    class Inner  
    {  
        int foo()  
        {  
            return m; // Ok to access local variable m of func()  
        }  
    }  
}
```

If a nested class has the `static` attribute, then it can not access variables of the enclosing scope that are local to the stack or need a `this`:

```
class Outer  
{  
    int m;  
    static int n;
```



```

static class Inner
{
    int foo()
    {
        return m;    // Error, Inner is static and m needs a this
        return n;    // Ok, n is static
    }
}

void func()
{
    int m;
    static int n;

    static class Inner
    {
        int foo()
        {
            return m;    // Error, Inner is static and m is local to the stack
            return n;    // Ok, n is static
        }
    }
}

```

Non-static nested classes work by containing an extra hidden member (called the context pointer) that is the frame pointer of the enclosing function if it is nested inside a function, or the `this` of the enclosing class's instance if it is nested inside a class.

When a non-static nested class is instantiated, the context pointer is assigned before the class's constructor is called, therefore the constructor has full access to the enclosing variables. A non-static nested class can only be instantiated when the necessary context pointer information is available:

```

class Outer
{
    class Inner { }

    static class SInner { }
}

void func()

```

```

{
    class Nested { }

    Outer o = new Outer;           // Ok
    Outer.Inner oi = new Outer.Inner; // Error, no 'this' for Outer
    Outer.SInner os = new Outer.SInner; // Ok

    Nested n = new Nested;        // Ok
}

```

A *this* can be supplied to the creation of an inner class instance by prefixing it to the *NewExpression*:

```

class Outer
{
    int a;

    class Inner
    {
        int foo()
        {
            return a;
        }
    }
}

int bar()
{
    Outer o = new Outer;
    o.a = 3;
    Outer.Inner oi = o.new Inner;
    return oi.foo();    // returns 3
}

```

Here *o* supplies the *this* to the outer class instance of **Outer**.

The property **.outer** used in a nested class gives the **this** pointer to its enclosing class. If there is no enclosing class context, **.outer** would return a pointer to enclosing function frame with **void***.

```

class Outer
{
    class Inner1
    {

```

```
    Outer getOuter()
    {
        return this.outer;
    }
}

void foo()
{
    Inner1 i = new Inner1;
    assert(i.getOuter() is this);
}

void bar()
{
    // x is referenced from nested scope, so
    // bar makes a closure environment.
    int x = 1;

    class Inner2
    {
        Outer getOuter()
        {
            x = 2;
            // The Inner2 instance owns function frame of bar
            // as static frame pointer, but .outer yet returns
            // the enclosing Outer class instance property.
            return this.outer;
        }
    }

    Inner2 i = new Inner2;
    assert(i.getOuter() is this);
}

static void baz()
{
    // make a closure environment
    int x = 1;

    class Inner3
```

```

    {
        void* getOuter()
        {
            x = 2;
            // There's no accessible enclosing class instance, so
            // .outer property returns the function frame of bar.
            return this.outer;
        }
    }

    Inner3 i = new Inner3;
    assert(i.getOuter() != null);
}

```

Anonymous Nested Classes

An anonymous nested class is both defined and instantiated with a *NewAnonClassExpression*:

NewAnonClassExpression:

```
new AllocatorArgumentsopt class ClassArgumentsopt SuperClassopt Interfacesopt AggregateBody
```

ClassArguments:

```
( ArgumentListopt )
```

which is equivalent to:

```
\lstineline|class| \textit{Identifier} \lstineline|:| \textit{SuperClass} \textit{Interface}
↳ Interfaces} \textit{AggregateBody}
```

```
\lstineline|new| \lstineline|(|\textit{ArgumentList}\lstineline|)| \textit{Interface}
↳ Identifier} \lstineline|(|\textit{ArgumentList}\lstineline|)|;
```

where *Identifier* is the name generated for the anonymous nested class.

Const, Immutable and Shared Classes

If a *ClassDeclaration* has a `const`, `immutable` or `shared` storage class, then it is as if each member of the class was declared with that storage class. If a base class is `const`, `immutable` or `shared`, then all classes derived from it are also `const`, `immutable` or `shared`.

Chapter 16

Interfaces

```
InterfaceDeclaration:  
    interface Identifier ;  
    interface Identifier BaseInterfaceListopt AggregateBody  
    InterfaceTemplateDeclaration
```

```
BaseInterfaceList:  
    : Interfaces
```

Interfaces describe a list of functions that a class that inherits from the interface must implement. A class that implements an interface can be converted to a reference to that interface.

Some operating system objects, like COM/OLE/ActiveX for Win32, have specialized interfaces. D interfaces that are compatible with COM/OLE/ActiveX are called *COM Interfaces*.

C++ Interfaces are another form of interfaces, meant to be binary compatible with C++.

Interfaces cannot derive from classes; only from other interfaces. Classes cannot derive from an interface multiple times.

```
interface D  
{  
    void foo();  
}  
  
class A : D, D // error, duplicate interface  
{  
}  
}
```

An instance of an interface cannot be created.

```
interface D
```

```
{
    void foo();
}
```

```
...
```

```
D d = new D(); // error, cannot create instance of interface
```

Virtual interface member functions do not have implementations. Interfaces are expected to implement static or final functions.

```
interface D
{
    void bar() { } // error, implementation not allowed
    static void foo() { } // ok
    final void abc() { } // ok
}
```

Classes that inherit from an interface may not override final or static interface member functions.

```
interface D
{
    void bar();
    static void foo() { }
    final void abc() { }
}

class C : D
{
    void bar() { } // ok
    void foo() { } // error, cannot override static D.foo()
    void abc() { } // error, cannot override final D.abc()
}
```

All interface functions must be defined in a class that inherits from that interface:

```
interface D
{
    void foo();
}

class A : D
{
```

```
    void foo() { } // ok, provides implementation
}
```

```
class B : D
{
    int foo() { } // error, no void foo() implementation
}
```

Interfaces can be inherited and functions overridden:

```
interface D
{
    int foo();
}
```

```
class A : D
{
    int foo() { return 1; }
}
```

```
class B : A
{
    int foo() { return 2; }
}
```

...

```
B b = new B();
b.foo(); // returns 2
D d = cast(D) b; // ok since B inherits A's D implementation
d.foo(); // returns 2;
```

Interfaces can be reimplemented in derived classes:

```
interface D
{
    int foo();
}
```

```
class A : D
{
    int foo() { return 1; }
}
```

```

}

class B : A, D
{
    int foo() { return 2; }
}

...

B b = new B();
b.foo();           // returns 2
D d = cast(D) b;
d.foo();           // returns 2
A a = cast(A) b;
D d2 = cast(D) a;
d2.foo();          // returns 2, even though it is A's D, not B's D

```

A reimplemented interface must implement all the interface functions, it does not inherit them from a super class:

```

interface D
{
    int foo();
}

class A : D
{
    int foo() { return 1; }
}

class B : A, D
{
    // error, no foo() for interface D
}

```

16.1 Interfaces with Contracts

Interface member functions can have contracts even though there is no body for the function. The contracts are inherited by any class member function that implements that interface member function.

```

interface I

```



```

{
    int foo(int i)
    in { assert(i > 7); }
    out (result) { assert(result & 1); }

    void bar();
}

```

16.2 Const and Immutable Interfaces

If an interface has `const` or `immutable` storage class, then all members of the interface are `const` or `immutable`. This storage class is not inherited.

16.3 COM Interfaces

A variant on interfaces is the COM interface. A COM interface is designed to map directly onto a Windows COM object. Any COM object can be represented by a COM interface, and any D object with a COM interface can be used by external COM clients.

A COM interface is defined as one that derives from the interface `core.stdc.windows.com.IUnknown`. ↵ `IUnknown`. A COM interface differs from a regular D interface in that:

- It derives from the interface `core.stdc.windows.com.IUnknown`.
- It cannot be the argument of a *DeleteExpression*.
- References cannot be upcast to the enclosing class object, nor can they be downcast to a derived interface. To accomplish this, an appropriate `QueryInterface()` would have to be implemented for that interface in standard COM fashion.
- Classes derived from COM interfaces are COM classes.
- The default linkage for member functions of COM classes is `extern(System)`.

Note that if you want to implement or override any base-class methods of D interfaces or classes (ones which do not inherit from `IUnknown`), you have to explicitly mark them as having the `extern(D)` linkage:

```

import core.sys.windows.windows;
import core.stdc.windows.com;

interface IText
{
    void write();
}

```

```

abstract class Printer : IText
{
    void print() { }
}

class C : Printer, IUnknown
{
    // Implements the IText \$(D write) class method.
    extern(D) void write() { }

    // Overrides the Printer \$(D print) class method.
    extern(D) override void print() { }

    // Overrides the Object base class \$(D toString) method.
    extern(D) override string toString() { return "Class_C"; }

    // Methods of class implementing the IUnknown interface have
    // the extern(System) calling convention by default.
    HRESULT QueryInterface(const(IID)*, void**);
    uint AddRef();
    uint Release();
}

```

The same applies to other `Object` methods such as `opCmp`, `toHash`, etc.

- The first member of the `Vtbl[]` is not the pointer to the `InterfaceInfo`, but the first virtual function pointer.

For more information, see *Modern COM Programming in D*

16.4 C++ Interfaces

C++ interfaces are interfaces declared with C++ linkage:

```

extern (C++) interface Ifoo
{
    void foo();
    void bar();
}

```

which is meant to correspond with the following C++ declaration:

```
class Ifoo
{
    virtual void foo();
    virtual void bar();
};
```

Any interface that derives from a C++ interface is also a C++ interface. A C++ interface differs from a D interface in that:

- It cannot be the argument of a *DeleteExpression*.
- References cannot be upcast to the enclosing class object, nor can they be downcast to a derived interface.
- The C++ calling convention is the default convention for its member functions, rather than the D calling convention.
- The first member of the `vtbl[]` is not the pointer to the **Interface**, but the first virtual function pointer.

Chapter 17

Enums

```
EnumDeclaration:  
    enum Identifier EnumBody  
    enum Identifier : EnumBaseType EnumBody  
    AnonymousEnumDeclaration
```

```
EnumBaseType:  
    Type
```

```
EnumBody:  
    { EnumMembers }  
    ;
```

```
EnumMembers:  
    EnumMember  
    EnumMember ,  
    EnumMember , EnumMembers
```

```
EnumMember:  
    Identifier  
    Identifier = AssignExpression
```

```
AnonymousEnumDeclaration:  
    enum : EnumBaseType { EnumMembers }  
    enum { EnumMembers }  
    enum { AnonymousEnumMembers }
```

```
AnonymousEnumMembers:
    AnonymousEnumMember
    AnonymousEnumMember ,
    AnonymousEnumMember , AnonymousEnumMembers
```

```
AnonymousEnumMember:
    EnumMember
    Type Identifier = AssignExpression
```

Enum declarations are used to define a group of constants. They come in these forms:

1. Named enums, which have a name.
2. Anonymous enums, which do not have a name.
3. Manifest constants.

17.1 Named Enums

Named enums are used to declare related constants and group them by giving them a unique type. The *EnumMembers* are declared in the scope of the named enum. The named enum declares a new type, and all the *EnumMembers* have that type.

This defines a new type *X* which has values *X.A=0*, *X.B=1*, *X.C=2*:

```
enum X { A, B, C } // named enum
```

If the *EnumBaseType* is not explicitly set, and the first *EnumMember* has an *AssignExpression*, it is set to the type of that *AssignExpression*. Otherwise, it defaults to type *int*.

Named enum members may not have individual *Types*.

A named enum member can be implicitly cast to its *EnumBaseType*, but *EnumBaseType* types cannot be implicitly cast to an enum type.

The value of an *EnumMember* is given by its *AssignExpression*. If there is no *AssignExpression* and it is the first *EnumMember*, its value is `\textit{EnumBaseType}.init`.

If there is no *AssignExpression* and it is not the first *EnumMember*, it is given the value of the previous `\textit{EnumMember}+1`. If the value of the previous *EnumMember* is `\textit{EnumBaseType}.max`, it is an error. If the value of the previous `\textit{EnumMember}+1` is the same as the value of the previous *EnumMember*, it is an error. (This can happen with floating point types.)

All *EnumMembers* are in scope for the *AssignExpressions*.

```
enum A = 3;
enum B
```

```

{
    A = A // error, circular reference
}
enum C
{
    A = B,    // A = 4
    B = D,    // B = 4
    C = 3,    // C = 3
    D         // D = 4
}
enum E : C
{
    E1 = C.D,
    E2      // error, C.D is C.max
}

```

An empty enum body (For example `enum E;`) signifies an opaque enum - the enum members are unknown.

Enum Default Initializer

The `.init` property of an enum type is the value of the first member of that enum. This is also the default initializer for the enum type.

```

enum X { A=3, B, C }
X x;      // x is initialized to 3

```

Enum Properties

Enum properties only exist for named enums.

<caption>Named Enum Properties</caption>	
<code>.init</code>	First enum member value
<code>.min</code>	Smallest value of enum
<code>.max</code>	Largest value of enum
<code>.sizeof</code>	Size of storage for an enumerated value

For example:

```

enum X { A=3, B, C }
X.min    // is X.A
X.max    // is X.C
X.sizeof // is same as int.sizeof

```

The *EnumBaseType* of named enums must support comparison in order to compute the `.max` and `.min` properties.

If the enum *Identifier* is not present, then the enum is an *anonymous enum*, and the *EnumMembers* are declared in the scope the *EnumDeclaration* appears in. No new type is created.

1. The *Type*, if present. Types are not permitted when an *EnumBaseType* is present.
2. The *EnumBaseType*, if present.
3. The type of the *AssignExpression*, if present.
4. The type of the previous *EnumMember*, if present.
5. `int`

Defines the constants A=0, B=1, C=2, all of type int.
Enums must have at least one member.

If there is no *AssignExpression* and it is not the first *EnumMember*, it is given the value of the previous `\textit{EnumMember}+1`. If the value of the previous *EnumMember* is the `.max` property of the previous *EnumMember*'s type, it is an error. If the value of the previous `\textit{EnumMember} \hookrightarrow +1` is the same as the value of the previous *EnumMember*, it is an error. (This can happen with floating point types.)

Sets A=0, B=12, C=13, D=21, and E=22, all of type int.

Sets A=3, B=4 all of type long.

```
enum
{
    A = 1.2f, // A is 1.2f of type float
}
```



```
B,          // B is 2.2f of type float
int C = 3,  // C is 3 of type int
D          // D is 4 of type int
}
```

17.3 Manifest Constants

If there is only one member of an anonymous enum, the `enum` can be omitted. Gramatically speaking, this is an *AutoDeclaration*.

```
enum i = 4;      // i is 4 of type int
enum long l = 3; // l is 3 of type long
```

Such declarations are not lvalues, meaning their address cannot be taken. They exist only in the memory of the compiler.

```
enum size = __traits(classInstanceSize, Foo); // evaluated at compile-time
```

Using manifest constants is an idiomatic D method to force compile-time evaluation of an expression.

Chapter 18

Type Qualifiers

Type qualifiers modify a type by applying a *TypeCtor*. *TypeCtors* are: `const`, `immutable`, `shared`, and `inout`. Each applies transitively to all subtypes.

18.1 Const and Immutable

When examining a data structure or interface, it is very helpful to be able to easily tell which data can be expected to not change, which data might change, and who may change that data. This is done with the aid of the language typing system. Data can be marked as `const` or `immutable`, with the default being changeable (or *mutable*).

`immutable` applies to data that cannot change. Immutable data values, once constructed, remain the same for the duration of the program's execution. Immutable data can be placed in ROM (Read Only Memory) or in memory pages marked by the hardware as read only. Since immutable data does not change, it enables many opportunities for program optimization, and has applications in functional style programming.

`const` applies to data that cannot be changed by the `const` reference to that data. It may, however, be changed by another reference to that same data. `Const` finds applications in passing data through interfaces that promise not to modify them.

Both `immutable` and `const` are *transitive*, which means that any data reachable through an `immutable` reference is also `immutable`, and likewise for `const`.

18.2 Immutable Storage Class

The simplest `immutable` declarations use it as a storage class. It can be used to declare manifest constants.

```
immutable int x = 3;  // x is set to 3
```

```
x = 4;           // error, x is immutable
char[x] s;       // s is an array of 3 char's
```

The type can be inferred from the initializer:

```
immutable y = 4; // y is of type int
y = 5;           // error, y is immutable
```

If the initializer is not present, the immutable can be initialized from the corresponding constructor:

```
immutable int z;
void test()
{
    z = 3; // error, z is immutable
}
static this()
{
    z = 3; // ok, can set immutable that doesn't
           // have static initializer
}
```

The initializer for a non-local immutable declaration must be evaluable at compile time:

```
int foo(int f) { return f * 3; }
int i = 5;
immutable x = 3 * 4; // ok, 12
immutable y = i + 1; // error, cannot evaluate at compile time
immutable z = foo(2) + 1; // ok, foo(2) can be evaluated at compile time, 7
```

The initializer for a non-static local immutable declaration is evaluated at run time:

```
int foo(int f)
{
    immutable x = f + 1; // evaluated at run time
    x = 3;               // error, x is immutable
}
```

Because immutable is transitive, data referred to by an immutable is also immutable:

```
immutable char[] s = "foo";
s[0] = 'a'; // error, s refers to immutable data
s = "bar";  // error, s is immutable
```

Immutable declarations can appear as lvalues, i.e. they can have their address taken, and occupy storage.

18.3 Const Storage Class

A const declaration is exactly like an immutable declaration, with the following differences:

- Any data referenced by the const declaration cannot be changed from the const declaration, but it might be changed by other references to the same data.
- The type of a const declaration is itself const.

18.4 Immutable Type

Data that will never change its value can be typed as immutable. The immutable keyword can be used as a *type qualifier*:

```
immutable(char)[] s = "hello";
```

The immutable applies to the type within the following parentheses. So, while `s` can be assigned new values, the contents of `s[]` cannot be:

```
s[0] = 'b'; // error, s[] is immutable
s = null;   // ok, s itself is not immutable
```

Immutability is transitive, meaning it applies to anything that can be referenced from the immutable type:

```
immutable(char)** p = ...;
p = ...;           // ok, p is not immutable
*p = ...;          // ok, *p is not immutable
**p = ...;         // error, **p is immutable
***p = ...;        // error, ***p is immutable
```

Immutable used as a storage class is equivalent to using immutable as a type qualifier for the entire type of a declaration:

```
immutable int x = 3; // x is typed as immutable(int)
immutable(int) y = 3; // y is immutable
```

18.5 Creating Immutable Data

The first way is to use a literal that is already immutable, such as string literals. String literals are always immutable.

```
auto s = "hello"; // s is immutable(char)[5]
char[] p = "world"; // error, cannot implicitly convert immutable
                      // to mutable
```

The second way is to cast data to immutable. When doing so, it is up to the programmer to ensure that no other mutable references to the same data exist.

```
char[] s = ...;
immutable(char)[] p = cast(immutable)s;    // undefined behavior
immutable(char)[] p = cast(immutable)s.dup; // ok, unique reference
```

The `.idup` property is a convenient way to create an immutable copy of an array:

```
auto p = s.idup;
p[0] = ...;    // error, p[] is immutable
```

18.6 Removing Immutable or Const with a Cast

An immutable or const type qualifier can be removed with a cast:

```
immutable int* p = ...;
int* q = cast(int*)p;
```

This does not mean, however, that one can change the data:

```
*q = 3; // allowed by compiler, but result is undefined behavior
```

The ability to cast away immutable-correctness is necessary in some cases where the static typing is incorrect and not fixable, such as when referencing code in a library one cannot change. Casting is, as always, a blunt and effective instrument, and when using it to cast away immutable-correctness, one must assume the responsibility to ensure the immutability of the data, as the compiler will no longer be able to statically do so.

Note that casting away a const qualifier and then mutating is undefined behavior, too, even when the referenced data is mutable. This is so that compilers and programmers can make assumptions based on const alone. For example, here it may be assumed that `f` does not alter `x`:

```
void f(const int* a);
void main()
{
    int x = 1;
    f(&x);
    assert(x == 1); // guaranteed to hold
}
```

18.7 Immutable Member Functions

Immutable member functions are guaranteed that the object and anything referred to by the `this` reference is immutable. They are declared as:

```

struct S
{
    int x;

    void foo() immutable
    {
        x = 4;          // error, x is immutable
        this.x = 4;     // error, x is immutable
    }
}

```

Note that using `immutable` on the left hand side of a method does not apply to the return type:

```

struct S
{
    immutable int[] bar() // bar is still immutable, return type is not!
    {
    }
}

```

To make the return type immutable, you need to surround the return type with parentheses:

```

struct S
{
    immutable(int[]) bar() // bar is now mutable, return type is immutable.
    {
    }
}

```

To make both the return type and the method immutable, you can write:

```

struct S
{
    immutable(int[]) bar() immutable
    {
    }
}

```

18.8 Const Type

Const types are like immutable types, except that `const` forms a read-only *view* of data. Other aliases to that same data may change it at any time.

18.9 Const Member Functions

Const member functions are functions that are not allowed to change any part of the object through the member function's this reference.

18.10 Implicit Conversions

Values can be implicitly converted between *mutable*, *const*, *immutable*, *shared const*, *inout* and *inout shared*.

References can be converted according to the following rules:

Implicit Conversion of Reference Types									
from/to	<i>mutable</i>	<i>const</i>	<i>shared</i>	<i>shared const</i>	<i>inout</i>	<i>inout const</i>	<i>inout shared</i>	<i>inout shared const</i>	<i>immutable</i>
<i>mutable</i>	✓	✓							
<i>const</i>		✓							
<i>shared</i>			✓	✓					
<i>shared const</i>				✓					
<i>inout</i>		✓			✓	✓			
<i>inout const</i>		✓				✓			
<i>inout shared</i>				✓			✓	✓	
<i>inout shared const</i>				✓				✓	
<i>immutable</i>		✓		✓		✓		✓	✓

If an implicit conversion is disallowed by the table, an *Expression* may be converted if:

An expression may be converted from *mutable* or *shared* to *immutable* if the expression is unique and all expressions it transitively refers to are either unique or *immutable*.

An expression may be converted from *mutable* to *shared* if the expression is unique and all expressions it transitively refers to are either unique, *immutable*, or *shared*.

An expression may be converted from *immutable* to *mutable* if the expression is unique.

An expression may be converted from *shared* to *mutable* if the expression is unique.

A *Unique Expression* is one for which there are no other references to the value of the expression and all expressions it transitively refers to are either also unique or are *immutable*. For example:


```
void main()
{
    immutable int** p = new int*(null); // ok, unique

    int x;
    immutable int** q = new int*(&x); // error, there may be other references to ↵
                                     ↵ x

    immutable int y;
    immutable int** r = new immutable(int)*(&y); // ok, y is immutable
}
```

Otherwise, a *CastExpression* can be used to force a conversion when an implicit version is disallowed, but this cannot be done in **@safe** code, and the correctness of it must be verified by the user.

Chapter 19

Functions

Grammar

FuncDeclaration:

*StorageClasses*_{opt} *BasicType* *FuncDeclarator* *FunctionBody*
AutoFuncDeclaration

AutoFuncDeclaration:

StorageClasses *Identifier* *FuncDeclaratorSuffix* *FunctionBody*

FuncDeclarator:

*BasicType2*_{opt} *Identifier* *FuncDeclaratorSuffix*

FuncDeclaratorSuffix:

Parameters *MemberFunctionAttributes*_{opt}
TemplateParameters *Parameters* *MemberFunctionAttributes*_{opt} *Constraint*_{opt}

Parameters:

(*ParameterList*_{opt})

ParameterList:

Parameter
Parameter , *ParameterList*
...

Parameter:

*InOut*_{opt} *BasicType* *Declarator*

```

InOutopt BasicType Declarator ...
InOutopt BasicType Declarator = AssignExpression
InOutopt Type
InOutopt Type ...

```

InOut:

```

InOutX
InOut InOutX

```

InOutX:

```

auto
TypeCtor
final
in
lazy
out
ref
return ref
scope

```

FunctionAttributes:

```

FunctionAttribute
FunctionAttribute FunctionAttributes

```

FunctionAttribute:

```

nothrow
pure
Property

```

MemberFunctionAttributes:

```

MemberFunctionAttribute
MemberFunctionAttribute MemberFunctionAttributes

```

MemberFunctionAttribute:

```

const
immutable
inout
return
shared
FunctionAttribute

```

```

FunctionBody:
    BlockStatement
    FunctionContractsopt BodyStatement
    FunctionContracts

```

```

FunctionContracts:
    InStatement OutStatementopt
    OutStatement InStatementopt

```

```

InStatement:
    in BlockStatement

```

```

OutStatement:
    out BlockStatement
    out ( Identifier ) BlockStatement

```

```

BodyStatement:
    body BlockStatement

```

Contracts

The **in** and **out** blocks of a function declaration specify the pre- and post-conditions of the function. They are used in Contract Programming. The code inside these blocks should not have any side-effects, including modifying function parameters and/or return values.

Function Return Values

Function return values are considered to be rvalues. This means they cannot be passed by reference to other functions.

Functions Without Bodies

Functions without bodies:

```
int foo();
```

that are not declared as **abstract** are expected to have their implementations elsewhere, and that implementation will be provided at the link step. This enables an implementation of a function to be completely hidden from the user of it, and the implementation may be in another language such as C, assembler, etc.

Pure Functions

Pure functions are functions that cannot access global/static mutable state, except if their arguments contain pointers to such. This enables optimizations based on the fact that a pure function may at most mutate state reachable through its parameters. To that end, a pure function:

- does not read or write any global or static mutable state
- cannot call functions that are not pure
- can override an impure function, but cannot be overridden by an impure function
- is covariant with an impure function
- cannot perform I/O

This definition of mutable functions is more general than the one traditionally employed by pure functional languages because it allows a D pure function to use state mutation, as long as all state is created internally or reachable through its arguments. In particular, a pure function may allocate memory by means of e.g. `new` or `malloc` without these being special cases. A pure function is allowed to loop indefinitely or terminate the program.

As a concession to practicality, a pure function can also:

- read and write the floating point exception flags
- read and write the floating point mode flags, as long as those flags are restored to their initial state upon function entry
- perform impure operations in statements that are in a *ConditionalStatement* controlled by a *DebugCondition*.

A pure function can throw exceptions.

```
import std.stdio;
int x;
immutable int y;
const int* pz;

pure int foo(int i,
             char* p,
             const char* q,
             immutable int* s)
{
    debug writeln("in foo()"); // ok, impure code allowed in debug statement
    x = i;    // error, modifying global state
    i = x;    // error, reading mutable global state
    i = y;    // ok, reading immutable global state
```

```

    i = *pz; // error, reading const global state
    return i;
}

```

An implementation may assume that a **pure** function that (a) accepts only parameters without mutable indirections, and (b) returns a result without mutable indirections, will have the same effect for all invocation with equivalent arguments, and is allowed to memoize the result of the function under the assumption that equivalent parameters always produce equivalent results. Such functions are termed *strongly pure* functions in this document. Note that a strongly pure function may still have behavior inconsistent with memoization by e.g. using **casts** or by changing behavior depending on the address of its parameters. An implementation is currently not required to enforce validity of memoization in all cases.

A pure function that accepts only parameters without mutable indirections and returns a result that has mutable indirections is called a *pure factory function*. An implementation may assume that all mutable memory returned by the call is not referenced by any other part of the program, i.e. it is newly allocated by the function. Conversely, the mutable references of the result may be assumed to not refer to any object that existed before the function call. For example:

```

struct List { int payload; List* next; }
pure List* make(int a, int b)
{
    auto result = new List(a, null);
    result.next = new List(b, result);
    return result;
}

```

Here, an implementation may assume (without having knowledge of the body of **make**) that all references in **make**'s result refer to other **List** objects created by **make**, and that no other part of the program refers to any of these objects.

Any **pure** function that is not strongly pure cannot be assumed to be memoizable, and calls to it may not be elided even if it returns **void** (save for compiler optimizations that prove the function has no effect). Function calls may still be elided, or results be memoized, by means of traditional inlining and optimization techniques available for all functions.

If a strongly pure function throws an exception or an error, the assumptions related to memoization and references do not carry to the thrown exception.

Pure destructors do not benefit of special elision.

Nothrow Functions

Nothrow functions can only throw exceptions derived from class *Error*.

Nothrow functions are covariant with throwing ones.

Ref Functions

Ref functions allow functions to return by reference. This is analogous to ref function parameters.

```
ref int foo()
{
    auto p = new int;
    return *p;
}
...
foo() = 3;  // reference returns can be lvalues
```

Auto Functions

Auto functions have their return type inferred from any *ReturnStatements* in the function body.

An auto function is declared without a return type. If it does not already have a storage class, use the auto storage class.

If there are multiple *ReturnStatements*, the types of them must be implicitly convertible to a common type. If there are no *ReturnStatements*, the return type is inferred to be void.

```
auto foo(int x) { return x + 3; }  // inferred to be int
auto foo(int x) { return x; return 2.5; }  // inferred to be double
```

Auto Ref Functions

Auto ref functions infer their return type just as auto functions do. In addition, they become ref functions if all return expressions are lvalues, and it would not be a reference to a local or a parameter.

```
auto ref foo(int x)    { return x; }  // value return
auto ref foo()        { return 3; }  // value return
auto ref foo(ref int x) { return x; } // ref return
auto ref foo(out int x) { return x; } // ref return
auto ref foo() { static int x; return x; } // ref return
```

The ref-ness of a function is determined from all *ReturnStatements* in the function body:

```
auto ref foo(ref int x) { return 3; return x; }  // ok, value return
auto ref foo(ref int x) { return x; return 3; }  // ok, value return
auto ref foo(ref int x, ref double y)
{
    return x; return y;
    // The return type is deduced to double, but cast(double)x is not an lvalue,
```



```

    // then become a value return.
}

```

Auto ref function can have explicit return type.

```

auto ref int foo(ref int x) { return x; } // ok, ref return
auto ref int foo(double x) { return x; } // error, cannot convert double to int

```

Inout Functions

Functions that deal with mutable, const, or immutable types with equanimity often need to transmit their type to the return value:

```

int[] foo(int[] a, int x, int y) { return a[x .. y]; }

const(int)[] foo(const(int)[] a, int x, int y) { return a[x .. y]; }

immutable(int)[] foo(immutable(int)[] a, int x, int y) { return a[x .. y]; }

```

The code generated by these three functions is identical. To indicate that these can be one function, the inout type constructor is employed:

```

inout(int)[] foo(inout(int)[] a, int x, int y) { return a[x .. y]; }

```

The inout forms a wildcard that stands in for any of mutable, const, immutable, inout, or inout const. When the function is called, the inout of the return type is changed to whatever the mutable, const, immutable, inout, or inout const status of the argument type to the parameter inout was.

Inout types can be implicitly converted to const or inout const, but to nothing else. Other types cannot be implicitly converted to inout. Casting to or from inout is not allowed in @safe functions.

A set of arguments to a function with inout parameters is considered a match if any inout argument types match exactly, or:

1. No argument types are composed of inout types.
2. A mutable, const or immutable argument type can be matched against each corresponding parameter inout type.

If such a match occurs, the inout is considered the common qualifier of the matched qualifiers. If more than two parameters exist, the common qualifier calculation is recursively applied.

Common qualifier of the two type qualifiers

	<i>mutable</i>	<i>const</i>	<i>immutable</i>	<i>inout</i>	<i>inout const</i>
<i>mutable</i> (= m)	m	c	c	c	c
<i>const</i> (= c)	c	c	c	c	c
<i>immutable</i> (= i)	c	c	i	wc	wc
<i>inout</i> (= w)	c	c	wc	w	wc
<i>inout const</i> (= wc)	c	c	wc	wc	wc

The *inout* in the return type is then rewritten to be the *inout* matched qualifiers:

```
int[] ma;
const(int)[] ca;
immutable(int)[] ia;

inout(int)[] foo(inout(int)[] a) { return a; }
void test1()
{
    // inout matches to mutable, so inout(int)[] is
    // rewritten to int[]
    int[] x = foo(ma);

    // inout matches to const, so inout(int)[] is
    // rewritten to const(int)[]
    const(int)[] y = foo(ca);

    // inout matches to immutable, so inout(int)[] is
    // rewritten to immutable(int)[]
    immutable(int)[] z = foo(ia);
}

inout(const(int))[] bar(inout(int)[] a) { return a; }
void test2()
{
    // inout matches to mutable, so inout(const(int))[] is
    // rewritten to const(int)[]
    const(int)[] x = bar(ma);

    // inout matches to const, so inout(const(int))[] is
    // rewritten to const(int)[]
    const(int)[] y = bar(ca);
}
```

```

    // inout matches to immutable, so inout(int)[] is
    // rewritten to immutable(int)[]
    immutable(int)[] z = bar(ia);
}

```

Note: Shared types are not overlooked. Shared types cannot be matched with inout.

Optional Parentheses

If a function call passes no explicit argument, i.e. it would syntactically use (), then these parentheses may be omitted, similar to a getter invocation of a property function.

```

void foo() {}    // no arguments
void fun(int x = 10) { }
void bar(int[] arr) {} // for UFCS

void main()
{
    foo();        // OK
    foo;          // also OK
    fun;          // OK

    int[] arr;
    arr.bar();    // OK
    arr.bar;      // also OK
}

```

Optional parentheses are not applied to delegates or function pointers.

```

void main()
{
    int function() fp;

    assert(fp == 6);    // Error, incompatible types int function() and int
    assert(*fp == 6);   // Error, incompatible types int() and int

    int delegate() dg;
    assert(dg == 6);    // Error, incompatible types int delegate() and int
}

```

If a function returns a delegate or function pointer, the parentheses are required if the returned value is to be called.

```
struct S {
    int function() callfp() { return &numfp; }
    int delegate() calldg() { return &numdg; }
    int numdg() { return 6; }
}

int numfp() { return 6; }

void main()
{
    S s;
    int function() fp;

    fp = s.callfp;
    assert(fp() == 6);

    fp = s.callfp();
    assert(fp() == 6);

    int x = s.callfp();
    assert(x == 6);

    int delegate() dg;

    dg = s.calldg;
    assert(dg() == 6);

    dg = s.calldg();
    assert(dg() == 6);

    int y = s.calldg();
    assert(y == 6);
}
```

Property Functions

Properties are functions that can be syntactically treated as if they were fields or variables. Properties can be read from or written to. A property is read by calling a method or function with no arguments; a property is written by calling a method or function with its argument being the value it is set to.

Simple getter and setter properties can be written using UFCS. These can be enhanced with the addition of the `@property` attribute to the function, which adds the following behaviors:

- `@property` functions cannot be overloaded with non-`@property` functions with the same name.
- `@property` functions can only have zero, one or two parameters.
- `@property` functions cannot have variadic parameters.
- For the expression `typeof(exp)` where `exp` is an `@property` function, the type is the return type of the function, rather than the type of the function.
- For the expression `__traits(compiles, exp)` where `exp` is an `@property` function, a further check is made to see if the function can be called.
- `@property` are mangled differently, meaning that `@property` must be consistently used across different compilation units.
- The ObjectiveC interface recognizes `@property` setter functions as special and modifies them accordingly.

A simple property would be:

```
struct Foo
{
    @property int data() { return m_data; } // read property

    @property int data(int value) { return m_data = value; } // write property

private:
    int m_data;
}
```

To use it:

```
int test()
{
    Foo f;

    f.data = 3;           // same as f.data(3);
    return f.data + 3; // same as return f.data() + 3;
}
```

The absence of a read method means that the property is write-only. The absence of a write method means that the property is read-only. Multiple write methods can exist; the correct one is selected using the usual function overloading rules.

In all the other respects, these methods are like any other methods. They can be static, have different linkages, have their address taken, etc.

The built in properties `.sizeof`, `.alignof`, and `.mangleof` may not be declared as fields or methods in structs, unions, classes or enums.

If a property function has no parameters, it works as a getter. If has exactly one parameter, it works as a setter.

Virtual Functions

Virtual functions are functions that are called indirectly through a function pointer table, called a `vtbl`, rather than directly. All `public` and `protected` member functions which are non-static and are not templated are virtual unless the compiler can determine that they will never be overridden (e.g. they are marked with `final` and do not override any functions in a base class), in which case, it will make them non-virtual. This results in fewer bugs caused by not declaring a function virtual and then overriding it anyway.

Member functions which are `private` or `package` are never virtual, and hence cannot be overridden.

Functions with non-D linkage cannot be virtual and hence cannot be overridden.

Member template functions cannot be virtual and hence cannot be overridden.

Functions marked as `final` may not be overridden in a derived class, unless they are also `private`. For example:

```
class A
{
    int def() { ... }
    final int foo() { ... }
    final private int bar() { ... }
    private int abc() { ... }
}

class B : A
{
    override int def() { ... } // ok, overrides A.def
    override int foo() { ... } // error, A.foo is final
    int bar() { ... } // ok, A.bar is final private, but not virtual
    int abc() { ... } // ok, A.abc is not virtual, B.abc is virtual
}

void test(A a)
{
    a.def(); // calls B.def
    a.foo(); // calls A.foo
    a.bar(); // calls A.bar
```

```

    a.abc();    // calls A.abc
}

void func()
{
    B b = new B();
    test(b);
}

```

Covariant return types are supported, which means that the overriding function in a derived class can return a type that is derived from the type returned by the overridden function:

```

class A { }
class B : A { }

class Foo
{
    A test() { return null; }
}

class Bar : Foo
{
    override B test() { return null; } // overrides and is covariant with Foo.↵
    ↵ test()
}

```

Virtual functions all have a hidden parameter called the *this* reference, which refers to the class object for which the function is called.

To avoid dynamic binding on member function call, insert base class name before the member function name. For example:

```

class B
{
    int foo() { return 1; }
}
class C : B
{
    override int foo() { return 2; }

    void test()
    {
        assert(B.foo() == 1); // translated to this.B.foo(), and

```

```

        // calls B.foo statically.
    assert(C.foo() == 2); // calls C.foo statically, even if
                        // the actual instance of 'this' is D.
    }
}
class D : C
{
    override int foo() { return 3; }
}
void main()
{
    auto d = new D();
    assert(d.foo() == 3); // calls D.foo
    assert(d.B.foo() == 1); // calls B.foo
    assert(d.C.foo() == 2); // calls C.foo
    d.test();
}

```

Function Inheritance and Overriding

A function in a derived class with the same name and parameter types as a function in a base class overrides that function:

```

class A
{
    int foo(int x) { ... }
}

class B : A
{
    override int foo(int x) { ... }
}

void test()
{
    B b = new B();
    bar(b);
}

void bar(A a)

```



```
{
    a.foo(1);    // calls B.foo(int)
}
```

However, when doing overload resolution, the functions in the base class are not considered:

```
class A
{
    int foo(int x) { ... }
    int foo(long y) { ... }
}

class B : A
{
    override int foo(long x) { ... }
}

void test()
{
    B b = new B();
    b.foo(1);    // calls B.foo(long), since A.foo(int) not considered
    A a = b;

    a.foo(1);    // issues runtime error (instead of calling A.foo(int))
}
```

To consider the base class's functions in the overload resolution process, use an *AliasDeclaration*:

```
class A
{
    int foo(int x) { ... }
    int foo(long y) { ... }
}

class B : A
{
    alias foo = A.foo;
    override int foo(long x) { ... }
}

void test()
{
```

```

    B b = new B();
    bar(b);
}

void bar(A a)
{
    a.foo(1);          // calls A.foo(int)
    B b = new B();
    b.foo(1);          // calls A.foo(int)
}

```

If such an *AliasDeclaration* is not used, the derived class's functions completely override all the functions of the same name in the base class, even if the types of the parameters in the base class functions are different. If, through implicit conversions to the base class, those other functions do get called, a compile-time error will be given:

```

class A
{
    void set(long i) { }
    void set(int i)  { }
}
class B : A
{
    void set(long i) { }
}

void foo(A a)
{
    int i;
    a.set(3);    // error, use of A.set(int) is hidden by B
                // use 'alias set = A.set;' to introduce base class overload set.
    assert(i == 1);
}

void main()
{
    foo(new B);
}

```

If an error occurs during the compilation of your program, the use of overloads and overrides needs to be reexamined in the relevant classes.

The compiler will not give an error if the hidden function is disjoint, as far as overloading is concerned, from all the other virtual functions in the inheritance hierarchy.

A function parameter's default value is not inherited:

```
class A
{
    void foo(int x = 5) { ... }
}

class B : A
{
    void foo(int x = 7) { ... }
}

class C : B
{
    void foo(int x) { ... }
}

void test()
{
    A a = new A();
    a.foo();           // calls A.foo(5)

    B b = new B();
    b.foo();           // calls B.foo(7)

    C c = new C();
    c.foo();           // error, need an argument for C.foo
}
```

If a derived class overrides a base class member function with different *FunctionAttributes*, the missing attributes will be automatically compensated by the compiler.

```
class B
{
    void foo() pure nothrow @safe {}
}

class D : B
{
    override void foo() {}
}
```

```
void main()
{
    auto d = new D();
    pragma(msg, typeof(&d.foo));
    // prints "void delegate() pure nothrow @safe" in compile time
}
```

Inline Functions

The compiler makes the decision whether to inline a function or not. This decision may be controlled by `pragma(inline)`, assuming that the compiler implements it, which is not mandatory.

Note that any *FunctionLiteral* should be inlined when used in its declaration scope.

Function Overloading

Functions are overloaded based on how well the arguments to a function can match up with the parameters. The function with the *best* match is selected. The levels of matching are:

1. no match
2. match with implicit conversions
3. match with conversion to `const`
4. exact match

Each argument (including any `this` pointer) is compared against the function's corresponding parameter, to determine the match level for that argument. The match level for a function is the *worst* match level of each of its arguments.

Literals do not match `ref` or `out` parameters.

If two or more functions have the same match level, then *partial ordering* is used to try to find the best match. Partial ordering finds the most specialized function. If neither function is more specialized than the other, then it is an ambiguity error. Partial ordering is determined for functions `f()` and `g()` by taking the parameter types of `f()`, constructing a list of arguments by taking the default values of those types, and attempting to match them against `g()`. If it succeeds, then `g()` is at least as specialized as `f()`. For example:

```
class A { }
class B : A { }
class C : B { }
void foo(A);
void foo(B);

void test()
```

```

{
    C c;
    /* Both foo(A) and foo(B) match with implicit conversion rules.
     * Applying partial ordering rules,
     * foo(B) cannot be called with an A, and foo(A) can be called
     * with a B. Therefore, foo(B) is more specialized, and is selected.
     */
    foo(c); // calls foo(B)
}

```

A function with a variadic argument is considered less specialized than a function without.

Functions defined with non-D linkage cannot be overloaded. This is because the name mangling might not take the parameter types into account.

Overload Sets

Functions declared at the same scope overload against each other, and are called an *Overload Set*. A typical example of an overload set are functions defined at module level:

```

module A;
void foo() { }
void foo(long i) { }

```

A.foo() and A.foo(long) form an overload set. A different module can also define functions with the same name:

```

module B;
class C { }
void foo(C) { }
void foo(int i) { }

```

and A and B can be imported by a third module, C. Both overload sets, the A.foo overload set and the B.foo overload set, are found. An instance of foo is selected based on it matching in exactly one overload set:

```

import A;
import B;

void bar(C c)
{
    foo();    // calls A.foo()
    foo(1L);  // calls A.foo(long)
    foo(c);   // calls B.foo(C)
    foo(1,2); // error, does not match any foo
}

```

```

    foo(1);    // error, matches A.foo(long) and B.foo(int)
    A.foo(1); // calls A.foo(long)
}

```

Even though `B.foo(int)` is a better match than `A.foo(long)` for `foo(1)`, it is an error because the two matches are in different overload sets.

Overload sets can be merged with an alias declaration:

```

import A;
import B;

alias foo = A.foo;
alias foo = B.foo;

void bar(C c)
{
    foo();    // calls A.foo()
    foo(1L);  // calls A.foo(long)
    foo(c);   // calls B.foo(C)
    foo(1,2); // error, does not match any foo
    foo(1);   // calls B.foo(int)
    A.foo(1); // calls A.foo(long)
}

```

Function Parameters

Parameter storage classes are `in`, `out`, `ref`, `lazy`, `const`, `immutable`, `shared`, `inout` or `scope`. For example:

```
int foo(in int x, out int y, ref int z, int q);
```

`x` is `in`, `y` is `out`, `z` is `ref`, and `q` is none.

- The function declaration makes it clear what the inputs and outputs to the function are.
- It eliminates the need for IDL (interface description language) as a separate language.
- It provides more information to the compiler, enabling more error checking and possibly better code generation.

Parameter Storage Classes

Storage Class	Description
<i>none</i>	parameter becomes a mutable copy of its argument
in	equivalent to const scope
out	parameter is initialized upon function entry with the default value for its type
ref	parameter is passed by reference
scope	references in the parameter cannot be escaped (e.g. assigned to a global variable). Ignored for parameters with no references
lazy	argument is evaluated by the called function and not by the caller
const	argument is implicitly converted to a const type
immutable	argument is implicitly converted to an immutable type
shared	argument is implicitly converted to a shared type
inout	argument is implicitly converted to an inout type

```
void foo(out int x)
{
    // x is set to int.init,
    // which is 0, at start of foo()
}
```

```
int a = 3;
foo(a);
// a is now 0
```

```
void abc(out int x)
{
    x = 2;
}
```

```
int y = 3;
abc(y);
// y is now 2
```

```
void def(ref int x)
{
    x += 1;
}
```

```
int z = 3;
```

```
def(z);
// z is now 4
```

For dynamic array and object parameters, which are passed by reference, in/out/ref apply only to the reference and not the contents.

lazy arguments are evaluated not when the function is called, but when the parameter is evaluated within the function. Hence, a **lazy** argument can be executed 0 or more times. A **lazy** parameter cannot be an lvalue.

```
void dotimes(int n, lazy void exp)
{
    while (n--)
        exp();
}

void test()
{
    int x;
    dotimes(3, writeln(x++));
}
```

prints to the console:

```
0 1 2
```

A **lazy** parameter of type **void** can accept an argument of any type.

Function Default Arguments

Function parameter declarations can have default values:

```
void foo(int x, int y = 3)
{
    ...
}
...
foo(4);    // same as foo(4, 3);
```

Default parameters are evaluated in the context of the function declaration. If the default value for a parameter is given, all following parameters must also have default values.

Return Ref Parameters

Note: The **return** attribute is currently only enforced by **dmd** when the **-dip25** switch is passed.

Return ref parameters are used with ref functions to ensure that the returned reference will not outlive the matching argument's lifetime.


```

ref int identity(return ref int x) {
    return x; // pass-through function that does nothing
}

ref int fun() {
    int x;
    return identity(x); // Error: escaping reference to local variable x
}

ref int gun(return ref int x) {
    return identity(x); // OK
}

```

Ref methods marked with the `return` attribute ensure the returned reference will not outlive the respective aggregate instance.

```

struct S
{
    private int x;
    ref int get() return { return x; }
}

ref int escape()
{
    S s;
    return s.get(); // Error: escaping reference to local variable s
}

```

Template functions and lambdas can deduce the `return` attribute.

Variadic Functions

Functions taking a variable number of arguments are called variadic functions. A variadic function can take one of three forms:

1. C-style variadic functions
2. Variadic functions with type info
3. Typesafe variadic functions

C-style Variadic Functions

A C-style variadic function is declared as taking a parameter of `...` after the required function parameters. It has non-D linkage, such as `extern (C)`:

```
extern (C) void foo(int x, int y, ...);

foo(3, 4);           // ok
foo(3, 4, 6.8);     // ok, one variadic argument
foo(2);             // error, y is a required argument
```

There must be at least one non-variadic parameter declared.

```
extern (C) int def(...); // error, must have at least one parameter
```

C-style variadic functions match the C calling convention for variadic functions, and is most useful for calling C library functions like `printf`.

C-style variadic functions cannot be marked as `@safe`.

Access to variadic arguments is done using the standard library module `core.stdc.stdarg`.

```
import core.stdc.stdarg;

void test()
{
    foo(3, 4, 5);    // first variadic argument is 5
}

void foo(int x, int y, ...)
{
    va_list args;

    va_start(args, y); // y is the last named parameter

    int z;
    va_arg(args, z);   // z is set to 5
}
```

D-style Variadic Functions

Variadic functions with argument and type info are declared as taking a parameter of `...` after the required function parameters. It has D linkage, and need not have any non-variadic parameters declared:

```
int abc(char c, ...); // one required parameter: c
int def(...);         // ok
```

To access them, the following import is required:

```
import core.vararg;
```

These variadic functions have a special local variable declared for them, `_argptr`, which is a `core.vararg` reference to the first of the variadic arguments. To access the arguments, `_argptr` must be used in conjunction with `va_arg`:

```
import core.vararg;

void test()
{
    foo(3, 4, 5);    // first variadic argument is 5
}

void foo(int x, int y, ...)
{
    int z;

    z = va_arg!int(_argptr); // z is set to 5
}
```

An additional hidden argument with the name `_arguments` and type `TypeInfo[]` is passed to the function. `_arguments` gives the number of arguments and the type of each, enabling type safety to be checked at run time.

```
import std.stdio;
import core.vararg;

class Foo { int x = 3; }
class Bar { long y = 4; }

void printargs(int x, ...)
{
    writeln("%d_arguments", _arguments.length);
    for (int i = 0; i < _arguments.length; i++)
    {
        writeln(_arguments[i]);

        if (_arguments[i] == typeid(int))
        {
            int j = va_arg!(int)(_argptr);
            writeln("\t%d", j);
        }
        else if (_arguments[i] == typeid(long))
        {

```

```

        long j = va_arg!(long)(_argptr);
        writefln("\t%d", j);
    }
    else if (_arguments[i] == typeid(double))
    {
        double d = va_arg!(double)(_argptr);
        writefln("\t%g", d);
    }
    else if (_arguments[i] == typeid(Foo))
    {
        Foo f = va_arg!(Foo)(_argptr);
        writefln("\t%s", f);
    }
    else if (_arguments[i] == typeid(Bar))
    {
        Bar b = va_arg!(Bar)(_argptr);
        writefln("\t%s", b);
    }
    else
        assert(0);
}

void main()
{
    Foo f = new Foo();
    Bar b = new Bar();

    writefln("%s", f);
    printargs(1, 2, 3L, 4.5, f, b);
}

```

which prints:

```

0x00870FE0
5 arguments
int          2
long
double      3

```

```

        4.5
Foo      0x00870FE0
Bar      0x00870FD0

```

D-style variadic functions cannot be marked as `@safe`.

Typesafe Variadic Functions

Typesafe variadic functions are used when the variable argument portion of the arguments are used to construct an array or class object.

For arrays:

```

int test()
{
    return sum(1, 2, 3) + sum(); // returns 6+0
}

int func()
{
    int[3] ii = [4, 5, 6];
    return sum(ii);              // returns 15
}

int sum(int[] ar ...)
{
    int s;
    foreach (int x; ar)
        s += x;
    return s;
}

```

For static arrays:

```

int test()
{
    return sum(2, 3); // error, need 3 values for array
    return sum(1, 2, 3); // returns 6
}

int func()

```

```

{
    int[3] ii = [4, 5, 6];
    int[] jj = ii;
    return sum(ii); // returns 15
    return sum(jj); // error, type mismatch
}

int sum(int[3] ar ...)
{
    int s;
    foreach (int x; ar)
        s += x;
    return s;
}

```

For class objects:

```

class Foo
{
    int x;
    string s;

    this(int x, string s)
    {
        this.x = x;
        this.s = s;
    }
}

void test(int x, Foo f ...);

...

Foo g = new Foo(3, "abc");
test(1, g);           // ok, since g is an instance of Foo
test(1, 4, "def");    // ok
test(1, 5);           // error, no matching constructor for Foo

```

An implementation may construct the object or array instance on the stack. Therefore, it is an error to refer to that instance after the variadic function has returned:

```

Foo test(Foo f ...)

```

```

{
    return f;    // error, f instance contents invalid after return
}

int[] test(int[] a ...)
{
    return a;        // error, array contents invalid after return
    return a[0..1];  // error, array contents invalid after return
    return a.dup;    // ok, since copy is made
}

```

For other types, the argument is built with itself, as in:

```

int test(int i ...)
{
    return i;
}

...
test(3);    // returns 3
test(3, 4); // error, too many arguments
int[] x;
test(x);    // error, type mismatch

```

Lazy Variadic Functions

If the variadic parameter is an array of delegates with no parameters:

```
void foo(int delegate()[] dgs ...);
```

Then each of the arguments whose type does not match that of the delegate is converted to a delegate.

```
int delegate() dg;
foo(1, 3+x, dg, cast(int delegate())null);
```

is the same as:

```
foo( { return 1; }, { return 3+x; }, dg, null );
```

Local Variables

It is an error to use a local variable without first assigning it a value. The implementation may not always be able to detect these cases. Other language compilers sometimes issue a warning for this, but since it is always a bug, it should be an error.

It is an error to declare a local variable that hides another local variable in the same function:

```
void func(int x)
{
    int x;          // error, hides previous definition of x
    double y;
    ...
    {
        char y;    // error, hides previous definition of y
        int z;
    }
    {
        wchar z;  // legal, previous z is out of scope
    }
}
```

While this might look unreasonable, in practice whenever this is done it either is a bug or at least looks like a bug.

It is an error to return the address of or a reference to a local variable.

It is an error to have a local variable and a label with the same name.

Local Static Variables

Local variables in functions can be declared as static or `__gshared` in which case they are statically allocated rather than being allocated on the stack. As such, their value persists beyond the exit of the function.

```
void foo()
{
    static int n;
    if (++n == 100)
        writeln("called_100_times");
}
```

The initializer for a static variable must be evaluatable at compile time, and they are initialized upon the start of the thread (or the start of the program for `__gshared`). There are no static constructors or static destructors for static local variables.

Although static variable name visibility follows the usual scoping rules, the names of them must be unique within a particular function.

```
void main()
{
    { static int x; }
```



```

    { static int x; } // error
    { int i; }
    { int i; } // ok
}

```

Nested Functions

Functions may be nested within other functions:

```

int bar(int a)
{
    int foo(int b)
    {
        int abc() { return 1; }

        return b + abc();
    }
    return foo(a);
}

void test()
{
    int i = bar(3); // i is assigned 4
}

```

Nested functions can be accessed only if the name is in scope.

```

void foo()
{
    void A()
    {
        B(); // error, B() is forward referenced
        C(); // error, C undefined
    }
    void B()
    {
        A(); // ok, in scope
        void C()
        {
            void D()
            {

```

```

        A();      // ok
        B();      // ok
        C();      // ok
        D();      // ok
    }
}
A(); // ok
B(); // ok
C(); // error, C undefined
}

```

and:

```

int bar(int a)
{
    int foo(int b) { return b + 1; }
    int abc(int b) { return foo(b); } // ok
    return foo(a);
}

void test()
{
    int i = bar(3); // ok
    int j = bar.foo(3); // error, bar.foo not visible
}

```

Nested functions have access to the variables and other symbols defined by the lexically enclosing function. This access includes both the ability to read and write them.

```

int bar(int a)
{
    int c = 3;

    int foo(int b)
    {
        b += c; // 4 is added to b
        c++;    // bar.c is now 5
        return b + c; // 12 is returned
    }
    c = 4;
    int i = foo(a); // i is set to 12
}

```

```

    return i + c;    // returns 17
}

void test()
{
    int i = bar(3); // i is assigned 17
}

```

This access can span multiple nesting levels:

```

int bar(int a)
{
    int c = 3;

    int foo(int b)
    {
        int abc()
        {
            return c;    // access bar.c
        }
        return b + c + abc();
    }
    return foo(3);
}

```

Static nested functions cannot access any stack variables of any lexically enclosing function, but can access static variables. This is analogous to how static member functions behave.

```

int bar(int a)
{
    int c;
    static int d;

    static int foo(int b)
    {
        b = d;           // ok
        b = c;           // error, foo() cannot access frame of bar()
        return b + 1;
    }
    return foo(a);
}

```

Functions can be nested within member functions:

```

struct Foo
{
    int a;

    int bar()
    {
        int c;

        int foo()
        {
            return c + a;
        }
        return 0;
    }
}

```

Nested functions always have the D function linkage type.

Unlike module level declarations, declarations within function scope are processed in order. This means that two nested functions cannot mutually call each other:

```

void test()
{
    void foo() { bar(); } // error, bar not defined
    void bar() { foo(); } // ok
}

```

There are several workarounds for this limitation:

- Declare the functions to be static members of a nested struct:

```

void test()
{
    static struct S
    {
        static void foo() { bar(); } // ok
        static void bar() { foo(); } // ok
    }

    S.foo(); // compiles (but note the infinite runtime loop)
}

```

- Declare one or more of the functions to be function templates even if they take no specific template arguments:

```

void test()
{
    void foo()() { bar(); } // ok (foo is a function template)
    void bar()   { foo(); } // ok
}

```

- Declare the functions inside of a mixin template:

```

mixin template T()
{
    void foo() { bar(); } // ok
    void bar() { foo(); } // ok
}

```

```

void test()
{
    mixin T!();
}

```

- Use a delegate:

```

void test()
{
    void delegate() fp;
    void foo() { fp(); }
    void bar() { foo(); }
    fp = &bar;
}

```

Nested functions cannot be overloaded.

Delegates, Function Pointers, and Closures

A function pointer can point to a static nested function:

```

int function() fp;

void test()
{
    static int a = 7;
    static int foo() { return a + 3; }

    fp = &foo;
}

```

```

}

void bar()
{
    test();
    int i = fp();      // i is set to 10
}

```

Note: Two functions with identical bodies, or two functions that compile to identical assembly code, are not guaranteed to have distinct function pointer values. The compiler is free to merge functions bodies into one if they compile to identical code.

```

int abc(int x) { return x + 1; }
int def(int y) { return y + 1; }

int function() fp1 = &abc;
int function() fp2 = &def;
// Do not rely on fp1 and fp2 being different values; the compiler may merge
// them.

```

A delegate can be set to a non-static nested function:

```

int delegate() dg;

void test()
{
    int a = 7;
    int foo() { return a + 3; }

    dg = &foo;
    int i = dg(); // i is set to 10
}

```

The stack variables referenced by a nested function are still valid even after the function exits (this is different from D 1.0). This is called a *closure*. Returning addresses of stack variables, however, is not a closure and is an error.

```

int* bar()
{
    int b;
    test();
    int i = dg(); // ok, test.a is in a closure and still exists
    return &b;    // error, bar.b not valid after bar() exits
}

```

Delegates to non-static nested functions contain two pieces of data: the pointer to the stack frame of the lexically enclosing function (called the *frame pointer*) and the address of the function. This is analogous to struct/class non-static member function delegates consisting of a *this* pointer and the address of the member function. Both forms of delegates are interchangeable, and are actually the same type:

```
struct Foo
{
    int a = 7;
    int bar() { return a; }
}

int foo(int delegate() dg)
{
    return dg() + 1;
}

void test()
{
    int x = 27;
    int abc() { return x; }
    Foo f;
    int i;

    i = foo(&abc);    // i is set to 28
    i = foo(&f.bar);  // i is set to 8
}
```

This combining of the environment and the function is called a *dynamic closure*.

The `.ptr` property of a delegate will return the *frame pointer* value as a `void*`.

The `.funcptr` property of a delegate will return the *function pointer* value as a function type.

Future directions: Function pointers and delegates may merge into a common syntax and be interchangeable with each other.

Anonymous Functions and Anonymous Delegates

See *FunctionLiterals*.

main() Function

For console programs, `main()` serves as the entry point. It gets called after all the module initializers are run, and after any unittests are run. After it returns, all the module destructors are run. `main` ↯

↳ `()` must be declared using one of the following forms:

```
void main() { ... }
void main(string[] args) { ... }
int main() { ... }
int main(string[] args) { ... }
```

Function Templates

Template functions are useful for avoiding code duplication - instead of writing several copies of a function, each with a different parameter type, a single function template can be sufficient. For example:

```
// Only one copy of func needs to be written
void func(T)(T x)
{
    writeln(x);
}

void main()
{
    func!(int)(1); // pass an int
    func(1);       // pass an int, inferring T = int
    func("x");     // pass a string
    func(1.0);     // pass a float

    struct S {}
    S s;
    func(s);       // pass a struct
}
```

`func` takes a template parameter `T` and a runtime parameter, `x`. `T` is a placeholder identifier that can accept any type. In this case `T` can be inferred from the runtime argument type.

Note: Using the name `T` is just a convention. The name `TypeOfX` could have been used instead. For more information, see function templates.

Compile Time Function Execution (CTFE)

Functions which are both portable and free of side-effects can be executed at compile time. This is useful when constant folding algorithms need to include recursion and looping. Compile time function execution is subject to the following restrictions:

1. The function source code must be available to the compiler. Functions which exist in the source code only as extern declarations cannot be executed at compile time.
2. Executed expressions may not reference any global or local static variables.
3. asm statements are not permitted
4. Non-portable casts (eg, from `int[]` to `float[]`), including casts which depend on endianness, are not permitted. Casts between signed and unsigned types are permitted

Pointers are permitted in CTFE, provided they are used safely:

- C-style semantics on pointer arithmetic are strictly enforced. Pointer arithmetic is permitted only on pointers which point to static or dynamic array elements. Such pointers must point to an element of the array, or to the first element past the array. Pointer arithmetic is completely forbidden on pointers which are null, or which point to a non-array.
- The memory location of different memory blocks is not defined. Ordered comparison (`<`, `<=`, `>`, `>=`) between two pointers is permitted when both pointers point to the same array, or when at least one pointer is `null`.
- Pointer comparisons between independent memory blocks will generate a compile-time error, unless two such comparisons are combined using `&&` or `||` to yield a result which is independent of the ordering of memory blocks. Each comparison must consist of two pointer expressions compared with `<`, `<=`, `>`, or `>=`, and may optionally be negated with `!`.
For example, the expression `(p1 > q1 && p2 <= q2)` is permitted when `p1`, `p2` are expressions yielding pointers to memory block *P*, and `q1`, `q2` are expressions yielding pointers to memory block *Q*, even when *P* and *Q* are unrelated memory blocks. It returns true if `[p1..p2]` lies inside `[q1..q2]`, and false otherwise. Similarly, the expression `(p1 < q1 | p2 > q2)` is true if `[p1..p2]` lies outside `[q1..q2]`, and false otherwise.
- Equality comparisons (`==`, `!=`, `is`, `!is`) are permitted between all pointers, without restriction.
- Any pointer may be cast to `void *` and from `void *` back to its original type. Casting between pointer and non-pointer types is prohibited.

Note that the above restrictions apply only to expressions which are actually executed. For example:

```
static int y = 0;

int countTen(int x)
{
    if (x > 10)
        ++y;
    return x;
}
```

```
static assert(countTen(6) == 6); // OK
static assert(countTen(12) == 12); // invalid, modifies y.
```

The `__ctfe` boolean pseudo-variable, which evaluates to true at compile time, but false at run time, can be used to provide an alternative execution path to avoid operations which are forbidden at compile time. Every usage of `__ctfe` is evaluated before code generation and therefore has no run-time cost, even if no optimizer is used.

In order to be executed at compile time, the function must appear in a context where it must be so executed, for example:

- initialization of a static variable
- dimension of a static array
- argument for a template value parameter

```
template eval( A... )
{
    const typeof(A[0]) eval = A[0];
}

int square(int i)
{
    return i * i;
}

void foo()
{
    static j = square(3);    // compile time
    writeln(j);
    writeln(square(4));     // run time
    writeln(eval!(square(5))); // compile time
}
```

Executing functions at compile time can take considerably longer than executing it at run time. If the function goes into an infinite loop, it will hang at compile time (rather than hanging at run time).

Non-recoverable errors (such as assert failures) do not throw exceptions; instead, they end interpretation immediately.

Functions executed at compile time can give different results from run time in the following scenarios:

- floating point computations may be done at a higher precision than run time
- dependency on implementation defined order of evaluation
- use of uninitialized variables

These are the same kinds of scenarios where different optimization settings affect the results.

String Mixins and Compile Time Function Execution

Any functions that execute at compile time must also be executable at run time. The compile time evaluation of a function does the equivalent of running the function at run time. This means that the semantics of a function cannot depend on compile time values of the function. For example:

```
int foo(char[] s)
{
    return mixin(s);
}

const int x = foo("1");
```

is illegal, because the runtime code for `foo()` cannot be generated. A function template would be the appropriate method to implement this sort of thing.

Function Safety

Safe functions are functions that are statically checked to exhibit no possibility of *undefined behavior*. Undefined behavior is often used as a vector for malicious attacks.

Safe Functions

Safe functions are marked with the `@safe` attribute.

The following operations are not allowed in safe functions:

- No casting from a pointer type to any type other than `void*`.
- No casting from any non-pointer type to a pointer type.
- No modification of pointer values.
- Cannot access unions that have pointers or references overlapping with other types.
- Calling any system functions.
- No catching of exceptions that are not derived from `class Exception`.
- No inline assembler.
- No explicit casting of mutable objects to immutable.
- No explicit casting of immutable objects to mutable.

- No explicit casting of thread local objects to shared.
- No explicit casting of shared objects to thread local.
- No taking the address of a local variable or function parameter.
- Cannot access `__gshared` variables.

Functions nested inside safe functions default to being safe functions.

Safe functions are covariant with trusted or system functions.

Note: The verifiable safety of functions may be compromised by bugs in the compiler and specification. Please report all such errors so they can be corrected.

Trusted Functions

Trusted functions are marked with the `@trusted` attribute.

Trusted functions are guaranteed by the programmer to not exhibit any undefined behavior if called by a safe function. Generally, trusted functions should be kept small so that they are easier to manually verify.

Trusted functions may call safe, trusted, or system functions.

Trusted functions are covariant with safe or system functions.

System Functions

System functions are functions not marked with `@safe` or `@trusted` and are not nested inside `@safe` functions. System functions may be marked with the `@system` attribute. A function being system does not mean it actually is unsafe, it just means that the compiler is unable to verify that it cannot exhibit undefined behavior.

System functions are **not** covariant with trusted or safe functions.

Function Attribute Inference

FunctionLiterals and function templates, since their function bodies are always present, infer the `pure`, `nothrow`, `@safe`, and `@nogc` attributes unless specifically overridden.

Attribute inference is not done for other functions, even if the function body is present.

The inference is done by determining if the function body follows the rules of the particular attribute.

Cyclic functions (i.e. functions that wind up directly or indirectly calling themselves) are inferred as being impure, throwing, and `@system`.

If a function attempts to test itself for those attributes, then the function is inferred as not having those attributes.

Uniform Function Call Syntax (UFCS)

A free function can be called with a syntax that looks as if the function were a member function of its first parameter type.

```
void func(X thisObj);

X obj;
obj.func();
// If 'obj' does not have regular member 'func',
// it's automatically rewritten to 'func(obj)'
```

This provides a way to add functions to a class externally as if they were public final member functions, which enables function chaining and component programming.

```
stdin.byLine(KeepTerminator.yes)
    .map!(a => a.idup)
    .array
    .sort
    .copy(stdout.lockingTextWriter());
```

It also works with `@property` functions:

```
@property prop(X thisObj);
@property prop(X thisObj, int value);

X obj;
obj.prop;      // Rewrites to: prop(obj);
obj.prop = 1;  // Rewrites to: prop(obj, 1);
```

Syntactically parenthesis-less check for `@property` functions is done at the same time as UFCS rewrite.

When UFCS rewrite is necessary, compiler searches the name on accessible module level scope, in order from the innermost scope.

```
module a;
void foo(X);
alias boo = foo;
void main()
{
    void bar(X);
    import b : baz; // void baz(X);

    X obj;
```

```

    obj.foo();    // OK, calls a.foo;
    //obj.bar();  // NG, UFCS does not see nested functions
    obj.baz();    // OK, calls b.baz, because it is declared at the
                  // top level scope of module b

import b : boo = baz;
obj.boo();       // OK, calls aliased b.baz instead of a.boo (== a.foo),
                  // because the declared alias name 'boo' in local scope
                  // overrides module scope name
}
class C
{
    void mfoo(X);
    static void sbar(X);
    import b : ibaz = baz;  // void baz(X);
    void test()
    {
        X obj;
        //obj.mfoo();  // NG, UFCS does not see member functions
        //obj.sbar();  // NG, UFCS does not see static member functions
        obj.ibaz();     // OK, ibaz is an alias of baz which declared at
                        // the top level scope of module b
    }
}

```

The reason why local symbols are not considered by UFCS, is to avoid unexpected name conflicts. See below problematic examples.

```

int front(int[] arr) { return arr[0]; }

void main()
{
    int[] a = [1,2,3];
    auto x = a.front();  // call .front by UFCS

    auto front = x;      // front is now a variable
    auto y = a.front();  // Error, front is not a function
}

class C
{

```

```
int[] arr;  
int front()  
{  
    return arr.front(); // Error, C.front is not callable  
                        // using argument types (int[])  
}  
}
```


Chapter 20

Operator Overloading

Operator overloading is accomplished by rewriting operators whose operands are class or struct objects into calls to specially named member functions. No additional syntax is used.

- Unary Operator Overloading
- Cast Operator Overloading
- Binary Operator Overloading
- Overloading the Comparison Operators
 - Overloading == and !=
 - Overloading <, <=, >, and >=
- Function Call Operator Overloading
- Assignment Operator Overloading
- Op Assignment Operator Overloading
- Array Indexing and Slicing Operators Overloading
 - Index Operator Overloading
 - Slice Operator Overloading
 - Dollar Operator Overloading
- Forwarding
- D1 style operator overloading

20.1 Unary Operator Overloading

Overloadable Unary Operators

<i>op</i>	<i>rewrite</i>
$-e$	<code>e.opUnary!("-")()</code>
$+e$	<code>e.opUnary!("+")()</code>
$\sim e$	<code>e.opUnary!("~")()</code>
$*e$	<code>e.opUnary!("*")()</code>
$++e$	<code>e.opUnary!("++")()</code>
$--e$	<code>e.opUnary!("--")()</code>

For example, in order to overload the $-$ (negation) operator for struct `S`, and no other operator:

```
struct S
{
    int m;

    int opUnary(string s)() if (s == "-")
    {
        return -m;
    }
}

int foo(S s)
{
    return -s;
}
```

Postincrement $e++$ and Postdecrement $e--$ Operators

These are not directly overloadable, but instead are rewritten in terms of the $++e$ and $-e$ prefix operators:

Postfix Operator Rewrites

<i>op</i>	<i>rewrite</i>
$e--$	<code>(auto t = e, --e, t)</code>
$e++$	<code>(auto t = e, ++e, t)</code>

Overloading Index Unary Operators

Overloadable Index Unary Operators	
<i>op</i>	<i>rewrite</i>
$-a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("-")(a.\textit{opSlice}(i, j))$
$+a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("+")(a.\textit{opSlice}(i, j))$
$\sim a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("~")(a.\textit{opSlice}(i, j))$
$*a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("*")(a.\textit{opSlice}(i, j))$
$++a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("++")(a.\textit{opSlice}(i, j))$
$--a[\textit{b}<\textit{sub}>1</\textit{sub}>, \textit{b}<\textit{sub}>2</\textit{sub}>, \dots \textit{b}<\textit{sub}>n</\textit{sub}>]$	$a.\textit{opIndexUnary}!("--")(a.\textit{opSlice}(i, j))$

Overloading Slice Unary Operators

Overloadable Slice Unary Operators	
<i>op</i>	<i>rewrite</i>
$-a[i..j]$	$a.\textit{opIndexUnary}!("-")(a.\textit{opSlice}(i, j))$
$+a[i..j]$	$a.\textit{opIndexUnary}!("+")(a.\textit{opSlice}(i, j))$
$\sim a[i..j]$	$a.\textit{opIndexUnary}!("~")(a.\textit{opSlice}(i, j))$
$*a[i..j]$	$a.\textit{opIndexUnary}!("*")(a.\textit{opSlice}(i, j))$
$++a[i..j]$	$a.\textit{opIndexUnary}!("++")(a.\textit{opSlice}(i, j))$
$--a[i..j]$	$a.\textit{opIndexUnary}!("--")(a.\textit{opSlice}(i, j))$
$-a[]$	$a.\textit{opIndexUnary}!("-")()$
$+a[]$	$a.\textit{opIndexUnary}!("+")()$
$\sim a[]$	$a.\textit{opIndexUnary}!("~")()$
$*a[]$	$a.\textit{opIndexUnary}!("*")()$
$++a[]$	$a.\textit{opIndexUnary}!("++")()$
$--a[]$	$a.\textit{opIndexUnary}!("--")()$

For backward compatibility, if the above rewrites fail and `opSliceUnary` is defined, then the rewrites `a.opSliceUnary!(op)(a, i, j)` and `a.opSliceUnary!(op)` are tried instead, respectively.

20.2 Cast Operator Overloading

To define how one type can be cast to another, define the `opCast` template method, which is used as follows:

Cast Operators	
<i>op</i>	<i>rewrite</i>
<code>cast(<i>type</i>)</code>	<code>\textit{e}.opCast!(<math>\textit{type}</math>)()</code>

Note that `opCast` is only ever used with an explicit `cast` expression, except in the case of boolean operations (see next section)

Boolean Operations

Notably absent from the list of overloaded unary operators is the `!` logical negation operator. More obscurely absent is a unary operator to convert to a `bool` result. Instead, these are covered by a rewrite to:

`opCast!(bool)(e)`

So,

```
if (e)    =>  if (e.opCast!(bool))
if (!e)   =>  if (!e.opCast!(bool))
```

etc., whenever a `bool` result is expected. This only happens, however, for instances of structs. Class references are converted to `bool` by checking to see if the class reference is null or not.

20.3 Binary Operator Overloading

The following binary operators are overloadable:

Overloadable Binary Operators							
<code>+</code>	<code>-</code>	<code>*</code>	<code>/</code>	<code>%</code>	<code>^^</code>	<code>&</code>	
<code> </code>	<code>^</code>	<code><<</code>	<code>>></code>	<code>>>></code>	<code>~</code>	<code>in</code>	

The expression:

`a op b`

is rewritten as both:

```
a.opBinary!(\textquotedblleft{<math>op</math>\textquotedblright{}})(b)
b.opBinaryRight!(\textquotedblleft{<math>op</math>\textquotedblright{}})(a)
```

and the one with the ‘better’ match is selected. It is an error for both to equally match.

Operator overloading for a number of operators can be done at the same time. For example, if only the `+` or `-` operators are supported:

```
T opBinary(string op)(T rhs)
{
    static if (op == "+") return data + rhs.data;
    else static if (op == "-") return data - rhs.data;
    else static assert(0, "Operator_~op~_not_implemented");
}
```

To do them all en masse:

```
T opBinary(string op)(T rhs)
{
    return mixin("data_~op~_rhs.data");
}
```

20.4 Overloading the Comparison Operators

D allows overloading of the comparison operators `==`, `!=`, `<`, `<=`, `>=`, `>` via two functions, `opEquals` and `opCmp`.

The equality and inequality operators are treated separately because while practically all user-defined types can be compared for equality, only a subset of types have a meaningful ordering. For example, while it makes sense to determine if two RGB color vectors are equal, it is not meaningful to say that one color is greater than another, because colors do not have an ordering. Thus, one would define `opEquals` for a `Color` type, but not `opCmp`.

Furthermore, even with orderable types, the order relation may not be linear. For example, one may define an ordering on sets via the subset relation, such that `x < y` is true if `x` is a (strict) subset of `y`. If `x` and `y` are disjoint sets, then neither `x < y` nor `y < x` holds, but that does not imply that `x == y`. Thus, it is insufficient to determine equality purely based on `opCmp` alone. For this reason, `opCmp` is only used for the inequality operators `<`, `<=`, `>=`, and `>`. The equality operators `==` and `!=` always employ `opEquals` instead.

Therefore, it is the programmer’s responsibility to ensure that `opCmp` and `opEquals` are consistent with each other. If `opEquals` is not specified, the compiler provides a default version that does member-wise comparison. If this suffices, one may define only `opCmp` to customize the behaviour of the inequality operators. But if not, then a custom version of `opEquals` should be defined as well, in order to preserve consistent semantics between the two kinds of comparison operators.

Finally, if the user-defined type is to be used as a key in the built-in associative arrays, then the programmer must ensure that the semantics of `opEquals` and `toHash` are consistent. If not, the associative array may not work in the expected manner.

Overloading `==` and `!=`

Expressions of the form `a != b` are rewritten as `!(a == b)`.

Given `a == b` :

1. If `a` and `b` are both class objects, then the expression is rewritten as:

```
.object.opEquals(a, b)
```

and that function is implemented as:

```
bool opEquals(Object a, Object b)
{
    if (a is b) return true;
    if (a is null || b is null) return false;
    if (typeid(a) == typeid(b)) return a.opEquals(b);
    return a.opEquals(b) && b.opEquals(a);
}
```

2. Otherwise the expressions `a.opEquals(b)` and `b.opEquals(a)` are tried. If both resolve to the same `opEquals` function, then the expression is rewritten to be `a.↗ ↘ opEquals(b)`.
3. If one is a better match than the other, or one compiles and the other does not, the first is selected.
4. Otherwise, an error results.

If overriding `Object.opEquals()` for classes, the class member function signature should look like:

```
class C
{
    override bool opEquals(Object o) { ... }
}
```

If structs declare an `opEquals` member function for the identity comparison, it could have several forms, such as:

```
struct S
{
    // lhs should be mutable object
```

```

bool opEquals(const S s) { ... }           // for r-values (e.g. temporaries)
bool opEquals(ref const S s) { ... }      // for l-values (e.g. variables)

// both hand side can be const object
bool opEquals(const S s) const { ... }    // for r-values (e.g. temporaries)
}

```

Alternatively, you can declare a single templated `opEquals` function with an auto ref parameter:

```

struct S
{
    // for l-values and r-values,
    // with converting both hand side implicitly to const
    bool opEquals()(auto ref const S s) const { ... }
}

```

Overloading <, <=, >, and >=

Comparison operations are rewritten as follows:

Rewriting of comparison operations		
comparison	rewrite 1	rewrite 2
<code>a < b</code>	<code>a.opCmp(b) < 0</code>	<code>b.opCmp(a) > 0</code>
<code>a <= b</code>	<code>a.opCmp(b) <= 0</code>	<code>b.opCmp(a) >= 0</code>
<code>a > b</code>	<code>a.opCmp(b) > 0</code>	<code>b.opCmp(a) < 0</code>
<code>a >= b</code>	<code>a.opCmp(b) >= 0</code>	<code>b.opCmp(a) <= 0</code>

Both rewrites are tried. If only one compiles, that one is taken. If they both resolve to the same function, the first rewrite is done. If they resolve to different functions, the best matching one is used. If they both match the same, but are different functions, an ambiguity error results.

If overriding `Object.opCmp()` for classes, the class member function signature should look like:

```

class C
{
    override int opCmp(Object o) { ... }
}

```

If structs declare an `opCmp` member function, it should have the following form:

```

struct S
{
    int opCmp(ref const S s) const { ... }
}

```

Note that `opCmp` is only used for the inequality operators; expressions like `a == b` always uses `opEquals`. If `opCmp` is defined but `opEquals` isn't, the compiler will supply a default version of `opEquals` that performs member-wise comparison. If this member-wise comparison is not consistent with the user-defined `opCmp`, then it is up to the programmer to supply an appropriate version of `opEquals`. Otherwise, inequalities like `a <= b` will behave inconsistently with equalities like `a == b`.

20.5 Function Call Operator Overloading `f()`

The function call operator, `()`, can be overloaded by declaring a function named `opCall`:

```
struct F
{
    int opCall();
    int opCall(int x, int y, int z);
}

void test()
{
    F f;
    int i;

    i = f();           // same as i = f.opCall();
    i = f(3,4,5);     // same as i = f.opCall(3,4,5);
}
```

In this way a struct or class object can behave as if it were a function.

Note that merely declaring `opCall` automatically disables struct literal syntax. To avoid the limitation, you need to also declare a constructor so that it takes priority over `opCall` in `Type(...)` syntax.

```
struct Multiplier
{
    int factor;
    this(int num) { factor = num; }
    int opCall(int value) { return value * factor; }
}

void test()
{
    Multiplier m = Multiplier(10); // invoke constructor
    int result = m(5);             // invoke opCall
}
```



```
    assert(result == 50);
}
```

Static opCall

`static opCall` also works as expected for a function call operator with type names.

```
struct Double
{
    static int opCall(int x) { return x * 2; }
}
void test()
{
    int i = Double(2);
    assert(i == 4);
}
```

Mixing struct constructors and `static opCall` is not allowed.

```
struct S
{
    this(int i) {}
    static S opCall() // disallowed due to constructor
    {
        return S.init;
    }
}
```

Note: `static opCall` can be used to simulate struct constructors with no arguments, but this is not recommended practice. Instead, the preferred solution is to use a factory function to create struct instances.

20.6 Assignment Operator Overloading

The assignment operator `=` can be overloaded if the left hand side is a struct aggregate, and `opAssign` is a member function of that aggregate.

For struct types, operator overloading for the identity assignment is allowed.

```
struct S
{
    // identity assignment, allowed.
    void opAssign(S rhs);
}
```

```

    // not identity assignment, also allowed.
    void opAssign(int);
}
S s;
s = S();      // Rewritten to s.opAssign(S());
s = 1;        // Rewritten to s.opAssign(1);

```

However for class types, identity assignment is not allowed. All class types have reference semantics, so identity assignment by default rebinds the left-hand-side to the argument at the right, and this is not overridable.

```

class C
{
    // If X is the same type as C or the type which is
    // implicitly convertible to C, then opAssign would
    // accept identity assignment, which is disallowed.
    // C opAssign(...);
    // C opAssign(X);
    // C opAssign(X, ...);
    // C opAssign(X ...);
    // C opAssign(X, U = defaultValue, etc.);

    // not an identity assignment - allowed
    void opAssign(int);
}
C c = new C();
c = new C(); // Rebinding referencee
c = 1;       // Rewritten to c.opAssign(1);

```

Index Assignment Operator Overloading

If the left hand side of an assignment is an index operation on a struct or class instance, it can be overloaded by providing an `opIndexAssign` member function. Expressions of the form `a[\textit{b}₁, \textit{b}₂, ... \textit{b}_n ↱ </sub>] = c` are rewritten as `a.opIndexAssign(c, \textit{b}₁, \textit{b}₂, ... \textit{b}_n ↱ </sub>)`.

```

struct A
{
    int opIndexAssign(int value, size_t i1, size_t i2);
}

```

```

}

void test()
{
    A a;
    a[i,3] = 7; // same as a.opIndexAssign(7,i,3);
}

```

Slice Assignment Operator Overloading

If the left hand side of an assignment is a slice operation on a struct or class instance, it can be overloaded by implementing an `opIndexAssign` member function that takes the return value of the `opSlice` function as parameter(s). Expressions of the form `a[i..j] = c` are rewritten as `a.opIndexAssign(c, a.opSlice(i, j))`, and `a[] = c` as `a.opIndexAssign(c)`.

See Array Indexing and Slicing Operators Overloading for more details.

```

struct A
{
    int opIndexAssign(int v); // overloads a[] = v
    int opIndexAssign(int v, size_t[2] x); // overloads a[i .. j] = v
    int[2] opSlice(size_t x, size_t y); // overloads i .. j
}

void test()
{
    A a;
    int v;

    a[] = v; // same as a.opIndexAssign(v);
    a[3..4] = v; // same as a.opIndexAssign(v, a.opSlice(3,4));
}

```

For backward compatibility, if rewriting `a[i..j]` as `a.opIndexAssign(a.opSlice(i, j))` fails to compile, the legacy rewrite `opSliceAssign(c, i, j)` is used instead.

20.7 Op Assignment Operator Overloading

The following op assignment operators are overloadable:

Overloadable Op Assignment Operators

+=	-=	*=	/=	%=	^^=	&=
=	^=	<<=	>>=	>>>=	~=	

The expression:

```
a op= b
```

is rewritten as:

```
a.opOpAssign!(\textquotedblleft{}op\textquotedblright{})(b)
```

Index Op Assignment Operator Overloading

If the left hand side of an `op=` is an index expression on a struct or class instance and `opIndexOpAssign` is a member:

```
a[b<sub>1</sub>, b<sub>2</sub>, ... b<sub>n</sub>] op= c
```

it is rewritten as:

```
a.opIndexOpAssign!(\textquotedblleft{}op\textquotedblright{})(c, b<sub>1</sub>,
    b<sub>2</sub>, ... b<sub>n</sub>)
```

Slice Op Assignment Operator Overloading

If the left hand side of an `op=` is a slice expression on a struct or class instance and `opIndexOpAssign` is a member:

```
a[i..j] op= c
```

it is rewritten as:

```
a.opIndexOpAssign!(\textquotedblleft{}op\textquotedblright{})(c, a.opSlice(i,
    j))
```

and

```
a[] op= c
```

it is rewritten as:

```
a.opIndexOpAssign!(\textquotedblleft{}op\textquotedblright{})(c)
```

For backward compatibility, if the above rewrites fail and `opSliceOpAssign` is defined, then the rewrites `a.opSliceOpAssign(c, i, j)` and `a.opSliceOpAssign(c)` are tried, respectively.

20.8 Array Indexing and Slicing Operators Overloading

The array indexing and slicing operators are overloaded by implementing the `opIndex`, `opSlice`, and `opDollar` methods. These may be combined to implement multidimensional arrays.

The code example below shows a simple implementation of a 2-dimensional array with overloaded indexing and slicing operators. The explanations of the various constructs employed are given in the sections following.

```
struct Array2D(E)
{
    E[] impl;
    int stride;
    int width, height;

    this(int width, int height, E[] initialData = [])
    {
        impl = initialData;
        this.stride = this.width = width;
        this.height = height;
        impl.length = width * height;
    }

    // Index a single element, e.g., arr[0, 1]
    ref E opIndex(int i, int j) { return impl[i + stride*j]; }

    // Array slicing, e.g., arr[1..2, 1..2], arr[2, 0..\$], arr[0..\$, 1].
    Array2D opIndex(int[2] r1, int[2] r2)
    {
        Array2D result;

        auto startOffset = r1[0] + r2[0]*stride;
        auto endOffset = r1[1] + (r2[1] - 1)*stride;
        result.impl = this.impl[startOffset .. endOffset];

        result.stride = this.stride;
        result.width = r1[1] - r1[0];
        result.height = r2[1] - r2[0];

        return result;
    }
}
```

```

auto opIndex(int[2] r1, int j) { return opIndex(r1, [j, j+1]); }
auto opIndex(int i, int[2] r2) { return opIndex([i, i+1], r2); }

// Support for 'x..y' notation in slicing operator for the given dimension.
int[2] opSlice(size_t dim)(int start, int end)
    if (dim >= 0 && dim < 2)
in { assert(start >= 0 && end <= this.opDollar!dim); }
body
{
    return [start, end];
}

// Support '\$' in slicing notation, e.g., arr[1..\$, 0..\$-1].
@property int opDollar(size_t dim : 0)() { return width; }
@property int opDollar(size_t dim : 1)() { return height; }
}

unittest
{
    auto arr = Array2D!int(4, 3, [
        0, 1, 2, 3,
        4, 5, 6, 7,
        8, 9, 10, 11
    ]);

    // Basic indexing
    assert(arr[0, 0] == 0);
    assert(arr[1, 0] == 1);
    assert(arr[0, 1] == 4);

    // Use of opDollar
    assert(arr[$-1, 0] == 3);
    assert(arr[0, $-1] == 8); // Note the value of \$ differs by dimension
    assert(arr[$-1, $-1] == 11);

    // Slicing
    auto slice1 = arr[1..\$, 0..\$];
    assert(slice1[0, 0] == 1 && slice1[1, 0] == 2 && slice1[2, 0] == 3 &&
        slice1[0, 1] == 5 && slice1[1, 1] == 6 && slice1[2, 1] == 7 &&
        slice1[0, 2] == 9 && slice1[1, 2] == 10 && slice1[2, 2] == 11);
}

```

```

    auto slice2 = slice1[0..2, 1..$];
    assert(slice2[0, 0] == 5 && slice2[1, 0] == 6 &&
           slice2[0, 1] == 9 && slice2[1, 1] == 10);

    // Thin slices
    auto slice3 = arr[2, 0..$];
    assert(slice3[0, 0] == 2 &&
           slice3[0, 1] == 6 &&
           slice3[0, 2] == 10);

    auto slice4 = arr[0..3, 2];
    assert(slice4[0, 0] == 8 && slice4[1, 0] == 9 && slice4[2, 0] == 10);
}

```

Index Operator Overloading

Expressions of the form `arr[b₁, b₂, ... b_n]` are translated into `arr.opIndex(b₁, b₂, ... b_n)`. For example:

```

struct A
{
    int opIndex(size_t i1, size_t i2, size_t i3);
}

void test()
{
    A a;
    int i;
    i = a[5,6,7]; // same as i = a.opIndex(5,6,7);
}

```

In this way a struct or class object can behave as if it were an array.

If an index expression can be rewritten using `opIndexAssign` or `opIndexOpAssign`, those are preferred over `opIndex`.

Slice Operator Overloading

Overloading the slicing operator means overloading expressions like `a[]` or `a[i..j]`, where the expressions inside the square brackets contain slice expressions of the form `i..j`.

To overload `a[]`, simply define `opIndex` with no parameters:

```
struct S
{
    int[] impl;
    int[] opIndex()
    {
        return impl[];
    }
}

void test()
{
    auto s = S([1,2,3]);
    auto t = s[]; // calls s.opIndex()
    assert(t == [1,2,3]);
}
```

To overload array indexing of the form `a[i..j, ...]`, two steps are needed. First, the expressions of the form `i..j` are translated via `opSlice` into user-defined objects that encapsulate the endpoints `i` and `j`. Then these user-defined objects are passed to `opIndex` to perform the actual slicing. This design was chosen in order to support mixed indexing and slicing in multidimensional arrays; for example, in translating expressions like `arr[1, 2..3, 4]`.

More precisely, an expression of the form `arr[b₁, b₂, ... b_n]` is translated into `arr.opIndex(c₁, c₂, ... c_n)`. Each argument `b_i` can be either a single expression, in which case it is passed directly as the corresponding argument `c_i` to `opIndex`; or it can be a slice expression of the form `x_i..y_i`, in which case the corresponding argument `c_i` to `opIndex` is `arr.opSlice!i(x_i, y_i)`. Namely:

<i>op</i>	<i>rewrite</i>
<code>arr[1, 2, 3]</code>	<code>arr.opIndex(1, 2, 3)</code>
<code>arr[1..2, 3..4, 5..6]</code>	<code>arr.opIndex(arr.opSlice!0(1,2), arr.opSlice!1(3,4), arr.opSlice!2(5,6))</code>
<code>arr[1, 2..3, 4]</code>	<code>arr.opIndex(1, arr.opSlice!1(2,3), 4)</code>

Similar translations are done for assignment operators involving slicing, for example:

<i>op</i>	<i>rewrite</i>
<code>arr[1, 2..3, 4] = c</code>	<code>arr.opIndexAssign(c, 1, arr.opSlice!1(2, 3), 4)</code>
<code>arr[2, 3..4] += c</code>	<code>arr.opIndexOpAssign!"+"(c, 2, arr.opSlice!1(2, 3))</code>

The intention is that `opSlice!i` should return a user-defined object that represents an interval of indices along the *i*'th dimension of the array. This object is then passed to `opIndex` to perform the actual slicing operation. If only one-dimensional slicing is desired, `opSlice` may be declared without the compile-time parameter *i*.

Note that in all cases, `arr` is only evaluated once. Thus, an expression like `getArray() ↪ [1, 2..3, $-1]=c` has the effect of:

```
auto __tmp = getArray();
__tmp.opIndexAssign(c, 1, __tmp.opSlice!1(2,3), __tmp.opDollar!2 - 1);
```

where the initial function call to `getArray` is only executed once.

For backward compatibility, `a[]` and `a[i..j]` can also be overloaded by implementing `opSlice ↪ ()` with no arguments and `opSlice(i, j)` with two arguments, respectively. This only applies for one-dimensional slicing, and dates from when D did not have full support for multidimensional arrays. This usage of `opSlice` is discouraged.

Dollar Operator Overloading

Within the arguments to array index and slicing operators, `$` gets translated to `opDollar!i`, where *i* is the position of the expression `$` appears in. For example:

<i>op</i>	<i>rewrite</i>
<code>arr[\$-1, \$-2, 3]</code>	<code>arr.opIndex(arr.opDollar!0 - 1, arr.opDollar!1 - 2, 3)</code>
<code>arr[1, 2, 3..\$]</code>	<code>arr.opIndex(1, 2, arr.opSlice!2(3, arr.opDollar!2))</code>

The intention is that `opDollar!i` should return the length of the array along its *i*'th dimension, or a user-defined object representing the end of the array along that dimension, that is understood by `opSlice` and `opIndex`.

```
struct Rectangle
{
    int width, height;
    int[] [] impl;
    this(int w, int h)
    {
        width = w;
        height = h;
        impl = new int[w][h];
    }
    int opIndex(size_t i1, size_t i2)
```

```

    {
        return impl[i1][i2];
    }
    int opDollar(size_t pos)()
    {
        static if (pos==0)
            return width;
        else
            return height;
    }
}

void test()
{
    auto r = Rectangle(10,20);
    int i = r[$-1, 0];    // same as: r.opIndex(r.opDollar!0, 0),
                        // which is r.opIndex(r.width-1, 0)
    int j = r[0, $-1];    // same as: r.opIndex(0, r.opDollar!1)
                        // which is r.opIndex(0, r.height-1)
}

```

As the above example shows, a different compile-time argument is passed to `opDollar` depending on which argument it appears in. A `$` appearing in the first argument gets translated to `opDollar!0`, a `$` appearing in the second argument gets translated to `opDollar!1`, and so on. Thus, the appropriate value for `$` can be returned to implement multidimensional arrays.

Note that `opDollar!i` is only evaluated once for each `i` where `$` occurs in the corresponding position in the indexing operation. Thus, an expression like `arr[$-sqrt($), 0, $-1]` has the effect of:

```

auto __tmp1 = arr.opDollar!0;
auto __tmp2 = arr.opDollar!2;
arr.opIndex(__tmp1 - sqrt(__tmp1), 0, __tmp2 - 1);

```

If `opIndex` is declared with only one argument, the compile-time argument to `opDollar` may be omitted. In this case, it is illegal to use `$` inside an array indexing expression with more than one argument.

20.9 Forwarding

Member names not found in a class or struct can be forwarded to a template function named `opDispatch` for resolution.

```
import std.stdio;

struct S
{
    void opDispatch(string s, T)(T i)
    {
        writeln("S.opDispatch('%s', %s)", s, i);
    }
}

class C
{
    void opDispatch(string s)(int i)
    {
        writeln("C.opDispatch('%s', %s)", s, i);
    }
}

struct D
{
    template opDispatch(string s)
    {
        enum int opDispatch = 8;
    }
}

void main()
{
    S s;
    s.opDispatch!("hello")(7);
    s.foo(7);

    auto c = new C();
    c.foo(8);

    D d;
    writeln("d.foo=%s", d.foo);
    assert(d.foo == 8);
}
```

20.10 D1 style operator overloading

While the preferred style for operator overloading is to use the above mechanisms, the D1 operator overload mechanisms are still allowed. There is no guarantee that these mechanisms will be supported in the future.

Chapter 21

Templates

I think that I can safely say that nobody understands C++ template mechanics.

– Richard Deyman

Templates are D's approach to generic programming. Templates are defined with a *TemplateDeclaration*:

TemplateDeclaration:

```
template Identifier TemplateParameters Constraintopt { DeclDefsopt }
```

TemplateParameters:

```
( TemplateParameterListopt )
```

TemplateParameterList:

```
TemplateParameter
TemplateParameter ,
TemplateParameter , TemplateParameterList
```

TemplateParameter:

```
TemplateTypeParameter
TemplateValueParameter
TemplateAliasParameter
TemplateSequenceParameter
TemplateThisParameter
```

The body of the *TemplateDeclaration* must be syntactically correct even if never instantiated. Semantic analysis is not done until instantiated. A template forms its own scope, and the template body can contain classes, structs, types, enums, variables, functions, and other templates.

Template parameters can be types, values, symbols, or sequences. Types can be any type. Value parameters must be of an integral type, floating point type, or string type and specializations for them must resolve to an integral constant, floating point constant, null, or a string literal. Symbols can be any non-local symbol. Sequences can contain zero or more types, values or symbols.

Template parameter specializations constrain the values or types the *TemplateParameter* can accept.

Template parameter defaults are the value or type to use for the *TemplateParameter* in case one is not supplied.

21.1 Explicit Template Instantiation

Templates are explicitly instantiated with:

TemplateInstance:

Identifier TemplateArguments

TemplateArguments:

! (*TemplateArgumentList*_{opt})

! *TemplateSingleArgument*

TemplateArgumentList:

TemplateArgument

TemplateArgument ,

TemplateArgument , *TemplateArgumentList*

TemplateArgument:

Type

AssignExpression

Symbol

Symbol:

SymbolTail

. *SymbolTail*

SymbolTail:

Identifier

Identifier . *SymbolTail*

```

    TemplateInstance
    TemplateInstance . SymbolTail

TemplateSingleArgument:
    Identifier
    BasicTypeX
    CharacterLiteral
    StringLiteral
    IntegerLiteral
    FloatLiteral
    true
    false
    null
    this
    SpecialKeyword

```

Once instantiated, the declarations inside the template, called the template members, are in the scope of the *TemplateInstance*:

```

template TFoo(T) { alias t = T*; }
...
TFoo!(int).t x; // declare x to be of type int*

```

If the *TemplateArgument* is one token long, the parentheses can be omitted:

```

TFoo!int.t x;    // same as TFoo!(int).t x;

```

A template instantiation can be aliased:

```

template TFoo(T) { alias t = T*; }
alias abc = TFoo!(int);
abc.t x;    // declare x to be of type int*

```

Multiple instantiations of a *TemplateDeclaration* with the same *TemplateArgumentList* all will refer to the same instantiation. For example:

```

template TFoo(T) { T f; }
alias a = TFoo!(int);
alias b = TFoo!(int);
...
a.f = 3;
assert(b.f == 3); // a and b refer to the same instance of TFoo

```

This is true even if the *TemplateInstances* are done in different modules.

Even if template arguments are implicitly converted to the same template parameter type, they still refer to the same instance:

```
struct TFoo(int x) { }
```

```
// 3 and 2+1 are both 3 of type int
static assert(is(TFoo!(3) == TFoo!(2 + 1)));
```

```
// 3u is implicitly converted to 3 to match int parameter,
// and refers to exactly the same instance as TFoo!(3)
static assert(is(TFoo!(3) == TFoo!(3u)));
```

If multiple templates with the same *Identifier* are declared, they are distinct if they have a different number of arguments or are differently specialized.

For example, a simple generic copy template would be:

```
template TCopy(T)
{
    void copy(out T to, T from)
    {
        to = from;
    }
}
```

To use the template, it must first be instantiated with a specific type:

```
int i;
TCopy!(int).copy(i, 3);
```

21.2 Instantiation Scope

TemplateInstances are always instantiated in the scope of where the *TemplateDeclaration* is declared, with the addition of the template parameters being declared as aliases for their deduced types.

For example:

```
module a
```

```
template TFoo(T) { void bar() { func(); } }
```

```
module b
```

```
import a;
```

```
void func() { }
```



```

alias f = TFoo!(int); // error: func not defined in module a

and:
module a

template TFoo(T) { void bar() { func(1); } }
void func(double d) { }

module b

import a;

void func(int i) { }
alias f = TFoo!(int);
...
f.bar(); // will call a.func(double)

```

TemplateParameter specializations and default values are evaluated in the scope of the *TemplateDeclaration*.

21.3 Argument Deduction

The types of template parameters are deduced for a particular template instantiation by comparing the template argument with the corresponding template parameter.

For each template parameter, the following rules are applied in order until a type is deduced for each parameter:

1. If there is no type specialization for the parameter, the type of the parameter is set to the template argument.
2. If the type specialization is dependent on a type parameter, the type of that parameter is set to be the corresponding part of the type argument.
3. If after all the type arguments are examined, there are any type parameters left with no type assigned, they are assigned types corresponding to the template argument in the same position in the *TemplateArgumentList*.
4. If applying the above rules does not result in exactly one type for each template parameter, then it is an error.

For example:

```

template TFoo(T) { }
alias Foo1 = TFoo!(int);    // (1) T is deduced to be int
alias Foo2 = TFoo!(char*); // (1) T is deduced to be char*

```

```

template TBar(T : T*) { }
alias Foo3 = TBar!(char*);    // (2) T is deduced to be char

template TAbc(D, U : D[]) { }
alias Bar1 = TAbc!(int, int[]); // (2) D is deduced to be int, U is int[]
alias Bar2 = TAbc!(char, int[]); // (4) error, D is both char and int

template TDef(D : E*, E) { }
alias Bar3 = TDef!(int*, int); // (1) E is int
                                // (3) D is int*

```

Deduction from a specialization can provide values for more than one parameter:

```

template Foo(T: T[U], U)
{
    ...
}

Foo!(int[long])    // instantiates Foo with T set to int, U set to long

```

When considering matches, a class is considered to be a match for any super classes or interfaces:

```

class A { }
class B : A { }

template TFoo(T : A) { }
alias Foo4 = TFoo!(B);    // (3) T is B

template TBar(T : U*, U : A) { }
alias Foo5 = TBar!(B*, B); // (2) T is B*
                           // (3) U is B

```

21.4 Template Type Parameters

TemplateTypeParameter:

Identifier

Identifier TemplateTypeParametersSpecialization

Identifier TemplateTypeParameterDefault

Identifier TemplateTypeParametersSpecialization TemplateTypeParameterDefault

TemplateTypeParametersSpecialization:

```

    : Type

    TemplateTypeParameterDefault:
        = Type

```

Specialization

Templates may be specialized for particular types of arguments by following the template parameter identifier with a `:` and the specialized type. For example:

```

template TFoo(T)          { ... } // #1
template TFoo(T : T[])    { ... } // #2
template TFoo(T : char)   { ... } // #3
template TFoo(T, U, V)    { ... } // #4

alias foo1 = TFoo!(int);           // instantiates #1
alias foo2 = TFoo!(double[]);      // instantiates #2 with T being double
alias foo3 = TFoo!(char);          // instantiates #3
alias fooe = TFoo!(char, int);      // error, number of arguments mismatch
alias foo4 = TFoo!(char, int, int); // instantiates #4

```

The template picked to instantiate is the one that is most specialized that fits the types of the *TemplateArgumentList*. Determining which is more specialized is done in the same way as the C++ partial ordering rules. If the result is ambiguous, it is an error.

21.5 Template This Parameters

```

TemplateThisParameter:
    this TemplateTypeParameter

```

TemplateThisParameters are used in member function templates to pick up the type of the *this* reference.

```

import std.stdio;

struct S
{
    const void foo(this T)(int i)
    {
        writeln(typeid(T));
    }
}

```

```

    }
}

void main()
{
    const(S) s;
    (&s).foo(1);
    S s2;
    s2.foo(2);
    immutable(S) s3;
    s3.foo(3);
}

```

Prints:

```
const(S) S immutable(S)
```

This is especially useful when used with inheritance. For example, you might want to implement a final base method which returns a derived class type. Typically you would return a base type, but this won't allow you to call or access derived properties of the type:

```

interface Addable(T)
{
    final auto add(T t)
    {
        return this;
    }
}

class List(T) : Addable!T
{
    List remove(T t)
    {
        return this;
    }
}

void main()
{
    auto list = new List!int;
    list.add(1).remove(1); // error: no 'remove' method for Addable!int
}

```

Here the method `add` returns the base type, which doesn't implement the `remove` method. The `template this` parameter can be used for this purpose:

```
interface Addable(T)
{
    final R add(this R)(T t)
    {
        return cast(R)this; // cast is necessary, but safe
    }
}

class List(T) : Addable!T
{
    List remove(T t)
    {
        return this;
    }
}

void main()
{
    auto list = new List!int;
    list.add(1).remove(1); // ok
}
```

21.6 Template Value Parameters

TemplateValueParameter:

BasicType Declarator

BasicType Declarator TemplateValueParameterSpecialization

BasicType Declarator TemplateValueParameterDefault

BasicType Declarator TemplateValueParameterSpecialization TemplateValueParameterDefault

TemplateValueParameterSpecialization:

: ConditionalExpression

TemplateValueParameterDefault:

= AssignExpression

= SpecialKeyword

Template value parameter types can be any type which can be statically initialized at compile time. Template value arguments can be integer values, floating point values, nulls, string values, array literals of template value arguments, associative array literals of template value arguments, or struct literals of template value arguments.

```
template foo(string s)
{
    string bar() { return s ~ "▯betty"; }
}

void main()
{
    writeln("%s", foo!("hello").bar()); // prints: hello betty
}
```

This example of template foo has a value parameter that is specialized for 10:

```
template foo(U : int, int T : 10)
{
    U x = T;
}

void main()
{
    assert(foo!(int, 10).x == 10);
}
```

21.7 Template Alias Parameters

TemplateAliasParameter:

```
alias Identifier TemplateAliasParameterSpecializationopt TemplateAliasParameterDefaultopt
alias BasicType Declarator TemplateAliasParameterSpecializationopt TemplateAliasParameterDefa
```

TemplateAliasParameterSpecialization:

```
: Type
: ConditionalExpression
```

TemplateAliasParameterDefault:

```
= Type
= ConditionalExpression
```

Alias parameters enable templates to be parameterized with almost any kind of D symbol, including user-defined type names, global names, local names, module names, template names, and template instance names. Literals can also be used as arguments to alias parameters.

Examples:

- User-defined type names

```
class Foo
{
    static int p;
}

template Bar(alias T)
{
    alias q = T.p;
}

void main()
{
    alias bar = Bar!(Foo);
    bar.q = 3; // sets Foo.p to 3
}
```

- Global names

```
shared int x;

template Foo(alias X)
{
    auto p = &X;
}

void main()
{
    alias bar = Foo!(x);
    *bar.p = 3; // set x to 3
    static shared int y;
    alias abc = Foo!(y);
    *abc.p = 3; // set y to 3
}
```

- Module names

```
import std.string;

template Foo(alias X)
{
    alias y = X.toString;
}

void main()
{
    alias bar = Foo!(std.string);
    bar.y(3);    // calls std.string.toString(3)
}
```

- Template names

```
shared int x;

template Foo(alias X)
{
    auto p = &X;
}

template Bar(alias T)
{
    alias abc = T!(x);
}

void main()
{
    alias bar = Bar!(Foo);
    *bar.abc.p = 3;    // sets x to 3
}
```

- Template instance names

```
shared int x;

template Foo(alias X)
{
    auto p = &X;
}
```



```

template Bar(alias T)
{
    alias q = T.p;
}

void main()
{
    alias foo = Foo!(x);
    alias bar = Bar!(foo);
    *bar.q = 3;  // sets x to 3
}

```

- Literals

```

template Foo(alias X, alias Y)
{
    static int i = X;
    static string s = Y;
}

void main()
{
    alias foo = Foo!(3, "bar");
    writeln(foo.i, foo.s);  // prints 3bar
}

```

Typed alias parameters

Alias parameters can also be typed. These parameters will accept symbols of that type:

```

template Foo(alias int x) { }
int x;
float f;

Foo!x;  // ok
Foo!f;  // fails to instantiate

```

Specialization

Alias parameters can accept both literals and user-defined type symbols, but they are less specialized than the matches to type parameters and value parameters:

```

template Foo(T)          { ... } // #1
template Foo(int n)      { ... } // #2
template Foo(alias sym) { ... } // #3

struct S {}
int var;

alias foo1 = Foo!(S);      // instantiates #1
alias foo2 = Foo!(1);      // instantiates #2
alias foo3a = Foo!([1,2]); // instantiates #3
alias foo3b = Foo!(var);   // instantiates #3

template Bar(alias A) { ... } // #4
template Bar(T : U!V, alias U, V...) { ... } // #5

class C(T) {}
alias bar = Bar!(C!int);    // instantiates #5

```

21.8 Template Sequence Parameters

TemplateSequenceParameter:
Identifier ...

If the last template parameter in the *TemplateParameterList* is declared as a *TemplateSequenceParameter*, it is a match with any trailing template arguments. Such a sequence of arguments can be defined using the template `std.meta.AliasSeq` and will thus henceforth be referred to by that name for clarity. An *AliasSeq* is not itself a type, value, or symbol. It is a compile-time sequence of any mix of types, values or symbols.

An *AliasSeq* whose elements consist entirely of types is called a type sequence or *TypeSeq*. An *AliasSeq* whose elements consist entirely of values is called a value sequence or *ValueSeq*.

An *AliasSeq* can be used as an argument list to instantiate another template, or as the list of parameters for a function.

```

template print(args...)
{
    void print()
    {
        writeln("args_are_", args); // args is a ValueSeq
    }
}

```

```

}

template write(Args...)
{
    void write(Args args) // Args is a TypeSeq
                        // args is a ValueSeq
    {
        writeln("args_are_", args);
    }
}

void main()
{
    print!(1, 'a', 6.8).print(); // prints: args are 1a6.8
    write!(int, char, double).write(1, 'a', 6.8); // prints: args are 1a6.8
}

```

The number of elements in an *AliasSeq* can be retrieved with the `.length` property. The n th element can be retrieved by indexing the *AliasSeq* with $[n]$, and sub-sequences are denoted by the slicing syntax.

AliasSeq-s are static compile-time entities, there is no way to dynamically change, add, or remove elements either at compile-time or run-time.

Type sequences can be deduced from the trailing parameters of an implicitly instantiated function template:

```

template print(T, Args...)
{
    void print(T first, Args args)
    {
        writeln(first);
        static if (args.length) // if more arguments
            print(args);        // recurse for remaining arguments
    }
}

void main()
{
    print(1, 'a', 6.8);
}

prints:
1 a 6.8

```

Type sequences can also be deduced from the type of a delegate or function parameter list passed as a function argument:

```
import std.stdio;

/* Partially applies a delegate by tying its first argument to a particular value
   ↳ .
   * R = return type
   * T = first argument type
   * Args = TypeSeq of remaining argument types
   */
R delegate(Args) partial(R, T, Args...)(R delegate(T, Args) dg, T first)
{
    // return a closure
    return (Args args) => dg(first, args);
}

void main()
{
    int plus(int x, int y, int z)
    {
        return x + y + z;
    }

    auto plus_two = partial(&plus, 2);
    writeln("%d", plus_two(6, 8)); // prints 16
}
```

See also: `std.functional.partial`

Specialization

If both a template with a sequence parameter and a template without a sequence parameter exactly match a template instantiation, the template without a *TemplateSequenceParameter* is selected.

```
template Foo(T)          { pragma(msg, "1"); } // #1
template Foo(int n)      { pragma(msg, "2"); } // #2
template Foo(alias sym) { pragma(msg, "3"); } // #3
template Foo(Args...)    { pragma(msg, "4"); } // #4

import std.stdio;
```

```
// Any sole template argument will never match to #4
alias foo1 = Foo!(int);           // instantiates #1
alias foo2 = Foo!(3);            // instantiates #2
alias foo3 = Foo!(std);          // instantiates #3

alias foo4 = Foo!(int, 3, std);   // instantiates #4
```

21.9 Template Parameter Default Values

Trailing template parameters can be given default values:

```
template Foo(T, U = int) { ... }
Foo!(uint, long); // instantiate Foo with T as uint, and U as long
Foo!(uint);       // instantiate Foo with T as uint, and U as int

template Foo(T, U = T*) { ... }
Foo!(uint);       // instantiate Foo with T as uint, and U as uint*
```

21.10 Eponymous Templates

If a template contains members whose name is the same as the template identifier and if the type or the parameters type of these members include at least all the template parameters then these members are assumed to be referred to in a template instantiation:

```
template foo(T)
{
    T foo; // declare variable foo of type T
}

void main()
{
    foo!(int) = 6; // instead of foo!(int).foo
}
```

Using functions and more types than the template:

```
template foo(S, T)
{
    // each member contains all the template parameters
    void foo(S s, T t) {}
    void foo(S s, T t, string) {}
}
```

```

}

void main()
{
    foo(1, 2, "test"); // foo!(int, int).foo(1, 2, "test")
    foo(1, 2); // foo!(int, int).foo(1, 2)
}

```

When the template parameters must be deduced, the eponymous members can't rely on a `static if` condition since the deduction relies on how the in members are used:

```

template foo(T)
{
    static if (is(T)) // T is not yet known...
        void foo(T t) {} // T is deduced from the member usage
}

void main()
{
    foo(0); // Error: cannot deduce function from argument types
    foo!int(0); // Ok since no deduction necessary
}

```

21.11 Template Constructors

ConstructorTemplate:

```

    this TemplateParameters Parameters MemberFunctionAttributesopt Constraintopt : ↵
↵
    this TemplateParameters Parameters MemberFunctionAttributesopt Constraintopt FunctionBody

```

Templates can be used to form constructors for classes and structs.

21.12 Aggregate Templates

ClassTemplateDeclaration:

```

class Identifier TemplateParameters Constraintopt BaseClassListopt AggregateBody
class Identifier TemplateParameters BaseClassListopt Constraintopt AggregateBody

```

InterfaceTemplateDeclaration:

```
interface Identifier TemplateParameters Constraintopt BaseInterfaceListopt AggregateBody
interface Identifier TemplateParameters BaseInterfaceList Constraint AggregateBody
```

StructTemplateDeclaration:

```
struct Identifier TemplateParameters Constraintopt AggregateBody
```

UnionTemplateDeclaration:

```
union Identifier TemplateParameters Constraintopt AggregateBody
```

If a template declares exactly one member, and that member is a class with the same name as the template:

```
template Bar(T)
{
    class Bar
    {
        T member;
    }
}
```

then the semantic equivalent, called a *ClassTemplateDeclaration* can be written as:

```
class Bar(T)
{
    T member;
}
```

Analogously to class templates, struct, union and interfaces can be transformed into templates by supplying a template parameter list.

21.13 Function Templates

If a template declares exactly one member, and that member is a function with the same name as the template, it is a function template declaration. Alternatively, a function template declaration is a function declaration with a *TemplateParameterList* immediately preceding the *Parameters*.

A function template to compute the square of type *T* is:

```
T Square(T)(T t)
{
    return t * t;
}
```

Function templates can be explicitly instantiated with a *!(TemplateArgumentList)*:

```
writeln("The square of %s is %s", 3, Square!(int)(3));
```

or implicitly, where the *TemplateArgumentList* is deduced from the types of the function arguments:

```
writeln("The square of %s is %s", 3, Square(3)); // T is deduced to be int
```

If there are fewer arguments supplied in the *TemplateArgumentList* than parameters in the *TemplateParameterList*, the arguments fulfill parameters from left to right, and the rest of the parameters are then deduced from the function arguments.

Function template type parameters that are to be implicitly deduced may not have specializations:

```
void Foo(T : T*)(T t) { ... }
```

```
int x,y;
```

```
Foo!(int*)(x); // ok, T is not deduced from function argument
```

```
Foo(&y); // error, T has specialization
```

Template arguments not implicitly deduced can have default values:

```
void Foo(T, U=T*)(T t) { U p; ... }
```

```
int x;
```

```
Foo(x); // T is int, U is int*
```

The deduced type parameter for dynamic array and pointer arguments has an unqualified head:

```
void foo(T)(T arg) { pragma(msg, T); }
```

```
int[] marr;
```

```
const(int[]) carr;
```

```
immutable(int[]) iarr;
```

```
foo(marr); // T == int[]
```

```
foo(carr); // T == const(int)[]
```

```
foo(iarr); // T == immutable(int)[]
```

```
int* mptr;
```

```
const(int*) cptr;
```

```
immutable(int*) iptr;
```

```
foo(mptr); // T == int*
```

```
foo(cptr); // T == const(int)*
```

```
foo(iptr); // T == immutable(int)*
```


Function templates can have their return types deduced based on the first *ReturnStatement* in the function:

```
auto Square(T)(T t)
{
    return t * t;
}
```

If there is more than one return statement, then the types of the return statement expressions must match. If there are no return statements, then the return type of the function template is `void`.

21.14 Variable Templates

Same as aggregates and functions, variable declarations with *Initializer* can have optional template parameters:

```
enum string constant(TL...) = TL.stringof;
ubyte[T.sizeof] storage(T) = 0;
auto array(alias a) = a;
```

These declarations are transformed into templates:

```
template constant(TL...)
{
    enum string constant = TL.stringof;
}
template storage(T)
{
    ubyte[T.sizeof] storage = 0;
}
template array(alias a)
{
    auto array = a;
}
```

21.15 Alias Templates

AliasDeclaration can also have optional template parameters:

```
alias Sequence(TL...) = TL;
```

It is lowered to:

```
template Sequence(TL...)
{
    alias Sequence = TL;
}
```

Function Templates with Auto Ref Parameters

An auto ref function template parameter becomes a ref parameter if its corresponding argument is an lvalue, otherwise it becomes a value parameter:

```
int foo(Args...)(auto ref Args args)
{
    int result;

    foreach (i, v; args)
    {
        if (v == 10)
            assert(__traits(isRef, args[i]));
        else
            assert(!__traits(isRef, args[i]));
        result += v;
    }
    return result;
}

void main()
{
    int y = 10;
    int r;
    r = foo(8);           // returns 8
    r = foo(y);           // returns 10
    r = foo(3, 4, y);     // returns 17
    r = foo(4, 5, y);     // returns 19
    r = foo(y, 6, y);     // returns 26
}
```

Auto ref parameters can be combined with auto ref return attributes:

```
auto ref min(T, U)(auto ref T lhs, auto ref U rhs)
{
    return lhs > rhs ? rhs : lhs;
}
```

```

}

void main()
{
    int x = 7, y = 8;
    int i;

    i = min(4, 3);      // returns 3
    i = min(x, y);      // returns 7
    min(x, y) = 10;     // sets x to 10
    static assert(!_traits(compiles, min(3, y) = 10));
    static assert(!_traits(compiles, min(y, 3) = 10));
}

```

21.16 Nested Templates

If a template is declared in aggregate or function local scope, the instantiated functions will implicitly capture the context of the enclosing scope.

```

class C
{
    int num;

    this(int n) { num = n; }

    template Foo()
    {
        // 'foo' can access 'this' reference of class C object.
        void foo(int n) { this.num = n; }
    }
}

void main()
{
    auto c = new C(1);
    assert(c.num == 1);

    c.Foo!().foo(5);
    assert(c.num == 5);
}

```

```

template Bar()
{
    // 'bar' can access local variable of 'main' function.
    void bar(int n) { c.num = n; }
}
Bar!().bar(10);
assert(c.num == 10);
}

```

Above, `Foo!().foo` will work just the same as a member function of class `C`, and `Bar!().bar` will work just the same as a nested function within function `main()`.

If a template has a template alias parameter, and is instantiated with a local symbol, the instantiated function will implicitly become nested in order to access runtime data of the given local symbol.

```

template Foo(alias sym)
{
    void foo() { sym = 10; }
}

class C
{
    int num;

    this(int n) { num = n; }

    void main()
    {
        assert(this.num == 1);

        alias fooX = Foo!(C.num).foo;

        // fooX will become member function implicitly, so &fooX returns delegate'
        //   ↳ object.
        static assert(is(typeof(&fooX) == delegate));

        fooX(); // called by using valid 'this' reference.
        assert(this.num == 10); // OK
    }
}

```

```

void main()
{
    new C(1).main();

    int num;
    alias fooX = Foo!num.foo;

    // fooX will become nested function implicitly, so &fooX returns delegate ↗
    ↘ object.
    static assert(is(typeof(&fooX) == delegate));

    fooX();
    assert(num == 10); // OK
}

```

Not only functions, but also instantiated class and struct types can become nested via implicitly captured context.

```

class C
{
    int num;
    this(int n) { num = n; }

    class N(T)
    {
        // instantiated class N!T can become nested in C
        T foo() { return num * 2; }
    }
}

void main()
{
    auto c = new C(10);
    auto n = c.new N!int();
    assert(n.foo() == 20);
}

void main()
{
    int num = 10;
    struct S(T)

```

```

{
    // instantiated struct S!T can become nested in main()
    T foo() { return num * 2; }
}
S!int s;
assert(s.foo() == 20);
}

```

A templated `struct` can become a nested `struct` if it is instantiated with a local symbol passed as an aliased argument:

```

struct A(alias F)
{
    int fun(int i) { return F(i); }
}

A!F makeA(alias F)() { return A!F(); }

void main()
{
    int x = 40;
    int fun(int i) { return x + i; }
    A!fun a = makeA!fun();
    assert(a.fun(2) == 42);
}

```

Limitation:

Currently nested templates can capture at most one context. As a typical example, non-static template member functions cannot take local symbol by using template alias parameter.

```

class C
{
    int num;
    void foo(alias sym)() { num = sym * 2; }
}

void main()
{
    auto c = new C();
    int var = 10;
}

```

```
    c.foo!var();    // NG, foo!var requires two contexts, 'this' and 'main()'
}
```

But, if one context is indirectly accessible from other context, it is allowed.

```
int sum(alias x, alias y)() { return x + y; }

void main()
{
    int a = 10;
    void nested()
    {
        int b = 20;
        assert(sum!(a, b)() == 30);
    }
    nested();
}
```

Two local variables `a` and `b` are in different contexts, but outer context is indirectly accessible from inner context, so nested template instance `sum!(a, b)` will capture only inner context.

21.17 Recursive Templates

Template features can be combined to produce some interesting effects, such as compile time evaluation of non-trivial functions. For example, a factorial template can be written:

```
template factorial(int n : 1)
{
    enum { factorial = 1 }
}

template factorial(int n)
{
    enum { factorial = n* factorial!(n-1) }
}

void test()
{
    writeln("%S", factorial!(4)); // prints 24
}
```

21.18 Template Constraints

Constraint:

```
if ( Expression )
```

Constraints are used to impose additional constraints on matching arguments to a template beyond what is possible in the *TemplateParameterList*. The *Expression* is computed at compile time and returns a result that is converted to a boolean value. If that value is true, then the template is matched, otherwise the template is not matched.

For example, the following function template only matches with odd values of N:

```
void foo(int N)()
    if (N & 1)
{
    ...
}
...
foo!(3)(); // OK, matches
foo!(4)(); // Error, no match
```

Template constraints can be used with aggregate types (structs, classes, unions). Constraints are effectively used with library module "std.traits":

```
import std.traits;

struct Bar(T)
    if (isIntegral!T)
{
    ...
}
...
auto x = Bar!int; // OK, int is an integral type
auto y = Bar!double; // Error, double does not satisfy constraint
```

21.19 Limitations

Templates cannot be used to add non-static fields or virtual functions to classes or interfaces. For example:

```
class Foo
{
```



```
template TBar(T)
{
    T xx;                // becomes a static field of Foo
    int func(T) { ... } // non-virtual

    static T yy;          // Ok
    static int func(T t, int y) { ... } // Ok
}
}
```


Chapter 22

Template Mixins

A *TemplateMixin* takes an arbitrary set of declarations from the body of a *TemplateDeclaration* and inserts them into the current context.

TemplateMixinDeclaration:

```
mixin template Identifier TemplateParameters Constraintopt { DeclDefsopt }
```

TemplateMixin:

```
mixin MixinTemplateName TemplateArgumentsopt Identifieropt ;
```

MixinTemplateName:

```
. QualifiedIdentifierList  
QualifiedIdentifierList  
Typeof . QualifiedIdentifierList
```

QualifiedIdentifierList:

```
Identifier  
Identifier . QualifiedIdentifierList  
TemplateInstance . QualifiedIdentifierList
```

A *TemplateMixin* can occur in declaration lists of modules, classes, structs, unions, and as a statement. The *MixinTemplateName* refers to a *TemplateDeclaration*. If the *TemplateDeclaration* has no parameters, the mixin form that has no *!(TemplateArgumentList)* can be used.

Unlike a template instantiation, a template mixin's body is evaluated within the scope where the mixin appears, not where the template declaration is defined. It is analogous to cutting and pasting the body of the template into the location of the mixin into a nested scope. It is useful for

injecting parameterized ‘boilerplate’ code, as well as for creating templated nested functions, which is not possible with template instantiations.

```

mixin template Foo()
{
    int x = 5;
}

mixin Foo;

struct Bar
{
    mixin Foo;
}

void test()
{
    writefln("x=%d", x);           // prints 5
    {
        Bar b;
        int x = 3;

        writefln("b.x=%d", b.x); // prints 5
        writefln("x=%d", x);      // prints 3
        {
            mixin Foo;
            writefln("x=%d", x); // prints 5
            x = 4;
            writefln("x=%d", x); // prints 4
        }
        writefln("x=%d", x);      // prints 3
    }
    writefln("x=%d", x);          // prints 5
}

```

Mixins can be parameterized:

```

mixin template Foo(T)
{
    T x = 5;
}

```

```
mixin Foo!(int);           // create x of type int
```

Mixins can add virtual functions to a class:

```
mixin template Foo()
{
    void func() { writeln("Foo.func()"); }
}

class Bar
{
    mixin Foo;
}

class Code : Bar
{
    override void func() { writeln("Code.func()"); }
}

void test()
{
    Bar b = new Bar();
    b.func();           // calls Foo.func()

    b = new Code();
    b.func();           // calls Code.func()
}
```

Mixins are evaluated in the scope of where they appear, not the scope of the template declaration:

```
int y = 3;

mixin template Foo()
{
    int abc() { return y; }
}

void test()
{
    int y = 8;
    mixin Foo; // local y is picked up, not global y
    assert(abc() == 8);
}
```

```
}
```

Mixins can parameterize symbols using alias parameters:

```
mixin template Foo(alias b)
{
    int abc() { return b; }
}
```

```
void test()
{
    int y = 8;
    mixin Foo!(y);
    assert(abc() == 8);
}
```

This example uses a mixin to implement a generic Duff's device for an arbitrary statement (in this case, the arbitrary statement is in bold). A nested function is generated as well as a delegate literal, these can be inlined by the compiler:

```
mixin template duffs_device(alias id1, alias id2, alias s)
{
    void duff_loop()
    {
        if (id1 < id2)
        {
            typeof(id1) n = (id2 - id1 + 7) / 8;
            switch ((id2 - id1) % 8)
            {
                case 0: do { s(); goto case;
                case 7:      s(); goto case;
                case 6:      s(); goto case;
                case 5:      s(); goto case;
                case 4:      s(); goto case;
                case 3:      s(); goto case;
                case 2:      s(); goto case;
                case 1:      s(); continue;
                default:      assert(0, "Impossible");
                            } while (--n > 0);
            }
        }
    }
}
```

```

}

void foo() { writeln("foo"); }

void test()
{
    int i = 1;
    int j = 11;

    mixin duffs_device!(i, j, delegate { foo(); });
    duff_loop(); // executes foo() 10 times
}

```

22.1 Mixin Scope

The declarations in a mixin are placed in a nested scope and then ‘imported’ into the surrounding scope. If the name of a declaration in a mixin is the same as a declaration in the surrounding scope, the surrounding declaration overrides the mixin one:

```

int x = 3;

mixin template Foo()
{
    int x = 5;
    int y = 5;
}

mixin Foo;
int y = 3;

void test()
{
    writefln("x=%d", x); // prints 3
    writefln("y=%d", y); // prints 3
}

```

If two different mixins are put in the same scope, and each define a declaration with the same name, there is an ambiguity error when the declaration is referenced:

```

mixin template Foo()
{

```

```

    int x = 5;
    void func(int x) { }
}

mixin template Bar()
{
    int x = 4;
    void func(long x) { }
}

mixin Foo;
mixin Bar;

void test()
{
    import std.stdio : writeln;
    writeln("x_ = %d", x); // error, x is ambiguous
    func(1);               // error, func is ambiguous
}

```

The call to `func()` is ambiguous because `Foo.func` and `Bar.func` are in different scopes.

If a mixin has an *Identifier*, it can be used to disambiguate between conflicting symbols:

```

int x = 6;

mixin template Foo()
{
    int x = 5;
    int y = 7;
    void func() { }
}

mixin template Bar()
{
    int x = 4;
    void func() { }
}

mixin Foo F;
mixin Bar B;

```



```

void test()
{
    writefln("y=%d", y);    // prints 7
    writefln("x=%d", x);    // prints 6
    writefln("F.x=%d", F.x); // prints 5
    writefln("B.x=%d", B.x); // prints 4
    F.func();               // calls Foo.func
    B.func();               // calls Bar.func
}

```

Alias declarations can be used to overload together functions declared in different mixins:

```

mixin template Foo()
{
    void func(int x) { }
}

mixin template Bar()
{
    void func(long x) { }
}

mixin Foo!() F;
mixin Bar!() B;

alias func = F.func;
alias func = B.func;

void main()
{
    func(1); // calls B.func
    func(1L); // calls F.func
}

```

A mixin has its own scope, even if a declaration is overridden by the enclosing one:

```

int x = 4;

mixin template Foo()
{
    int x = 5;
    int bar() { return x; }
}

```

```
}

mixin Foo;

void test()
{
    writefln("x_=%d", x);           // prints 4
    writefln("bar()_=%d", bar());   // prints 5
}
```

Chapter 23

Contract Programming

Contracts are a breakthrough technique to reduce the programming effort for large projects. Contracts are the concept of preconditions, postconditions, errors, and invariants. Contracts can be done in C++ without modification to the language, but the result is clumsy and inconsistent.

Building contract support into the language makes for:

1. a consistent look and feel for the contracts
2. tool support
3. it's possible the compiler can generate better code using information gathered from the contracts
4. easier management and enforcement of contracts
5. handling of contract inheritance

in</code>` or the `<code>out</code>` clause can be omitted. If the `<code>out</code>` clause is for a function body, the variable `<code>result</code>` is declared and assigned the return value of the function. For example, let's implement a square root function:

```
long square_root(long x)
in
{
    assert(x >= 0);
}
```

```

out (result)
{
    assert((result * result) <= x && (result+1) * (result+1) > x);
}
body
{
    return cast(long)std.math.sqrt(cast(real)x);
}

```

The assert's in the in and out bodies are called `<dfn>contracts</dfn>`. Any other D statement or expression is allowed in the bodies, but it is important to ensure that the code has no side effects, and that the release version of the code will not depend on any effects of the code. For a release build of the code, the in and out code is not inserted.

If the function returns a void, there is no result, and so there can be no result declaration in the out clause. In that case, use:

```

void func()
out
{
    ...contracts...
}
body
{
    ...
}

```

In an out statement, *result* is initialized and set to the return value of the function.

23.3 In, Out and Inheritance

If a function in a derived class overrides a function in its super class, then only one of the `in` contracts of the function and its base functions must be satisfied. Overriding functions then becomes a process of *loosening* the `in` contracts.

A function without an `in` contract means that any values of the function parameters are allowed. This implies that if any function in an inheritance hierarchy has no `in` contract, then `in` contracts on functions overriding it have no useful effect.

Conversely, all of the `out` contracts need to be satisfied, so overriding functions becomes a processes of *tightening* the `out` contracts.

23.4 Invariants

Invariants are used to specify characteristics of a class or struct that always must be true (except while executing a member function). For example, a class representing a date might have an invariant that the day must be 1..31 and the hour must be 0..23:

```
class Date
{
    int day;
    int hour;

    this(int d, int h)
    {
        day = d;
        hour = h;
    }

    invariant
    {
        assert(1 <= day && day <= 31);
        assert(0 <= hour && hour < 24);
    }
}
```

The invariant is a contract saying that the asserts must hold true. The invariant is checked when a class or struct constructor completes, at the start of the class or struct destructor. For public or exported functions, the order of execution is:

1. preconditions
2. invariant
3. body
4. invariant
5. postconditions

The invariant is not checked if the class or struct is implicitly constructed using the default `.init` value.

The code in the invariant may not call any public non-static members of the class or struct, either directly or indirectly. Doing so will result in a stack overflow, as the invariant will wind up being called in an infinitely recursive manner.

Invariants are implicitly `const`.

Since the invariant is called at the start of public or exported members, such members should not be called from constructors.

```

class Foo
{
    public void f() { }
    private void g() { }

    invariant
    {
        f(); // error, cannot call public member function from invariant
        g(); // ok, g() is not public
    }
}

```

The invariant can be checked with an `assert()` expression:

1. classes need to pass a class object
2. structs need to pass the address of an instance

```

auto mydate = new Date(); //class
auto s = S(); //struct
...
assert(mydate); // check that class Date invariant holds
assert(&s); // check that struct S invariant holds

```

Invariants contain `assert` expressions, and so when they fail, they throw a `AssertErrors`. Class invariants are inherited, that is, any class invariant is implicitly added with the invariants of its base classes.

There can be more than one `invariant` per class or struct.

When compiling for release, the invariant code is not generated, and the compiled program runs at maximum speed. The compiler is free to assume the invariant holds true, regardless of whether code is generated for it or not, and may optimize code accordingly.

23.5 References

- Contracts Reading List
- Adding Contracts to Java

Chapter 24

Conditional Compilation

Conditional compilation is the process of selecting which code to compile and which code to not compile. (In C and C++, conditional compilation is done with the preprocessor directives `#if` / `#else` / `#endif`.)

ConditionalDeclaration:

```
Condition DeclarationBlock
Condition DeclarationBlock else DeclarationBlock
Condition : DeclDefsopt
Condition DeclarationBlock else : DeclDefsopt
```

ConditionalStatement:

```
Condition NoScopeNonEmptyStatement
Condition NoScopeNonEmptyStatement else NoScopeNonEmptyStatement
```

If the *Condition* is satisfied, then the following *DeclarationBlock* or *Statement* is compiled in. If it is not satisfied, the *DeclarationBlock* or *Statement* after the optional **else** is compiled in.

Any *DeclarationBlock* or *Statement* that is not compiled in still must be syntactically correct.

No new scope is introduced, even if the *DeclarationBlock* or *Statement* is enclosed by { }.

ConditionalDeclarations and *ConditionalStatements* can be nested.

The *StaticAssert* can be used to issue errors at compilation time for branches of the conditional compilation that are errors.

Condition comes in the following forms:

Condition:

```
VersionCondition
DebugCondition
```

StaticIfCondition

Version Condition

```
VersionCondition:
    version ( IntegerLiteral )
    version ( Identifier )
    version ( unittest )
    version ( assert )
```

Versions enable multiple versions of a module to be implemented with a single source file.

The *VersionCondition* is satisfied if the *IntegerLiteral* is greater than or equal to the current *version level*, or if *Identifier* matches a *version identifier*.

The *version level* and *version identifier* can be set on the command line by the **-version** switch or in the module itself with a *VersionSpecification*, or they can be predefined by the compiler.

Version identifiers are in their own unique name space, they do not conflict with debug identifiers or other symbols in the module. Version identifiers defined in one module have no influence over other imported modules.

```
int k;
version (Demo) // compile in this code block for the demo version
{
    int i;
    int k;      // error, k already defined

    i = 3;
}
x = i;         // uses the i declared above

version (X86)
{
    ... // implement custom inline assembler version
}
else
{
    ... // use default, but slow, version
}
```

The `version(unittest)` is satisfied if and only if the code is compiled with unit tests enabled (the `-unittest` option on `dmd`).

Version Specification

```
VersionSpecification:
    version = Identifier ;
    version = IntegerLiteral ;
```

The version specification makes it straightforward to group a set of features under one major version, for example:

```
version (ProfessionalEdition)
{
    version = FeatureA;
    version = FeatureB;
    version = FeatureC;
}
version (HomeEdition)
{
    version = FeatureA;
}
...
version (FeatureB)
{
    ... implement Feature B ...
}
```

Version identifiers or levels may not be forward referenced:

```
version (Foo)
{
    int x;
}
version = Foo; // error, Foo already used
```

VersionSpecifications may only appear at module scope.

While the debug and version conditions superficially behave the same, they are intended for very different purposes. Debug statements are for adding debug code that is removed for the release version. Version statements are to aid in portability and multiple release versions.

Here's an example of a *full* version as opposed to a *demo* version:

```
class Foo
{
    int a, b;
```

```

version(full)
{
    int extrafunctionality()
    {
        ...
        return 1; // extra functionality is supported
    }
}
else // demo
{
    int extrafunctionality()
    {
        return 0; // extra functionality is not supported
    }
}
}

```

Various different version builds can be built with a parameter to version:

```

version(n) // add in version code if version level is >= n
{
    ... version code ...
}

version(identifier) // add in version code if version
                    // keyword is identifier
{
    ... version code ...
}

```

These are presumably set by the command line as `-version=n` and `-version=identifier`.

Predefined Versions

Several environmental version identifiers and identifier name spaces are predefined for consistent usage. Version identifiers do not conflict with other identifiers in the code, they are in a separate name space. Predefined version identifiers are global, i.e. they apply to all modules being compiled and imported.

Table 24.1: Predefined Version Identifiers

Version Identifier	Description
DigitalMars	DMD (Digital Mars D) is the compiler
GNU	GDC (GNU D Compiler) is the compiler
LDC	LDC (LLVM D Compiler) is the compiler
SDC	SDC (Stupid D Compiler) is the compiler
Windows	Microsoft Windows systems
Win32	Microsoft 32-bit Windows systems
Win64	Microsoft 64-bit Windows systems
linux	All Linux systems
OSX	Mac OS X
FreeBSD	FreeBSD
OpenBSD	OpenBSD
NetBSD	NetBSD
DragonFlyBSD	DragonFlyBSD
BSD	All other BSDs
Solaris	Solaris
Posix	All POSIX systems (includes Linux, FreeBSD, OS X, Solaris, etc.)
AIX	IBM Advanced Interactive eXecutive OS
Haiku	The Haiku operating system
SkyOS	The SkyOS operating system
SysV3	System V Release 3
SysV4	System V Release 4
Hurd	GNU Hurd
Android	The Android platform
PlayStation	The PlayStation platform
PlayStation4	The PlayStation 4 platform
Cygwin	The Cygwin environment
MinGW	The MinGW environment
FreeStanding	An environment without an operating system (such as Bare-metal targets)
CRuntime_Bionic	Bionic C runtime
CRuntime_DigitalMars	DigitalMars C runtime
CRuntime_Glibc	Glibc C runtime
CRuntime_Microsoft	Microsoft C runtime
X86	Intel and AMD 32-bit processors
X86_64	Intel and AMD 64-bit processors

(continued)

Version Identifier	Description
ARM	The ARM architecture (32-bit) (AArch32 et al)
ARM_Thumb	ARM in any Thumb mode
ARM_SoftFloat	The ARM soft floating point ABI
ARM_SoftFP	The ARM softfp floating point ABI
ARM_HardFloat	The ARM hardfp floating point ABI
AArch64	The Advanced RISC Machine architecture (64-bit)
Epiphany	The Epiphany architecture
PPC	The PowerPC architecture, 32-bit
PPC_SoftFloat	The PowerPC soft float ABI
PPC_HardFloat	The PowerPC hard float ABI
PPC64	The PowerPC architecture, 64-bit
IA64	The Itanium architecture (64-bit)
MIPS32	The MIPS architecture, 32-bit
MIPS64	The MIPS architecture, 64-bit
MIPS_O32	The MIPS O32 ABI
MIPS_N32	The MIPS N32 ABI
MIPS_O64	The MIPS O64 ABI
MIPS_N64	The MIPS N64 ABI
MIPS_EABI	The MIPS EABI
MIPS_SoftFloat	The MIPS soft-float ABI
MIPS_HardFloat	The MIPS hard-float ABI
NVPTX	The Nvidia Parallel Thread Execution (PTX) architecture, 32-bit
NVPTX64	The Nvidia Parallel Thread Execution (PTX) architecture, 64-bit
SPARC	The SPARC architecture, 32-bit
SPARC_V8Plus	The SPARC v8+ ABI
SPARC_SoftFloat	The SPARC soft float ABI
SPARC_HardFloat	The SPARC hard float ABI
SPARC64	The SPARC architecture, 64-bit
S390	The System/390 architecture, 32-bit
SystemZ	The System Z architecture, 64-bit
HPPA	The HP PA-RISC architecture, 32-bit
HPPA64	The HP PA-RISC architecture, 64-bit
SH	The SuperH architecture, 32-bit
SH64	The SuperH architecture, 64-bit

↩

(continued)

Version Identifier	Description
<code>Alpha</code>	The Alpha architecture
<code>Alpha_SoftFloat</code>	The Alpha soft float ABI
<code>Alpha_HardFloat</code>	The Alpha hard float ABI
<code>LittleEndian</code>	Byte order, least significant first
<code>BigEndian</code>	Byte order, most significant first
<code>ELFv1</code>	The Executable and Linkable Format v1
<code>ELFv2</code>	The Executable and Linkable Format v2
<code>D_Coverage</code>	Code coverage analysis instrumentation (command line switch <code>-cov</code>) is being generated
<code>D_Ddoc</code>	Ddoc documentation (command line switch <code>-D</code>) is being generated
<code>D_InlineAsm_X86</code>	Inline assembler for X86 is implemented
<code>D_InlineAsm_X86_64</code>	Inline assembler for X86-64 is implemented
<code>D_LP64</code>	Pointers are 64 bits (command line switch <code>-m64</code>). (Do not confuse this with C's LP64 model)
<code>D_X32</code>	Pointers are 32 bits, but words are still 64 bits (x32 ABI) (This can be defined in parallel to <code>X86_64</code>)
<code>D_HardFloat</code>	The target hardware has a floating point unit
<code>D_SoftFloat</code>	The target hardware does not have a floating point unit
<code>D_PIC</code>	Position Independent Code (command line switch <code>-fPIC</code>) is being generated
<code>D_SIMD</code>	Vector extensions (via <code>__simd</code>) are supported
<code>D_Version2</code>	This is a D version 2 compiler
<code>D_NoBoundsChecks</code>	Array bounds checks are disabled (command line switch <code>-boundscheck=off</code>)
<code>D_ObjectiveC</code>	The target supports interfacing with Objective-C
<code>unittest</code>	Unit tests are enabled (command line switch <code>-unittest</code>)
<code>assert</code>	Checks are being emitted for <i>AssertExpressions</i>
<code>none</code>	Never defined; used to just disable a section of code
<code>all</code>	Always defined; used as the opposite of <code>none</code>

The following identifiers are defined, but are deprecated:

Predefined Version Identifiers (deprecated)

Version Identifier	Description
<code>darwin</code>	The Darwin operating system; use <code>OSX</code> instead
<code>Thumb</code>	ARM in Thumb mode; use <code>ARM_Thumb</code> instead
<code>S390X</code>	The System/390X architecture 64-bit; use <code>SystemZ</code> instead

Others will be added as they make sense and new implementations appear.

It is inevitable that the D language will evolve over time. Therefore, the version identifier namespace beginning with "D_" is reserved for identifiers indicating D language specification or new feature conformance. Further, all identifiers derived from the ones listed above by appending any character(s) are reserved. This means that e.g. `ARM_foo` and `Windows_bar` are reserved while `foo_ARM` and `bar_Windows` are not.

Furthermore, predefined version identifiers from this list cannot be set from the command line or from version statements. (This prevents things like both `Windows` and `linux` being simultaneously set.)

Compiler vendor specific versions can be predefined if the trademarked vendor identifier prefixes it, as in:

```
version(DigitalMars_funky_extension)
{
    ...
}
```

It is important to use the right version identifier for the right purpose. For example, use the vendor identifier when using a vendor specific feature. Use the operating system identifier when using an operating system specific feature, etc.

24.1 Debug Condition

DebugCondition:

```
debug
debug ( IntegerLiteral )
debug ( Identifier )
```

Two versions of programs are commonly built, a release build and a debug build. The debug build includes extra error checking code, test harnesses, pretty-printing code, etc. The `debug` statement conditionally compiles in its statement body. It is D's way of what in C is done with `#ifdef DEBUG` / `#endif` pairs.

The `debug` condition is satisfied when the `-debug` switch is passed to the compiler or when the debug level is ≥ 1 .

The `debug` (*IntegerLiteral*) condition is satisfied when the debug level is \geq *IntegerLiteral*.

The `debug` (*Identifier*) condition is satisfied when the debug identifier matches *Identifier*.

```
class Foo
{
    int a, b;
debug:
    int flag;
}
```

Debug Specification

```
DebugSpecification:
    debug = Identifier ;
    debug = IntegerLiteral ;
```

Debug identifiers and levels are set either by the command line switch `-debug` or by a *DebugSpecification*.

Debug specifications only affect the module they appear in, they do not affect any imported modules. Debug identifiers are in their own namespace, independent from version identifiers and other symbols.

It is illegal to forward reference a debug specification:

```
debug(foo) writeln("Foo");
debug = foo;    // error, foo used before set
```

DebugSpecifications may only appear at module scope.

Various different debug builds can be built with a parameter to debug:

```
debug(IntegerLiteral) { } // add in debug code if debug level is >= ↵
    ↵ IntegerLiteral
debug(identifier) { } // add in debug code if debug keyword is identifier
```

These are presumably set by the command line as `-debug=n` and `-debug=identifier`.

24.2 Static If Condition

```
StaticIfCondition:
    static if ( AssignExpression )
```

AssignExpression is implicitly converted to a boolean type, and is evaluated at compile time. The condition is satisfied if it evaluates to `true`. It is not satisfied if it evaluates to `false`.

It is an error if *AssignExpression* cannot be implicitly converted to a boolean type or if it cannot be evaluated at compile time.

StaticIfConditions can appear in module, class, template, struct, union, or function scope. In function scope, the symbols referred to in the *AssignExpression* can be any that can normally be referenced by an expression at that point.

```
const int i = 3;
int j = 4;

static if (i == 3)    // ok, at module scope
    int x;

class C
{
    const int k = 5;

    static if (i == 3) // ok
        int x;
    else
        long x;

    static if (j == 3) // error, j is not a constant
        int y;

    static if (k == 5) // ok, k is in current scope
        int z;
}

template INT(int i)
{
    static if (i == 32)
        alias INT = int;
    else static if (i == 16)
        alias INT = short;
    else
        static assert(0); // not supported
}
```

```

INT!(32) a;  // a is an int
INT!(16) b;  // b is a short
INT!(17) c;  // error, static assert trips

```

A *StaticIfConditional* condition differs from an *IfStatement* in the following ways:

1. It can be used to conditionally compile declarations, not just statements.
2. It does not introduce a new scope even if { } are used for conditionally compiled statements.
3. For unsatisfied conditions, the conditionally compiled code need only be syntactically correct. It does not have to be semantically correct.
4. It must be evaluatable at compile time.

24.3 Static Assert

StaticAssert:

```

static assert ( AssignExpression ,opt );
static assert ( AssignExpression , AssignExpression ,opt );

```

AssignExpression is evaluated at compile time, and converted to a boolean value. If the value is true, the static assert is ignored. If the value is false, an error diagnostic is issued and the compile fails.

Unlike *AssertExpressions*, *StaticAsserts* are always checked and evaluated by the compiler unless they appear in an unsatisfied conditional.

```

void foo()
{
    if (0)
    {
        assert(0);  // never trips
        static assert(0);  // always trips
    }
    version (BAR)
    {
    }
    else
    {
        static assert(0);  // trips when version BAR is not defined
    }
}

```

StaticAssert is useful tool for drawing attention to conditional configurations not supported in the code.

The optional second *AssignExpression* can be used to supply additional information, such as a text string, that will be printed out along with the error diagnostic.

Chapter 25

Traits

Traits are extensions to the language to enable programs, at compile time, to get at information internal to the compiler. This is also known as compile time reflection. It is done as a special, easily extended syntax (similar to Pragmas) so that new capabilities can be added as required.

TraitsExpression:

```
__traits ( TraitsKeyword , TraitsArguments )
```

TraitsKeyword:

```
isAbstractClass  
isArithmetic  
isAssociativeArray  
isFinalClass  
isPOD  
isNested  
isFloating  
isIntegral  
isScalar  
isStaticArray  
isUnsigned  
isVirtualFunction  
isVirtualMethod  
isAbstractFunction  
isFinalFunction  
isStaticFunction  
isOverrideFunction  
isTemplate  
isRef
```

```

isOut
isLazy
hasMember
identifier
getAliasThis
getAttributes
getFunctionAttributes
getMember
getOverloads
getPointerBitmap
getProtection
getVirtualFunctions
getVirtualMethods
getUnitTests
parent
classInstanceSize
getVirtualIndex
allMembers
derivedMembers
isSame
compiles

```

```

TraitsArguments:
  TraitsArgument
  TraitsArgument , TraitsArguments

```

```

TraitsArgument:
  AssignExpression
  Type

```

Additionally special keywords are provided for debugging purposes:

```

SpecialKeyword:
  __FILE__
  __FILE_FULL_PATH__
  __MODULE__
  __LINE__
  __FUNCTION__
  __PRETTY_FUNCTION__

```

25.1 *isArithmetic*

If the arguments are all either types that are arithmetic types, or expressions that are typed as arithmetic types, then `true` is returned. Otherwise, `false` is returned. If there are no arguments, `false` is returned.

```
import std.stdio;

void main()
{
    int i;
    writeln(__traits(isArithmetic, int));
    writeln(__traits(isArithmetic, i, i+1, int));
    writeln(__traits(isArithmetic));
    writeln(__traits(isArithmetic, int*));
}
```

Prints:

```
true true false false
```

25.2 *isFloating*

Works like `isArithmetic`, except it's for floating point types (including imaginary and complex types).

25.3 *isIntegral*

Works like `isArithmetic`, except it's for integral types (including character types).

25.4 *isScalar*

Works like `isArithmetic`, except it's for scalar types.

25.5 *isUnsigned*

Works like `isArithmetic`, except it's for unsigned types.

25.6 *isStaticArray*

Works like `isArithmetic`, except it's for static array types.

25.7 *isAssociativeArray*

Works like `isArithmetic`, except it's for associative array types.

25.8 *isAbstractClass*

If the arguments are all either types that are abstract classes, or expressions that are typed as abstract classes, then `true` is returned. Otherwise, `false` is returned. If there are no arguments, `false` is returned.

```
import std.stdio;

abstract class C { int foo(); }

void main()
{
    C c;
    writeln(__traits(isAbstractClass, C));
    writeln(__traits(isAbstractClass, c, C));
    writeln(__traits(isAbstractClass));
    writeln(__traits(isAbstractClass, int*));
}

Prints:
true true false false
```

25.9 *isFinalClass*

Works like `isAbstractClass`, except it's for final classes.

25.10 *isPOD*

Takes one argument, which must be a type. It returns `true` if the type is a POD type, otherwise `false`.

25.11 *isNested*

Takes one argument. It returns `true` if the argument is a nested type which internally stores a context pointer, otherwise it returns `false`. Nested types can be classes, structs, and functions.

25.12 *isVirtualFunction*

The same as *isVirtualMethod*, except that final functions that don't override anything return true.

25.13 *isVirtualMethod*

Takes one argument. If that argument is a virtual function, `true` is returned, otherwise `false`. Final functions that don't override anything return false.

```
import std.stdio;

struct S
{
    void bar() { }
}

class C
{
    void bar() { }
}

void main()
{
    writeln(__traits(isVirtualMethod, C.bar)); // true
    writeln(__traits(isVirtualMethod, S.bar)); // false
}
```

25.14 *isAbstractFunction*

Takes one argument. If that argument is an abstract function, `true` is returned, otherwise `false`.

```
import std.stdio;

struct S
{
```

```

    void bar() { }
}

class C
{
    void bar() { }
}

class AC
{
    abstract void foo();
}

void main()
{
    writeln(__traits(isAbstractFunction, C.bar));    // false
    writeln(__traits(isAbstractFunction, S.bar));    // false
    writeln(__traits(isAbstractFunction, AC.foo));    // true
}

```

25.15 *isFinalFunction*

Takes one argument. If that argument is a final function, `true` is returned, otherwise `false`.

```

import std.stdio;

struct S
{
    void bar() { }
}

class C
{
    void bar() { }
    final void foo();
}

final class FC
{
    void foo();
}

```

```

}

void main()
{
    writeln(__traits(isFinalFunction, C.bar)); // false
    writeln(__traits(isFinalFunction, S.bar)); // false
    writeln(__traits(isFinalFunction, C.foo)); // true
    writeln(__traits(isFinalFunction, FC.foo)); // true
}

```

25.16 *isOverrideFunction*

Takes one argument. If that argument is a function marked with `override`, `true` is returned, otherwise `false`.

```

import std.stdio;

class Base
{
    void foo() { }
}

class Foo : Base
{
    override void foo() { }
    void bar() { }
}

void main()
{
    writeln(__traits(isOverrideFunction, Base.foo)); // false
    writeln(__traits(isOverrideFunction, Foo.foo)); // true
    writeln(__traits(isOverrideFunction, Foo.bar)); // false
}

```

25.17 *isStaticFunction*

Takes one argument. If that argument is a static function, meaning it has no context pointer, `true` is returned, otherwise `false`.

25.18 *isRef, isOut, isLazy*

Takes one argument. If that argument is a declaration, `true` is returned if it is ref, out, or lazy, otherwise `false`.

```
void fooref(ref int x)
{
    static assert(__traits(isRef, x));
    static assert(!__traits(isOut, x));
    static assert(!__traits(isLazy, x));
}

void fooout(out int x)
{
    static assert(!__traits(isRef, x));
    static assert(__traits(isOut, x));
    static assert(!__traits(isLazy, x));
}

void foolazy(lazy int x)
{
    static assert(!__traits(isRef, x));
    static assert(!__traits(isOut, x));
    static assert(__traits(isLazy, x));
}
```

25.19 *isTemplate*

Takes one argument. If that argument is a template then `true` is returned, otherwise `false`.

```
void foo(T){}
static assert(__traits(isTemplate,foo));
static assert(!__traits(isTemplate,foo!int()));
static assert(!__traits(isTemplate,"string"));
```

25.20 *hasMember*

The first argument is a type that has members, or is an expression of a type that has members. The second argument is a string. If the string is a valid property of the type, `true` is returned, otherwise `false`.

```
import std.stdio;

struct S
{
    int m;
}

void main()
{
    S s;

    writeln(__traits(hasMember, S, "m")); // true
    writeln(__traits(hasMember, s, "m")); // true
    writeln(__traits(hasMember, S, "y")); // false
    writeln(__traits(hasMember, int, "sizeof")); // true
}
```

25.21 *identifier*

Takes one argument, a symbol. Returns the identifier for that symbol as a string literal.

25.22 *getAliasThis*

Takes one argument, a symbol of aggregate type. If the given aggregate type has alias `this`, returns a list of alias `this` names, by a tuple of `strings`. Otherwise returns an empty tuple.

25.23 *getAttributes*

Takes one argument, a symbol. Returns a tuple of all attached user defined attributes. If no UDA's exist it will return an empty tuple.

For more information, see: User Defined Attributes

```
@(3) int a;
@("string", 7) int b;

enum Foo;
@Foo int c;

pragma(msg, __traits(getAttributes, a));
```

```
pragma(msg, __traits(getAttributes, b));
pragma(msg, __traits(getAttributes, c));
```

Prints:

```
tuple(3) tuple("string", 7) tuple((Foo))
```

25.24 *getFunctionAttributes*

Takes one argument which must either be a function symbol, function literal, or a function pointer. It returns a string tuple of all the attributes of that function **excluding** any user defined attributes (UDAs can be retrieved with the `getAttributes` trait). If no attributes exist it will return an empty tuple.

Note: The order of the attributes in the returned tuple is implementation-defined and should not be relied upon.

A list of currently supported attributes are:

- `pure`, `nothrow`, `@nogc`, `@property`, `@system`, `@trusted`, `@safe`, and `ref`

Note: `ref` is a function attribute even though it applies to the return type.

Additionally the following attributes are only valid for non-static member functions:

- `const`, `immutable`, `inout`, `shared`

For example:

```
int sum(int x, int y) pure nothrow { return x + y; }
```

```
// prints ("pure", "nothrow", "@system")
pragma(msg, __traits(getFunctionAttributes, sum));
```

```
struct S
{
    void test() const @system { }
```

```
// prints ("const", "@system")
pragma(msg, __traits(getFunctionAttributes, S.test));
```

Note that some attributes can be inferred. For example:

```
// prints ("pure", "nothrow", "@nogc", "@trusted")
pragma(msg, __traits(getFunctionAttributes, (int x) @trusted { return x * 2; }));
```

25.25 *getMember*

Takes two arguments, the second must be a string. The result is an expression formed from the first argument, followed by a '.', followed by the second argument as an identifier.

```
import std.stdio;

struct S
{
    int mx;
    static int my;
}

void main()
{
    S s;

    __traits(getMember, s, "mx") = 1;  // same as s.mx=1;
    writeln(__traits(getMember, s, "m" ~ "x")); // 1

    __traits(getMember, S, "mx") = 1;  // error, no this for S.mx
    __traits(getMember, S, "my") = 2;  // ok
}
```

25.26 *getOverloads*

The first argument is an aggregate (e.g. struct/class/module). The second argument is a string that matches the name of one of the functions in that aggregate. The result is a tuple of all the overloads of that function.

```
import std.stdio;

class D
{
    this() { }
    ~this() { }
    void foo() { }
    int foo(int) { return 2; }
}
```

```

void main()
{
    D d = new D();

    foreach (t; __traits(getOverloads, D, "foo"))
        writeln(typeid(typeof(t)));

    alias b = typeof(__traits(getOverloads, D, "foo"));
    foreach (t; b)
        writeln(typeid(t));

    auto i = __traits(getOverloads, d, "foo")[1](1);
    writeln(i);
}

```

Prints:

```
void() int() void() int() 2
```

25.27 *getPointerBitmap*

The argument is a type. The result is an array of `size_t` describing the memory used by an instance of the given type.

The first element of the array is the size of the type (for classes it is the `classInstanceSize`).

The following elements describe the locations of GC managed pointers within the memory occupied by an instance of the type. For type `T`, there are `T.sizeof / size_t.sizeof` possible pointers represented by the bits of the array values.

This array can be used by a precise GC to avoid false pointers.

```

class C
{
    // implicit virtual function table pointer not marked
    // implicit monitor field not marked, usually managed manually
    C next;
    size_t sz;
    void* p;
    void function () fn; // not a GC managed pointer
}

struct S
{
    size_t val1;
}

```



```

    void* p;
    C c;
    byte[] arr;           // { length, ptr }
    void delegate () dg; // { context, func }
}

static assert (__traits(getPointerBitmap, C) == [6*size_t.sizeof, 0b010100]);
static assert (__traits(getPointerBitmap, S) == [7*size_t.sizeof, 0b0110110]);

```

25.28 *getProtection*

The argument is a symbol. The result is a string giving its protection level: "public", "private", "protected", "export", or "package".

```

import std.stdio;

class D
{
    export void foo() { }
    public int bar;
}

void main()
{
    D d = new D();

    auto i = __traits(getProtection, d.foo);
    writeln(i);

    auto j = __traits(getProtection, d.bar);
    writeln(j);
}

Prints:
export public

```

25.29 *getVirtualFunctions*

The same as *getVirtualMethods*, except that final functions that do not override anything are included.

25.30 *getVirtualMethods*

The first argument is a class type or an expression of class type. The second argument is a string that matches the name of one of the functions of that class. The result is a tuple of the virtual overloads of that function. It does not include final functions that do not override anything.

```
import std.stdio;

class D
{
    this() { }
    ~this() { }
    void foo() { }
    int foo(int) { return 2; }
}

void main()
{
    D d = new D();

    foreach (t; __traits(getVirtualMethods, D, "foo"))
        writeln(typeid(t));

    alias b = typeid(__traits(getVirtualMethods, D, "foo"));
    foreach (t; b)
        writeln(typeid(t));

    auto i = __traits(getVirtualMethods, d, "foo")[1](1);
    writeln(i);
}
```

Prints:

```
void() int() void() int() 2
```

25.31 *getUnitTests*

Takes one argument, a symbol of an aggregate (e.g. struct/class/module). The result is a tuple of all the unit test functions of that aggregate. The functions returned are like normal nested static functions, CTFE will work and UDA's will be accessible.

Note:

The `-unittest` flag needs to be passed to the compiler. If the flag is not passed `__traits(↵ getUnitTests)` will always return an empty tuple.

```
module foo;

import core.runtime;
import std.stdio;

struct name { string name; }

class Foo
{
    unittest
    {
        writeln("foo.Foo.unittest");
    }
}

@name("foo") unittest
{
    writeln("foo.unittest");
}

template Tuple (T...)
{
    alias Tuple = T;
}

shared static this()
{
    // Override the default unit test runner to do nothing. After that, "main" will
    // be called.
    Runtime.moduleUnitTester = { return true; };
}

void main()
{
    writeln("start_↵main");
}
```

```

alias tests = Tuple!(__traits(getUnitTests, foo));
static assert(tests.length == 1);

alias attributes = Tuple!(__traits(getAttributes, tests[0]));
static assert(attributes.length == 1);

foreach (test; tests)
    test();

foreach (test; __traits(getUnitTests, Foo))
    test();
}

```

By default, the above will print:

```
start main foo.unittest foo.Foo.unittest
```

25.32 *parent*

Takes a single argument which must evaluate to a symbol. The result is the symbol that is the parent of it.

25.33 *classInstanceSize*

Takes a single argument, which must evaluate to either a class type or an expression of class type. The result is of type `size_t`, and the value is the number of bytes in the runtime instance of the class type. It is based on the static type of a class, not the polymorphic type.

25.34 *getVirtualIndex*

Takes a single argument which must evaluate to a function. The result is a `ptrdiff_t` containing the index of that function within the vtable of the parent type. If the function passed in is final and does not override a virtual function, `-1` is returned instead.

25.35 *allMembers*

Takes a single argument, which must evaluate to either a type or an expression of type. A tuple of string literals is returned, each of which is the name of a member of that type combined with all of the members of the base classes (if the type is a class). No name is repeated. Builtin properties are not included.

```
import std.stdio;

class D
{
    this() { }
    ~this() { }
    void foo() { }
    int foo(int) { return 0; }
}

void main()
{
    auto b = [ __traits(allMembers, D) ];
    writeln(b);
    // ["__ctor", "__dtor", "foo", "toString", "toHash", "opCmp", "opEquals", "✓
    ↪ Monitor", "factory"]
}
```

The order in which the strings appear in the result is not defined.

25.36 *derivedMembers*

Takes a single argument, which must evaluate to either a type or an expression of type. A tuple of string literals is returned, each of which is the name of a member of that type. No name is repeated. Base class member names are not included. Builtin properties are not included.

```
import std.stdio;

class D
{
    this() { }
    ~this() { }
    void foo() { }
    int foo(int) { return 0; }
}

void main()
{
    auto a = [__traits(derivedMembers, D)];
    writeln(a);    // ["__ctor", "__dtor", "foo"]
}
```

The order in which the strings appear in the result is not defined.

25.37 *isSame*

Takes two arguments and returns bool `true` if they are the same symbol, `false` if not.

```
import std.stdio;

struct S { }

int foo();
int bar();

void main()
{
    writeln(__traits(isSame, foo, foo)); // true
    writeln(__traits(isSame, foo, bar)); // false
    writeln(__traits(isSame, foo, S));   // false
    writeln(__traits(isSame, S, S));     // true
    writeln(__traits(isSame, std, S));   // false
    writeln(__traits(isSame, std, std)); // true
}
```

If the two arguments are expressions made up of literals or enums that evaluate to the same value, true is returned.

25.38 *compiles*

Returns a bool `true` if all of the arguments compile (are semantically correct). The arguments can be symbols, types, or expressions that are syntactically correct. The arguments cannot be statements or declarations.

If there are no arguments, the result is `false`.

```
import std.stdio;

struct S
{
    static int s1;
    int s2;
}
```

```

int foo();
int bar();

void main()
{
    writeln(__traits(compiles));           // false
    writeln(__traits(compiles, foo));      // true
    writeln(__traits(compiles, foo + 1));  // true
    writeln(__traits(compiles, &foo + 1)); // false
    writeln(__traits(compiles, typeof(1))); // true
    writeln(__traits(compiles, S.s1));      // true
    writeln(__traits(compiles, S.s3));      // false
    writeln(__traits(compiles, 1,2,3,int,long,std)); // true
    writeln(__traits(compiles, 3[1]));      // false
    writeln(__traits(compiles, 1,2,3,int,long,3[1])); // false
}

```

This is useful for:

- Giving better error messages inside generic code than the sometimes hard to follow compiler ones.
- Doing a finer grained specialization than template partial specialization allows for.

25.39 Special Keywords

`__FILE__` and `__LINE__` expand to the source file name and line number at the point of instantiation. The path of the source file is left up to the compiler.

`__FILE_FULL_PATH__` expands to the absolute source file name at the point of instantiation.

`__MODULE__` expands to the module name at the point of instantiation.

`__FUNCTION__` expands to the fully qualified name of the function at the point of instantiation.

`__PRETTY_FUNCTION__` is similar to `__FUNCTION__`, but also expands the function return type, its parameter types, and its attributes.

Example usage:

```

module test;
import std.stdio;

void test(string file = __FILE__, size_t line = __LINE__, string mod = __MODULE__
    ↪ ,
        string func = __FUNCTION__, string pretty = __PRETTY_FUNCTION__, string
    ↪ fileFullPath = __FILE_FULL_PATH__)

```

```
{
    writefln("file:_%'s',_line:_%'s',_module:_%'s',\nfunction:_%'s',_pretty_↵
        ↵ function:_%'s',\nfile_full_path:_%'s'",
            file, line, mod, func, pretty, fileFullPath);
}

int main(string[] args)
{
    test();
    return 0;
}
```

Assuming the file was at /example/test.d, this will output:

```
file:  'test.d', line:  '13', module:  'test', function:  'test.main', pretty
function:  'int test.main(string[] args)', file full path:  '/example/test.d'
```


Chapter 26

Error Handling

I came, I coded, I crashed.

– *Julius C'ster*

All programs have to deal with errors. Errors are unexpected conditions that are not part of the normal operation of a program. Examples of common errors are:

- Out of memory.
- Out of disk space.
- Invalid file name.
- Attempting to write to a read-only file.
- Attempting to read a non-existent file.
- Requesting a system service that is not supported.

26.1 The Error Handling Problem

The traditional C way of detecting and reporting errors is not traditional, it is ad-hoc and varies from function to function, including:

- Returning a NULL pointer.
- Returning a 0 value.
- Returning a non-zero error code.
- Requiring `errno` to be checked.
- Requiring that a function be called to check if the previous function failed.

To deal with these possible errors, tedious error handling code must be added to each function call. If an error happened, code must be written to recover from the error, and the error must be

reported to the user in some user friendly fashion. If an error cannot be handled locally, it must be explicitly propagated back to its caller. The long list of `errno` values needs to be converted into appropriate text to be displayed. Adding all the code to do this can consume a large part of the time spent coding a project - and still, if a new `errno` value is added to the runtime system, the old code can not properly display a meaningful error message.

Good error handling code tends to clutter up what otherwise would be a neat and clean looking implementation.

Even worse, good error handling code is itself error prone, tends to be the least tested (and therefore buggy) part of the project, and is frequently simply omitted. The end result is likely a "blue screen of death" as the program failed to deal with some unanticipated error.

Quick and dirty programs are not worth writing tedious error handling code for, and so such utilities tend to be like using a table saw with no blade guards.

What's needed is an error handling philosophy and methodology such that:

- It is standardized - consistent usage makes it more useful.
- The result is reasonable even if the programmer fails to check for errors.
- Old code can be reused with new code without having to modify the old code to be compatible with new error types.
- No errors get inadvertently ignored.
- 'Quick and dirty' utilities can be written that still correctly handle errors.
- It is easy to make the error handling source code look good.

26.2 The D Error Handling Solution

Let's first make some observations and assumptions about errors:

- Errors are not part of the normal flow of a program. Errors are exceptional, unusual, and unexpected.
- Because errors are unusual, execution of error handling code is not performance critical.
- The normal flow of program logic is performance critical.
- All errors must be dealt with in some way, either by code explicitly written to handle them, or by some system default handling.
- The code that detects an error knows more about the error than the code that must recover from the error.

The solution is to use exception handling to report errors. All errors are objects derived from abstract class **Error**. **Error** has a pure virtual function called `toString()` which produces a `char[]` with a human readable description of the error.

If code detects an error like "out of memory," then an Error is thrown with a message saying "Out of memory". The function call stack is unwound, looking for a handler for the Error. Finally blocks are executed as the stack is unwound. If an error handler is found, execution resumes there. If not, the default Error handler is run, which displays the message and terminates the program.

How does this meet our criteria?

It is standardized - consistent usage makes it more useful. This is the D way, and is used consistently in the D runtime library and examples.

The result is reasonable result even if the programmer fails to check for errors. If no catch handlers are there for the errors, then the program gracefully exits through the default error handler with an appropriate message.

Old code can be reused with new code without having to modify the old code to be compatible with new error types. Old code can decide to catch all errors, or only specific ones, propagating the rest upwards. In any case, there is no more need to correlate error numbers with messages, the correct message is always supplied.

No errors get inadvertently ignored. Error exceptions get handled one way or another. There is nothing like a NULL pointer return indicating an error, followed by trying to use that NULL pointer.

'Quick and dirty' utilities can be written that still correctly handle errors. Quick and dirty code need not write any error handling code at all, and don't need to check for errors. The errors will be caught, an appropriate message displayed, and the program gracefully shut down all by default.

It is easy to make the error handling source code look good. The try/catch/finally statements look a lot nicer than endless if (error) goto errorhandler; statements.

How does this meet our assumptions about errors?

Errors are not part of the normal flow of a program. Errors are exceptional, unusual, and unexpected. D exception handling fits right in with that.

Because errors are unusual, execution of error handling code is not performance critical. Exception handling stack unwinding is a relatively slow process.

The normal flow of program logic is performance critical. Since the normal flow code does not have to check every function call for error returns, it can be realistically faster to use exception handling for the errors.

All errors must be dealt with in some way, either by code explicitly written to handle them, or by some system default handling. If there's no handler for a particular error, it is handled by the runtime library default handler. If an error is ignored, it is because the programmer specifically added code to ignore an error, which presumably means it was intentional.

The code that detects an error knows more about the error than the code that must recover from the error. There is no more need to translate error codes into human readable strings, the correct string is generated by the error detection code, not the error recovery code. This also leads to consistent error messages for the same error between applications.

Using exceptions to handle errors leads to another issue - how to write exception safe programs. Here's how.

Chapter 27

Unit Tests

```
UnitTest:  
    unittest BlockStatement
```

Unit tests are a series of test cases applied to a module to determine if it is working properly. Ideally, unit tests should be run every time a program is compiled.

Unit tests are a special function defined like:

```
unittest  
{  
    ...test code...  
}
```

There can be any number of unit test functions in a module, including within struct, union and class declarations. They are executed in lexical order. The order in which the modules are called to run their unit tests is implementation defined. Stylistically, a unit test for a function should appear immediately following it.

A compiler switch, such as **-unittest** for **dmd**, will cause the unittest test code to be compiled and incorporated into the resulting executable. The unittest code gets run after static initialization is run and before the **main()** function is called.

UnitTests must be grammatically correct even if **-unittest** is not used.

For example, given a class **Sum** that is used to add two values, a unit test can be given:

```
class Sum  
{  
    int add(int x, int y) { return x + y; }  
  
    unittest  
    {
```

```

        Sum sum = new Sum;
        assert(sum.add(3,4) == 7);
        assert(sum.add(-2,0) == -2);
    }
}

```

Attributed Unittests

A unittest may be attributed with any of the global function attributes. Such unittests are useful in verifying the given attribute(s) on a template function:

```

void myFunc(T)(T[] data)
{
    if (data.length > 2)
        data[0] = data[1];
}

@safe nothrow unittest
{
    auto arr = [1,2,3];
    myFunc(arr);
    assert(arr == [2,2,3]);
}

```

This unittest verifies that `myFunc` contains only `@safe`, `nothrow` code. Although this can also be accomplished by attaching these attributes to `myFunc` itself, that would prevent `myFunc` from being instantiated with types `T` that have `@system` or throwing code in their `opAssign` method, or other methods that `myFunc` may call. The above idiom allows `myFunc` to be instantiated with such types, yet at the same time verify that the `@system` and throwing behaviour is not introduced by the code within `myFunc` itself.

Documented Unittests

Documented unittests allow the developer to deliver code examples to the user, while at the same time automatically verifying that the examples are valid. This avoids the frequent problem of having outdated documentation for some piece of code.

If a declaration is followed by a documented unittest, the code in the unittest will be inserted in the **example** section of the declaration:

```

/// Math class
class Math
{

```

```

    /// add function
    static int add(int x, int y) { return x + y; }

    ///
    unittest
    {
        assert(add(2, 2) == 4);
    }
}

///
unittest
{
    auto math = new Math();
    auto result = math.add(2, 2);
}

```

The above will generate the following documentation:

```

<dl><dt><big><a name="Math"></a>class <u>Math</u>;
</big></dt>    <dd><u>Math</u>    class<br><br>    <b>Example:</b><pre
class="d_code"><font color="blue">auto</font> math = <font color="blue">new</font>
<u>Math</u>; <font color="blue">auto</font> result = math.add(2, 2); </pre><br>
    <dl><dt><big><a name="Math.add"></a>int <u>add</u>(int <i>x</i>, int
<i>y</i>); </big></dt>    <dd><u>add</u> function<br><br>    <b>Example:</b><pre
class="d_code"><font color="blue">assert</font>(<u>add</u>(2, 2) == 4); </pre> </dd>
</dl> </dd> </dl>

```

A unittest which is not documented, or is marked as private will not be used to generate code samples.

There can be multiple documented unittests and they can appear in any order. They will be attached to the last non-unittest declaration:

```

/// add function
int add(int x, int y) { return x + y; }

/// code sample generated
unittest
{
    assert(add(1, 1) == 2);
}

/// code sample not generated because the unittest is private

```

```

private unittest
{
    assert(add(2, 2) == 4);
}

unittest
{
    /// code sample not generated because the unittest isn't documented
    assert(add(3, 3) == 6);
}

/// code sample generated, even if it only includes comments (or is empty)
unittest
{
    /** assert(add(4, 4) == 8); */
}

```

The above will generate the following documentation:

```

<dl><dt><big><a name="add"></a>int <u>add</u>(int <i>x</i>, int <i>y</i>);
</big></dt> <dd><u>add</u> function<br><br> <b>Examples:</b><br> code sam-
ple generated <pre class="d_code"> <font color="blue">assert</font>(<u>add</u>(1, 1) ==
2); </pre> <br><br><b>Examples:</b><br> code sample generated, even if it is empty or only
includes comments
<pre class="d_code"> <font color="green">/** assert(add(4, 4) == 8); */</font> </pre>
<br><br> </dd> </dl>

```

Versioning

The version identifier `unittest` is predefined if the compilation is done with unit tests enabled.

Chapter 28

Garbage Collection

D is a systems programming language with support for garbage collection. Usually it is not necessary to free memory explicitly. Just allocate as needed, and the garbage collector will periodically return all unused memory to the pool of available memory.

D also provides the mechanisms to write code where the garbage collector is **not involved**. More information is provided below.

C and C++ programmers accustomed to explicitly managing memory allocation and deallocation will likely be skeptical of the benefits and efficacy of garbage collection. Experience both with new projects written with garbage collection in mind, and converting existing projects to garbage collection shows that:

- Garbage collected programs are often faster. This is counterintuitive, but the reasons are:
 - Reference counting is a common solution to solve explicit memory allocation problems. The code to implement the increment and decrement operations whenever assignments are made is one source of slowdown. Hiding it behind smart pointer classes doesn't help the speed. (Reference counting methods are not a general solution anyway, as circular references never get deleted.)
 - Destructors are used to deallocate resources acquired by an object. For most classes, this resource is allocated memory. With garbage collection, most destructors then become empty and can be discarded entirely.
 - All those destructors freeing memory can become significant when objects are allocated on the stack. For each one, some mechanism must be established so that if an exception happens, the destructors all get called in each frame to release any memory they hold. If the destructors become irrelevant, then there's no need to set up special stack frames to handle exceptions, and the code runs faster.
 - Garbage collection kicks in only when memory gets tight. When memory is not tight, the program runs at full speed and does not spend any time tracing and freeing memory.

- Garbage collected programs do not suffer from gradual deterioration due to an accumulation of memory leaks.
- Garbage collectors reclaim unused memory, therefore they do not suffer from "memory leaks" which can cause long running applications to gradually consume more and more memory until they bring down the system. GC programs have longer term stability.
- Garbage collected programs have fewer hard-to-find pointer bugs. This is because there are no dangling references to freed memory. There is no code to explicitly manage memory, hence no bugs in such code.
- Garbage collected programs are faster to develop and debug, because there's no need for developing, debugging, testing, or maintaining the explicit deallocation code.

Garbage collection is not a panacea. There are some downsides:

- It is not always obvious when the GC allocates memory, which in turn can trigger a collection, so the program can pause unexpectedly.
- The time it takes for a collection to complete is not bounded. While in practice it is very quick, this cannot normally be guaranteed.
- Normally, all threads other than the collector thread must be halted while the collection is in progress.
- Garbage collectors can keep around some memory that an explicit deallocator would not.
- Garbage collection should be implemented as a basic operating system kernel service. But since it is not, garbage collecting programs must carry around with them the garbage collection implementation. While this can be a shared library, it is still there.

These constraints are addressed by techniques outlined in Memory Management, including the mechanisms provided by D to control allocations outside the GC heap.

There is currently work in progress to make the runtime library free of GC heap allocations, to allow its use in scenarios where the use of GC infrastructure is not possible.

28.1 How Garbage Collection Works

The GC works by:

1. Stopping all other threads than the thread currently trying to allocate GC memory.
2. 'Hijacking' the current thread for GC work.
3. Scanning all 'root' memory ranges for pointers into GC allocated memory.
4. Recursively scanning all allocated memory pointed to by roots looking for more pointers into GC allocated memory.

5. Freeing all GC allocated memory that has no active pointers to it and do not need destructors to run.
6. Queueing all unreachable memory that needs destructors to run.
7. Resuming all other threads.
8. Running destructors for all queued memory.
9. Freeing any remaining unreachable memory.
10. Returning the current thread to whatever work it was doing.

28.2 Interfacing Garbage Collected Objects With Foreign Code

The garbage collector looks for roots in:

1. the static data segment
2. the stacks and register contents of each thread
3. the TLS (thread-local storage) areas of each thread
4. any roots added by `core.memory.GC.addRoot()` or `core.memory.GC.addRange()`

If the only pointer to an object is held outside of these areas, then the collector will miss it and free the memory.

To avoid this from happening, either

- maintain a pointer to the object in an area the collector does scan for pointers;
- add a root where a pointer to the object is stored using `core.memory.GC.addRoot()` or `core.memory.GC.addRange()`.
- reallocate and copy the object using the foreign code's storage allocator or using the C runtime library's `malloc/free`.

28.3 Pointers and the Garbage Collector

Pointers in D can be broadly divided into two categories: Those that point to garbage collected memory, and those that do not. Examples of the latter are pointers created by calls to C's `malloc()`, pointers received from C library routines, pointers to static data, pointers to objects on the stack, etc. For those pointers, anything that is legal in C can be done with them.

For garbage collected pointers and references, however, there are some restrictions. These restrictions are minor, but they are intended to enable the maximum flexibility in garbage collector design.

Undefined behavior:

- Do not xor pointers with other values, like the xor pointer linked list trick used in C.

- Do not use the xor trick to swap two pointer values.
- Do not store pointers into non-pointer variables using casts and other tricks.

```
void* p;
...
int x = cast(int)p;    // error: undefined behavior
```

The garbage collector does not scan non-pointer fields for GC pointers.

- Do not take advantage of alignment of pointers to store bit flags in the low order bits:

```
p = cast(void*)(cast(int)p | 1); // error: undefined behavior
```

- Do not store into pointers values that may point into the garbage collected heap:

```
p = cast(void*)12345678;    // error: undefined behavior
```

A copying garbage collector may change this value.

- Do not store magic values into pointers, other than `null`.
- Do not write pointer values out to disk and read them back in again.
- Do not use pointer values to compute a hash function. A copying garbage collector can arbitrarily move objects around in memory, thus invalidating the computed hash value.
- Do not depend on the ordering of pointers:

```
if (p1 < p2) // error: undefined behavior
...

```

since, again, the garbage collector can move objects around in memory.

- Do not add or subtract an offset to a pointer such that the result points outside of the bounds of the garbage collected object originally allocated.

```
char* p = new char[10];
char* q = p + 6; // ok
q = p + 11;      // error: undefined behavior
q = p - 1;      // error: undefined behavior
```

- Do not misalign pointers if those pointers may point into the GC heap, such as:

```
struct Foo
{
    align (1):
        byte b;
        char* p; // misaligned pointer
}
```

Misaligned pointers may be used if the underlying hardware supports them **and** the pointer is never used to point into the GC heap.

- Do not use byte-by-byte memory copies to copy pointer values. This may result in intermediate conditions where there is not a valid pointer, and if the gc pauses the thread in such a condition, it can corrupt memory. Most implementations of `memcpy()` will work since the internal implementation of it does the copy in aligned chunks greater than or equal to the pointer size, but since this kind of implementation is not guaranteed by the C standard, use `memcpy()` only with extreme caution.
- Do not have pointers in a struct instance that point back to the same instance. The trouble with this is if the instance gets moved in memory, the pointer will point back to where it came from, with likely disastrous results.

Things that are reliable and can be done:

- Use a union to share storage with a pointer:

```
union U { void* ptr; int value }
```

- A pointer to the start of a garbage collected object need not be maintained if a pointer to the interior of the object exists.

```
char[] p = new char[10];
char[] q = p[3..6];
// q is enough to hold on to the object, don't need to keep
// p as well.
```

One can avoid using pointers anyway for most tasks. D provides features rendering most explicit pointer uses obsolete, such as reference objects, dynamic arrays, and garbage collection. Pointers are provided in order to interface successfully with C APIs and for some low level work.

28.4 Working with the Garbage Collector

Garbage collection doesn't solve every memory deallocation problem. For example, if a pointer to a large data structure is kept, the garbage collector cannot reclaim it, even if it is never referred to again. To eliminate this problem, it is good practice to set a reference or pointer to an object to null when no longer needed.

This advice applies only to static references or references embedded inside other objects. There is not much point for such stored on the stack to be nulled because new stack frames are initialized anyway.

28.5 Object Pinning and a Moving Garbage Collector

Although D does not currently use a moving garbage collector, by following the rules listed above one can be implemented. No special action is required to pin objects. A moving collector will

only move objects for which there are no ambiguous references, and for which it can update those references. All other objects will be automatically pinned.

28.6 D Operations That Involve the Garbage Collector

Some sections of code may need to avoid using the garbage collector. The following constructs may allocate memory using the garbage collector:

- *NewExpression*
- Array appending
- Array concatenation
- Array literals (except when used to initialize static data)
- Associative array literals
- Any insertion, removal, or lookups in an associative array
- Extracting keys or values from an associative array
- Taking the address of (i.e. making a delegate to) a nested function that accesses variables in an outer scope
- A function literal that accesses variables in an outer scope
- An *AssertExpression* that fails its condition

28.7 Configuring the Garbage Collector

Since version 2.067, The garbage collector can now be configured through the command line, the environment or by options embedded into the executable.

By default, GC options can only be passed on the command line of the program to run, e.g.

```
app "--DRT-gcopt=profile:1_minPoolSize:16" arguments to app
```

Available GC options are:

- disable:0|1 - start disabled
- profile:0|1 - enable profiling with summary when terminating program
- gc:conservative|manual - select gc implementation (default = conservative)
- initReserve:N - initial memory to reserve in MB
- minPoolSize:N - initial and minimum pool size in MB
- maxPoolSize:N - maximum pool size in MB
- incPoolSize:N - pool size increment MB
- heapSizeFactor:N - targeted heap size to used memory ratio

In addition, `-DRT-gcopt=help` will show the list of options and their current settings.

Command line options starting with `"-DRT-"` are filtered out before calling `main`, so the program will not see them. They are still available via `rt_args`.

Configuration via the command line can be disabled by declaring a variable for the linker to pick up before using its default from the runtime:

```
extern(C) __gshared bool rt_cmdline_enabled = false;
```

Likewise, declare a boolean `rt_envvars_enabled` to enable configuration via the environment variable `DRT_GCOPT`:

```
extern(C) __gshared bool rt_envvars_enabled = true;
```

Setting default configuration properties in the executable can be done by specifying an array of options named `rt_options`:

```
extern(C) __gshared string[] rt_options = [ "gcopt=initReserve:100▯profile:1" ];
```

Evaluation order of options is `rt_options`, then environment variables, then command line arguments, i.e. if command line arguments are not disabled, they can override options specified through the environment or embedded in the executable.

28.8 References

- Wikipedia
- GC FAQ
- Uniprocessor Garbage Collector Techniques
- Garbage Collection: Algorithms for Automatic Dynamic Memory Management

Chapter 29

Floating Point

Floating Point Intermediate Values

On many computers, greater precision operations do not take any longer than lesser precision operations, so it makes numerical sense to use the greatest precision available for internal temporaries. The philosophy is not to dumb down the language to the lowest common hardware denominator, but to enable the exploitation of the best capabilities of target hardware.

For floating point operations and expression intermediate values, a greater precision can be used than the type of the expression. Only the minimum precision is set by the types of the operands, not the maximum. **Implementation Note:** On Intel x86 machines, for example, it is expected (but not required) that the intermediate calculations be done to the full 80 bits of precision implemented by the hardware.

It's possible that, due to greater use of temporaries and common subexpressions, optimized code may produce a more accurate answer than unoptimized code.

Algorithms should be written to work based on the minimum precision of the calculation. They should not degrade or fail if the actual precision is greater. Float or double types, as opposed to the real (extended) type, should only be used for:

- reducing memory consumption for large arrays
- when speed is more important than accuracy
- data and function argument compatibility with C

Floating Point Constant Folding

Regardless of the type of the operands, floating point constant folding is done in **real** or greater precision. It is always done following IEEE 754 rules and round-to-nearest is used.

Floating point constants are internally represented in the implementation in at least **real** ↯ precision, regardless of the constant's type. The extra precision is available for constant folding.

Committing to the precision of the result is done as late as possible in the compilation process. For example:

```
const float f = 0.2f;
writeln(f - 0.2);
```

will print 0. A non-const static variable's value cannot be propagated at compile time, so:

```
static float f = 0.2f;
writeln(f - 0.2);
```

will print 2.98023e-09. Hex floating point constants can also be used when specific floating point bit patterns are needed that are unaffected by rounding. To find the hex value of 0.2f:

```
import std.stdio;

void main()
{
    writefln("%a", 0.2f);
}
```

which is 0x1.99999ap-3. Using the hex constant:

```
const float f = 0x1.99999ap-3f;
writeln(f - 0.2);
```

prints 2.98023e-09.

Different compiler settings, optimization settings, and inlining settings can affect opportunities for constant folding, therefore the results of floating point calculations may differ depending on those settings.

Rounding Control

IEEE 754 floating point arithmetic includes the ability to set 4 different rounding modes. These are accessible via the functions in `core.stdc.fenv`.

If the floating-point rounding mode is changed within a function, it must be restored before the function exits. If this rule is violated (for example, by the use of inline asm), the rounding mode used for subsequent calculations is undefined.

Exception Flags

IEEE 754 floating point arithmetic can set several flags based on what happened with a computation:

```

FE_INVALID
FE_DENORMAL
FE_DIVBYZERO
FE_OVERFLOW
FE_UNDERFLOW
FE_INEXACT

```

These flags can be set/reset via the functions in `core.stdc.fenv`.

Floating Point Transformations

An implementation may perform transformations on floating point computations in order to reduce their strength, i.e. their runtime computation time. Because floating point math does not precisely follow mathematical rules, some transformations are not valid, even though some other programming languages still allow them.

The following transformations of floating point expressions are not allowed because under IEEE rules they could produce different results.

Disallowed Floating Point Transformations	
transformation	comments
$x + 0 \rightarrow x$	not valid if x is -0
$x - 0 \rightarrow x$	not valid if x is ± 0 and rounding is towards $-\infty$
$-x \leftrightarrow 0 - x$	not valid if x is $+0$
$x - x \rightarrow 0$	not valid if x is NaN or $\pm\infty$
$x - y \leftrightarrow -(y - x)$	not valid because $(1-1=+0)$ whereas $-(1-1)=-0$
$x * 0 \rightarrow 0$	not valid if x is NaN or $\pm\infty$
$x / c \leftrightarrow x * (1/c)$	valid if $(1/c)$ yields an exact result
$x != x \rightarrow \text{false}$	not valid if x is a NaN
$x == x \rightarrow \text{true}$	not valid if x is a NaN
$x !op y \leftrightarrow !(x op y)$	not valid if x or y is a NaN

Of course, transformations that would alter side effects are also invalid.

Chapter 30

D x86 Inline Assembler

D, being a systems programming language, provides an inline assembler. The inline assembler is standardized for D implementations across the same CPU family, for example, the Intel Pentium inline assembler for a Win32 D compiler will be syntax compatible with the inline assembler for Linux running on an Intel Pentium.

Implementations of D on different architectures, however, are free to innovate upon the memory model, function call/return conventions, argument passing conventions, etc.

This document describes the `x86` and `x86_64` implementations of the inline assembler. The inline assembler platform support that a compiler provides is indicated by the `D_InlineAsm_X86` and `D_InlineAsm_X86_64` version identifiers, respectively.

AsmInstruction:

```
Identifier : AsmInstruction
align IntegerExpression
even
naked
db Operands
ds Operands
di Operands
dl Operands
df Operands
dd Operands
de Operands
db StringLiteral
ds StringLiteral
di StringLiteral
dl StringLiteral
dw StringLiteral
```

```

    dq StringLiteral
    Opcode
    Opcode Operands

Operands:
    Operand
    Operand , Operands

```

30.1 Labels

Assembler instructions can be labeled just like other statements. They can be the target of goto statements. For example:

```

void *pc;
asm
{
    call L1          ;
L1:                  ;
    pop  EBX         ;
    mov  pc[EBP],EBX ; // pc now points to code at L1
}

```

30.2 align *IntegerExpression*

```

IntegerExpression:
    IntegerLiteral
    Identifier

```

Causes the assembler to emit NOP instructions to align the next assembler instruction on an *IntegerExpression* boundary. *IntegerExpression* must evaluate at compile time to an integer that is a power of 2.

Aligning the start of a loop body can sometimes have a dramatic effect on the execution speed.

30.3 even

Causes the assembler to emit NOP instructions to align the next assembler instruction on an even boundary.

30.4 naked

Causes the compiler to not generate the function prolog and epilog sequences. This means such is the responsibility of inline assembly programmer, and is normally used when the entire function is to be written in assembler.

30.5 db, ds, di, dl, df, dd, de

These pseudo ops are for inserting raw data directly into the code. **db** is for bytes, **ds** is for 16 bit words, **di** is for 32 bit words, **dl** is for 64 bit words, **df** is for 32 bit floats, **dd** is for 64 bit doubles, and **de** is for 80 bit extended reals. Each can have multiple operands. If an operand is a string literal, it is as if there were *length* operands, where *length* is the number of characters in the string. One character is used per operand. For example:

```
asm
{
    db 5,6,0x83;    // insert bytes 0x05, 0x06, and 0x83 into code
    ds 0x1234;      // insert bytes 0x34, 0x12
    di 0x1234;      // insert bytes 0x34, 0x12, 0x00, 0x00
    dl 0x1234;      // insert bytes 0x34, 0x12, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00
    df 1.234;       // insert float 1.234
    dd 1.234;       // insert double 1.234
    de 1.234;       // insert real 1.234
    db "abc";       // insert bytes 0x61, 0x62, and 0x63
    ds "abc";       // insert bytes 0x61, 0x00, 0x62, 0x00, 0x63, 0x00
}
```

30.6 Opcodes

A list of supported opcodes is at the end.

The following registers are supported. Register names are always in upper case.

Register:

```
AL AH AX EAX
BL BH BX EBX
CL CH CX ECX
DL DH DX EDX
BP EBP
SP ESP
DI EDI
```

```

SI ESI
ES CS SS DS GS FS
CR0 CR2 CR3 CR4
DR0 DR1 DR2 DR3 DR6 DR7
TR3 TR4 TR5 TR6 TR7
ST
ST(0) ST(1) ST(2) ST(3) ST(4) ST(5) ST(6) ST(7)
MM0 MM1 MM2 MM3 MM4 MM5 MM6 MM7
XMM0 XMM1 XMM2 XMM3 XMM4 XMM5 XMM6 XMM7

```

x86_64 adds these additional registers.

Register64:

```

RAX RBX RCX RDX
BPL RBP
SPL RSP
DIL RDI
SIL RSI
R8B R8W R8D R8
R9B R9W R9D R9
R10B R10W R10D R10
R11B R11W R11D R11
R12B R12W R12D R12
R13B R13W R13D R13
R14B R14W R14D R14
R15B R15W R15D R15
XMM8 XMM9 XMM10 XMM11 XMM12 XMM13 XMM14 XMM15
YMM0 YMM1 YMM2 YMM3 YMM4 YMM5 YMM6 YMM7
YMM8 YMM9 YMM10 YMM11 YMM12 YMM13 YMM14 YMM15

```

Special Cases

lock, rep, repe, repne, repnz, repz These prefix instructions do not appear in the same statement as the instructions they prefix; they appear in their own statement. For example:

```

asm
{
    rep    ;
    movsb ;

```



```
}
```

pause This opcode is not supported by the assembler, instead use

```
asm
{
    rep  ;
    nop  ;
}
```

which produces the same result.

floating point ops Use the two operand form of the instruction format;

```
fdiv ST(1);    // wrong
fmul ST;       // wrong
fdiv ST,ST(1); // right
fmul ST,ST(0); // right
```

30.7 Operands

Operand:

AsmExp

AsmExp:

AsmLogOrExp

AsmLogOrExp ? *AsmExp* : *AsmExp*

AsmLogOrExp:

AsmLogAndExp

AsmLogOrExp || *AsmLogAndExp*

AsmLogAndExp:

AsmOrExp

AsmLogAndExp && *AsmOrExp*

AsmOrExp:

AsmXorExp

AsmOrExp | *AsmXorExp*

AsmXorExp:

```

AsmAndExp
AsmXorExp ^ AsmAndExp

```

```

AsmAndExp:
  AsmEqualExp
  AsmAndExp & AsmEqualExp

```

```

AsmEqualExp:
  AsmRelExp
  AsmEqualExp == AsmRelExp
  AsmEqualExp != AsmRelExp

```

```

AsmRelExp:
  AsmShiftExp
  AsmRelExp < AsmShiftExp
  AsmRelExp <= AsmShiftExp
  AsmRelExp > AsmShiftExp
  AsmRelExp >= AsmShiftExp

```

```

AsmShiftExp:
  AsmAddExp
  AsmShiftExp << AsmAddExp
  AsmShiftExp >> AsmAddExp
  AsmShiftExp >>> AsmAddExp

```

```

AsmAddExp:
  AsmMulExp
  AsmAddExp + AsmMulExp
  AsmAddExp - AsmMulExp

```

```

AsmMulExp:
  AsmBrExp
  AsmMulExp * AsmBrExp
  AsmMulExp / AsmBrExp
  AsmMulExp % AsmBrExp

```

```

AsmBrExp:
  AsmUnExp
  AsmBrExp [ AsmExp ]

```

AsmUnaExp:

AsmTypePrefix AsmExp
offsetof AsmExp
seg AsmExp
+ AsmUnaExp
- AsmUnaExp
! AsmUnaExp
~ AsmUnaExp
AsmPrimaryExp

AsmPrimaryExp:

IntegerLiteral
FloatLiteral
__LOCAL_SIZE
\$
Register
Register : AsmExp
Register64
Register64 : AsmExp
DotIdentifier
this

DotIdentifier:

Identifier
Identifier . DotIdentifier

The operand syntax more or less follows the Intel CPU documentation conventions. In particular, the convention is that for two operand instructions the source is the right operand and the destination is the left operand. The syntax differs from that of Intel's in order to be compatible with the D language tokenizer and to simplify parsing.

The **seg** means load the segment number that the symbol is in. This is not relevant for flat model code. Instead, do a move from the relevant segment register.

Operand Types

AsmTypePrefix:

near ptr
far ptr
byte ptr

```

short ptr
int ptr
word ptr
dword ptr
qword ptr
float ptr
double ptr
real ptr

```

In cases where the operand size is ambiguous, as in:

```
add [EAX],3 ;
```

it can be disambiguated by using an *AsmTypePrefix*:

```
add byte ptr [EAX],3 ;
add int ptr [EAX],7 ;
```

far ptr is not relevant for flat model code.

Struct/Union/Class Member Offsets

To access members of an aggregate, given a pointer to the aggregate is in a register, use the `.offsetof` property of the qualified name of the member:

```

struct Foo { int a,b,c; }
int bar(Foo *f)
{
    asm
    {
        mov EBX,f ;
        mov EAX,Foo.b.offsetof[EBX] ;
    }
}
void main()
{
    Foo f = Foo(0, 2, 0);
    assert(bar(&f) == 2);
}

```

Alternatively, inside the scope of an aggregate, only the member name is needed:

```
struct Foo // or class
```

```

{
    int a,b,c;
    int bar()
    {
        asm
        {
            mov EBX, this ;
            mov EAX, b[EBX] ;
        }
    }
}
void main()
{
    Foo f = Foo(0, 2, 0);
    assert(f.bar() == 2);
}

```

Stack Variables

Stack variables (variables local to a function and allocated on the stack) are accessed via the name of the variable indexed by EBP:

```

int foo(int x)
{
    asm
    {
        mov EAX,x[EBP] ; // loads value of parameter x into EAX
        mov EAX,x      ; // does the same thing
    }
}

```

If the [EBP] is omitted, it is assumed for local variables. If **naked** is used, this no longer holds.

Special Symbols

\$ Represents the program counter of the start of the next instruction. So,

```
jmp $ ;
```

branches to the instruction following the jmp instruction. The \$ can only appear as the target of a jmp or call instruction.

`__LOCAL_SIZE` This gets replaced by the number of local bytes in the local stack frame. It is most handy when the `naked` is invoked and a custom stack frame is programmed.

30.8 Opcodes Supported

Table 30.1: Opcodes

aaa	aad	aam	aas	adc
add	addpd	addps	addsd	addss
and	andnpd	andnps	andpd	andps
arpl	bound	bsf	bsr	bswap
bt	btc	btr	bts	call
cbw	cdq	clc	cld	clflush
cli	clts	cmc	cmova	cmovae
cmovb	cmovbe	cmovc	cmove	cmovg
cmovge	cmovl	cmovle	cmovna	cmovnae
cmovnb	cmovnbe	cmovnc	cmovne	cmovng
cmovnge	cmovnl	cmovnle	cmovno	cmovnp
cmovns	cmovnz	cmovo	cmovp	cmovpe
cmovpo	cmovs	cmovz	cmp	cmppd
cmpps	cmps	cmpsb	cmpsd	cmpss
cmpsw	cmpxchg	cmpxchg8b	cmpxchg16b	
comisd	comiss			
cpuid	cvtdq2pd	cvtdq2ps	cvtpd2dq	cvtpd2pi
cvtpd2ps	cvtpi2pd	cvtpi2ps	cvtps2dq	cvtps2pd
cvtps2pi	cvtsd2si	cvtsd2ss	cvtsi2sd	cvtsi2ss
cvtss2sd	cvtss2si	cvttpd2dq	cvttpd2pi	cvttps2dq
cvttps2pi	cvttss2si	cvttss2si	cwd	cwde
da	daa	das	db	dd
de	dec	df	di	div
divpd	divps	divsd	divss	dl
dq	ds	dt	dw	emms
enter	f2xm1	fabs	fadd	faddp
fbld	fbstp	fchs	fclex	fcmovb
fcmovbe	fcmove	fcmovnb	fcmovnbe	fcmovne
fcmovnu	fcmovu	fcom	fcomi	fcomip
fcomp	fcompp	fcos	fdecstp	fdisi
fdiv	fdivp	fdivr	fdivrp	feni

←

(continued)

ffree	fiadd	ficom	ficomp	fidiv
fidivr	fild	fimul	fincstp	finit
fist	fistp	fisub	fisubr	fld
fld1	fldcw	fldenv	fldl2e	fldl2t
fldlg2	fldln2	fldpi	fldz	fmul
fmulp	fnclex	fndisi	fneni	fninit
fnop	fnsave	fnstcw	fnstenv	fnstsw
fpatan	fprem	fprem1	fptan	frndint
frstor	fsave	fscale	fsetpm	fsin
fsincos	fsqrt	fst	fstcw	fstenv
fstp	fstsw	fsub	fsubp	fsubr
fsubrp	ftst	fucom	fucomi	fucomip
fucomp	fucompp	fwait	fxam	fxch
fxrstor	fxsave	fxtract	fyl2x	fyl2xp1
hlt	idiv	imul	in	inc
ins	insb	insd	insw	int
into	invd	invlpg	iret	iretd
ja	jae	jb	jbe	jc
jcxz	je	jecxz	jg	jge
jl	jle	jmp	jna	jnae
jnb	jnb	jnc	jne	jng
jnge	jnl	jnle	jno	jnp
jns	jnz	jo	jp	jpe
jpo	js	jz	lahf	lar
ldmxcsr	lds	lea	leave	les
lfence	lfs	lgdt	lgs	lidt
lldt	lmsw	lock	lods	lodsrb
lodsd	lodsw	loop	loope	loopne
loopnz	loopz	lsl	lss	ltr
maskmovdqu	maskmovq	maxpd	maxps	maxsd
maxss	mfence	minpd	minps	minsd
minss	mov	movapd	movaps	movd
movd2q	movdqa	movdqu	movhlps	movhpd
movhps	movlhps	movlpd	movlps	movmskpd
movmskps	movntdq	movnti	movntpd	movntps
movntq	movq	movq2dq	movs	movsb
movsd	movss	movsw	movsx	movupd
movups	movzx	mul	mulpd	mulps

←

(continued)

mulsd	mulss	neg	nop	not
or	orpd	orps	out	outs
outsb	outsd	outsw	packssdw	packsswb
packuswb	paddb	padd	paddq	paddsb
paddsw	paddusb	paddusw	paddw	pand
pandn	pavgb	pavgw	pcmpeqb	pcmpeqd
pcmpeqw	pcmpgtb	pcmpgtd	pcmpgtw	pextrw
pinsrw	pmaddwd	pmaxsw	pmaxub	pminsw
pminub	pmovmskb	pmulhuw	pmulhw	pmullw
pmuludq	pop	popa	popad	popf
popfd	por	prefetchnta	prefetcht0	prefetcht1
prefetcht2	psadbw	pshufd	pshufhw	pshufw
pshufw	pslld	pslldq	psllq	psllw
psrad	psraw	psrld	psrldq	psrlq
psrlw	psubb	psubd	psubq	psubsb
psubsw	psubusb	psubusw	psubw	punpckhbw
punpckhdq	punpckhqdq	punpckhwd	punpcklbw	punpckldq
punpcklqdq	punpcklwd	push	pusha	pushad
pushf	pushfd	pxor	rcl	rcpps
rcpss	rcr	rdmsr	rdpmc	rdtsc
rep	repe	repne	repnz	repz
ret	retf	rol	ror	rsm
rsqrtps	rsqrtss	sahf	sal	sar
sbb	scas	scasb	scasd	scasw
seta	setae	setb	setbe	setc
sete	setg	setge	setl	setle
setna	setnae	setnb	setnbe	setnc
setne	setng	setnge	setnl	setnle
setno	setnp	setns	setnz	seto
setp	setpe	setpo	sets	setz
sfence	sgdt	shl	shld	shr
shrd	shufpd	shufps	sidt	sldt
smsw	sqrtpd	sqrtps	sqrtsd	sqrtss
stc	std	sti	stmxcscr	stos
stosb	stosd	stosw	str	sub
subpd	subps	subsd	subss	syscall
sysenter	sysexit	sysret	test	ucomisd
ucomiss	ud2	unpckhpd	unpckhps	unpcklpd

←

(continued)

unpcklps	verr	verw	wait	wbinvd
wrmsr	xadd	xchg	xlat	xlatb
xor	xorpd	xorps		

Pentium 4 (Prescott) Opcodes Supported

Table 30.2: Pentium 4 Opcodes

addsubpd	addsubps	fisttp	haddpd	haddps
hsubpd	hsubps	lddqu	monitor	movddup
movshdup	movsldup	mwait		

AMD Opcodes Supported

Table 30.3: AMD Opcodes

pavgusb	pf2id	pfacc	pfadd	pfcmpaq
pfcmpge	pfcmpgt	pfmax	pfmin	pfmul
pfnacc	pfpnacc	pfrcp	pfrcpit1	pfrcpit2
pfrsqit1	pfrsqr	pfsb	pfsbr	pi2fd
pmulhrw	pswapd			

SIMD

SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2 and AVX are supported.

Chapter 31

Embedded Documentation

The D programming language enables embedding both contracts and test code along side the actual code, which helps to keep them all consistent with each other. One thing lacking is the documentation, as ordinary comments are usually unsuitable for automated extraction and formatting into manual pages. Embedding the user documentation into the source code has important advantages, such as not having to write the documentation twice, and the likelihood of the documentation staying consistent with the code.

Some existing approaches to this are:

- Doxygen which already has some support for D
- Java's Javadoc, probably the most well-known
- C#'s embedded XML
- Other documentation tools

D's goals for embedded documentation are:

1. It looks good as embedded documentation, not just after it is extracted and processed.
2. It's easy and natural to write, i.e. minimal reliance on `<tags>` and other clumsy forms one would never see in a finished document.
3. It does not repeat information that the compiler already knows from parsing the code.
4. It doesn't rely on embedded HTML, as such will impede extraction and formatting for other purposes.
5. It's based on existing D comment forms, so it is completely independent of parsers only interested in D code.
6. It should look and feel different from code, so it won't be visually confused with code.
7. It should be possible for the user to use Doxygen or other documentation extractor if desired.

31.1 Specification

The specification for the form of embedded documentation comments only specifies how information is to be presented to the compiler. It is implementation-defined how that information is used and the form of the final presentation. Whether the final presentation form is an HTML web page, a man page, a PDF file, etc. is not specified as part of the D Programming Language.

Phases of Processing

Embedded documentation comments are processed in a series of phases:

1. Lexical - documentation comments are identified and attached to tokens.
2. Parsing - documentation comments are associated with specific declarations and combined.
3. Sections - each documentation comment is divided up into a sequence of sections.
4. Special sections are processed.
5. Highlighting of non-special sections is done.
6. All sections for the module are combined.
7. Macro and Escape text substitution is performed to produce the final result.

Lexical

Embedded documentation comments are one of the following forms:

1. `/** ... */` The two `*`'s after the opening `/`
2. `/++ ... +/` The two `+`'s after the opening `/`
3. `///` The three slashes

The following are all embedded documentation comments:

```
/// This is a one line documentation comment.
```

```
/** So is this. */
```

```
/++ And this. +/
```

```
/**  
    This is a brief documentation comment.  
*/
```

```
/**  
 * The leading * on this line is not part of the documentation comment.
```

```

*/

/*****
    The extra *'s immediately following the /** are not
    part of the documentation comment.
*/

/++
    This is a brief documentation comment.
+/

/++
+ The leading + on this line is not part of the documentation comment.
+/

/+++++
    The extra +'s immediately following the /++ are not
    part of the documentation comment.
    +/

/*****Closing*'s are not part *****/

```

The extra *'s and +'s on the comment opening, closing and left margin are ignored and are not part of the embedded documentation. Comments not following one of those forms are not documentation comments.

Parsing

Each documentation comment is associated with a declaration. If the documentation comment is on a line by itself or with only whitespace to the left, it refers to the next declaration. Multiple documentation comments applying to the same declaration are concatenated. Documentation comments not associated with a declaration are ignored. Documentation comments preceding the *ModuleDeclaration* apply to the entire module. If the documentation comment appears on the same line to the right of a declaration, it applies to that.

If a documentation comment for a declaration consists only of the identifier *ditto* then the documentation comment for the previous declaration at the same declaration scope is applied to this declaration as well.

If there is no documentation comment for a declaration, that declaration may not appear in the output. To ensure it does appear in the output, put an empty declaration comment for it.

```
int a;   /// documentation for a; b has no documentation
```

```

int b;

/** documentation for c and d */
/** more documentation for c and d */
int c;
/** ditto */
int d;

/** documentation for e and f */ int e;
int f;  /// ditto

/** documentation for g */
int g; /// more documentation for g

/// documentation for C and D
class C
{
    int x; /// documentation for C.x

    /** documentation for C.y and C.z */
    int y;
    int z; /// ditto
}

/// ditto
class D { }

```

Sections

The document comment is a series of *Sections*. A *Section* is a name that is the first non-blank character on a line immediately followed by a `:'`. This name forms the section name. The section name is not case sensitive.

Summary

The first section is the *Summary*, and does not have a section name. It is first paragraph, up to a blank line or a section name. While the summary can be any length, try to keep it to one line. The *Summary* section is optional.

Description

The next unnamed section is the *Description*. It consists of all the paragraphs following the *Summary* until a section name is encountered or the end of the comment.

While the *Description* section is optional, there cannot be a *Description* without a *Summary* section.

```

/*****
 * Brief summary of what
 * myfunc does, forming the summary section.
 *
 * First paragraph of synopsis description.
 *
 * Second paragraph of
 * synopsis description.
 */

```

```
void myfunc() { }
```

Named sections follow the *Summary* and *Description* unnamed sections.

Standard Sections

For consistency and predictability, there are several standard sections. None of these are required to be present.

Authors: Lists the author(s) of the declaration.

```

/**
 * Authors: Melvin D. Nerd, melvin@mailinator.com
 */

```

Bugs: Lists any known bugs.

```

/**
 * Bugs: Doesn't work for negative values.
 */

```

Date: Specifies the date of the current revision. The date should be in a form parseable by `std.date`.

```

/**
 * Date: March 14, 2003
 */

```

Deprecated: Provides an explanation for and corrective action to take if the associated declaration is marked as deprecated.

```
/**
 * Deprecated: superseded by function bar().
 */

deprecated void foo() { ... }
```

Examples: Any usage examples

```
/**
 * Examples:
 * -----
 * writeln("3"); // writes '3' to stdout
 * -----
 */
```

History: Revision history.

```
/**
 * History:
 *      V1 is initial version
 *
 *      V2 added feature X
 */
```

License: Any license information for copyrighted code.

```
/**
 * License: use freely for any purpose
 */

void bar() { ... }
```

Returns: Explains the return value of the function. If the function returns **void**, don't redundantly document it.

```
/**
 * Read the file.
 * Returns: The contents of the file.
 */
```



```
void[] readFile(char[] filename) { ... }
```

See _Also: List of other symbols and URL's to related items.

```
/**
 * See_Also:
 *   foo, bar, http://www.digitalmars.com/d/phobos/index.html
 */
```

Standards: If this declaration is compliant with any particular standard, the description of it goes here.

```
/**
 * Standards: Conforms to DSPEC-1234
 */
```

Throws: Lists exceptions thrown and under what circumstances they are thrown.

```
/**
 * Write the file.
 * Throws: WriteException on failure.
 */
```

```
void writeFile(char[] filename) { ... }
```

Version: Specifies the current version of the declaration.

```
/**
 * Version: 1.6a
 */
```

Special Sections

Some sections have specialized meanings and syntax.

Copyright: This contains the copyright notice. The macro COPYRIGHT is set to the contents of the section when it documents the module declaration. The copyright section only gets this special treatment when it is for the module declaration.

```
/** Copyright: Public Domain */
```

```
module foo;
```

Params: Function parameters can be documented by listing them in a params section. Each line that starts with an identifier followed by an '=' starts a new parameter description. A description can span multiple lines.

```

/*****
 * foo does this.
 * Params:
 *      x =      is for this
 *              and not for that
 *      y =      is for that
 */

void foo(int x, int y)
{
}

```

Macros: The macros section follows the same syntax as the **Params:** section. It's a series of *NAME=value* pairs. The *NAME* is the macro name, and *value* is the replacement text.

```

/**
 * Macros:
 *      FOO =    now is the time for
 *              all good men
 *      BAR =    bar
 *      MAGENTA = <font color="magenta">\$(DOLLAR)0</font>
 */

```

Escapes= The escapes section is a series of substitutions which replace special characters with a string. It's useful when the output format requires escaping of certain characters, for example in HTML & should be escaped with &. The syntax is */c/string/*, where *c* is either a single character, or multiple characters separated by whitespace or commas, and *string* is the replacement text.

```

/**
 * ESCAPES = /&/AddressOf!/
 *           /!/Exclamation/
 *           /?/QuestionMark/
 *           /,/Comma/
 *           /{ }/Parens/
 *           /<, >/Arrows/
 */

```

31.2 Highlighting

Embedded Comments

The documentation comments can themselves be commented using the `$(DDOC_COMMENT comment ↵ text)` syntax. These comments do not nest.

Embedded Code

D code can be embedded using lines beginning with at least three hyphens (ignoring whitespace) to delineate the code section:

```
/++
+ Our function.
+
+ Example:
+ ---
+ import std.stdio;
+
+ void foo()
+ {
+     writeln("foo!"); /* print the string */
+ }
+ ---
+ /
```

Note that the documentation comment uses the `/++ ... +/` form so that `/* ... */` can be used inside the code section.

Inline Code

Inline code can be written between backtick characters (```), similarly to the syntax used on GitHub, Reddit, Stack Overflow, and other websites. Both the opening and closing ``` character must appear on the same line to trigger this behavior.

Text inside these sections will be escaped according to the rules described above, then wrapped in a `\$(DDOC_BACKQUOTED)` macro. By default, this macro expands to be displayed as an inline text span, formatted as code.

A literal backtick character can be output either as a non-paired ``` on a single line or by using the `\$(BACKTICK)` macro.

```
/// Returns 'true' if 'a == b'.
void foo() {}
```

```

/// Backquoted '<html>' will be displayed to the user instead
/// of passed through as embedded HTML (see below).
void bar() {}

```

Embedded HTML

HTML can be embedded into the documentation comments, and it will be passed through to the HTML output unchanged. However, since it is not necessarily true that HTML will be the desired output format of the embedded documentation comment extractor, it is best to avoid using it where practical.

```

/**
 * Example of embedded HTML:
 *
 * <ol>
 *   <li><a href="http://www.digitalmars.com">Digital Mars</a></li>
 *   <li><a href="http://www.classicempire.com">Empire</a></li>
 * </ol>
 */

```

Emphasis

Identifiers in documentation comments that are function parameters or are names that are in scope at the associated declaration are emphasized in the output. This emphasis can take the form of italics, boldface, a hyperlink, etc. How it is emphasized depends on what it is - a function parameter, type, D keyword, etc. To prevent unintended emphasis of an identifier, it can be preceded by an underscore (`_`). The underscore will be stripped from the output.

Character Entities

Some characters have special meaning to the documentation processor, to avoid confusion it can be best to replace them with their corresponding character entities:

Characters and Entities

Character	Entity
<	<
>	>
&	&

It is not necessary to do this inside a code section, or if the special character is not immediately followed by a `#` or a letter.

No Documentation

No documentation is generated for the following constructs, even if they have a documentation comment:

- Invariants
- Postblits
- Destructors
- Static constructors and static destructors
- Class info, type info, and module info

31.3 Macros

The documentation comment processor includes a simple macro text preprocessor. When a $\$(NAME)$ appears in section text it is replaced with $NAME$'s corresponding replacement text.

For example:

```
/**
Macros:
  PARAM = <u>\$1</u>
  MATH_DOCS = <a href="https://dlang.org/phobos/std_math.html">Math Docs</a>
*/
module math;

/**
 * This function returns the sum of \$(PARAM a) and \$(PARAM b).
 * See also the \$(MATH_DOCS).
 */
int sum(int a, int b) { return a + b; }
```

The above would generate the following output:

```
<h1>test</h1>
<dl><dt><big><a name="sum"></a>int <u>sum</u>(int <i>a</i>, int <i>b</i>);
</big></dt>
<dd>This function returns the <u>sum</u> of <u><i>a</i></u> and <u><i>b</i></u>.
  See also the <a href="https://dlang.org/phobos/std_math.html">Math Docs</a>.
</dd>
</dl>
```

The replacement text is recursively scanned for more macros. If a macro is recursively encountered, with no argument or with the same argument text as the enclosing macro, it is replaced with

no text. Macro invocations that cut across replacement text boundaries are not expanded. If the macro name is undefined, the replacement text has no characters in it. If a \$(NAME) is desired to exist in the output without being macro expanded, the \$ should be replaced with `$`.

Macros can have arguments. Any text from the end of the identifier to the closing ‘)’ is the \$0 argument. A \$0 in the replacement text is replaced with the argument text. If there are commas in the argument text, \$1 will represent the argument text up to the first comma, \$2 from the first comma to the second comma, etc., up to \$9. \$+ represents the text from the first comma to the closing ‘)’. The argument text can contain nested parentheses, “ or ” strings, `<!-- ... -->` comments, or tags. If stray, unnested parentheses are used, they can be replaced with the entity `(` for (and `)` for).

Macro definitions come from the following sources, in the specified order:

1. Predefined macros.
2. Definitions from file specified by `sc.ini`’s or `dmd.conf` `DDOCFILE` setting.
3. Definitions from *.ddoc files specified on the command line.
4. Runtime definitions generated by Ddoc.
5. Definitions from any Macros: sections.

Macro redefinitions replace previous definitions of the same name. This means that the sequence of macro definitions from the various sources forms a hierarchy.

Macro names beginning with “D_” and “DDOC_” are reserved.

Predefined Macros

These are hardwired into Ddoc, and represent the minimal definitions needed by Ddoc to format and highlight the presentation. The definitions are for simple HTML.

B	=	<code>\\$0</code>
I	=	<code><i>\\$0</i></code>
U	=	<code><u>\\$0</u></code>
P	=	<code><p>\\$0</p></code>
DL	=	<code><dl>\\$0</dl></code>
DT	=	<code><dt>\\$0</dt></code>
DD	=	<code><dd>\\$0</dd></code>
TABLE	=	<code><table>\\$0</table></code>
TR	=	<code><tr>\\$0</tr></code>
TH	=	<code><th>\\$0</th></code>
TD	=	<code><td>\\$0</td></code>
OL	=	<code>\\$0</code>
UL	=	<code>\\$0</code>
LI	=	<code>\\$0</code>

```

BIG          = <big>\$0</big>
SMALL        = <small>\$0</small>
BR           = <br>
LINK         = <a href="\$0">\$0</a>
LINK2        = <a href="\$1">\$+</a>
LPAREN       = (
RPAREN       = )
BACKTICK     = `
DOLLAR       = \$
DEPRECATED   = \$0

RED          = <font color=red>\$0</font>
BLUE         = <font color=blue>\$0</font>
GREEN        = <font color=green>\$0</font>
YELLOW       = <font color=yellow>\$0</font>
BLACK        = <font color=black>\$0</font>
WHITE        = <font color=white>\$0</font>

D_CODE       = <pre class="d_code">\$0</pre>
D_INLINECODE = <pre style="display:inline;" class="d_inline_code">\$0</pre>
D_COMMENT    = \$(GREEN \$0)
D_STRING     = \$(RED \$0)
D_KEYWORD    = \$(BLUE \$0)
D_PSYMBOL    = \$(U \$0)
D_PARAM      = \$(I \$0)

DDOC = <html><head>
      <META http-equiv="content-type" content="text/html; charset=utf-8">
      <title>\$(TITLE)</title>
    </head><body>
      <h1>\$(TITLE)</h1>
      \$(BODY)
      <hr>\$(SMALL Page generated by \$(LINK2 https://dlang.org/ddoc.html, Ddoc)
        ↪ . \$(COPYRIGHT))
    </body></html>

DDOC_COMMENT    = <!-- \$0 -->
DDOC_DECL       = \$(DT \$(BIG \$0))
DDOC_DECL_DD    = \$(DD \$0)
DDOC_DITTO      = \$(BR)\$0

```

```

DDOC_SECTIONS           = \$0
DDOC_SUMMARY             = \$0\$(BR)\$(BR)
DDOC_DESCRIPTION         = \$0\$(BR)\$(BR)
DDOC_AUTHORS             = \$(B Authors:)\$(BR)\$0\$(BR)\$(BR)
DDOC_BUGS                = \$(RED BUGS:)\$(BR)\$0\$(BR)\$(BR)
DDOC_COPYRIGHT           = \$(B Copyright:)\$(BR)\$0\$(BR)\$(BR)
DDOC_DATE                = \$(B Date:)\$(BR)\$0\$(BR)\$(BR)
DDOC_DEPRECATED          = \$(RED Deprecated:)\$(BR)\$0\$(BR)\$(BR)
DDOC_EXAMPLES            = \$(B Examples:)\$(BR)\$0\$(BR)\$(BR)
DDOC_HISTORY             = \$(B History:)\$(BR)\$0\$(BR)\$(BR)
DDOC_LICENSE             = \$(B License:)\$(BR)\$0\$(BR)\$(BR)
DDOC_RETURNS             = \$(B Returns:)\$(BR)\$0\$(BR)\$(BR)
DDOC_SEE_ALSO            = \$(B See Also:)\$(BR)\$0\$(BR)\$(BR)
DDOC_STANDARDS           = \$(B Standards:)\$(BR)\$0\$(BR)\$(BR)
DDOC_THROWS              = \$(B Throws:)\$(BR)\$0\$(BR)\$(BR)
DDOC_VERSION             = \$(B Version:)\$(BR)\$0\$(BR)\$(BR)
DDOC_SECTION_H           = \$(B \$0)\$(BR)\$(BR)
DDOC_SECTION             = \$0\$(BR)\$(BR)
DDOC_MEMBERS             = \$(DL \$0)
DDOC_MODULE_MEMBERS      = \$(DDOC_MEMBERS \$0)
DDOC_CLASS_MEMBERS       = \$(DDOC_MEMBERS \$0)
DDOC_STRUCT_MEMBERS      = \$(DDOC_MEMBERS \$0)
DDOC_ENUM_MEMBERS        = \$(DDOC_MEMBERS \$0)
DDOC_TEMPLATE_MEMBERS    = \$(DDOC_MEMBERS \$0)
DDOC_ENUM_BASETYPE       = \$0
DDOC_PARAMS              = \$(B Params:)\$(BR)\n\$(TABLE \$0)\$(BR)
DDOC_PARAM_ROW           = \$(TR \$0)
DDOC_PARAM_ID            = \$(TD \$0)
DDOC_PARAM_DESC          = \$(TD \$0)
DDOC_BLANKLINE           = \$(BR)\$(BR)

DDOC_ANCHOR              = <a name="\$1"></a>
DDOC_PSYMBOL             = \$(U \$0)
DDOC_PSUPER_SYMBOL       = \$(U \$0)
DDOC_KEYWORD             = \$(B \$0)
DDOC_PARAM               = \$(I \$0)
DDOC_CONSTRAINT          = <span class="constraint">if ()</span> if ()
DDOC_OVERLOAD_SEPARATOR =
DDOC_TEMPLATE_PARAM_LIST=
DDOC_TEMPLATE_PARAM      =

```



```
DDOC_BACKQUOTED      = \$(D_INLINECODE \$0)
```

```
ESCAPES = /</\&lt;/
          />/\&gt;/
          /\&/\&amp;/
```

Ddoc does not generate HTML code. It formats into the basic formatting macros, which (in their predefined form) are then expanded into HTML. If output other than HTML is desired, then these macros need to be redefined.

Table 31.1: Basic Formatting Macros

Name	Description
B	boldface the argument
I	italicize the argument
U	underline the argument
P	argument is a paragraph
DL	argument is a definition list
DT	argument is a definition in a definition list
DD	argument is a description of a definition
TABLE	argument is a table
TR	argument is a row in a table
TH	argument is a header entry in a row
TD	argument is a data entry in a row
OL	argument is an ordered list
UL	argument is an unordered list
LI	argument is an item in a list
BIG	argument is one font size bigger
SMALL	argument is one font size smaller
BR	start new line
LINK	generate clickable link on argument
LINK2	generate clickable link, first arg is address
RED	argument is set to be red
BLUE	argument is set to be blue
GREEN	argument is set to be green
YELLOW	argument is set to be yellow
BLACK	argument is set to be black
WHITE	argument is set to be white
D_CODE	argument is D code

(continued)

Name	Description
D_INLINECODE	argument is inline D code
DDOC	overall template for output

DDOC is special in that it specifies the boilerplate into which the entire generated text is inserted (represented by the Ddoc generated macro **BODY**). For example, in order to use a style sheet, **DDOC** would be redefined as:

```
DDOC = <!DOCTYPE HTML PUBLIC "-//W3C//DTD_HTML_4.01//EN" "http://www.w3.org/TR/2001/11/html4/strict.dtd">
<html><head>
<META http-equiv="content-type" content="text/html; charset=utf-8">
<title>\$(TITLE)</title>
<link rel="stylesheet" type="text/css" href="style.css">
</head><body>
<h1>\$(TITLE)</h1>
\$(BODY)
</body></html>
```

DDOC_COMMENT is used to insert comments into the output file.
Highlighting of D code is performed by the following macros:

Table 31.2: D Code Formatting Macros

Name	Description
D_COMMENT	Highlighting of comments
D_STRING	Highlighting of string literals
D_KEYWORD	Highlighting of D keywords
D_PSYMBOL	Highlighting of current declaration name
D_PARAM	Highlighting of current function declaration parameters

The highlighting macros start with **DDOC_**. They control the formatting of individual parts of the presentation.

Table 31.3: Ddoc Section Formatting Macros

Name	Description
DDOC_COMMENT	Inserts a HTML comment.
DDOC_DECL	Highlighting of the declaration.
DDOC_DECL_DD	Highlighting of the description of a declaration.
DDOC_DITTO	Highlighting of ditto declarations.
DDOC_SECTIONS	Highlighting of all the sections.
DDOC_SUMMARY	Highlighting of the summary section.
DDOC_DESCRIPTION	Highlighting of the description section.
DDOC_AUTHORS	Highlighting of the authors section.
DDOC_BUGS	Highlighting of the bugs section.
DDOC_COPYRIGHT	Highlighting of the copyright section.
DDOC_DATE	Highlighting of the date section.
DDOC_DEPRECATED	Highlighting of the deprecated section.
DDOC_EXAMPLES	Highlighting of the examples section.
DDOC_HISTORY	Highlighting of the history section.
DDOC_LICENSE	Highlighting of the license section.
DDOC_RETURNS	Highlighting of the returns section.
DDOC_SEE_ALSO	Highlighting of the see-also section.
DDOC_STANDARDS	Highlighting of the standards section.
DDOC_THROWS	Highlighting of the throws section.
DDOC_VERSION	Highlighting of the version section.
DDOC_SECTION_H	Highlighting of the section name of a non-standard section.
DDOC_SECTION	Highlighting of the contents of a non-standard section.
DDOC_MEMBERS	Default highlighting of all the members of a class, struct, etc.
DDOC_MODULE_MEMBERS	Highlighting of all the members of a module.
DDOC_CLASS_MEMBERS	Highlighting of all the members of a class.
DDOC_STRUCT_MEMBERS	Highlighting of all the members of a struct.
DDOC_ENUM_MEMBERS	Highlighting of all the members of an enum.
DDOC_TEMPLATE_MEMBERS	Highlighting of all the members of a template.
DDOC_ENUM_BASETYPE	Highlighting of the type an enum is based upon
DDOC_PARAMS	Highlighting of a function parameter section.
DDOC_PARAM_ROW	Highlighting of a name=value function parameter.
DDOC_PARAM_ID	Highlighting of the parameter name.
DDOC_PARAM_DESC	Highlighting of the parameter value.
DDOC_BLANKLINE	Inserts a blank line.

(continued)

Name	Description
<code>DDOC_ANCHOR</code>	Expands to a named anchor used for hyperlinking to a particular declaration section. Argument <code>\$1</code> expands to the qualified declaration name.
<code>DDOC_PSYMBOL</code>	Highlighting of declaration name to which a particular section is referring.
<code>DDOC_PSUPER_SYMBOL</code>	Highlighting of the base type of a class.
<code>DDOC_KEYWORD</code>	Highlighting of D keywords.
<code>DDOC_PARAM</code>	Highlighting of function parameters.
<code>DDOC_BACKQUOTED</code>	Inserts inline code.

For example, one could redefine `DDOC_SUMMARY`:

```
DDOC_SUMMARY = \$(GREEN \ $0)
```

And all the summary sections will now be green.

Macro Definitions from `sc.ini`'s `DDOCFILE`

A text file of macro definitions can be created, and specified in `sc.ini`:

```
DDOCFILE=myproject.ddoc
```

Macro Definitions from `.ddoc` Files on the Command Line

File names on the DMD command line with the extension `.ddoc` are text files that are read and processed in order.

Macro Definitions Generated by Ddoc

Table 31.4: Generated Macro Definitions

Macro Name	Content
BODY	Set to the generated document text.
TITLE	Set to the module name.
DATETIME	Set to the current date and time.
YEAR	Set to the current year.
COPYRIGHT	Set to the contents of any Copyright: section that is part of the module comment.

←

(continued)

Macro Name	Content
DOCFILENAME	Set to the name of the generated output file.
SRCFILENAME	Set to the name of the source file the documentation is being generated from.

31.4 Using Ddoc to generate examples from unit tests

Ddoc can automatically generate usage examples for declarations using unit tests. If a declaration is followed by a documented unit test, the code from the test will be inserted into the example section of the declaration. This avoids the frequent problem of having outdated documentation for pieces of code.

To create a documented unit test just add three forward slashes before the unittest block, like this:

```
///
unittest
{
    ...
}
```

For more information please see the full section on documented unit tests.

31.5 Using Ddoc for other Documentation

Ddoc is primarily designed for use in producing documentation from embedded comments. It can also, however, be used for processing other general documentation. The reason for doing this would be to take advantage of the macro capability of Ddoc and the D code syntax highlighting capability.

If the .d source file starts with the string "Ddoc" then it is treated as general purpose documentation, not as a D code source file. From immediately after the "Ddoc" string to the end of the file or any "Macros:" section forms the document. No automatic highlighting is done to that text, other than highlighting of D code embedded between lines delineated with — lines. Only macro processing is done.

Much of the D documentation itself is generated this way, including this page. Such documentation is marked at the bottom as being generated by Ddoc.

31.6 Links to D documentation generators

A list of current D documentation generators which use Ddoc can be found on our wiki page.

Chapter 32

Interfacing to C

D is designed to fit comfortably with a C compiler for the target system. D makes up for not having its own VM by relying on the target environment's C runtime library. It would be senseless to attempt to port to D or write D wrappers for the vast array of C APIs available. How much easier it is to just call them directly.

This is done by matching the C compiler's data types, layouts, and function call/return sequences.

32.1 Calling C Functions

C functions can be called directly from D. There is no need for wrapper functions, argument swizzling, and the C functions do not need to be put into a separate DLL.

The C function must be declared and given a calling convention, most likely the "C" calling convention, for example:

```
extern (C) int strcmp(const char* string1, const char* string2);
```

and then it can be called within D code in the obvious way:

```
import std.string;
int myDfunction(char[] s)
{
    return strcmp(std.string.toStringz(s), "foo");
}
```

There are several things going on here:

- D understands how C function names are "mangled" and the correct C function call/return sequence.
- C functions cannot be overloaded with another C function with the same name.

- There are no `__cdecl`, `__far`, `__stdcall`, `__declspec`, or other such C extended type modifiers in D. These are handled by linkage attributes, such as `extern (C)`.
- There is no volatile type modifier in D. To declare a C function that uses volatile, just drop the keyword from the declaration.
- Strings are not 0 terminated in D. See "Data Type Compatibility" for more information about this. However, string literals in D are 0 terminated.

C code can correspondingly call D functions, if the D functions use an attribute that is compatible with the C compiler, most likely the `extern (C)`:

```
// myfunc() can be called from any C function
extern (C)
{
    void myfunc(int a, int b)
    {
        ...
    }
}
```

32.2 Storage Allocation

C code explicitly manages memory with calls to `malloc()` and `free()`. D allocates memory using the D garbage collector, so no explicit `free`'s are necessary.

D can still explicitly allocate memory using `core.stdc.stdlib.malloc()` and `core.stdc.stdlib.free()`, these are useful for connecting to C functions that expect `malloc`'d buffers, etc.

If pointers to D garbage collector allocated memory are passed to C functions, it's critical to ensure that that memory will not be collected by the garbage collector before the C function is done with it. This is accomplished by:

- Making a copy of the data using `core.stdc.stdlib.malloc()` and passing the copy instead.
- Leaving a pointer to it on the stack (as a parameter or automatic variable), as the garbage collector will scan the stack.
- Leaving a pointer to it in the static data segment, as the garbage collector will scan the static data segment.
- Registering the pointer with the garbage collector with the `std.gc.addRoot()` or `std.gc.addRange()` calls.

An interior pointer to the allocated memory block is sufficient to let the GC know the object is in use; i.e. it is not necessary to maintain a pointer to the beginning of the allocated memory.

The garbage collector does not scan the stacks of threads not created by the D Thread interface. Nor does it scan the data segments of other DLL's, etc.

32.3 Data Type Compatibility

D And C Type Equivalence		
D	C	
	32 bit	64 bit
<code>void</code>	<code>void</code>	
<code>byte</code>	<code>signed char</code>	
<code>ubyte</code>	<code>unsigned char</code>	
<code>char</code>	<code>char</code> (chars are unsigned in D)	
<code>wchar</code>	<code>wchar_t</code> (when <code>sizeof(wchar_t)</code> is 2)	
<code>dchar</code>	<code>wchar_t</code> (when <code>sizeof(wchar_t)</code> is 4)	
<code>short</code>	<code>short</code>	
<code>ushort</code>	<code>unsigned short</code>	
<code>int</code>	<code>int</code>	
<code>uint</code>	<code>unsigned</code>	
<code>ulong</code>	<code>unsigned long long</code>	<code>unsigned long</code>
<code>core.stdc.config.c_long</code>	<code>long</code>	<code>long</code>
<code>core.stdc.config.c_ulong</code>	<code>unsigned long</code>	<code>unsigned long</code>
<code>long</code>	<code>long long</code>	<code>long</code> (or <code>long long</code>)
<code>ulong</code>	<code>unsigned long long</code>	<code>unsigned long</code> (or <code>unsigned long long</code>)
<code>float</code>	<code>float</code>	
<code>double</code>	<code>double</code>	
<code>real</code>	<code>long double</code>	
<code>cdouble</code>	<code>double _Complex</code>	
<code>creal</code>	<code>long double _Complex</code>	
<code>struct</code>	<code>struct</code>	
<code>union</code>	<code>union</code>	
<code>enum</code>	<code>enum</code>	
<code>class</code>	no equivalent	
<code>type *</code>	<code>type *</code>	
<code>type[dim]</code>	<code>type[dim]</code>	
<code>type[dim]*</code>	<code>type(*)[dim]</code>	
<code>type[]</code>	no equivalent	
<code>type1[type2]</code>	no equivalent	
<code>type function(params)</code>	<code>type(*) (params)</code>	
<code>type delegate(params)</code>	no equivalent	
<code>size_t</code>	<code>size_t</code>	
<code>ptrdiff_t</code>	<code>ptrdiff_t</code>	

These equivalents hold for most C compilers. The C standard does not pin down the sizes of the types, so some care is needed.

32.4 Passing D Array Arguments to C Functions

In C, arrays are passed to functions as pointers even if the function prototype says its an array. In D, static arrays are passed by value, not by reference. Thus, the function prototype must be adjusted to match what C expects.

D And C Function Prototype Equivalence

D type	C type
T^*	$T[]$
<code>ref T[dim]</code>	$T[dim]$

For example:

```
void foo(int a[3]) { ... } // C code

extern (C)
{
    void foo(ref int[3] a); // D prototype
}
```

32.5 Calling printf()

This mostly means checking that the printf format specifier matches the corresponding D data type. Although printf is designed to handle 0 terminated strings, not D dynamic arrays of chars, it turns out that since D dynamic arrays are a length followed by a pointer to the data, the `%.s` format works:

```
void foo(char[] string)
{
    printf("my_string is: %.s\n", string.length, string.ptr);
}
```

The `printf` format string literal in the example doesn't end with `'\0'`. This is because string literals, when they are not part of an initializer to a larger data structure, have a `'\0'` character helpfully stored after the end of them.

An improved D function for formatted output is `std.stdio.writef()`.

32.6 Structs and Unions

D structs and unions are analogous to C's.

C code often adjusts the alignment and packing of struct members with a command line switch or with various implementation specific `#pragma`'s. D supports explicit alignment attributes that correspond to the C compiler's rules. Check what alignment the C code is using, and explicitly set it for the D struct declaration.

D does not support bit fields. If needed, they can be emulated with shift and mask operations, or use the `std.bitmanip.bitfields` library type. `htod` will convert bit fields to inline functions that do the right shift and masks.

D does not support declaring variables of anonymous struct types. In such a case you can define a named struct in D and make it private:

```
union Info // C code
{
    struct
    {
        char *name;
    } file;
};

union Info // D code
{
    private struct File
    {
        char* name;
    }
    File file;
}
```

32.7 Callbacks

D can easily call C callbacks (function pointers), and C can call callbacks provided by D code if the callback is an `extern(C)` function, or some other linkage that both sides have agreed to (e.g. `extern(Windows)`).

Here's an example of C code providing a callback to D code:

```
void someFunc(void *arg) { printf("Called_someFunc!\n"); } // C code
typedef void (*Callback)(void *);
extern "C" Callback getCallback(void)
{
```

```

    return someFunc;
}

extern(C) alias Callback = int function(int, int); // D code
extern(C) Callback getCallback();
void main()
{
    Callback cb = getCallback();
    cb(); // invokes the callback
}

```

And an example of D code providing a callback to C code:

```

extern "C" void printer(int (*callback)(int, int)) // C code
{
    printf("calling_callback_with_2_and_4_returns:%d\n", callback(2, 4));
}

extern(C) alias Callback = int function(int, int); // D code
extern(C) void printer(Callback callback);
extern(C) int sum(int x, int y) { return x + y; }
void main()
{
    printer(&sum);
}

```

For more info about callbacks read the closures section.

32.8 Using Existing C Libraries

Since D can call C code directly, it can also call any C library functions, giving D access to the smorgasbord of existing C libraries. To do so, however, one needs to write a D interface (.di) file, which is a translation of the C .h header file for the C library into D.

For popular C libraries, the first place to look for the corresponding D interface file is the Deimos Project. If it isn't there already, and you write one, please contribute it to the Deimos Project.

32.9 Accessing C Globals

C globals can be accessed directly from D. C globals have the C naming convention, and so must be in an `extern` (C) block. Use the `extern` storage class to indicate that the global is allocated in the C code, not the D code. C globals default to being in global, not thread local, storage. To reference global storage from D, use the `__gshared` storage class.

```
extern (C) extern __shared int x;
```

Chapter 33

Interfacing to C++

This document specifies how to interface with C++ directly.

It is also possible to indirectly interface with C++ code, either through a C interface or a COM interface.

33.1 The General Idea

Being 100% a fully functional C++ compiler front end to D. Anecdotal evidence suggests that writing such is a minimum of a 10 man-year project, essentially making a D compiler with such capability unimplementable. Other languages looking to hook up to C++ face the same problem, and the solutions have been:

1. Support the COM interface (but that only works for Windows).
2. Laboriously construct a C wrapper around the C++ code.
3. Use an automated tool such as SWIG to construct a C wrapper.
4. Reimplement the C++ code in the other language.
5. Give up.

D takes a pragmatic approach that assumes a couple modest accommodations can solve a significant chunk of the problem:

- matching C++ name mangling conventions
- matching C++ function calling conventions
- matching C++ virtual function table layout for single inheritance

33.2 Global Functions

C++ global functions, including those in namespaces, can be declared and called in D, or defined in D and called in C++.

Calling C++ Global Functions from D

Given a C++ function in a C++ source file:

```
#include <iostream>

using namespace std;

int foo(int i, int j, int k)
{
    cout << "i_=" << i << endl;
    cout << "j_=" << j << endl;
    cout << "k_=" << k << endl;

    return 7;
}
```

In the corresponding D code, `foo` is declared as having C++ linkage and function calling conventions:

```
extern (C++) int foo(int i, int j, int k);
```

and then it can be called within the D code:

```
extern (C++) int foo(int i, int j, int k);
```

```
void main()
{
    foo(1, 2, 3);
}
```

Compiling the two files, the first with a C++ compiler, the second with a D compiler, linking them together, and then running it yields:

```
i = 1 j = 2 k = 3
```

There are several things going on here:

- D understands how C++ function names are "mangled" and the correct C++ function call/return sequence.

- Because modules are not part of C++, each function with C++ linkage in the global namespace must be globally unique within the program.
- There are no `__cdecl`, `__far`, `__stdcall`, `__declspec`, or other such nonstandard C++ extensions in D.
- There are no volatile type modifiers in D.
- Strings are not 0 terminated in D. See "Data Type Compatibility" for more information about this. However, string literals in D are 0 terminated.

Calling Global D Functions From C++

To make a D function accessible from C++, give it C++ linkage:

```
import std.stdio;

extern (C++) int foo(int i, int j, int k)
{
    writeln("i=%s", i);
    writeln("j=%s", j);
    writeln("k=%s", k);
    return 1;
}

extern (C++) void bar();

void main()
{
    bar();
}
```

The C++ end looks like:

```
int foo(int i, int j, int k);

void bar()
{
    foo(6, 7, 8);
}
```

Compiling, linking, and running produces the output:

```
i = 6 j = 7 k = 8
```

33.3 C++ Namespaces

C++ symbols that reside in namespaces can be accessed from D. A namespace can be added to the `extern (C++) LinkageAttribute`:

```
extern (C++, N) int foo(int i, int j, int k);

void main()
{
    N.foo(1, 2, 3);    // foo is in C++ namespace 'N'
}
```

33.4 Classes

C++ classes can be declared in D by using the `extern (C++)` attribute on `class`, `struct` and `interface` declarations. `extern (C++)` interfaces have the same restrictions as D interfaces, which means that Multiple Inheritance is supported to the extent that only one base class can have member fields.

`extern (C++)` structs do not support virtual functions but can be used to map C++ value types.

Unlike classes and interfaces with D linkage, `extern (C++)` classes and interfaces are not rooted in `Object` and cannot be used with `typeid`.

D structs and classes have different semantics whereas C++ structs and classes are basically the same. The use of a D struct or class depends on the C++ implementation and not on the used C++ keyword. When mapping a D `class` onto a C++ `struct`, use `extern(C++, struct)` to avoid linking problems with C++ compilers (notably MSVC) that distinguish between C++'s `class` and `struct` when mangling. Conversely, use `extern(C++, class)` to map a D `struct` onto a C++ `class`.

`extern(C++, class)` and `extern(C++, struct)` can be combined with C++ namespaces:

```
extern (C++, struct) extern (C++, foo)
class Bar
{
}
```

Using C++ Classes From D

The following example shows binding of a pure virtual function, its implementation in a derived class, a non-virtual member function, and a member field:

```
#include <iostream>

using namespace std;

class Base
{
    public:
        virtual void print3i(int a, int b, int c) = 0;
};

class Derived : public Base
{
    public:
        int field;
        Derived(int field) : field(field) {}

        void print3i(int a, int b, int c)
        {
            cout << "a_=" << a << endl;
            cout << "b_=" << b << endl;
            cout << "c_=" << c << endl;
        }

        int mul(int factor);
};

int Derived::mul(int factor)
{
    return field * factor;
}

Derived *createInstance(int i)
{
    return new Derived(i);
}

void deleteInstance(Derived *&d)
{
    delete d;
    d = 0;
}
```

```
}
```

We can use it in D code like:

```
extern(C++)
{
    interface Base
    {
        void print3i(int a, int b, int c);
    }

    class Derived : Base
    {
        int field;
        @disable this();
        override void print3i(int a, int b, int c);
        final int mul(int factor);
    }

    Derived createInstance(int i);
    void deleteInstance(ref Derived d);
}

void main()
{
    import std.stdio;

    auto d1 = createInstance(5);
    writeln(d1.field);
    writeln(d1.mul(4));

    Base b1 = d1;
    b1.print3i(1, 2, 3);

    deleteInstance(d1);
    assert(d1 is null);

    auto d2 = createInstance(42);
    writeln(d2.field);

    deleteInstance(d2);
}
```

```
    assert(d2 is null);
}
```

Compiling, linking, and running produces the output:

```
5 20 a = 1 b = 2 c = 3 42
```

Note how in the above example, the constructor is not bindable and is instead disabled on the D side; an alternative would be to reimplement the constructor in D. See the section below on lifetime management for more information.

Using D Classes From C++

Given D code like:

```
extern (C++) int callE(E);

extern (C++) interface E
{
    int bar(int i, int j, int k);
}

class F : E
{
    extern (C++) int bar(int i, int j, int k)
    {
        writeln("i=%s", i);
        writeln("j=%s", j);
        writeln("k=%s", k);
        return 8;
    }
}

void main()
{
    F f = new F();
    callE(f);
}
```

The C++ code to access it looks like:

```
class E
{
public:
```

```
    virtual int bar(int i, int j, int k);  
};  
  
int callE(E *e)  
{  
    return e->bar(11, 12, 13);  
}
```

33.5 C++ Templates

C++ function and type templates can be bound by using the `extern` (C++) attribute on a function or type template declaration.

Note that all instantiations used in D code must be provided by linking to C++ object code or shared libraries containing the instantiations.

For example:

```
#include <iostream>  
  
template<class T>  
struct Foo  
{  
    private:  
        T field;  
  
    public:  
        Foo(T t) : field(t) {}  
        T get();  
        void set(T t);  
};  
  
template<class T>  
T Foo<T>::get()  
{  
    return field;  
}  
  
template<class T>  
void Foo<T>::set(T t)  
{
```

```
        field = t;
    }

    Foo<int> makeIntFoo(int i)
    {
        return Foo<int>(i);
    }

    Foo<char> makeCharFoo(char c)
    {
        return Foo<char>(c);
    }

    template<class T>
    void increment(Foo<T> &foo)
    {
        foo.set(foo.get() + 1);
    }

    template<class T>
    void printThreeNext(Foo<T> foo)
    {
        for(size_t i = 0; i < 3; ++i)
        {
            std::cout << foo.get() << std::endl;
            increment(foo);
        }
    }

    // The following two functions ensure that the required instantiations of
    // printThreeNext are provided by this code module
    void printThreeNexti(Foo<int> foo)
    {
        printThreeNext(foo);
    }

    void printThreeNextc(Foo<char> foo)
    {
        printThreeNext(foo);
    }
```

```
extern(C++):
struct Foo(T)
{
    private:
        T field;

    public:
        @disable this();
        T get();
        void set(T t);
}

Foo!int makeIntFoo(int i);
Foo!char makeCharFoo(char c);
void increment(T)(ref Foo!T foo);
void printThreeNext(T)(Foo!T foo);

extern(D) void main()
{
    auto i = makeIntFoo(42);
    assert(i.get() == 42);
    i.set(1);
    increment(i);
    assert(i.get() == 2);

    auto c = makeCharFoo('a');
    increment(c);
    assert(c.get() == 'b');

    c.set('A');
    printThreeNext(c);
}
```

Compiling, linking, and running produces the output:

```
A B C
```

33.6 Function Overloading

C++ and D follow different rules for function overloading. D source code, even when calling `extern` (C++) functions, will still follow D overloading rules.

33.7 Memory Allocation

C++ code explicitly manages memory with calls to `::operator new()` and `::operator delete` (↪ `()`). D's `new` operator allocates memory using the D garbage collector, so no explicit delete is necessary. D's `new` operator is not compatible with C++'s `::operator new` and `::operator delete` (↪ `delete`). Attempting to allocate memory with D's `new` and deallocate with C++ `::operator delete` (↪ `delete`) will result in miserable failure.

D can explicitly manage memory using a variety of library tools, such as with `std.experimental allocator`. Additionally, `core.stdc.stdlib.malloc` and `core.stdc.stdlib.free` (↪ `.free`) can be used directly for connecting to C++ functions that expect `malloc`'d buffers.

If pointers to memory allocated on the D garbage collector heap are passed to C++ functions, it's critical to ensure that the referenced memory will not be collected by the D garbage collector before the C++ function is done with it. This is accomplished by:

- Making a copy of the data using `std.experimental allocator` or `core.stdc.stdlib.malloc` (↪ `malloc`) and passing the copy instead.
- Leaving a pointer to it on the stack (as a parameter or automatic variable), as the garbage collector will scan the stack.
- Leaving a pointer to it in the static data segment, as the garbage collector will scan the static data segment.
- Registering the pointer with the garbage collector using the `core.memory.GC.addRoot` or `core.memory.GC.addRange` functions.

An interior pointer to the allocated memory block is sufficient to let the GC know the object is in use; i.e. it is not necessary to maintain a pointer to the *beginning* of the allocated memory.

The garbage collector does not scan the stacks of threads not registered with the D runtime, nor does it scan the data segments of shared libraries that aren't registered with the D runtime.

33.8 Data Type Compatibility

D And C++ Type Equivalence	
D type	C++ type
void	void
byte	signed char
ubyte	unsigned char
char	char (chars are unsigned in D)
<code>core.stdc.stddef.wchar_t</code>	<code>wchar_t</code>
short	short
ushort	unsigned short
int	int
uint	unsigned
long	long long
ulong	unsigned long long
<code>core.stdc.config.cpp_long</code>	long
<code>core.stdc.config.cpp_ulong</code>	unsigned long
float	float
double	double
real	long double
<code>extern</code> (C++) struct	struct or class
<code>extern</code> (C++) class	struct or class
<code>extern</code> (C++) interface	struct or class with no member fields
union	union
enum	enum
<i>type</i> *	<i>type</i> *
ref <i>type</i> (in parameter lists only)	<i>type</i> &
<i>type</i> [<i>dim</i>]	<i>type</i> [<i>dim</i>]
<i>type</i> [<i>dim</i>]*	<i>type</i> (*)[<i>dim</i>]
<i>type</i> []	no equivalent
<i>type</i> [<i>type</i>]	no equivalent
<i>type</i> function (<i>parameters</i>)	<i>type</i> (*)(<i>parameters</i>)
<i>type</i> delegate (<i>parameters</i>)	no equivalent

These equivalents hold when the D and C++ compilers used are companions on the host platform.

33.9 Packing and Alignment

D structs and unions are analogous to C's.

C code often adjusts the alignment and packing of struct members with a command line switch or with various implementation specific `#pragma`'s. D supports explicit alignment attributes that correspond to the C compiler's rules. Check what alignment the C code is using, and explicitly set it for the D struct declaration.

D supports bitfields in the standard library: see `std.bitmanip.bitfields`.

33.10 Lifetime Management

C++ constructors, copy constructors, move constructors and destructors cannot be called directly in D code, and D constructors, postblit operators and destructors cannot be directly exported to C++ code. Interoperation of types with these special operators is possible by either 1) disabling the operator in the client language and only using it in the host language, or 2) faithfully reimplementing the operator in the client language. With the latter approach, care needs to be taken to ensure observable semantics remain the same with both implementations, which can be difficult, or in some edge cases impossible, due to differences in how the operators work in the two languages. For example, in D all objects are movable and there is no move constructor.

33.11 Special Member Functions

D cannot directly call C++ special member functions, and vice versa. These include constructors, destructors, conversion operators, operator overloading, and allocators.

33.12 Runtime Type Identification

D runtime type identification uses completely different techniques than C++. The two are incompatible.

33.13 Exception Handling

Exception interoperability is a work in progress.

At present, C++ exceptions cannot be caught in or thrown from D, and D exceptions cannot be caught in or thrown from C++. Additionally, objects in C++ stack frames are not guaranteed to be destroyed when unwinding the stack due to a D exception, and vice versa.

The plan is to support all of the above except throwing D exceptions directly in C++ code (but they will be throwable indirectly by calling into a D function with C++ linkage).

33.14 Comparing D Immutable and Const with C++ Const

Const, Immutable Comparison		
Feature	D	C++98
<code>const</code> keyword	Yes	Yes
<code>immutable</code> keyword	Yes	No
const notation	<pre>// Functional: //ptr to const ptr to const int const(int*)* p; // Yes: //const ptr to const ptr to const int const int** p; **p = 3; // error // Yes: // ptr to const int const(int)* p; int* q = cast(int*)p; // ok // No: // ptr to const int const(int)* p; int* q = cast(int*)p; *q = 3; // undefined behavior</pre>	<pre>// Postfix: //ptr to const ptr to const int const int *const *p; // No: // const ptr to ptr to int int** const p; **p = 3; // ok // Yes: // ptr to const int const int* p; int* q = const_cast<int*>p; // Yes: // ptr to const int const int* p; int* q = const_cast<int*>p; *q = 3; // ok // No: void foo(int x); void foo(const int x); //error // Yes: void foo(const int* x, int* y) { bar(*x); // bar(3) *y = 4; bar(*x); // bar(4) } ... int i = 3; foo(&i, &i); // No: void foo(immutable int* x, int* y) { bar(*x); // bar(3) *y = 4; // undefined behavior bar(*x); // bar(??) } ... int i = 3; foo(cast(immutable)&i, &i);</pre>
transitive const		
cast away const		
cast+mutate		
overloading		
const/mutable aliasing		
immutable/mutable aliasing		No immutables
type of string literal	<code>immutable(char) []</code>	<code>const char*</code>
string literal to non-	not allowed	allowed but deprecated

Chapter 34

Interfacing to Objective-C

D has limited support for interfacing with Objective-C. It supports external classes and calling instance methods. It is only available on OS X, compiling for 64bit.

Fully working example is available at the bottom.

34.1 Declaring an External Class

```
extern (Objective-C)
interface NSString
{
    const(char)* UTF8String() @selector("UTF8String");
}
```

Currently all Objective-C classes need to be declared as interfaces in D. All Objective-C classes that should be accessible from within D need to be declared with the Objective-C linkage.

The `@selector` attribute indicates which Objective-C selector should be used when calling this method from D. This attribute needs to be attached to all methods.

34.2 Calling an Instance Method

Calling an Objective-C instance method uses the same syntax as calling regular D methods:

```
const(char)* result = object.UTF8String();
```

When the compiler sees a call to a method with Objective-C linkage it will generate a call similar to how an Objective-C compiler would call the method.

34.3 The @selector Attribute

The `@selector` attribute is a compiler recognized UDA. It is used to tell the compiler which selector to use when calling an Objective-C method.

Selectors in Objective-C can contain the colon character, which is not valid in D identifiers. D supports method overloading while Objective-C achieves something similar by using different selectors. For these two reasons it is better to be able to specify the selectors manually in D, instead of trying to infer it. This allows to have a more natural names for the methods in D. Example:

```
extern (Objective-C)
interface NSString
{
    NSString initWith(in char*) @selector("initWithUTF8String:");
    NSString initWith(NSString) @selector("initWithString:");
}
```

Here the method `initWith` is overloaded with two versions, one accepting in `char*`, the other one `NSString`. These two methods are mapped to two different Objective-C selectors, `initWithUTF8String:` and `initWithString:`.

The attribute is defined in druntime in `core.attribute` and aliased in `object`, meaning it will be implicitly imported. The attribute is only defined when the version identifier `D_ObjectiveC` is enabled.

Compiler Checks

The compiler performs the following checks to enforce the correct usage of the `@selector` attribute:

- The attribute can only be attached to methods with Objective-C linkage
- The attribute can only be attached once to a method
- The attribute cannot be attached to a template method
- The number of colons in the selector needs to match the number of parameters the method is declared with

If any of the checks fail, a compile error will occur.

34.4 The D_ObjectiveC Version Identifier

The `D_ObjectiveC` version identifier is a predefined version identifier. It is enabled if Objective-C support is available for the target.

34.5 Objective-C Linkage

Objective-C linkage is achieved by attaching the `extern (Objective-C)` attribute to an interface. Example:

```
extern (Objective-C)
interface NSObject
{
    NSObject init() @selector("init");
}
```

All methods inside an interface declared as `extern (Objective-C)` will get implicit Objective-C linkage.

The linkage is recognized on all platforms but will issue a compile error if it is used on a platform where Objective-C support is not available. This allows to easily hide Objective-C declarations from platforms where it is not available using the `version` statement, without resorting to string mixins or other workarounds.

34.6 Memory Management

The preferred way to do memory management in Objective-C is to use Automatic Reference Counting, ARC. This is not supported in D, therefore manual memory management is required to be used instead. This is achieved by calling `release` on an Objective-C instance, like in the old days of Objective-C.

34.7 Frameworks

Most Objective-C code is bundled in something called a "Framework". This is basically a regular directory, with the `.framework` extension and a specific directory layout. A framework contains a dynamic library, all public header files and any resources (images, sounds and so on) required by the framework.

These directories are recognized by some tools, like the Objective-C compiler and linker, to be frameworks. To link with a framework from DMD, use the following flags:

```
-L-framework -L<Framework>
```

Where `<Framework>` is the name of the framework to link with, without the `.framework` extension. The two `-L` flags are required because the linker expects a space between the `-framework` flag and the name of the framework. DMD cannot handle this and will instead interpret the name of the framework as a separate flag.

Framework Paths

Using the above flag, the linker will search in the standard framework paths. The standard search paths for frameworks are:

- /System/Library/Frameworks
- /Library/Frameworks

The following flag from DMD can be used to add a new path in which to search for frameworks:

`-L-F<framework_path>`

For more information see the reference documentation and the `ld` man page.

34.8 Full Usage Example

Since the only parts of Objective-C that is currently supported is calling instance methods, this example demonstrates how the Objective-C runtime can be used to achieve a running example.

This example will create an Objective-C string, `NSString`, and log the message using `NSLog` to `stderr`.

```
extern (Objective-C)
interface Class
{
    NSString alloc() @selector("alloc");
}
```

This interface is used to emulate the `Class` type available in the Objective-C runtime. The instance method `alloc` will be used to emulate a class method in `NSObject`, `alloc`.

```
extern (Objective-C)
interface NSString
{
    NSString initWithUTF8String(in char* str) @selector("initWithUTF8String:");
    void release() @selector("release");
}
```

This is a simplified declaration of the `NSString` class. The `initWithUTF8String:` method will be used to convert a C string in UTF-8 to an Objective-C string, `NSString`. The `release` method is used to release and deallocate the string. Since D doesn't support ARC it's needed to manually release Objective-C instances.

```
extern (C) Class objc_lookupClass(in char* name)
```

The `objc_lookupClass` function is used to get the class definition of the class with the given name.

```
extern (C) void NSLog(NSString, ...);
```

This `NSLog` function prints a message to the System Log facility, i.e. to `stderr` and Console.

```
auto cls = objc_lookupClass("NSString");
```

Get the class definition of `NSString`.

```
auto str = cls.alloc();
```

Allocate an instance of the class, `NSString`.

```
str = str.initWithUTF8String("Hello_World!");
```

Initialize the Objective-C string using a C string.

```
NSLog(str);
```

Log the string to `stderr`, this will print something like this in the terminal:

```
2015-07-18 13:14:27.978 main[11045:2934950] Hello World!
```

```
str.release();
```

Release and deallocate the string.

All steps combined look like this:

```
module main;
```

```
extern (Objective-C)
interface Class
{
    NSString alloc() @selector("alloc");
}
```

```
extern (Objective-C)
interface NSString
{
    NSString initWithUTF8String(in char* str) @selector("initWithUTF8String:");
    void release() @selector("release");
}
```

```
extern (C) void NSLog(NSString, ...);
extern (C) Class objc_lookupClass(in char* name);
```

```
void main()
{
    auto cls = objc_lookUpClass("NSString");
    auto str = cls.alloc().initWithUTF8String("Hello_!World!");
    NSLog(str);
    str.release();
}
```

When compiling the application remember to link with the required libraries, in this case the Foundation framework. Example:

```
dmd -L-framework -LFoundation main.d
```

Chapter 35

Portability Guide

It's good software engineering practice to minimize gratuitous portability problems in the code. Techniques to minimize potential portability problems are:

- The integral and floating type sizes should be considered as minimums. Algorithms should be designed to continue to work properly if the type size increases.
- Floating point computations can be carried out at a higher precision than the size of the floating point variable can hold. Floating point algorithms should continue to work properly if precision is arbitrarily increased.
- Avoid depending on the order of side effects in a computation that may get reordered by the compiler. For example:

`a + b + c`

can be evaluated as $(a + b) + c$, $a + (b + c)$, $(a + c) + b$, $(c + b) + a$, etc. Parentheses control operator precedence, parentheses do not control order of evaluation.

If the operands of an associative operator `+` or `*` are floating point values, the expression is not reordered.

- Avoid dependence on byte order; i.e. whether the CPU is big-endian or little-endian.
- Avoid dependence on the size of a pointer or reference being the same size as a particular integral type.
- If size dependencies are inevitable, put a `static assert` in the code to verify it:

```
static assert(int.sizeof == (int*).sizeof);
```

35.1 32 to 64 Bit Portability

64 bit processors and operating systems are here. With that in mind:

- Integral types will remain the same sizes between 32 and 64 bit code.
- Pointers and object references will increase in size from 4 bytes to 8 bytes going from 32 to 64 bit code.
- Use `size_t` as an alias for an unsigned integral type that can span the address space. Array indices should be of type `size_t`.
- Use `ptrdiff_t` as an alias for a signed integral type that can span the address space. A type representing the difference between two pointers should be of type `ptrdiff_t`.
- The `.length`, `.size`, `.sizeof`, `.offsetof` and `.alignof` properties will be of type `size_t`.

35.2 Endianness

Endianness refers to the order in which multibyte types are stored. The two main orders are *big endian* and *little endian*. The compiler predefines the version identifier `BigEndian` or `LittleEndian` depending on the order of the target system. The x86 systems are all little endian.

The times when endianness matters are:

- When reading data from an external source (like a file) written in a different endian format.
- When reading or writing individual bytes of a multibyte type like `longs` or `doubles`.

35.3 OS Specific Code

System specific code is handled by isolating the differences into separate modules. At compile time, the correct system specific module is imported.

Minor differences can be handled by constant defined in a system specific import, and then using that constant in an *IfStatement* or *StaticIfStatement*.

Chapter 36

Named Character Entities

NamedCharacterEntity:
& Identifier ;

D supports the full list of named character entities from HTML 5. However, dmd does not yet support named entities which contain multiple code points. Below is a *partial* list of the named character entities. See the HTML 5 Spec for the full list.

Note: Not all will display properly in the **Symbol** column in all browsers.

Table 36.1: Named Character Entities

Name	Value	Symbol
quot	34	"
amp	38	&
lt	60	<
gt	62	>
OElig	338	Œ
oelig	339	œ
Scaron	352	Š
scaron	353	š
Yuml	376	ÿ
circ	710	^
tilde	732	~
ensp	8194	
emsp	8195	
thinsp	8201	

←

(continued)

Name	Value	Symbol
zwnj	8204	
zwj	8205	
lrm	8206	
rlm	8207	
ndash	8211	—
mdash	8212	—
lsquo	8216	‘
rsquo	8217	’
sbquo	8218	
ldquo	8220	“
rdquo	8221	”
bdquo	8222	”
dagger	8224	†
Dagger	8225	‡
permil	8240	‰
lsaquo	8249	‹
rsaquo	8250	›
euro	8364	€

Table 36.2: Latin-1 (ISO-8859-1) Entities

Name	Value	Symbol
nbs	160	
iexcl	161	¡
cent	162	¢
pound	163	£
curren	164	¤
yen	165	¥
brvbar	166	¸
sect	167	§
uml	168	¨
copy	169	©
ordf	170	ª
laquo	171	«
not	172	¬
shy	173	¸

(continued)

Name	Value	Symbol
reg	174	®
macr	175	—
deg	176	°
plusmn	177	±
sup2	178	²
sup3	179	³
acute	180	′
micro	181	μ
para	182	¶
middot	183	·
cedil	184	¸
sup1	185	¹
ordm	186	^o
raquo	187	»
frac14	188	$\frac{1}{4}$
frac12	189	$\frac{1}{2}$
frac34	190	$\frac{3}{4}$
iquest	191	¿
Agrave	192	À
Aacute	193	Á
Acirc	194	Â
Atilde	195	Ã
Auml	196	Ä
Aring	197	Å
AElig	198	Æ
Ccedil	199	Ç
Egrave	200	È
Eacute	201	É
Ecirc	202	Ê
Euml	203	Ë
Igrave	204	Ì
Iacute	205	Í
Icirc	206	Î
Iuml	207	Ï
ETH	208	Ð
Ntilde	209	Ñ

↩

(continued)

Name	Value	Symbol
Ograve	210	Ò
Oacute	211	Ó
Ocirc	212	Ô
Otilde	213	Õ
Ouml	214	Ö
times	215	×
Oslash	216	Ø
Ugrave	217	Û
Uacute	218	Ú
Ucirc	219	Û
Uuml	220	Ü
Yacute	221	Ý
THORN	222	Þ
szlig	223	ß
agrave	224	à
aacute	225	á
acirc	226	â
atilde	227	ã
auml	228	ä
aring	229	å
aelig	230	æ
ccedil	231	ç
egrave	232	è
eacute	233	é
ecirc	234	ê
euml	235	ë
igrave	236	ì
iacute	237	í
icirc	238	î
iuml	239	ï
eth	240	ð
ntilde	241	ñ
ograve	242	ò
oacute	243	ó
ocirc	244	ô
otilde	245	õ

↩

(continued)

Name	Value	Symbol
ouml	246	ö
divide	247	÷
oslash	248	ø
ugrave	249	ù
uacute	250	ú
ucirc	251	û
uuml	252	ü
yacute	253	ý
thorn	254	þ
yuml	255	ÿ

Table 36.3: Symbols and Greek letter entities

Name	Value	Symbol
fnof	402	<i>f</i>
Alpha	913	Α
Beta	914	Β
Gamma	915	Γ
Delta	916	Δ
Epsilon	917	Ε
Zeta	918	Ζ
Eta	919	Η
Theta	920	Θ
Iota	921	Ι
Kappa	922	Κ
Lambda	923	Λ
Mu	924	Μ
Nu	925	Ν
Xi	926	Ξ
Omicron	927	Ο
Pi	928	Π
Rho	929	Ρ
Sigma	931	Σ
Tau	932	Τ
Upsilon	933	Υ
Phi	934	Φ

↩

(continued)

Name	Value	Symbol
Chi	935	X
Psi	936	Ψ
Omega	937	Ω
alpha	945	α
beta	946	β
gamma	947	γ
delta	948	δ
epsilon	949	ϵ
zeta	950	ζ
eta	951	η
theta	952	θ
iota	953	ι
kappa	954	κ
lambda	955	λ
mu	956	μ
nu	957	ν
xi	958	ξ
omicron	959	o
pi	960	π
rho	961	ρ
sigmaf	962	ς
sigma	963	σ
tau	964	τ
upsilon	965	v
phi	966	ϕ
chi	967	χ
psi	968	ψ
omega	969	ω
thetasym	977	ϑ
upsih	978	Υ
piv	982	ϖ
bull	8226	\bullet
hellip	8230	\dots
prime	8242	$'$
Prime	8243	$''$
oline	8254	

↩

(continued)

Name	Value	Symbol
frac1	8260	/
weierp	8472	\wp
image	8465	\Im
real	8476	\Re
trade	8482	TM
alefsym	8501	\aleph
larr	8592	\leftarrow
uarr	8593	\uparrow
rarr	8594	\rightarrow
darr	8595	\downarrow
harr	8596	\leftrightarrow
crarr	8629	\curvearrowright
lArr	8656	\Leftarrow
uArr	8657	\Uparrow
rArr	8658	\Rightarrow
dArr	8659	\Downarrow
hArr	8660	\Leftrightarrow
forall	8704	\forall
part	8706	∂
exist	8707	\exists
empty	8709	\emptyset
nabla	8711	∇
isin	8712	\in
notin	8713	\notin
ni	8715	\ni
prod	8719	\prod
sum	8721	\sum
minus	8722	$-$
lowast	8727	$*$
radic	8730	$\sqrt{\quad}$
prop	8733	\propto
infin	8734	∞
ang	8736	\angle
and	8743	\wedge
or	8744	\vee
cap	8745	\cap

↩

(continued)

Name	Value	Symbol
cup	8746	\cup
int	8747	\int
there4	8756	\therefore
sim	8764	\sim
cong	8773	\cong
asymp	8776	\asymp
ne	8800	\neq
equiv	8801	\equiv
le	8804	\leq
ge	8805	\geq
sub	8834	\subset
sup	8835	\supset
nsub	8836	$\not\subset$
sube	8838	\subseteq
supe	8839	\supseteq
oplus	8853	\oplus
otimes	8855	\otimes
perp	8869	\perp
sdot	8901	\cdot
lceil	8968	\lceil
rceil	8969	\rceil
lfloor	8970	\lfloor
rfloor	8971	\rfloor
loz	9674	\diamond
spades	9824	\spadesuit
clubs	9827	\clubsuit
hearts	9829	\heartsuit
diams	9830	\diamondsuit
lang	10216	\langle
rang	10217	\rangle

Chapter 37

Memory Safety

Memory Safety for a program is defined as it being impossible for the program to corrupt memory. Therefore, the safe subset of D consists only of programming language features that are guaranteed to never result in memory corruption. See this article for a rationale.

Memory-safe code cannot use certain language features, such as:

- Casts that break the type system.
- Modification of pointer values.
- Taking the address of a local variable or function parameter.

Usage

Memory safety can be enabled on a per-function basis using the `@safe` attribute. This can be inferred when the compiler has the function body available. The `@trusted` attribute can be used when a function has a safe interface, but uses unsafe code internally. These functions can be called from `@safe` code.

Array bounds checks are necessary to enforce memory safety, so these are enabled (by default) for `@safe` code even in `-release` mode.

Limitations

Memory safety does not imply that code is portable, uses only sound programming practices, is free of byte order dependencies, or other bugs. It is focussed only on eliminating memory corruption possibilities.

Chapter 38

Application Binary Interface

A D implementation that conforms to the D ABI (Application Binary Interface) will be able to generate libraries, DLL's, etc., that can interoperate with D binaries built by other implementations.

C ABI

The C ABI referred to in this specification means the C Application Binary Interface of the target system. C and D code should be freely linkable together, in particular, D code shall have access to the entire C ABI runtime library.

Endianness

The endianness (byte order) of the layout of the data will conform to the endianness of the target machine. The Intel x86 CPUs are *little endian* meaning that the value 0x0A0B0C0D is stored in memory as: 0D 0C 0B 0A.

Basic Types

bool	8 bit byte with the values 0 for false and 1 for true
byte	8 bit signed value
ubyte	8 bit unsigned value
short	16 bit signed value
ushort	16 bit unsigned value
int	32 bit signed value
uint	32 bit unsigned value
long	64 bit signed value
ulong	64 bit unsigned value
cent	128 bit signed value
ucent	128 bit unsigned value
float	32 bit IEEE 754 floating point value
double	64 bit IEEE 754 floating point value
real	implementation defined floating point value, for x86 it is 80 bit IEEE 754 extended real

Delegates

Delegates are *fat pointers* with two parts:

Delegate Layout		
offset	property	contents
0	.ptr	context pointer
<i>ptrsize</i>	.funcptr	pointer to function

The *context pointer* can be a class *this* reference, a struct *this* pointer, a pointer to a closure (nested functions) or a pointer to an enclosing function's stack frame (nested functions).

Structs

Conforms to the target's C ABI struct layout.

Classes

An object consists of:

Class Object Layout

size	property	contents
<i>ptrsize</i>	<code>._vptr</code>	pointer to vtable
<i>ptrsize</i>	<code>._monitor</code>	monitor
<i>ptrsize...</i>		vptr's for any interfaces implemented by this class in left to right, most to least derived, order
...	...	super's non-static fields and super's interface vptrs, from least to most derived
...	named fields	non-static fields

The vtable consists of:

Virtual Function Pointer Table Layout

size	contents
<i>ptrsize</i>	pointer to instance of TypeInfo
<i>ptrsize...</i>	pointers to virtual member functions

Casting a class object to an interface consists of adding the offset of the interface's corresponding vptr to the address of the base of the object. Casting an interface ptr back to the class type it came from involves getting the correct offset to subtract from it from the object.Interface entry at `vtbl[0]`. Adjustor thunks are created and pointers to them stored in the method entries in the `vtbl[]` in order to set the this pointer to the start of the object instance corresponding to the implementing method.

An adjustor thunk looks like:

```
ADD EAX,offset
JMP method
```

The leftmost side of the inheritance graph of the interfaces all share their vptrs, this is the single inheritance model. Every time the inheritance graph forks (for multiple inheritance) a new vptr is created and stored in the class' instance. Every time a virtual method is overridden, a new `vtbl[]` must be created with the updated method pointers in it.

The class definition:

```
class XXXX
{
    ....
};
```

Generates the following:

- An instance of Class called ClassXXXX.

- A type called `StaticClassXXXX` which defines all the static members.
- An instance of `StaticClassXXXX` called `StaticXXXX` for the static members.

Interfaces

An interface is a pointer to a pointer to a `Vtbl[]`. The `Vtbl[0]` entry is a pointer to the corresponding instance of the `object.Interface` class. The rest of the `Vtbl[1..$]` entries are pointers to the virtual functions implemented by that interface, in the order that they were declared.

A COM interface differs from a regular interface in that there is no `object.Interface` entry in `Vtbl[0]`; the entries `Vtbl[0..$]` are all the virtual function pointers, in the order that they were declared. This matches the COM object layout used by Windows.

A C++ interface differs from a regular interface in that it matches the layout of a C++ class using single inheritance on the target machine.

Arrays

A dynamic array consists of:

Dynamic Array Layout		
offset	property	contents
0	<code>.length</code>	array dimension
<code>size_t</code>	<code>.ptr</code>	pointer to array data

A dynamic array is declared as:

```
type[] array;
```

whereas a static array is declared as:

```
type[dimension] array;
```

Thus, a static array always has the dimension statically available as part of the type, and so it is implemented like in C. Static array's and Dynamic arrays can be easily converted back and forth to each other.

Associative Arrays

Associative arrays consist of a pointer to an opaque, implementation defined type.

The current implementation is contained in and defined by `rt/aaA.d`.

Reference Types

D has reference types, but they are implicit. For example, classes are always referred to by reference; this means that class instances can never reside on the stack or be passed as function parameters.

Name Mangling

D accomplishes typesafe linking by *mangling* a D identifier to include scope and type information.

```
MangledName:
  _D QualifiedName Type
  _D QualifiedName M Type
  _D QualifiedName Z          // Internal
```

```
QualifiedName:
  SymbolName
  SymbolName QualifiedName
```

```
SymbolName:
  LName
  TemplateInstanceName
  0          // anonymous symbols
```

The **M** means that the symbol is a function that requires a `this` pointer.

Template Instance Names have the types and values of its parameters encoded into it:

```
TemplateInstanceName:
  Number TemplateID LName TemplateArgs Z

TemplateID:
  __T
  __U          // for symbols declared inside template constraint

TemplateArgs:
  TemplateArg
  TemplateArg TemplateArgs

TemplateArg:
  TemplateArgX
  H TemplateArgX

TemplateArgX:
  T Type
  V Type Value
  S SymbolName
```

S MangledName

Value:

n
i Number
N Number
e HexFloat
c HexFloat c HexFloat
CharWidth Number _ HexDigits
A Number Value...
S Number Value...

HexFloat:

NAN
INF
NINF
N HexDigits P Exponent
HexDigits P Exponent

Exponent:

N Number
Number

HexDigits:

HexDigit
HexDigit HexDigits

HexDigit:

Digit
A
B
C
D
E
F

CharWidth:

a
w
d

n is for **null** arguments.

Number is for positive numeric literals (including character literals).

N Number is for negative numeric literals.

e HexFloat is for real and imaginary floating point literals.

c HexFloat c HexFloat is for complex floating point literals.

CharWidth Number _ HexDigits *CharWidth* is whether the characters are 1 byte (**a**), 2 bytes (**w**) or 4 bytes (**d**) in size. *Number* is the number of characters in the string. The *HexDigits* are the hex data for the string.

A Number Value... An array or associative array literal. *Number* is the length of the array. *Value* is repeated *Number* times for a normal array, and $2 * \textit{Number}$ times for an associative array.

S Number Value... A struct literal. *Value* is repeated *Number* times.

Name:

Namestart

Namestart Namechars

Namestart:

-

Alpha

Namechar:

Namestart

Digit

Namechars:

Namechar

Namechar Namechars

A *Name* is a standard D identifier.

```

LName:
    Number Name

Number:
    Digit
    Digit Number

Digit:
    0
    1
    2
    3
    4
    5
    6
    7
    8
    9

```

An *LName* is a name preceded by a *Number* giving the number of characters in the *Name*.

Type Mangling

Types are mangled using a simple linear scheme:

```

Type:
    TypeModifiersopt TypeX

TypeX:
    TypeArray
    TypeStaticArray
    TypeAssocArray
    TypePointer
    TypeFunction
    TypeIdent
    TypeClass
    TypeStruct
    TypeEnum
    TypeTypedef

```


TypeDelegate
TypeVoid
TypeByte
TypeUbyte
TypeShort
TypeUshort
TypeInt
TypeUint
TypeLong
TypeUlong
TypeCent
TypeUcent
TypeFloat
TypeDouble
TypeReal
TypeIfloat
TypeIdouble
TypeIreal
TypeCfloat
TypeCdouble
TypeCreal
TypeBool
TypeChar
TypeWchar
TypeDchar
TypeNull
TypeTuple
TypeVector

TypeModifiers:

Const
Wild
Wild Const
Shared
Shared Const
Shared Wild
Shared Wild Const
Immutable

Shared:

```

0

Const:
    x

Immutable:
    y

Wild:
    Ng

TypeArray:
    A Type

TypeStaticArray:
    G Number Type

TypeAssocArray:
    H Type Type

TypePointer:
    P Type

TypeVector:
    Nh Type

TypeFunction:
    CallConvention FuncAttrsopt Parameters ParamClose Type

CallConvention:
    F      // D
    U      // C
    W      // Windows
    V      // Pascal
    R      // C++
    Y      // Objective-C

FuncAttrs:
    FuncAttr
    FuncAttr FuncAttrs

```

FuncAttr:

empty

FuncAttrPure

FuncAttrNogc

FuncAttrNothrow

FuncAttrProperty

FuncAttrRef

FuncAttrReturn

FuncAttrScope

FuncAttrTrusted

FuncAttrSafe

FuncAttrPure:

Na

FuncAttrNogc:

Ni

FuncAttrNothrow:

Nb

FuncAttrProperty:

Nd

FuncAttrRef:

Nc

FuncAttrReturn:

Nj

FuncAttrScope:

M

FuncAttrTrusted:

Ne

FuncAttrSafe:

Nf

```

Parameters:
    Parameter
    Parameter Parameters

Parameter:
    Parameter2
    M Parameter2      // scope

Parameter2:
    Type
    J Type            // out
    K Type            // ref
    L Type            // lazy

ParamClose
    X      // variadic T t...) style
    Y      // variadic T t,...) style
    Z      // not variadic

TypeIdent:
    I QualifiedName

TypeClass:
    C QualifiedName

TypeStruct:
    S QualifiedName

TypeEnum:
    E QualifiedName

TypeTypedef:
    T QualifiedName

TypeDelegate:
    D TypeModifiersopt TypeFunction

TypeVoid:
    v

```

TypeByte:
g

TypeUbyte:
h

TypeShort:
s

TypeUshort:
t

TypeInt:
i

TypeUint:
k

TypeLong:
l

TypeUlong:
m

TypeCent:
zi

TypeUcent:
zk

TypeFloat:
f

TypeDouble:
d

TypeReal:
e

TypeIfloat:

```
o
TypeIdouble:
p
TypeIreal:
j
TypeCfloat:
q
TypeCdouble:
r
TypeCreal:
c
TypeBool:
b
TypeChar:
a
TypeWchar:
u
TypeDchar:
w
TypeNull:
n
TypeTuple:
B Number Parameters
```

Function Calling Conventions

The `extern (C)` and `extern (D)` calling convention matches the C calling convention used by the supported C compiler on the host system. Except that the `extern (D)` calling convention for

Windows x86 is described here.

Register Conventions

- EAX, ECX, EDX are scratch registers and can be destroyed by a function.
- EBX, ESI, EDI, EBP must be preserved across function calls.
- EFLAGS is assumed destroyed across function calls, except for the direction flag which must be forward.
- The FPU stack must be empty when calling a function.
- The FPU control word must be preserved across function calls.
- Floating point return values are returned on the FPU stack. These must be cleaned off by the caller, even if they are not used.

Return Value

- The types bool, byte, ubyte, short, ushort, int, uint, pointer, Object, and interfaces are returned in EAX.
- long and ulong are returned in EDX,EAX, where EDX gets the most significant half.
- float, double, real, ifloat, idouble, ireal are returned in ST0.
- cfloat, cdouble, creal are returned in ST1,ST0 where ST1 is the real part and ST0 is the imaginary part.
- Dynamic arrays are returned with the pointer in EDX and the length in EAX.
- Associative arrays are returned in EAX.
- References are returned as pointers in EAX.
- Delegates are returned with the pointer to the function in EDX and the context pointer in EAX.
- 1, 2 and 4 byte structs and static arrays are returned in EAX.
- 8 byte structs and static arrays are returned in EDX,EAX, where EDX gets the most significant half.
- For other sized structs and static arrays, the return value is stored through a hidden pointer passed as an argument to the function.
- Constructors return the this pointer in EAX.

Parameters

The parameters to the non-variadic function:

```
foo(a1, a2, ..., an);
```

are passed as follows:

```

a1
a2
...
an
hidden
this

```

where *hidden* is present if needed to return a struct value, and *this* is present if needed as the this pointer for a member function or the context pointer for a nested function.

The last parameter is passed in EAX rather than being pushed on the stack if the following conditions are met:

- It fits in EAX.
- It is not a 3 byte struct.
- It is not a floating point type.

Parameters are always pushed as multiples of 4 bytes, rounding upwards, so the stack is always aligned on 4 byte boundaries. They are pushed most significant first. **out** and **ref** are passed as pointers. Static arrays are passed as pointers to their first element. On Windows, a real is pushed as a 10 byte quantity, a creal is pushed as a 20 byte quantity. On Linux, a real is pushed as a 12 byte quantity, a creal is pushed as two 12 byte quantities. The extra two bytes of pad occupy the ‘most significant’ position.

The callee cleans the stack.

The parameters to the variadic function:

```

void foo(int p1, int p2, int[] p3...)
foo(a1, a2, ..., an);

```

are passed as follows:

```

p1
p2
a3
hidden
this

```

The variadic part is converted to a dynamic array and the rest is the same as for non-variadic functions.

The parameters to the variadic function:

```

void foo(int p1, int p2, ...)
foo(a1, a2, a3, ..., an);

```

are passed as follows:


```

an
...
a3
a2
a1
_arguments
hidden
this

```

The caller is expected to clean the stack. `_argptr` is not passed, it is computed by the callee.

Exception Handling

Windows

Conforms to the Microsoft Windows Structured Exception Handling conventions.

Linux, FreeBSD and OS X

Uses static address range/handler tables. It is not compatible with the ELF/Mach-O exception handling tables. The stack is walked assuming it uses the EBP/RBP stack frame convention. The EBP/RBP convention must be used for every function that has an associated EH (Exception Handler) table.

For each function that has exception handlers, an EH table entry is generated.

EH Table Entry	
field	description
<code>void*</code>	pointer to start of function
<code>DHandlerTable*</code>	pointer to corresponding EH data
<code>uint</code>	size in bytes of the function

The EH table entries are placed into the following special segments, which are concatenated by the linker.

EH Table Segment	
Operating System	Segment Name
Windows	<code>FI</code>
Linux	<code>.deh_eh</code>
FreeBSD	<code>.deh_eh</code>
OS X	<code>__deh_eh, __DATA</code>

The rest of the EH data can be placed anywhere, it is immutable.

DHandlerTable	
field	description
<code>void*</code>	pointer to start of function
<code>uint</code>	offset of ESP/RSP from EBP/RBP
<code>uint</code>	offset from start of function to return code
<code>uint</code>	number of entries in <code>DHandlerInfo[]</code>
<code>DHandlerInfo[]</code>	array of handler information

DHandlerInfo	
field	description
<code>uint</code>	offset from function address to start of guarded section
<code>uint</code>	offset of end of guarded section
<code>int</code>	previous table index
<code>uint</code>	if != 0 offset to <code>DCatchInfo</code> data from start of table
<code>void*</code>	if not null, pointer to finally code to execute

DCatchInfo	
field	description
<code>uint</code>	number of entries in <code>DCatchBlock[]</code>
<code>DCatchBlock[]</code>	array of catch information

DCatchBlock	
field	description
<code>ClassInfo</code>	catch type
<code>uint</code>	offset from EBP/RBP to catch variable
<code>void*</code>	catch handler code

Garbage Collection

The interface to this is found in `phobos/internal/gc`.

Runtime Helper Functions

These are found in `phobos/internal`.

Module Initialization and Termination

All the static constructors for a module are aggregated into a single function, and a pointer to that function is inserted into the ctor member of the ModuleInfo instance for that module.

All the static denstructors for a module are aggregated into a single function, and a pointer to that function is inserted into the dtor member of the ModuleInfo instance for that module.

Unit Testing

All the unit tests for a module are aggregated into a single function, and a pointer to that function is inserted into the unitTest member of the ModuleInfo instance for that module.

38.1 Symbolic Debugging

D has types that are not represented in existing C or C++ debuggers. These are dynamic arrays, associative arrays, and delegates. Representing these types as structs causes problems because function calling conventions for structs are often different than that for these types, which causes C/C++ debuggers to misrepresent things. For these debuggers, they are represented as a C type which does match the calling conventions for the type. The **dmd** compiler will generate only C symbolic type info with the **-gc** compiler switch.

Types for C Debuggers	
D type	C representation
dynamic array	<code>unsigned long long</code>
associative array	<code>void*</code>
delegate	<code>long long</code>
dchar	<code>unsigned long</code>

For debuggers that can be modified to accept new types, the following extensions help them fully support the types.

Codeview Debugger Extensions

The D **dchar** type is represented by the special primitive type 0x78.

D makes use of the Codeview OEM generic type record indicated by **LF_OEM** (0x0015). The format is:

Codeview OEM Extensions for D						
field size	2	2	2	2	2	2
D Type	Leaf Index	OEM Identifier	recOEM	num indices	type index	type index
dynamic array	LF_OEM	<i>OEM</i>	1	2	@ <i>index</i>	@ <i>element</i>
associative array	LF_OEM	<i>OEM</i>	2	2	@ <i>key</i>	@ <i>element</i>
delegate	LF_OEM	<i>OEM</i>	3	2	@ <i>this</i>	@ <i>function</i>

Where:

<i>OEM</i>	0x42
<i>index</i>	type index of array index
<i>key</i>	type index of key
<i>element</i>	type index of array element
<i>this</i>	type index of context pointer
<i>function</i>	type index of function

These extensions can be pretty-printed by obj2asm.

The Ddbg debugger supports them.

Chapter 39

Vector Extensions

Modern CPUs often support specialized vector types and vector operations, sometimes called "media instructions". Vector types are a fixed array of floating or integer types, and vector operations operate simultaneously on them, thus achieving great speed-ups.

When the compiler takes advantage of these instructions with standard D code to speed up loops over arithmetic data, this is called auto-vectorization. Auto-vectorization, however, has had only limited success and has not been able to really take advantage of the richness (and often quiriness) of the native vector instructions.

D has the array operation notation, such as:

```
int[] a,b;  
...  
a[] += b[];
```

which can be vectorized by the compiler, but again success is limited for the same reason auto-vectorization is.

The difficulties with trying to use vector instructions on regular arrays are:

1. The vector types have stringent alignment requirements that are not and cannot be met by conventional arrays.
2. C ABI's often have vector extensions and have special name mangling for them, call/return conventions, and symbolic debug support.
3. The only way to get at the full vector instruction set would be to use inline assembler - but the compiler cannot do register allocation across inline assembler blocks (or other optimizations), leading to poor code performance.
4. Interleaving conventional array code with vector operations on the same data can unwittingly lead to extremely poor runtime performance.

These issues are cleared up by using special vector types.

39.1 core.simd

Vector types and operations are introduced to D code by importing `core.simd`:

```
import core.simd;
```

These types and operations will be the ones defined for the architecture the compiler is targeting. If a particular CPU family has varying support for vector types, an additional runtime check may be necessary. The compiler does not emit runtime checks; those must be done by the programmer.

Depending on the architecture, compiler flags may be required to activate support for SIMD types.

The types defined will all follow the naming convention:

typeNN

where *type* is the vector element type and *NN* is the number of those elements in the vector type. The type names will not be keywords.

Properties

Vector types have the property:

Vector Type Properties	
Property	Description
.array	Returns static array representation

All the properties of the static array representation also work.

Conversions

Vector types of the same size can be implicitly converted among each other. Vector types can be cast to the static array representation.

Integers and floating point values can be implicitly converted to their vector equivalents:

```
int4 v = 7;
v = 3 + v;    // add 3 to each element in v
```

Accessing Individual Vector Elements

They cannot be accessed directly, but can be when converted to an array type:

```
int4 v;
(cast(int*)&v)[3] = 2;    // set 3rd element of the 4 int vector
(cast(int[4])v)[3] = 2;  // set 3rd element of the 4 int vector
v.array[3] = 2;          // set 3rd element of the 4 int vector
v.ptr[3] = 2;            // set 3rd element of the 4 int vector
```

Conditional Compilation

If vector extensions are implemented, the version identifier `D_SIMD` is set.

Whether a type exists or not can be tested at compile time with an *IsExpression*:

```
static if (is(typeNN))
    ... yes, it is supported ...
else
    ... nope, use workaround ...
```

Whether a particular operation on a type is supported can be tested at compile time with:

```
float4 a,b;
static if (__traits(compiles, a+b))
    ... yes, it is supported ...
else
    ... nope, use workaround ...
```

For runtime testing to see if certain vector instructions are available, see the functions in `core.cpuid`.

A typical workaround would be to use array vector operations instead:

```
float4 a,b;
static if (__traits(compiles, a/b))
    c = a / b;
else
    c[] = a[] / b[];
```

39.2 X86 And X86_64 Vector Extension Implementation

The rest of this document describes the specific implementation of the vector types for the X86 and X86_64 architectures.

The vector extensions are currently implemented for the OS X 32 bit target, and all 64 bit targets.

`core.simd` defines the following types:

Vector Types		
Type Name	Description	gcc Equivalent
<code>void16</code>	16 bytes of untyped data	<i>no equivalent</i>
<code>byte16</code>	16 <code>byte</code> 's	<code>signed char __attribute__((vector_size(16)))</code>
<code>ubyte16</code>	16 <code>ubyte</code> 's	<code>unsigned char __attribute__((vector_size(16)))</code>
<code>short8</code>	8 <code>short</code> 's	<code>short __attribute__((vector_size(16)))</code>
<code>ushort8</code>	8 <code>ushort</code> 's	<code>ushort __attribute__((vector_size(16)))</code>
<code>int4</code>	4 <code>int</code> 's	<code>int __attribute__((vector_size(16)))</code>
<code>uint4</code>	4 <code>uint</code> 's	<code>unsigned __attribute__((vector_size(16)))</code>
<code>long2</code>	2 <code>long</code> 's	<code>long __attribute__((vector_size(16)))</code>
<code>ulong2</code>	2 <code>ulong</code> 's	<code>unsigned long __attribute__((vector_size(16)))</code>
<code>float4</code>	4 <code>float</code> 's	<code>float __attribute__((vector_size(16)))</code>
<code>double2</code>	2 <code>double</code> 's	<code>double __attribute__((vector_size(16)))</code>
<code>void32</code>	32 bytes of untyped data	<i>no equivalent</i>
<code>byte32</code>	32 <code>byte</code> 's	<code>signed char __attribute__((vector_size(32)))</code>
<code>ubyte32</code>	32 <code>ubyte</code> 's	<code>unsigned char __attribute__((vector_size(32)))</code>
<code>short16</code>	16 <code>short</code> 's	<code>short __attribute__((vector_size(32)))</code>
<code>ushort16</code>	16 <code>ushort</code> 's	<code>ushort __attribute__((vector_size(32)))</code>
<code>int8</code>	8 <code>int</code> 's	<code>int __attribute__((vector_size(32)))</code>
<code>uint8</code>	8 <code>uint</code> 's	<code>unsigned __attribute__((vector_size(32)))</code>
<code>long4</code>	4 <code>long</code> 's	<code>long __attribute__((vector_size(32)))</code>
<code>ulong4</code>	4 <code>ulong</code> 's	<code>unsigned long __attribute__((vector_size(32)))</code>
<code>float8</code>	8 <code>float</code> 's	<code>float __attribute__((vector_size(32)))</code>
<code>double4</code>	4 <code>double</code> 's	<code>double __attribute__((vector_size(32)))</code>

Note: for 32 bit gcc, it's `long long` instead of `long`.

Supported 128-bit Vector Operators

[illegible]

Supported 256-bit Vector Operators

Operator	void32	byte32	ubyte32	short16	ushort16	int8	uint8	long4	ulong4	float8	double4
=	×	×	×	×	×	×	×	×	×	×	×
+	—	×	×	×	×	×	×	×	×	×	×
-	—	×	×	×	×	×	×	×	×	×	×
*	—	—	—	—	—	—	—	—	—	×	×
/	—	—	—	—	—	—	—	—	—	×	×
&	—	×	×	×	×	×	×	×	×	—	—
	—	×	×	×	×	×	×	×	×	—	—
^	—	×	×	×	×	×	×	×	×	—	—
+=	—	×	×	×	×	×	×	×	×	×	×
-=	—	×	×	×	×	×	×	×	×	×	×
*=	—	—	—	—	—	—	—	—	—	×	×
/=	—	—	—	—	—	—	—	—	—	×	×
&=	—	×	×	×	×	×	×	×	×	—	—
=	—	×	×	×	×	×	×	×	×	—	—
^=	—	×	×	×	×	×	×	×	×	—	—
<i>unary~</i>	—	×	×	×	×	×	×	×	×	—	—
<i>unary+</i>	—	×	×	×	×	×	×	×	×	×	×
<i>unary-</i>	—	×	×	×	×	×	×	×	×	×	×

Operators not listed are not supported at all.

Vector Operation Intrinsics

See `core.simd` for the supported intrinsics.